

1 Sharing sunlight between trackers and crops: dynamic
2 light-sharing control for agrivoltaics under a
3 food-security floor on the world's staple croplands

4 Quang-Vinh Dang*

5 *British University Vietnam, Hung Yen, Vietnam*

6 **Abstract**

Single-axis trackers let agrivoltaic arrays shift the daily rotation from maximum electricity towards letting light through (light sharing). We quantify its value under a food-security floor of 90% of open-field yield with a differentiable simulator of the whole system. The crop response to shade is the least certain part. A hierarchical model fitted to 428 observations from 58 field studies predicts a new study of the four staple crop types only to about 0.22 in relative yield, no better than proportional loss, and we turned the learned curves and four field trials into an ensemble of plausible responses by gradient descent through the simulator. In a deployable protocol (density chosen from each site's history, held-out sites and years, 400 crop-site pairs), light-sharing rules allowed arrays three times as dense as static arrays at equal retention and raised the land-equivalent ratio (LER) from 1.07 to 1.18–1.24; a phenology rule needing no forecast was the best. A neural controller trained through the simulator exceeded the written rule by 0.044 LER but a rule tuned on the same training seasons by only 0.018 (95% interval 0.012–0.024); training on the ensemble or adding a farm-measurable biomass signal changed nothing. The design decision mattered more: a density chosen under the mean response left about 70% of crop-site pairs at the floor when another plausible response held (oracle 85–87%), and a floor raised by 0.02 under the mean response did better at equal compliance than a rule built on the ensemble's quantiles. The simulator is too mild for maize and soybean, which supply 53% of the upscaled electricity, so gains for these crops are probably overestimated. Under illustrative costs, light sharing breaks even with the open field at 73 (learned controller) to 92 USD MWh⁻¹, and a conventional plant replacing the crop stays more profitable, so agrivoltaics

*Corresponding author
Preprint submitted to *Artificial Intelligence in Agriculture*
Email address: vinh.dq4@buv.edu.vn (Quang-Vinh Dang)

needs a premium or land-use rules. Equipping 5% of staple cropland would give about 14 PWh yr^{-1} but needs about 10 TWp of modules.

7 *Keywords:* Agrivoltaics, Light sharing, Solar tracking, Food–energy nexus,
8 Crop modelling, Differentiable simulation, Hierarchical shrinkage model

9 1. Introduction

10 Two of the defining challenges of this century compete for the same re-
11 source: land. The energy transition requires terawatts of new solar pho-
12 tovoltaic (PV) capacity, and ground-mounted solar farms need roughly one
13 to two hectares per megawatt. Food production must rise while crops are
14 exposed to more frequent heat waves and droughts. Warming alone is pro-
15 jected to reduce global yields of maize, wheat, rice and soybean by several
16 per cent per degree (Zhao et al., 2017), and a new generation of climate and
17 crop models shows these impacts emerging earlier than previously expected
18 (Jägermeyr et al., 2021). Croplands are also where solar power potential is
19 greatest (Adeh et al., 2019). They are flat, sunny, accessible and close to
20 rural electricity demand, so converting farmland into solar farms is tempting
21 and controversial at the same time.

22 Agrivoltaic (AV) systems combine crops and PV modules on the same
23 land and aim to resolve this conflict (Dupraz et al., 2011; Dinesh and Pearce,
24 2016; Weselek et al., 2019). By sharing sunlight between modules and plants,
25 a hectare can deliver more food plus energy than the two uses side by side,
26 with land-equivalent ratios (LER) well above one in field trials and mod-
27 elling studies (Dupraz et al., 2011; Amaducci et al., 2018; Valle et al., 2017;
28 Trommsdorff et al., 2021). The partial shade also changes the crop micro-
29 climate. It lowers canopy temperature and evaporative demand, conserves
30 soil moisture and improves water-use efficiency, especially in drylands (Mar-
31 rou et al., 2013b; Hassanpour Adeh et al., 2018; Barron-Gafford et al., 2019;
32 Elamri et al., 2018; Weselek et al., 2021). Agrivoltaics is therefore a candi-
33 date adaptation measure against heat and drought stress, in addition to a
34 mitigation measure through renewable electricity. The same shade, however,
35 also reduces the light available for photosynthesis. Regulators increasingly
36 require AV systems to retain most of the reference crop yield. The German
37 pre-standard DIN SPEC 91434, for example, requires at least 66% of the
38 reference yield (DIN, 2021), and other schemes and many farmers set stricter

39 floors. How much electricity can be produced without sacrificing food secu-
40 rity depends on the crop, the local climate, the year and the growth stage.

41 Most AV designs fix the geometry of the array, i.e. the ground coverage
42 ratio (GCR, module area per ground area) and the tilt or tracking rule, from
43 generic rules of thumb. Single-axis trackers add a degree of freedom: the
44 modules can be rotated away from the sun to let light through when the crop
45 needs it, as demonstrated for lettuce by Valle et al. (2017), or towards the
46 sun to maximise electricity. Choosing the rotation every day is a sequential
47 decision problem. The value of light depends on the crop’s phenological stage,
48 soil water and upcoming weather, and the effect of a decision appears in the
49 harvest months later. Recent work in this journal has used evolutionary
50 multi-objective optimisation to lay out PV arrays on farmland (Zhang et al.,
51 2026). Machine learning for crop management has mainly used reinforcement
52 learning with black-box simulators (Gautron et al., 2022), which is sample-
53 inefficient and hard to scale to global, multi-decadal assessments.

54 Two questions follow. First, how much of the electricity that an agri-
55 voltaic system can produce under a food-security floor comes from the array
56 density, and how much from dynamic light-sharing control of the trackers?
57 Second, how robust is that answer to the weakest part of any such study, the
58 crop response to shade, and does a controller learned from a simulator add
59 anything to a well-designed hand-written rule? Answering them requires
60 a simulator that couples the physics of light, PV, microclimate and crop
61 growth, a protocol in which all controllers face the same design constraints,
62 a check of the shade response against field data, and enough sites and seasons
63 to separate real effects from noise.

64 This paper makes three contributions. The first is a quantification of
65 light-sharing control under a yield floor. Light sharing is parameterised by
66 one daily level between full backtracking and the rotation that minimises the
67 shadow on the crop, and we compare static arrays, a seasonal schedule, a rule
68 that follows crop phenology, a rule that reacts to heat and drought stress,
69 and a neural controller (SALS) trained through the simulator. The second
70 is a deployable evaluation protocol. The array density of every site is chosen
71 from its historical seasons only, with the same rule for every method, and
72 performance is measured on held-out sites and years, over three training seeds
73 and three spatial splits. The third is a differentiable agrivoltaic simulator
74 that supports both, coupling hourly solar geometry, single-axis trackers with
75 backtracking, row-to-row shading and PV electricity, the direct and diffuse
76 light reaching the crop, canopy temperature and reference evapotranspiration

under the modules, and the SIMPLE crop model (Zhao et al., 2019) extended with a flowering-stage shade penalty. All components are differentiable tensor operations (Karniadakis et al., 2021). The PV part is benchmarked against PVGIS. The crop response to shade is learned from the open data of a meta-analysis (Laub et al., 2022), checked against field trials on rice, wheat, potato and maize (Gonocruz et al., 2021; Weselek et al., 2021; Ramos-Fuentes et al., 2023), and propagated through the simulator as an ensemble.

We apply the protocol to 400 crop–site pairs that represent where maize, wheat, rice, soybean and potato are grown worldwide, sampled in proportion to their harvested area (Yu et al., 2020) with local crop calendars (Jägermeyr et al., 2021), and driven by NASA POWER daily weather for 2001–2019 (Sparks, 2018) under baseline and +1.5, +2 and +3 °C climates. The main findings are as follows.

1. Light sharing is the main lever. At the same mean yield retention of 90%, arrays with light-sharing control carry about three times the module area of static arrays and produce 1.8–2.4 times as much electricity (Section 3.1).
2. The best hand-written rule follows crop phenology and needs no weather forecast. Rules that react to heat or drought stress are not better and depend on forecast quality (Section 3.3).
3. The crop response to shade can be learned from data. A hierarchical model fitted to 428 observations from 58 field studies predicts a held-out study of the four staple crop types to about 0.22 in relative yield, no better than proportional loss (it is better only for crop types the simulator does not represent), so between-study spread is large. The simulator is close to these curves for cereals and potato but too mild for maize and soybean, so the gains for those two crops are probably overestimated (Section 3.7).
4. A learned controller adds a small gain over the best rule: 0.044 LER over the written phenology rule and 0.018 over a rule tuned on the same training seasons. Training it on an ensemble of shade responses, or giving it a farm-measurable biomass signal, does not change this (Section 3.8).
5. The bottleneck is the uncertainty of the crop response that enters the design. A density chosen under the mean response leaves about 70% of crop–site pairs at the floor when another plausible response holds (oracle 85–87%). A rule built on the quantiles of an ensemble of responses

114 raises this by 9 points for 0.04 LER, but a floor raised by 0.02 under
 115 the mean response does better at equal compliance, so the ensemble is
 116 useful to size a margin and not to replace it (Section 3.8).

- 117 6. Under illustrative costs, light sharing breaks even with the open field at
 118 73 to 92 USD MWh⁻¹, learning is worth about 2 kUSD ha⁻¹ yr⁻¹ over a
 119 tuned rule at 90 USD MWh⁻¹, and a conventional plant replacing the
 120 crop is more profitable at every price (Section 3.9).
- 121 7. At the yield floor, agrivoltaics on 5% of the world’s staple cropland
 122 could supply about 14 PWh yr⁻¹, but only with about 10 TWp of mod-
 123 ules, several times the installed PV capacity of the world (Section 3.11).

124 2. Materials and methods

125 2.1. Global cropland sites, weather and crop calendars

126 We studied five staple crops that together cover about 670 Mha of har-
 127 vested area and 580 Mha of physical cropland: wheat, rice, maize, soybean
 128 and potato. Harvested areas (rainfed and irrigated) were taken from the
 129 SPAM 2010 gridded production maps (Yu et al., 2020) and aggregated to 1°
 130 cells. For each crop, 80 cells with at least 5000 ha of the crop were drawn
 131 without replacement with probability proportional to harvested area. An un-
 132 weighted average over the sample therefore approximates an area-weighted
 133 average over global production. The 400 crop-site pairs lie in 381 distinct
 134 cells in 50 countries (Fig. 2). Sowing and maturity dates came from the
 135 GGCM Phase 3 crop calendar at 0.5° (Jägermeyr et al., 2021). For wheat
 136 we used the winter-wheat calendar unless the coldest-month mean temper-
 137 ature was below −8 °C, in which case the spring-wheat calendar was used.
 138 Daily weather for 2001–2020 was retrieved from NASA POWER (Sparks,
 139 2018): global and diffuse shortwave irradiance, maximum, minimum and
 140 dew-point temperature, precipitation and wind speed at 2 m. Each site con-
 141 tributes 19 seasons (sowing years 2001–2019), each simulated as the 365 days
 142 after sowing, which gives 7600 site-seasons.

143 Climate change was represented by the global warming levels +1.5, +2
 144 and +3 °C relative to 2001–2020. Daily maximum, minimum and dew-point
 145 temperatures were shifted uniformly, which keeps relative humidity approx-
 146 imately constant and raises the vapour-pressure deficit. The atmospheric
 147 CO₂ concentration was set to 390, 450, 500 and 600 ppm for the baseline and
 148 the three warming levels. Precipitation and irradiance were left unchanged.

149 This delta approach isolates the thermal signal of warming; its limitations
150 are discussed in Section 4.5.

151 2.2. A differentiable agrivoltaic simulator

152 The simulator advances all site-seasons in parallel with a daily time step
153 for the crop and an hourly time step for radiation and PV (Fig. 1). All equa-
154 tions are implemented as PyTorch tensor operations (Paszke et al., 2019), so
155 that yields and electricity are differentiable with respect to the design and
156 control variables.

157 *Hourly irradiance.* Daily global and diffuse irradiance were disaggregated
158 into hourly values with the Collares-Pereira and Rabl (1979) and Liu and
159 Jordan (1960) ratios and renormalised to the daily totals. Beam irradiance
160 on the horizontal is the difference between the two. Solar position follows
161 the standard declination and hour-angle equations.

162 *Array geometry and PV electricity.* We consider elevated horizontal single-
163 axis trackers with north-south axes. The axis height is 1.25 module widths,
164 e.g. 2.75 m for the common 2.2-m single-module-high (1P) tracker. The ro-
165 tation ϕ is limited to $\pm 60^\circ$. The ground coverage ratio g is the module width
166 divided by the row pitch and is restricted to $g \leq 0.5$. For 2.2-m modules this
167 keeps a pitch of at least 4.4 m, and a free strip of at least 2.2 m between the
168 rows when the modules are horizontal, which is the stow position during field
169 operations. Let ψ be the solar zenith angle projected onto the east-west ver-
170 tical plane. The energy-optimal rotation with backtracking (Lorenzo et al.,
171 2011) is $\phi_{\text{bt}} = \psi - \text{sgn}(\psi) \arccos(\cos \psi / g)$ when $\cos \psi < g$, and $\phi_{\text{bt}} = \psi$ oth-
172 erwise. The *light-sharing* rotation ϕ_{sh} is the feasible angle closest to edge-on
173 to the sun, which minimises the shadow on the crop. Given a light-sharing
174 level $u \in [0, 1]$ chosen once per day, the rotation in every hour is

$$\phi = (1 - u) \phi_{\text{bt}} + u \phi_{\text{sh}}. \quad (1)$$

175 Plane-of-array irradiance combines three terms: beam irradiance on the mod-
176 ule, reduced by the fraction of the module shaded by the neighbouring row,
177 $\max\{0, 1 - \cos \psi / (g |\cos(\phi - \psi)|)\}$; isotropic sky diffuse; and ground-reflected
178 light (albedo 0.2). The cell temperature follows the NOCT model (45 °C),
179 with an hourly air temperature interpolated sinusoidally between the daily
180 extremes. Electricity per unit land is plane-of-array irradiance $\times \eta \times g$, with
181 $\eta = 0.21 [1 - 0.0035(T_{\text{cell}} - 25)] (1 - 0.14)$ for module efficiency, temperature

182 losses and system losses. The installed capacity is 2.1 g MWp ha^{-1} . A con-
 183 ventional ground-mounted PV plant ($g = 0.4$, backtracking, no crop) is the
 184 energy reference E_{PV} .

185 *Validation of the PV model.* We compared the simulated specific yield of an
 186 isolated tracker row ($g = 0.05$, negligible row-to-row shading) with PVGIS 5.2
 187 (European Commission, Joint Research Centre, 2022), for a horizontal north-
 188 south single-axis tracker with 14% losses, at 30 randomly chosen sites (29
 189 returned data). Across sites, the simulated and PVGIS specific yields were
 190 correlated with $r = 0.971$. The simulator was on average 4.9% lower (mean
 191 absolute deviation 5.5%). This is consistent with its $\pm 60^\circ$ rotation limit and
 192 with the different irradiance databases (NASA POWER vs. SARA5 and NSRDB). Electricity estimates are therefore, if anything, conservative.

194 *Light and microclimate at the crop.* Averaged over the row pitch, the fraction
 195 of ground reached by beam light is $1 - \min\{1, g |\cos(\phi - \psi)| / \cos \psi\}$. The
 196 fraction of sky diffuse reaching the ground is a two-dimensional isotropic
 197 sky-view factor. We pre-computed it for 25 values of g and 13 rotation
 198 angles by ray casting over 31 neighbouring rows, and interpolate it bilinearly
 199 (differentiably). The daily ratio s of the irradiance at the crop to that of the
 200 open field defines the crop’s shading. The canopy maximum temperature is
 201 lowered by $\kappa(1 - s)$.

202 κ is uncertain. Field studies under PV modules report small reductions
 203 of daytime crop or canopy temperature and little change in air temperature
 204 (Marrou et al., 2013b; Barron-Gafford et al., 2019), and heat-stress thresholds
 205 calibrated on air temperature do not transfer directly to canopy temperature
 206 (Siebert et al., 2014; Webber et al., 2018). We therefore use a moderate
 207 $\kappa = 1^\circ\text{C}$ as the main case and report $\kappa = 0$ (no cooling) and $\kappa = 2^\circ\text{C}$
 208 for all methods (Section 3.2). Nocturnal warming under the modules is not
 209 represented. The reference evapotranspiration under the modules, ET_0 , is re-
 210 computed with the FAO-56 Penman–Monteith equation (Allen et al., 1998)
 211 from the reduced radiation and canopy temperature. Soil water conserva-
 212 tion under the modules (Hassanpour Adeg et al., 2018; Elamri et al., 2018)
 213 therefore emerges from the model, but rain redistribution by the modules
 214 (Marrou et al., 2013a) is ignored.

215 *Crop growth.* Crop growth follows the SIMPLE model (Zhao et al., 2019)
 216 with the species parameters of its Table 1a (Table 1). Daily biomass growth

217 is

$$\Delta B = R_{\text{crop}} \cdot \text{RUE} \cdot f_{\text{Solar}} \cdot f_{\text{CO}_2} \cdot f_{\text{Temp}} \cdot \min\{f_{\text{Heat}}, f_{\text{Water}}\}, \quad (2)$$

218 where $R_{\text{crop}} = sR$ is the radiation received by the crop. The terms are as
219 follows.

- 220 • f_{Solar} is the canopy’s radiation interception. In SIMPLE it depends
221 only on thermal time, with senescence accelerated by heat and drought
222 stress. Shade therefore reduces the daily biomass increment but not
223 leaf-area development.
- 224 • f_{Heat} decreases linearly between the heat threshold T_{heat} and the ex-
225 treme temperature T_{ext} of the canopy maximum temperature.
- 226 • $f_{\text{Water}} = 1 - S_{\text{water}} \text{ARID}$, where the agricultural reference index for
227 drought (Woli et al., 2012) is computed from a daily soil-water balance
228 with curve-number runoff, drainage and root water uptake. Irrigated
229 fields have $f_{\text{Water}} = 1$.

230 SIMPLE has a fixed harvest index and therefore ignores the high sen-
231 sitivity of grain number to shading around flowering (Fischer, 1985; An-
232 drade et al., 1999). It also ignores the differences in shade tolerance be-
233 tween crop types. In the meta-analysis of Laub et al. (2022), maize and
234 grain legumes lost yield strongly even at low shade levels, whereas C3 ce-
235 reals and tuber crops initially lost less than proportionally; shading trials
236 also show marked effects of shade on maize development (Ramos-Fuentes
237 et al., 2023). We added a stage-dependent penalty. During a critical window
238 ($0.45 \leq \text{TT}/T_{\text{sum}} \leq 0.65$), the harvest index is multiplied by $1 - h\bar{d}$, where $\bar{d}$
239 is the mean relative radiation deficit $1 - s$ in that window. The sensitivity h
240 is highest for maize (0.6) and soybean (0.5), the crop types that Laub et al.
241 (2022) found most shade-sensitive, lower for the C3 cereals (wheat 0.4; rice
242 0.5, for which no specific curve was available) and lowest for potato (0.2).
243 These values are assumptions that reproduce this ranking; they are not cal-
244 ibrated estimates. Section 3.6 shows the resulting shade-response curves.
245 Because RUE is linear, SIMPLE does not credit the higher light-use effi-
246 ciency of shaded canopies under more diffuse light. The penalty, conversely,
247 makes shading around flowering more damaging than proportional. The net
248 effect is closer to the field evidence than either assumption alone.

249 The thermal-time requirement to maturity was calibrated for every site as
250 the mean thermal time between the GGCM sowing and maturity dates under

the 2001–2020 climate, and the canopy parameters were scaled accordingly. The cultivar is thus adapted to the present climate, and warming shortens the season. Growth stops smoothly once the thermal-time requirement is met. Yield is biomass multiplied by the (shade-adjusted) harvest index. Each site is represented by a rainfed and an irrigated sub-field in proportion to its SPAM irrigated fraction. Each sub-field carries its own array, so both yield *and* electricity are area-weighted over the two sub-fields.

2.3. SALS: simulation-trained adaptive light sharing

SALS learns a *control network* $u_t = \pi_\omega(\mathbf{x}_t; g)$ (Algorithm 1). It is a multilayer perceptron with two hidden layers of 64 units and a bounded-logit output, $u = \sigma(5 \tanh(z/5)) \in [0.007, 0.993]$, which keeps gradients non-zero at the extremes. It chooses the daily light-sharing level of Eq. (1). Its inputs are the information available on day t :

- the crop and soil state: relative thermal time TT/T_{sum} , radiation interception f_{Solar} , relative soil water and the previous day’s ARID;
- a *forecast* of the day’s maximum temperature (relative to T_{heat} and T_{opt}), irradiance and ET_0 ;
- static information: the irrigation status of the sub-field, the array’s GCR g , latitude, day of season and the crop.

Forecasts are perturbed with random errors, independently every day: Gaussian with a standard deviation of 2°C for maximum temperature, and multiplicative with standard deviations of 20% for irradiance and 15% for ET_0 . The physics always uses the true weather. The same forecast errors are applied to the rule-based controllers. The controller is conditioned on g . During training, each season is simulated with a GCR close to the density that is sustainable at its training site. This is the largest value on a coarse grid $(0.05, 0.10, \dots, 0.50)$ at which seasonal sharing keeps the mean retention of the site’s training seasons at or above ρ , plus a uniform perturbation in $[-0.05, +0.10]$ truncated to $[0.05, 0.50]$. The network therefore learns to control the densities that design selection will later pick. Outside the cropping season the trackers backtrack for maximum electricity.

For site-season i with open-field yield Y_i^{open} , let $r_i = Y_i/Y_i^{\text{open}}$ be the yield retention and $e_i = E_i/E_{\text{PV},i}$ the electricity relative to a conventional

284 PV plant. Training minimises the penalised objective

$$\mathcal{L}(\omega) = \mathbb{E}_{i \sim \mathcal{B}} \left[-e_i + \mu_1 \max\{0, \rho - r_i\} + \mu_2 \max\{0, \rho - r_i\}^2 \right], \quad \rho = 0.9, \mu_1 = 5, \mu_2 = 30, \quad (3)$$

285 which combines an exact (hinge) and a quadratic penalty for the food-security
 286 floor ρ . Because r_i and e_i are outputs of the differentiable simulator, the
 287 exact gradient of Eq. (3) is obtained by back-propagation through the 365
 288 daily steps of every season (back-propagation through time). This is the key
 289 difference from reinforcement learning with black-box simulators (Gautron
 290 et al., 2022), which estimates gradients from sampled returns.

291 Each minibatch contains 512 site-seasons, each simulated as a rainfed and
 292 an irrigated sub-field, drawn from one of the four climates in turn. Seasons
 293 with a failed open-field crop ($Y^{\text{open}} < 0.2 \text{ t ha}^{-1}$) are excluded from training
 294 and evaluation. The network is fine-tuned through the simulator with Adam
 295 (Kingma and Ba, 2015) (learning rate 10^{-3} , cosine decay, gradient-norm
 296 clipping at 1) for 300 steps. Growth-stage logistic functions are evaluated in
 297 a numerically stable form, so that gradients remain finite even under extreme
 298 heat.

299 *Initialisation.* Trained from random weights, the controller converged to
 300 sharing light throughout the cropping window ($u > 0.99$ on all days), i.e. to
 301 the seasonal-sharing rule, because the gradient of electricity with respect to u
 302 is small and of the same sign for most days of a season. We therefore initialise
 303 the network by behaviour cloning of the phenology rule of Section 2.4: binary
 304 cross-entropy against the rule’s output (soft targets 0.95 and 0.05) on about
 305 1.7 million day-states visited by the rule in the training seasons, with 1500
 306 Adam steps at a learning rate of 3×10^{-3} . Fine-tuning then moves the
 307 policy away from this initialisation wherever doing so improves the objective.
 308 Consequently, the gain of SALS over the phenology rule measures what fine-
 309 tuning through the simulator adds to that rule, not what can be learned
 310 without it.

311 2.4. Deployable design selection and baselines

312 A real array has one fixed density, which must be chosen before it is built.
 313 Every method, SALS included, therefore uses the same *history-based design*
 314 *rule*. For a site, the GCR is the largest value on the grid $\{0.05, 0.075, \dots, 0.50\}$
 315 for which the method’s mean yield retention over the site’s 2001–2014 sea-
 316 sons is at least $\rho = 0.9$. If no value qualifies, no array is built and the site
 317 remains open field. A stricter *chance-constrained* variant requires $r \geq 0.9$ in

Algorithm 1 SALS training

Require: training site-seasons in four climates; floor ρ ; penalties μ_1, μ_2

- 1: initialise π_ω by behaviour cloning of the phenology rule
- 2: **for** step = $1, \dots, N$ **do**
- 3: sample a climate and a minibatch $\mathcal{B}$ of site-seasons (rainfed and irrigated sub-fields)
- 4: set g_i to the site’s sustainable density plus a perturbation; initialise crop and soil states
- 5: **for** day $t = 1, \dots, 365$ **do**
- 6: $u_{it} \leftarrow \pi_\omega(\mathbf{x}_{it}; g_i)$ from state and noisy forecast if the crop is in the field, else 0
- 7: hourly rotation (Eq. 1), PV electricity, crop light, canopy temperature, ET_0
- 8: advance the soil-water balance and the SIMPLE crop state (Eq. 2)
- 9: $r_i \leftarrow Y_i/Y_i^{\text{open}}, e_i \leftarrow E_i/E_{\text{PV},i}$ (area-weighted); loss $\mathcal{L}$ (Eq. 3)
- 10: back-propagate through the season; Adam update of ω

318 at least 80% of the historical seasons. The historical seasons of the test sites
319 are never used for training, and the weather of the evaluation years is never
320 used for design. We compare four controllers:

- 321 • *AV static*: backtracking trackers without light sharing ($u = 0$);
- 322 • *Seasonal sharing*: the light-sharing rotation ($u = 1$) throughout the
323 cropping season;
- 324 • *Stress rule*: an expert rule that shares light ($u = 1$) during the main
325 growth phase ($0.1 < TT/T_{\text{sum}} < 0.9$) except on days with a heat-stress
326 forecast ($T_{\text{max}} \geq T_{\text{heat}}$) or drought ($\text{ARID} > 0.5$), when the modules
327 track the sun and shade the crop;
- 328 • *SALS*: the learned controller.

329 The open field and a conventional PV plant without crops ($\text{LER} = 1$ each)
330 are the references. Because the design rule is identical for all controllers, dif-
331 ferences between methods reflect the quality of control and the array density
332 that each controller can sustain under the food-security floor. We also com-
333 pare the controllers at identical GCR over the whole grid (matched-density
334 frontiers).

335 2.5. A shade-response model learned from field data

336 The simulator needs a crop response to shade that is supported by ob-
 337 servations. We learned one from the open data of the meta-analysis of Laub
 338 et al. (2022) (Zenodo 10.5281/zenodo.5716091, CC BY 4.0): 428 observa-
 339 tions of yield relative to an unshaded control at relative shade reductions of
 340 0–90% from 58 studies, with crop and crop type, experiment type (artificial
 341 shading or intercropping), climate zone, irradiation and shade type. Rela-
 342 tive yield is modelled as $y(s) = \exp(-\beta s^\gamma)$ with Student- t errors ($\nu = 4$;
 343 maximum a posteriori estimation, so this is a hierarchical shrinkage regres-
 344 sion with bootstrap intervals and not a full Bayesian model). We compared
 345 six models: proportional loss $y = 1 - s$; one curve for all crops; one curve
 346 per crop type without pooling; one curve per crop type with partial pooling
 347 (Gaussian shrinkage of $\log \beta$ and $\log \gamma$ towards a common value, scale 0.5);
 348 partial pooling with covariates on $\log \beta$ (standardised irradiation, subtrop-
 349 ical zone, solar-panel shade and intercropping); and gradient-boosted trees
 350 on the same inputs. Fits use L-BFGS. Models were compared by leave-one-
 351 study-out cross-validation, in which every study is held out once, which tests
 352 the prediction of a new study and not of a new observation in a known study.
 353 Intervals for the crop-type curves and covariate effects come from a cluster
 354 bootstrap over studies (150 resamples).

355 2.6. Shade-response uncertainty: gradient calibration, robust controllers and 356 design

357 The crop’s response to shade is the least certain part of the simulator,
 358 so we treat it as an unknown with a distribution. Two parameters per crop
 359 enter the simulator: a radiation-use compensation c (growth is multiplied by
 360 $1 + c(1 - s)$, s being the share of open-field irradiance that reaches the crop)
 361 and a multiplier h on the flowering-window harvest-index penalty. Because
 362 the simulator is differentiable, both can be fitted by gradient descent through
 363 whole simulated seasons to published yield–shading observations. The ob-
 364 servations are the relative yields at 20% and 40% shading of the data-driven
 365 curves of Section 2.5 for the five crops (median, with an error equal to the
 366 width of the 80% bootstrap interval divided by 2.56 and at least 0.04; rice
 367 takes the curve of C3 cereals, which contains no rice), the field trials on rice,
 368 wheat, potato and maize (Gonocruz et al., 2021; Weselek et al., 2021; Ramos-
 369 Fuentes et al., 2023) (each season of each trial as a separate observation, error
 370 0.06–0.08) and the unaffected harvest index of shaded maize (Ramos-Fuentes
 371 et al., 2023). We call the resulting ensemble E2. A wider ensemble, E3, uses

the same procedure with a target error of 0.2 for the learned curves, which is the between-study spread of a new study (Section 3.7). An earlier ensemble, E1, was fitted in the same way to curves read by hand from the figures of Laub et al. (2022) and was used to train the robust and adaptive controllers below. Priors are $c \sim \mathcal{N}(0.4, 0.5^2)$ and $h \sim \mathcal{N}(1, 0.5^2)$, and the parameters are clamped to $[0, 2.5]$ (c) and $[0, 2]$ (h) after every step, which is why a response milder than the data cannot be lowered further. We fitted 12 ensemble members by randomised maximum a posteriori estimation: every member sees the observations perturbed by their errors and a prior draw perturbed by its width, and is optimised with Adam (40 steps, with the loss still falling slightly at the last step; 40 training-block seasons per crop and member; shading of the static arrays cached). The ensemble is a device for generating plausible responses and not a validated posterior: it has no coverage check, and for some parameters (the harvest-index multiplier of wheat and potato) its spread equals that of the prior. The design analyses use five fixed members (indices 0, 3, 5, 8 and 10 of 12), chosen in advance. The ensemble is used in three ways.

Robust and adaptive controllers. SALS-robust has the architecture and training of SALS (Section 2.3), but every training season is simulated with the response of a randomly drawn ensemble member that the controller does not observe, and the density jitter is widened to $[-0.05, +0.25]$ because feasible densities differ between members. SALS-adaptive additionally observes the realised biomass of the crop under the modules relative to an open reference plot, a quantity that can be measured on the farm by remote sensing or by a reference plot. A feedback rule that shares light whenever this ratio falls below a target serves as the hand-written counterpart; its target (1.0) was chosen on training-block seasons under the ensemble.

Uncertainty-aware design. The design rule of Section 2.4 chooses the largest GCR whose historical retention meets the floor. Under shade-response uncertainty we choose the largest GCR that meets the floor in at least 75% of five ensemble members, and compare it with a design under the ensemble mean and with an oracle design under the true response. Evaluation of this rule is leave-one-out: the truth is the held-out member and the design uses the other four.

Stress tests. Controllers and rules are also evaluated under four fixed responses, the proportional response, the response calibrated to the meta-analysis, the response fitted to the field trials and a harsh response (proportional growth and a 1.5 times larger harvest-index penalty). The proportional

410 and harsh responses lie outside the range of the ensemble.

411 2.7. Economic analysis

412 To judge whether the electricity gains matter economically, we computed
413 an annual net value per hectare for each method relative to the open field:
414 electricity revenue minus the annuity and operating cost of the installed ca-
415 pacity minus the loss of crop revenue (crop production costs are unchanged).
416 The installed capacity is 2.1 g MWp per hectare (modules cover a share g of
417 the land at 0.21 kWp m⁻²). Yield is dry matter, valued at assumed prices.
418 A conventional ground-mounted plant (GCR 0.4) that replaces the crop is
419 the comparator. All parameters (Table 13) are assumptions chosen to be
420 representative, not data, and their ranges drive the conclusions; we therefore
421 report break-even electricity prices and one-at-a-time sensitivities, and we do
422 not monetise food security, which market crop prices do not capture.

423 2.8. Evaluation protocol, statistics and upscaling

424 Sites were split spatially: 10°×10° blocks were randomly assigned to train-
425 ing (70%) or test (30%), so that no test site has a training site in the same
426 block. Training uses the 2001–2014 seasons of the training sites. Design
427 selection uses the 2001–2014 seasons of the test sites. Evaluation uses the
428 2015–2019 seasons of the test sites. The same site-seasons are analysed in all
429 four climates.

430 We report the following metrics:

- 431 • the mean electricity (MWh ha⁻¹ yr⁻¹) and the relative electricity e ;
- 432 • the mean yield retention r and the production-weighted retention $\sum Y / \sum Y^{\text{open}}$;
- 433 • season-level compliance (share of seasons with $r \geq \rho$) and site-level
434 compliance (share of sites whose mean test-season retention is $\geq \rho$);
- 435 • the lower tail (10th percentile of r , share of seasons with $r < 0.8$);
- 436 • the land-equivalent ratio $\text{LER} = r + e$ (Dupraz et al., 2011).

437 Statistical comparisons use site means as the unit of analysis. We report
438 paired Wilcoxon signed-rank tests and cluster (site) bootstrap 95% confi-
439 dence intervals of the mean difference. SALS was trained with three random
440 seeds on the main split, and with one seed on each of two further spatial
441 splits.

Table 1: SIMPLE crop parameters (Zhao et al., 2019, Table 1a) and the harvest-index shade sensitivity h added in this study. T_{sum} , I_{50A} and I_{50B} are rescaled per site to the local crop calendar (Section 2.2). Units: temperatures in $^{\circ}\text{C}$, thermal times in $^{\circ}\text{C d}$, RUE in gMJ^{-1} ; S_{CO_2} is the relative RUE increase per 100 ppm above 350 ppm.

Crop (cultivar)	h	T_{sum}	HI	I_{50A}	I_{50B}	T_{base}	T_{opt}	RUE	$I_{50 \text{ max } H}$	$I_{50 \text{ max } W}$	T_{heat}	T_{ext}	S_{CO_2}	S_{water}
Wheat (Batten)	0.4	2150	0.34	280	50	0	15	1.24	100	25	34	45	0.08	0.4
Rice (IR72)	0.5	2300	0.47	850	200	9	26	1.24	100	10	34	50	0.08	1.0
Maize (McCurdy 84aa)	0.6	2050	0.50	500	50	8	28	2.10	100	12	34	50	0.01	1.2
Soybean (Williams 82)	0.5	2350	0.40	600	200	6	27	0.86	120	20	36	50	0.07	0.9
Potato (Sebago)	0.2	2400	0.85	500	350	4	22	1.30	50	30	34	45	0.10	0.4

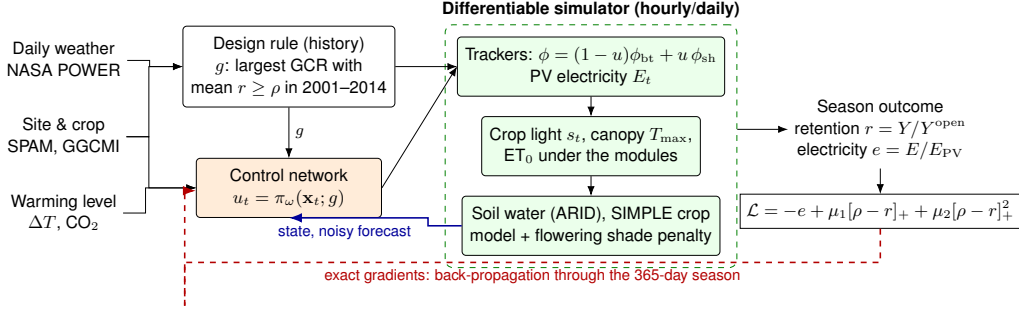


Figure 1: Simulator and controller. The array density g of a site is fixed from its historical seasons. A control network maps the daily crop and soil state, the density g and a noisy forecast to the light-sharing level u_t . The differentiable simulator converts these decisions into PV electricity, light and microclimate at the crop, soil water and crop growth. The season’s electricity and yield retention enter a penalised objective whose exact gradient is back-propagated through the 365-day season.

For the global estimate, every site is evaluated in the split in which it is held out. The mean electricity and yield per hectare of each crop are scaled to the *physical* area of that crop in SPAM 2010, which avoids double counting multiple-cropped land (Waha et al., 2020). We report the installed capacity (TWp), the capacity factor and a supply curve.

3. Results

3.1. Electricity and food on held-out sites and years

Table 2 compares the four controllers on the 92 held-out sites of the main split (seasons sown in 2015–2019), after each site’s density has been chosen from its 2001–2014 seasons with the same rule for all controllers. By

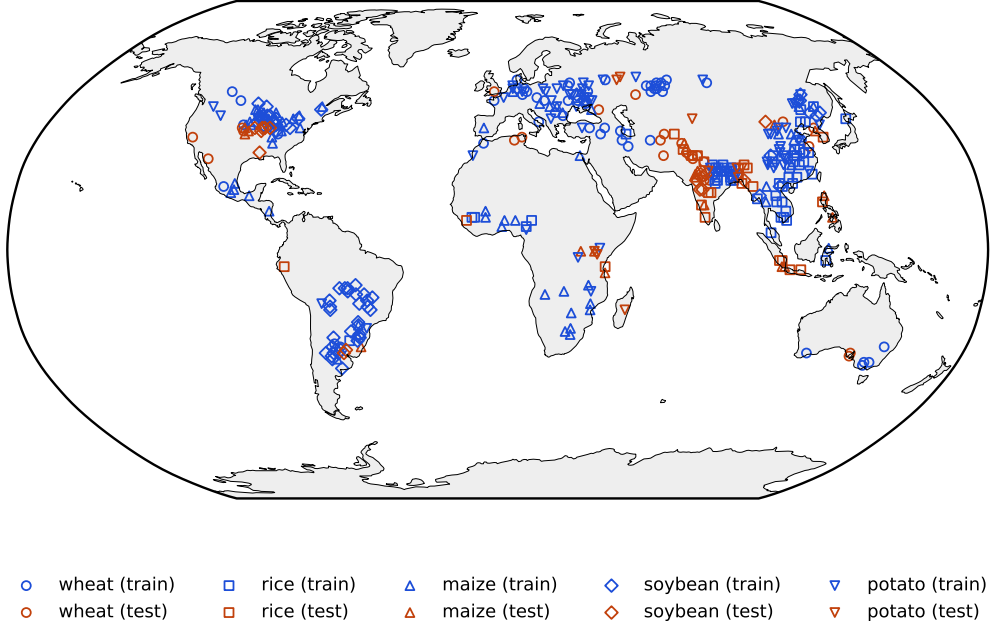


Figure 2: The 400 crop-site pairs, sampled in proportion to the harvested area of each crop (SPAM 2010). Blue: training sites; orange: geographically held-out test sites of the main split.

452 construction all methods keep the mean yield retention close to 90% (0.904–
 453 0.917), so differences in electricity are differences in the density that each
 454 controller can sustain under this floor.

455 **Static arrays are the weakest.** Without light sharing, the floor limits
 456 the ground coverage to 6% on average, and arrays produce $217 \text{ MWh ha}^{-1} \text{ yr}^{-1}$,
 457 16% of a conventional PV plant per hectare (LER 1.07). Rotating the
 458 modules to let light through during the cropping season (seasonal shar-
 459 ing) allows a coverage of 17% and $397 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ (LER 1.19). A rule
 460 that shares light only while the crop is growing (phenology rule) reaches
 461 $465 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ (LER 1.24). An expert rule that additionally shades the
 462 crop on forecast heat or drought days (stress rule) is worse than the phenol-
 463 ogy rule (LER 1.18): its density is limited to 13% because it reduces light at
 464 times when the floor is most binding.

465 **The learned controller is best by a modest, consistent margin.**
 466 SALS sustains a density of 17% (similar to the phenology rule), produces

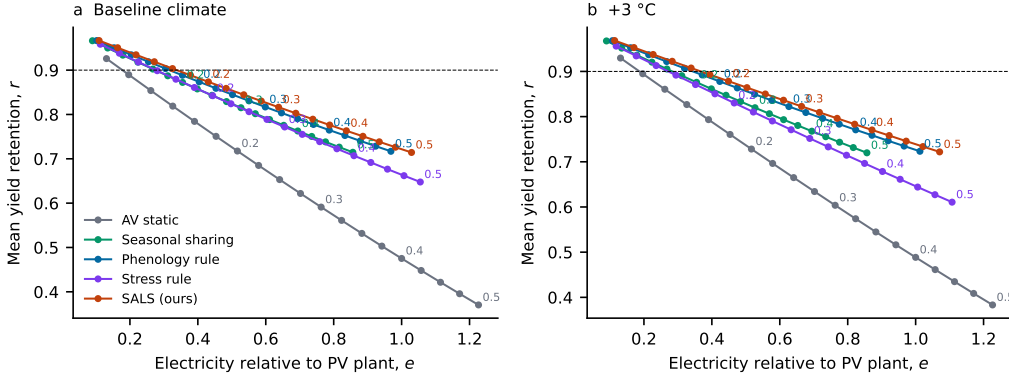


Figure 3: Electricity-retention frontiers at identical array densities (labels: GCR) for the held-out sites. Each point is the mean over test seasons at a fixed GCR; the dashed line is the 90% floor.

510 MWh ha⁻¹ yr⁻¹ and reaches an LER of 1.270 at the same retention. Compared with the phenology rule this is +0.038 in LER (95% CI over sites 0.025–0.053; +10% electricity) and it is higher at 84% of the 92 sites (Table A.15). Compared with static arrays it is +0.205 (0.176–0.236), with 2.35 times the electricity at 2.9 times the density. The frontier in Fig. 6 shows the same picture at matched densities: the learned policy lies on or slightly above the phenology rule over the whole density range, both lie well above seasonal sharing and the stress rule, and static arrays are far below.

What the floor does and does not guarantee. Enforcing the floor on the historical mean does not guarantee it in every season. In the test seasons, 70% of SALS seasons have $r \geq 0.9$ (phenology rule 69%, stress rule 76%, static arrays 93%), and 2.7% have $r < 0.8$ (1.3%, 0.9% and 0.9%). The lower tail of SALS is therefore slightly heavier than that of the rules: its 10th percentile of r is 0.883, against 0.887, 0.886 and 0.903. Requiring $r \geq 0.9$ in at least 80% of a site’s historical seasons (chance-constrained design) raises the season-level compliance to 90–94% for the light-sharing controllers. It costs about 0.065 in LER for SALS (1.205), which remains the best controller (phenology rule 1.192; Section 3.2).

3.2. Robustness

Seeds and splits. Three training seeds of SALS give LERs of 1.270, 1.268 and 1.269 under the baseline climate and 1.297, 1.298 and 1.297 at +3 °C. On two further spatial splits, in which different blocks of sites are

held out, SALS exceeds the phenology rule by 0.047 and 0.041 in LER under the baseline climate and by 0.052 and 0.039 at +3 °C (main split: 0.034 and 0.032). In every split SALS ranks first, the phenology rule second and static arrays last.

Canopy cooling. The canopy-cooling coefficient κ is the most uncertain microclimate parameter. We repeated the whole protocol, including the historical design step, with $\kappa = 0, 1$ and 2°C . Electricity changes with κ because cooler canopies tolerate denser arrays (SALS: 475, 510 and 549 MWh ha⁻¹ yr⁻¹ under the baseline climate), but the retention is held at 0.905 by the design rule and the ranking is unchanged. SALS exceeds the phenology rule by 0.029, 0.034 and 0.042 in LER under the baseline climate (0.022, 0.032 and 0.051 at +3 °C).

Which ingredients matter. Two ablations test the design of SALS at the same training budget. Masking the stress inputs (soil water, ARID, and the maximum temperature, irradiance and ET₀ forecasts) leaves the LER unchanged (1.269 vs 1.270 under the baseline climate, 1.297 vs 1.297 at +3 °C). Training only on the baseline climate gives 1.268 and 1.291. The improvement over the rules therefore does not depend on stress-aware inputs or on training in warmed climates.

3.3. Floor conditioning, chance-constrained training and forecast errors

One network for all floors. Regulations and farmers set different floors. The network of the main experiments was trained at $\rho = 0.9$ only. We evaluated it at floors of 0.80, 0.85, 0.90 and 0.95, re-running the historical design step for each floor (Table 3). Its LER exceeds that of the phenology rule by 0.037, 0.036, 0.034 and 0.033 under the baseline climate (0.040, 0.037, 0.032 and 0.033 at +3 °C), so the benefit carries over to other floors without retraining. A network that received the floor as an additional input and was trained over all four floors did not improve on it (differences in LER ≤ 0.001). Because the design step already adapts the density to the floor, conditioning the controller on the floor is not needed.

Chance-constrained training. We also trained a network whose objective penalises a batch share of seasons below the floor that exceeds 20%. It did not improve the tail. With the mean-based design step its LER was 1.263 (baseline) and 1.290 (+3 °C) against 1.270 and 1.297, and with chance-constrained design it reached season-level compliance of 92.6% against 92.0% and an LER of 1.198 against 1.205. The season-level guarantee is obtained by the design step (Section 3.1), not by the training objective.

526 **Forecast errors.** We raised the forecast errors seen by the controllers
527 to 0, 1, 2 and 3 times the errors used in training. The LER of SALS and
528 of the phenology rule was identical at all four error levels to three decimals
529 (densities 0.2 and 0.3: SALS 1.307 and 1.457, phenology rule 1.278 and
530 1.416). Neither controller uses the weather forecast: both rely on the crop’s
531 thermal time, which is observable from on-farm temperature records. The
532 stress rule, which depends on the forecast of maximum temperature, degrades
533 with forecast error: at density 0.3 its retention falls from 0.778 to 0.751 and
534 its LER from 1.432 to 1.418 as the errors triple.

535 **A tuned rule.** The phenology rule above was written, not tuned,
536 whereas SALS is trained. We therefore tuned the window of the rule on
537 the training-block seasons with the same objective (electricity at the density
538 where mean retention equals the floor); the best window starts at 0.15 in-
539 stead of 0.05 of the thermal-time requirement (Table 4). Pooled over three
540 splits and 76 distinct crop–site pairs, the tuned rule raises LER over the orig-
541 inal rule by 0.026 (95% cluster-bootstrap interval 0.016 to 0.037), and SALS
542 exceeds the tuned rule by 0.018 (0.012 to 0.024), or 5.7% more electricity at
543 equal retention, and the original rule by 0.044 (0.030 to 0.060). Tuning the
544 rule thus accounts for about 60% of the margin of SALS over the written
545 rule, and the remaining margin is small but not zero. SALS meets the floor
546 in fewer seasons than the rules (83% against 84–85%).

547 3.4. *What the learned policy does*

548 Averaged over the rainfed test seasons of the main split at a density of 0.3,
549 the light-sharing level of SALS is a smooth function of relative thermal time.
550 It is close to zero in the first 10% of the thermal time required for maturity,
551 rises through 0.55 at 10–20%, stays at 0.94–0.99 between 30% and 70%, and
552 falls through 0.58 at 80–90% to about zero at maturity. The plateau covers
553 the window in which our crop model penalises shading through the harvest
554 index. The ramps at both ends let the modules produce electricity when
555 additional light is worth least to the crop. Fig. 4 shows two example seasons.
556 The phenology rule, which is also the initialisation of the network, shares
557 light during the whole growing period. Fine-tuning through the simulator
558 found that sharing can be reduced at the edges of the season: SALS gains its
559 advantage by yielding light when it matters and collecting electricity when
560 it matters less.

561 The policy also shares less light on days with a forecast heat stress: the
562 mean level is 0.35 on such days against 0.60 (baseline) and 0.51 (+3 °C)

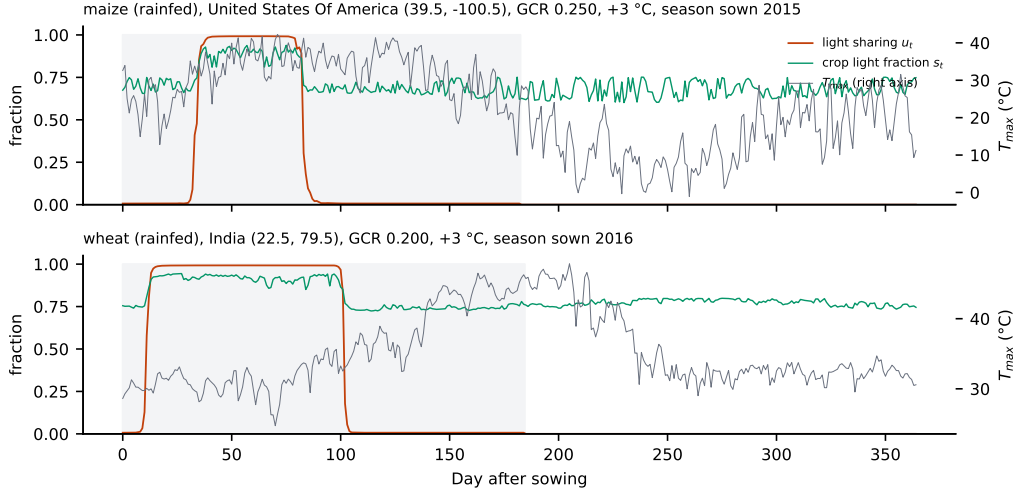


Figure 4: Learned daily control for two test seasons under $+3^{\circ}\text{C}$, at the density chosen for the site, simulated for the dominant sub-field of the site. Shading marks the cropping window. Orange: light-sharing level u_t ; green: fraction of open-field irradiance reaching the crop; grey: daily maximum temperature (right axis).

on other days, and 58–63% of sites show a lower level on heat days. This cannot be read as a stress response, however. Heat days occur late in the season, when the phenological ramp is already falling, and a policy without any stress input shows a similar contrast (0.37 vs 0.57 in a separate check). The ablation above confirms that stress information does not improve the result.

3.5. Crops and regions

The benefit of light sharing differs between crops (baseline climate, SALS). Soybean and maize tolerate the densest arrays (GCR 0.24 and 0.19; 691 and 577 $\text{MWh ha}^{-1}\text{yr}^{-1}$; LER 1.41 and 1.30), potato follows (0.19; 542; 1.31), and rice and wheat are the most restrictive (0.13 and 0.15; 395 and 428; 1.20 for both). The stress rule has a lower LER for every crop (1.26 for soybean, 1.14–1.21 for the others).

The floor is not met evenly. Production-weighted retention is 0.888 for soybean and 0.898 for maize, below 0.9, because high-yielding seasons lose more than low-yielding ones. Regionally, Latin America and Oceania (4 sites) combine the highest density (GCR 0.41) and electricity ($1090 \text{ MWh ha}^{-1}\text{yr}^{-1}$) with a production-weighted retention of only 0.862. Africa and the Middle

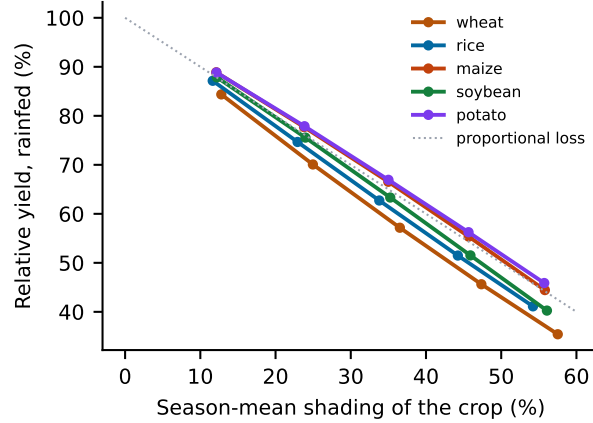


Figure 5: Simulated relative yield of rainfed crops under static arrays (GCR 0.1–0.5) against season-mean shading of the crop, baseline climate, held-out sites. The dotted line is proportional loss.

581 East (8 sites) reach 0.894, and only 63% of their sites keep a mean retention
 582 of at least 0.9. South Asia, the region with the most sites (39), reaches 0.902
 583 (67% of sites) at a density of 0.14. A regulatory floor enforced per site would
 584 therefore need a region- and crop-specific design step, not a global one.

585 3.6. Model behaviour and validation

586 **PV output.** The simulated specific yield of an isolated tracker row agrees
 587 with PVGIS 5.2 (European Commission, Joint Research Centre, 2022) at 29
 588 sites with $r = 0.971$, a mean bias of -4.9% and a mean absolute deviation
 589 of 5.5% (Section 2.2).

590 **Shade response.** Fig. 5 shows the simulated relative yield of rainfed
 591 crops under static arrays as a function of the season-mean shading of the crop.
 592 Yield declines close to proportionally with shading: 84–89% relative yield at
 593 12% shading, 70–78% at 25%, 57–67% at 37% and 35–46% at 57%. Wheat
 594 loses most and potato and maize least. This response is steeper than the *less*
 595 *than proportional* losses that Laub et al. (2022) report for small shade levels
 596 of C3 cereals and tubers, and it does not reproduce their finding that maize
 597 and grain legumes are more shade sensitive than C3 cereals. In the simulator,
 598 shading also relieves heat and drought in rainfed maize and soybean, which
 599 offsets the higher harvest-index penalty we assigned to them. The absolute
 600 effect of shading on each crop is therefore uncertain (Section 4.5).

601 **Comparison with field trials.** We fitted one radiation-use compen-
 602 sation coefficient per crop to the yield–shading curves of Laub et al. (2022)
 603 using training-block sites only (Methods, calibration script in the repository).
 604 The fit changed wheat and rice only (coefficients 0.4 and 0.1) because the
 605 proportional model already lies within about 0.02 of the C3-cereal and tu-
 606 ber curves, and it cannot lower the maize and legume responses below the
 607 proportional case (root-mean-square error 0.16 for maize and 0.08 for soy-
 608 bean). Against independent field trials (Table 5) the model underestimates
 609 retained yield: rice 0.79 against 0.84 (Gonocruz et al., 2021), wheat and
 610 potato 0.74 and 0.72 against 0.81–0.82 in the worse season of Weselek et al.
 611 (2021) (and above 1 in the better one), and maize 0.70–0.75 against 0.83–0.91
 612 at 30–35% shading (Ramos-Fuentes et al., 2023), where harvest index was
 613 not affected by shade. The model is therefore conservative for food reten-
 614 tion with respect to these trials, except for the stronger maize losses of the
 615 meta-analysis. Table 6 shows the consequence for the design results: with
 616 the calibrated response the allowed density and the LER of every method
 617 rise (phenology rule 1.28 against 1.23), and with a response fitted to the
 618 field trials the LER of the phenology rule is 1.62. The phenology rule is the
 619 best method in all three cases, and light sharing adds LER +0.17, +0.19 and
 620 +0.30 over static arrays (conservative, calibrated, field-fitted); only the order
 621 of the seasonal and stress rules changes. The absolute gains for rice, wheat
 622 and potato are thus conservative with respect to these trials; for maize and
 623 soybean the meta-analysis points the other way (Section 3.7).

624 3.7. *A shade-response model learned from field data*

625 Cross-validation by study shows what the data can and cannot support
 626 (Tables 7 and 8). Over all 339 shaded observations the root-mean-square
 627 error of a held-out study falls from 0.43 for proportional loss $y = 1 - s$ to 0.37
 628 for one curve for all crops and to 0.31 for crop-type curves, but this gain comes
 629 from crop types that the simulator does not represent (berries, forages, fruits
 630 and vegetables: 0.53 to 0.36). For the four crop types that the simulator uses,
 631 the learned curves are not better than proportional loss: the error is 0.222
 632 against 0.220 pooled (difference +0.002, 95% interval -0.023 to $+0.030$),
 633 0.187 against 0.214 for C3 cereals, and slightly worse for maize, grain legumes
 634 and tubers. Partial pooling is as good as separate curves (0.312 against 0.313)
 635 and the covariates add almost nothing (0.311). Gradient-boosted trees on the
 636 same inputs are worse (0.321) than the parametric model. The error of 0.22
 637 for the four crop types is large: yields of a new study differ from the learned

638 curve by about 20 percentage points, which is the between-study spread
639 that a design under uncertainty has to cover. Crop type is the information
640 that transfers between studies, and none of the covariate effects on $\log \beta$ is
641 clearly different from zero (irradiation -0.10 , 90% interval -0.32 to 0.09 ;
642 subtropical zone -0.37 , -0.73 to 0.07 ; intercropping 0.16 , -0.48 to 0.45).
643 The estimated effect of solar-panel shade is not interpretable, because all 46
644 panel observations are of leafy vegetables, mainly lettuce, and the effect is
645 confounded with crop type.

646 The learned curves show a clear ordering of crop types (Table 9). At
647 20% shading C3 cereals and tubers lose less than proportionally (relative
648 yield 0.93), whereas maize (0.73) and grain legumes (0.71) lose more; at 40%
649 shading the values are 0.72 and 0.71 against 0.42 and 0.40 . Intervals are wide
650 for maize (0.27 – 0.59 at 40%) and for tubers, which rest on two studies. The
651 simulator’s responses are close to the data for cereals and potato but, under
652 both calibrated ensembles, too mild for maize and soybean at 40% shading
653 (0.75 and 0.48 for E2 against 0.42 and 0.40 in the data). The calibration
654 cannot lower the maize response further, because the simulated shading also
655 relieves heat and drought stress and the compensation and harvest-index pa-
656 rameters are bounded below at zero. The field trials of maize (Ramos-Fuentes
657 et al., 2023) are also milder than the meta-analysis. For maize and soybean
658 the gains reported in this paper are therefore probably overestimated.

659 3.8. Learning under shade-response uncertainty

660 **Calibration.** Gradient descent through the simulator fitted the five crops
661 to the data-driven curves and the field trials (ensemble E2, Table 10). The
662 data identify two features: the harvest index of maize is hardly affected by
663 shade (multiplier 0.16 ± 0.13 , $c \approx 0.14$), as observed by Ramos-Fuentes et al.
664 (2023), and soybean is more shade sensitive than the cereals (compensation
665 0.05 , harvest-index multiplier 1.66), as in Laub et al. (2022). Wheat has the
666 largest compensation (0.91 ± 0.31) and the largest uncertainty; the compen-
667 sation is an effective parameter that absorbs diffuse-light and light-saturation
668 effects that the simulator does not model explicitly, not a physical value.

669 **Learned controllers.** Controllers were trained on E1 and evaluated
670 under four fixed responses on 21 crop–site pairs of split 0 with one training
671 seed, so the intervals below, which are cluster bootstraps over pairs, do not
672 include seed or split variation. Under every response SALS stays better than
673 the written phenology rule (Table 11): by $+0.019$ (0.007 to 0.039), $+0.024$
674 (0.009 to 0.049), $+0.038$ (0.017 to 0.062) and $+0.015$ (0.001 to 0.038) LER

under the proportional, calibrated, field-fitted and harsh responses. Training on the ensemble did not help: SALS-robust differs from SALS, which was trained on the proportional response only, by -0.005 to $+0.003$ LER with every interval containing zero, and SALS-adaptive, which sees the measured biomass ratio, by -0.010 to $+0.007$ with every interval containing zero. The training ensemble E1 is narrow, so this null result may reflect its low diversity. The feedback rule is worse than the phenology rule by 0.026 – 0.058 LER under every response (intervals exclude zero). A plausible reason is that the biomass ratio does not reveal the harvest-index loss from shading around flowering, but we did not test this explanation.

Design under uncertainty. The larger risk is in the design step (Table 12, Fig. 6). Pooled over the three spatial splits (76 distinct crop–site pairs), a density chosen under the mean response of E2 leaves 69% of the pairs at the floor when the true response is another plausible one, against 87% for an oracle design; with the wider ensemble E3 the figures are 70% and 85%. Choosing the density so that the floor is met in at least 75% of the plausible responses raises compliance by $+9.0$ points under E2 (95% cluster-bootstrap interval 5.3 to 12.9) and $+8.7$ under E3 (4.2 to 13.7), at a cost of 0.038 and 0.034 in LER; it recovers about half of the gap to the oracle. The ensemble-based rule is, however, not better than a simple margin. Keeping the density rule of the main analysis and raising the floor used in the design from 0.90 to 0.92 under the mean response gives 89% (E2) and 85% (E3) of the pairs at the floor at LER 1.434 and 1.468, whereas the ensemble rule at 90% reliability reaches 89% and 88% at 1.412 and 1.409: at equal compliance the margin keeps 0.02 (E2) to 0.04 (E3) more LER. The ensemble is thus useful to size the margin (about 0.02 to 0.04 above the floor for 85–95% of the pairs) but a rule built on its quantiles does not beat it. On the first split alone, with the earlier ensemble E1, the quantile rule restored the oracle level (76% and 74%, against 58% and 55% under the mean response), but without the margin comparison. The guarantee holds only within the range of responses that the ensemble spans: with the E2 design on split 0 (21 pairs), compliance under the proportional, calibrated and harsh responses is 19–24% (mean-response design 19%), and 81% under the field-fitted response (71%).

3.9. Economics of light sharing

Under the illustrative parameters of Table 13 the economics are dominated by the capacity installed, not by the crop (Table 14). The annuity of the modules is 11–32 kUSD per hectare and year, whereas the crop revenue

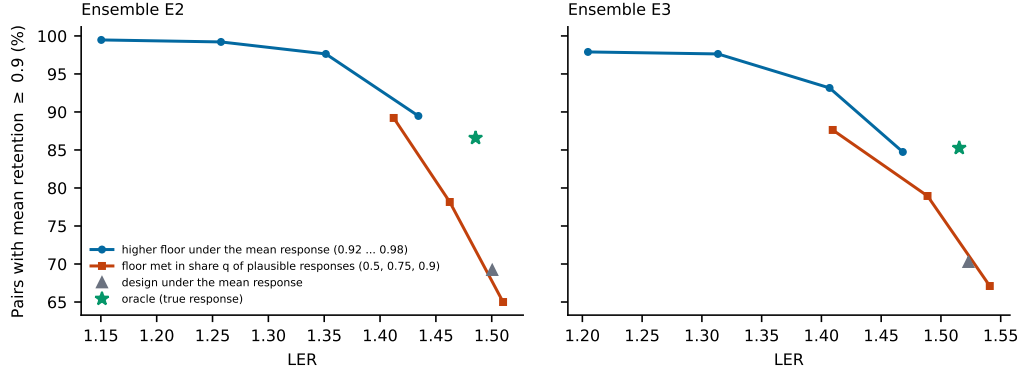


Figure 6: Share of crop-site pairs at the floor against LER for the phenology rule under design rules for shade-response uncertainty (leave-one-member-out, three splits, 76 pairs): a higher floor under the mean response (blue), a floor met in a share q of the plausible responses (orange), the design under the mean response and the oracle design.

lost is about 0.3 kUSD, so a method is worth what it earns per installed watt. At 60 USD MWh⁻¹ all light-sharing methods lose money relative to the open field (−3 to −10 kUSD ha⁻¹ yr⁻¹) and static arrays are close to break-even; break-even prices are 73 (SALS), 77 (tuned phenology rule), 80 (written rule), 70 (stress rule) and 92 (seasonal sharing) USD MWh⁻¹, against 48 for a conventional plant that replaces the crop. At 90 USD MWh⁻¹ SALS earns 7.6 kUSD ha⁻¹ yr⁻¹, against 5.5 for the tuned rule and 4.0 for the written rule, so learning is worth about 2.1 kUSD ha⁻¹ yr⁻¹ (39%) over the tuned rule at that price; the stress rule, which installs fewer modules, has a lower break-even price than SALS (70 against 73) although it earns less electricity per hectare. A conventional plant is far more profitable per hectare at every price (16 kUSD at 60, 94 kUSD at 120, against −6 and 21 for SALS), so agrivoltaics competes with it only if food production, land-use rules or a premium are valued; the net value of agrivoltaics relative to the plant stays negative across the whole price range.

Cost of the safety margin. At 90 USD MWh⁻¹, moving from a density designed under the mean response to a floor of 0.92 under that response costs 1.07 kUSD ha⁻¹ yr⁻¹ (18%) under E2 for 20 more points of compliance (69% to 89%) and 1.08 kUSD (18%) under E3 for 14 points (70% to 85%), about 50–75 USD ha⁻¹ yr⁻¹ per point. The ensemble rule at 90% reliability costs 1.03 and 1.47 kUSD for 20 and 18 points, so its price per point is similar to or higher than that of the margin. At 60 USD MWh⁻¹ a margin has a negative

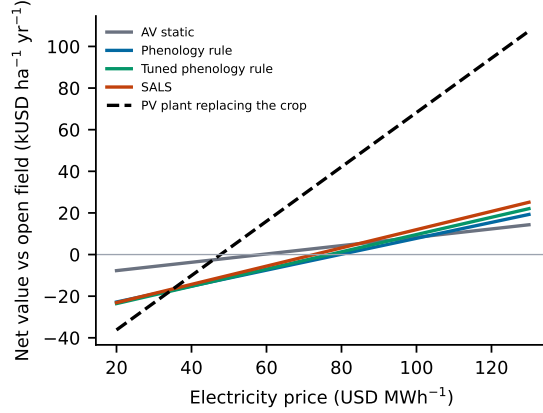


Figure 7: Net value per hectare relative to the open field against electricity price for static arrays, rules and SALS, and for a conventional plant replacing the crop (dashed). Illustrative parameters of Table 13.

734 cost, because every module beyond break-even loses money; the economics
 735 of the design margin therefore depend on the price level, whereas its effect
 736 on compliance does not.

737 3.10. Agrivoltaics under warming

738 Under warming, the open-field production of the held-out sites (Fig. 8)
 739 falls at +3 °C by 23% for maize and 18% for rice, is nearly unchanged for soy-
 740 bean, and rises for wheat (+3%) and potato (+10%). Agrivoltaic production
 741 with light-sharing control follows the open field at about 90% of its level, so
 742 it inherits the same climate impacts (maize −31%, rice −27% relative to the
 743 baseline open field at +3 °C). Agrivoltaics does not protect staple production
 744 against warming in this model. The LER of light-sharing systems increases
 745 slightly with warming (SALS 1.270 to 1.297; phenology rule 1.236 to 1.265),
 746 because the floor allows slightly denser arrays in warmer conditions (GCR
 747 0.173 to 0.177 and electricity 510 to 540 MWh ha^{−1} yr^{−1}). The margin of
 748 SALS over the rules does not grow with warming (Table A.15).

749 3.11. Global potential and what it would take

750 Every site is evaluated when it is held out (219 of the 400 crop-site
 751 pairs, in at least one split), and each crop's mean electricity per hectare
 752 is scaled to its physical area in SPAM 2010 (581 Mha for the five crops,
 753 which avoids counting multiple-cropped land twice). Equipping 5% of this

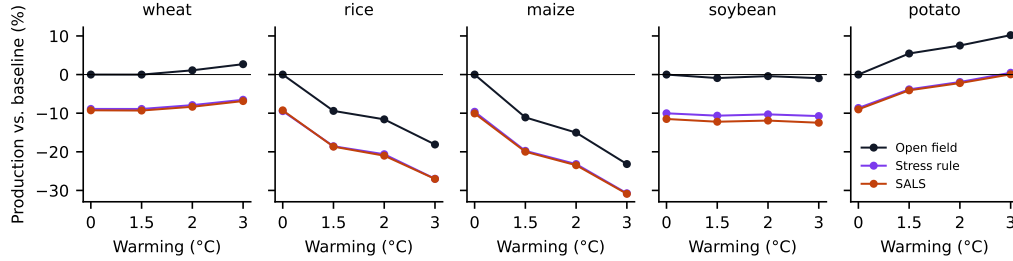


Figure 8: Production relative to the baseline open field (sum over held-out site-seasons) as a function of warming, for the open field, the stress rule and SALS.

area (29.1 Mha) with SALS arrays at the floor would produce 14.4 PWh yr^{-1} under the baseline climate and 16.0 PWh yr^{-1} at $+3^\circ\text{C}$, with a production-weighted retention of 0.903 and 0.900. The phenology rule would give 13.0 and 14.5 PWh yr^{-1} , and static arrays 6.4 and 6.9 PWh yr^{-1} . For comparison, world electricity generation was about 30 PWh in 2023 (Ember, 2024).

These numbers require large amounts of hardware. SALS at 5% coverage needs 10.6 TWp of modules (11.3 TWp at $+3^\circ\text{C}$), at a capacity factor of 15.5% (16.2%). World solar PV capacity at the end of 2024 was about 1.9 TW (IRENA, 2025), so the requirement is about 5.7 times the installed PV fleet. The electricity is also variable, peaks at midday and would be curtailed or stored. The figures are technical potentials, not projections. Ordering sites by electricity per hectare gives a supply curve (Fig. 9): the best 2% of the staple cropland would supply 16 PWh yr^{-1} and the best 5% 40 PWh yr^{-1} with SALS. The best tenth of the sites are mostly rainfed maize and soybean sites in the Americas and in southern and eastern Africa (7% mean irrigated fraction against 29% overall), which are also among the regions where the floor is hardest to meet (Section 3.5). Siting by electricity alone would therefore concentrate agrivoltaics where food-security risk is highest.

Maize and soybean, for which the simulated shade response is probably too mild (Section 3.7), account for 53% of the electricity of the upscaling (maize 30%, soybean 23%, wheat 29%, rice 15%, potato 3%). Counting only wheat, rice and potato would give about 7 PWh yr^{-1} at the same coverage, so the global figure should be read as an upper estimate for about half of its content. By crop, the LER of SALS under the baseline climate is 1.20 for wheat and rice, 1.31 for potato, 1.30 for maize and 1.41 for soybean.

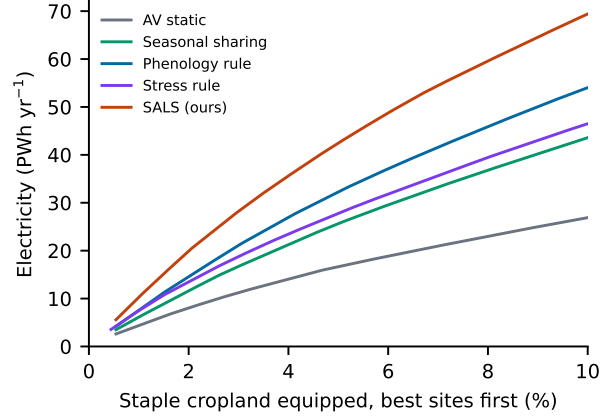


Figure 9: Supply curve under the baseline climate: cumulative electricity when sites are equipped in order of decreasing electricity per hectare, against the share of staple cropland (physical area) equipped. Each controller is evaluated at the density that keeps its mean retention at 0.9.

779 4. Discussion

780 4.1. What the results establish

781 Four findings stand on consistent evidence. First, sharing light with the
782 crop is the main lever on the electricity that tracker-based agrivoltaic arrays
783 can produce under a yield floor. With static trackers the floor limits arrays
784 to about 6% ground coverage and an LER of 1.07. Any light-sharing control
785 allows about three times the density, and the LER rises to 1.18–1.24 for hand-
786 written rules. Second, the best simple rule follows crop phenology. It shares
787 light across the middle of the season, around flowering and grain filling, and
788 produces electricity at the beginning and the end. It needs only the crop’s
789 thermal time, which is observable on the farm, and it is insensitive to forecast
790 errors, whereas the rule that reacts to heat and drought stress is worse in
791 every climate and degrades with forecast error. Third, the qualitative result
792 does not depend on the weakest part of the model. The simulated shade
793 response is conservative against the field trials we could use (Section 3.6);
794 calibrating it raises the allowed density, and the phenology rule stays the
795 best method, with a gain over static arrays of 0.17, 0.19 and 0.30 LER
796 under the conservative, calibrated and field-fitted responses. However, the
797 response learned from the meta-analysis shows steeper losses for maize and
798 grain legumes than the simulator reproduces (Section 3.7), so the absolute

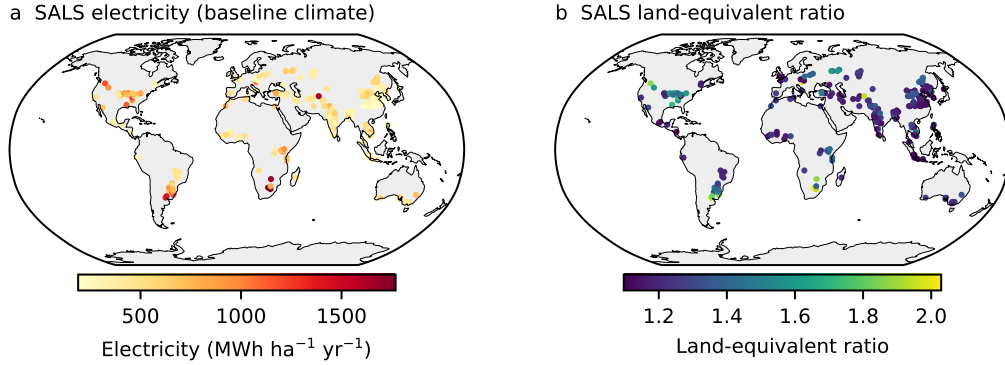


Figure 10: SALS electricity (a) and land-equivalent ratio (b) at the pooled held-out sites (baseline climate, mean of seasons 2015–2019).

gains are conservative for wheat, rice and potato and probably overestimates for maize and soybean (53% of the upscaled electricity). Fourth, a controller learned through the differentiable simulator is better than the best rule we could write, by 0.04 in LER against the written rule (0.018 against a rule whose window was tuned on the training seasons, 5.7% more electricity at equal retention), reproducibly over training seeds (range 0.002 in LER), three spatial splits, three canopy-cooling assumptions and four climates, and at 84–86% of the held-out sites. This gain is small compared with the gap between static and dynamic control.

4.2. What role does artificial intelligence play?

We tested three ways in which learning could matter, and the answers differ. The first is learning the crop’s response to shade from data. The learned model is useful mainly as a measurement of how little can be learned: a hierarchical model fitted to 428 observations from 58 studies predicts a held-out study of the four staple crop types with an error of 0.22 in relative yield, no better than proportional loss, and the gain over proportional loss (0.43 to 0.31 over all types) comes from crop types that the simulator does not represent. Crop type is the information that transfers, whereas irradiation, climate zone and shade type add nothing measurable and a non-parametric model (gradient-boosted trees) is worse with this number of studies (Section 3.7). Two things follow. The shade response of the staple crops is a distribution with a spread of about 20 percentage points between studies and not a curve, and the process-based simulator does not reproduce the steep losses of maize

822 and legumes seen in the data. Learning the response was limited by the data
823 and not by the model: two studies support the tuber curve, no rice study
824 exists, and 46 panel observations are all leafy vegetables.

825 The second is calibration through the differentiable simulator. Gradient
826 descent through whole seasons fitted the response of each crop to the learned
827 curves and to the field trials, and it produced an ensemble of plausible re-
828 sponses that reconciles them (Section 3.8). The ensemble informs the design
829 decision: a density chosen under the mean response leaves about 70% of
830 crop–site pairs at the floor if another plausible response holds, against 85–
831 87% for an oracle, and a density chosen to meet the floor in 75% of plausible
832 responses recovers about half of that gap at a cost of 0.04 LER. A simple mar-
833 gin, however, does better: raising the floor from 0.90 to 0.92 under the mean
834 response reaches 85–89% of the pairs and keeps 0.02–0.04 more LER than
835 the ensemble rule at equal compliance. What the ensemble adds is the size
836 of the margin, not a better rule. The quantile rule is a scenario-type design
837 (Calafiore and Campi, 2006; Campi and Garatti, 2008); with five members it
838 has none of the finite-sample guarantees of that theory, and the comparison
839 with the margin shows that it is not needed in practice. The bottleneck of
840 tracker-based agrivoltaics under a food-security floor is the uncertainty of
841 the crop response that enters the design, and it is data on that response, not
842 larger networks, that would improve the result; the learned curves show a
843 spread of about 20 percentage points between studies. The guarantee holds
844 only within the range of responses that the ensemble spans.

845 The third is control, and here learning contributes little. The learned
846 controller adds 0.018 LER over a tuned phenology rule (and 0.044 over the
847 written one), and 0.01–0.03 over the written rule under the alternative re-
848 sponses on one split; a reader who needs a simple, auditable controller can use
849 that rule and obtain most of the electricity gain over static arrays. Neither of
850 the two extensions that we expected to matter under uncertainty did: train-
851 ing on an ensemble of responses did not change the policy, and neither did
852 giving the controller a farm-measurable signal, the biomass of the crop rela-
853 tive to a reference plot, nor a feedback rule built on that signal. The learned
854 policy also has to be initialised from the phenology rule (Section 4.5). We
855 therefore do not claim that learned control is what makes agrivoltaics work:
856 in this model, the schedule that a learned policy converges to is a smooth
857 function of crop phenology, which a rule can already express.

858 4.3. *What the results do not establish*

859 The early hypothesis that stress-aware control, that is, shading the crop
860 on heat and drought days, would be a main source of benefit under warming
861 was not supported. The stress rule performs worse than the phenology rule
862 in every climate, and a learned controller without stress inputs performs as
863 well as one with them. A plausible reason is that in our crop model the yield
864 gain from shading on a hot or dry day is smaller than the loss of light on
865 that day, so stress information does not help. Whether this holds for crop
866 models with more detailed canopy-temperature and heat-stress responses is
867 an open question. Agrivoltaics does not protect staple production against
868 warming in our simulations: light-sharing systems produce about 90% of the
869 open-field yield in every climate and follow the same decline in maize and
870 rice (Fig. 8).

871 4.4. *Food security in practice*

872 Holding the historical mean retention at 90% does not hold it in each
873 season or at each site. 70% of SALS seasons meet the floor, the lower tail is
874 slightly heavier than for the rules, and production-weighted retention is below
875 0.9 for soybean (0.888) and maize (0.898) and in several regions. Chance-
876 constrained design raises season-level compliance to about 92% at a cost of
877 0.065 in LER, and the ranking of the controllers is preserved. Whichever for-
878 mulation is chosen, a per-site regulatory floor would need a region- and crop-
879 specific design step. The upscaled potential (Section 3.11) shows the trade-off
880 in magnitude. The 14 PWh yr⁻¹ that 5% of the staple cropland could supply
881 need about 10 TWp of modules, several times the present global fleet, and
882 the best sites are rainfed maize and soybean areas where production-weighted
883 retention is lowest.

884 4.5. *Limitations*

- 885 • *Simulation only.* All results come from a simulator. PV output was
886 benchmarked against PVGIS (29 sites, mean bias -4.9%), but the crop
887 response to shade was learned from one meta-analysis data set and
888 compared with four field trials (Sections 3.6 and 3.7). The simulated
889 response is conservative with respect to the trials (rice, wheat, potato,
890 maize) but too mild for maize and soybean relative to the meta-analysis,
891 where the 40%-shading relative yields are 0.42 and 0.40 against 0.75

892 and 0.48 in the simulator. The harvest-index sensitivities h are assump-
893 tions, and the sensitivity analysis uses rule controllers on a subsample
894 of 22 sites.

- 895 • *Canopy cooling.* The effect of shade on canopy temperature was repre-
896 sented by one coefficient κ , applied to a heat-stress function that was
897 calibrated on air temperature (Siebert et al., 2014; Webber et al., 2018).
898 The ranking of the controllers is robust to κ between 0 and 2 °C, but
899 the absolute electricity changes by about $\pm 8\%$. Nocturnal warming
900 under modules and rain redistribution are not modelled.
- 901 • *Crop model.* SIMPLE has a fixed harvest index and linear radiation-
902 use efficiency. More detailed models with grain-number formation and
903 canopy energy balance may favour different schedules.
- 904 • *Climate scenarios.* Warming was applied as uniform temperature shifts
905 with fixed precipitation, radiation and relative humidity, and with un-
906 changed cultivars and sowing dates. It is not a bias-adjusted climate
907 projection.
- 908 • *Initialisation and state information.* The network is initialised from the
909 phenology rule, because training from random weights collapsed to full
910 sharing. The gain of SALS is therefore the gain of fine-tuning that rule,
911 and we cannot say whether a better optimiser or a richer policy class
912 would find larger gains. The controller uses the simulated state of the
913 crop and soil (thermal time, canopy interception, soil water) without
914 measurement error, and the forecast errors we tested apply to weather
915 only.
- 916 • *Economics and deployment.* The economic analysis (Section 3.9) uses
917 assumed prices and costs, no grid, curtailment, storage or subsidy
918 model, and does not value food security; its break-even prices move
919 with the capex and the discount rate (Fig. 7). The value of midday
920 electricity, curtailment, storage, structural cost and farmer acceptance
921 (Trommsdorff et al., 2021) will determine the realisable share. The av-
922 erage row pitch at the chosen densities (about 13 m for 2.2-m modules)
923 leaves room for machinery, but the densest designs (GCR up to 0.5)
924 have a pitch of only 4.4 m.

5. Conclusions

Light sharing, the daily choice between maximum electricity and letting light through to the crop, is the main lever of tracker-based agrivoltaics under a food-security floor. In a differentiable simulator of trackers, PV output, crop light, microclimate and crop growth, evaluated fairly on 400 crop-site pairs with the array density of each site chosen from its own history, light-sharing rules allowed about three times the module density of static arrays at a mean yield retention of 90%. They raised the land-equivalent ratio from 1.07 to 1.18–1.24, and a rule that follows crop phenology and needs no weather forecast was the best of them. The simulated shade response is conservative against the field trials we could use for rice, wheat and potato, and calibrating it raised the gains without changing the ranking of methods; for maize and soybean, which supply 53% of the upscaled electricity, the meta-analysis indicates larger losses than the simulator reproduces, so the gains for these crops are probably overestimated. The crop’s response to shade, the least certain input, was learned from 428 field observations: a hierarchical model predicted a held-out study of the staple crop types no better than proportional loss (error 0.22), which shows the size of the between-study spread, and gradient descent through the simulator turned it into an ensemble of plausible responses. The simulator is too mild for maize and soybean, so the gains for those crops are probably overestimated. A learned controller exceeded the written phenology rule by 0.044 LER and a rule with a window tuned on the training seasons by 0.018 (95% interval 0.012–0.024), by following a smooth phenology-dependent schedule that a rule can express; training it on an ensemble of shade responses or giving it a farm-measurable biomass signal added nothing. The decision that mattered was the design. A density chosen under the mean response left about 70% of crop-site pairs at the floor when another plausible response held (oracle 85–87%); choosing the density so that the floor holds in 75% of the plausible responses recovered about half of the gap at a cost of 0.04 LER, but a floor raised by 0.02 under the mean response did better at equal compliance. The ensemble is therefore useful to size a safety margin, not to replace it. Under illustrative costs, light sharing breaks even with the open field at 73 (learned controller) to 92 USD MWh^{−1}; a conventional plant replacing the crop stays more profitable at every price, so agrivoltaics depends on a premium, land-use rules or a value for food production that markets do not give. The yield floor is met on average but not in every season or region, and the electricity potential of 5% of staple

cropland (about 14 PWh yr^{-1}) requires about 10 TWp of modules. Future work should validate the crop response to shade under trackers with more field data, notably for maize and legumes where the meta-analysis and the trials disagree, extend the crop model to grain-number formation and canopy energy balance, add costs and grid constraints, and pursue controllers that guarantee the food-security floor per site and per season.

CRedit authorship contribution statement

Quang-Vinh Dang: Conceptualization, Methodology, Software, Formal analysis, Investigation, Data curation, Visualization, Writing – original draft, Writing – review & editing.

Declaration of competing interest

The author has no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data and code availability

All data used are public: SPAM 2010 v2r0 harvested areas (Yu et al., 2020), the GGCM Phase 3 crop calendar (Jägermeyr et al., 2021), NASA POWER daily weather (Sparks, 2018) and the SIMPLE crop parameters (Zhao et al., 2019). The code of the simulator, of SALS and of all experiments, together with the site list, is available from the author’s public repository (<https://github.com/vinhqddang/Artificial-Intelligence-in-Agriculture>).

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Appendix A. Statistical comparison

References

Adeh, E.H., Good, S.P., Calaf, M., Higgins, C.W., 2019. Solar pv power potential is greatest over croplands. *Scientific Reports* 9, 11442. doi:10.1038/s41598-019-47803-3.

- 991 Allen, R.G., Pereira, L.S., Raes, D., Smith, M., 1998. Crop evapotranspira-
992 tion: Guidelines for computing crop water requirements. FAO Irrigation
993 and Drainage Paper 56. Food and Agriculture Organization of the United
994 Nations. Rome.
- 995 Amaducci, S., Yin, X., Colauzzi, M., 2018. Agrivoltaic systems to optimise
996 land use for electric energy production. *Applied Energy* 220, 545–561.
997 doi:10.1016/j.apenergy.2018.03.081.
- 998 Andrade, F.H., Vega, C., Uhart, S., Cirilo, A., Cantarero, M., Valentinuz, O.,
999 1999. Kernel number determination in maize. *Crop Science* 39, 453–459.
1000 doi:10.2135/cropsci1999.0011183X0039000200026x.
- 1001 Barron-Gafford, G.A., Pavao-Zuckerman, M.A., Minor, R.L., Sutter, L.F.,
1002 Barnett-Moreno, I., Blackett, D.T., Thompson, M., Dimond, K., Gerlak,
1003 A.K., Nabhan, G.P., Macknick, J.E., 2019. Agrivoltaics provide mutual
1004 benefits across the food–energy–water nexus in drylands. *Nature Sustain-*
1005 *ability* 2, 848–855. doi:10.1038/s41893-019-0364-5.
- 1006 Calafiore, G.C., Campi, M.C., 2006. The scenario approach to robust control
1007 design. *IEEE Transactions on Automatic Control* 51, 742–753. doi:10.
1008 1109/TAC.2006.875041.
- 1009 Campi, M.C., Garatti, S., 2008. The exact feasibility of randomized solutions
1010 of uncertain convex programs. *SIAM Journal on Optimization* 19, 1211–
1011 1230. doi:10.1137/07069821X.
- 1012 Collares-Pereira, M., Rabl, A., 1979. The average distribution of solar
1013 radiation-correlations between diffuse and hemispherical and between daily
1014 and hourly insolation values. *Solar Energy* 22, 155–164. doi:10.1016/
1015 0038-092x(79)90100-2.
- 1016 DIN, 2021. DIN SPEC 91434:2021-05: Agri-photovoltaic systems – require-
1017 ments for primary agricultural use. Beuth Verlag, Berlin.
- 1018 Dinesh, H., Pearce, J.M., 2016. The potential of agrivoltaic systems. *Renew-*
1019 *able and Sustainable Energy Reviews* 54, 299–308. doi:10.1016/j.rser.
1020 2015.10.024.
- 1021 Dupraz, C., Marrou, H., Talbot, G., Dufour, L., Nogier, A., Ferard, Y., 2011.
1022 Combining solar photovoltaic panels and food crops for optimising land

1023 use: Towards new agrivoltaic schemes. *Renewable Energy* 36, 2725–2732.
1024 doi:10.1016/j.renene.2011.03.005.

1025 Elamri, Y., Cheviron, B., Lopez, J.M., Dejean, C., Belaud, G., 2018. Water
1026 budget and crop modelling for agrivoltaic systems: Application to irrigated
1027 lettuces. *Agricultural Water Management* 208, 440–453. doi:10.1016/j.
1028 agwat.2018.07.001.

1029 Ember, 2024. *Global Electricity Review 2024*. Technical Report. Ember.
1030 London.

1031 European Commission, Joint Research Centre, 2022. Photovoltaic geo-
1032 graphical information system (PVGIS), version 5.2. https://re.jrc.ec.europa.eu/pvg_tools/.

1034 Fischer, R.A., 1985. Number of kernels in wheat crops and the influence of
1035 solar radiation and temperature. *The Journal of Agricultural Science* 105,
1036 447–461. doi:10.1017/S0021859600056495.

1037 Gautron, R., Maillard, O.A., Preux, P., Corbeels, M., Sabbadin, R., 2022.
1038 Reinforcement learning for crop management support: Review, prospects
1039 and challenges. *Computers and Electronics in Agriculture* 200, 107182.
1040 doi:10.1016/j.compag.2022.107182.

1041 Gonocruz, R.A., Nakamura, R., Yoshino, K., Homma, M., Doi, T., Yoshida,
1042 Y., Tani, A., 2021. Analysis of the rice yield under an agrivoltaic
1043 system: A case study in japan. *Environments* 8, 65. doi:10.3390/
1044 environments8070065.

1045 Hassanpour Adeg, E., Selker, J.S., Higgins, C.W., 2018. Remarkable agri-
1046 voltaic influence on soil moisture, micrometeorology and water-use effi-
1047 ciency. *PLOS ONE* 13, e0203256. doi:10.1371/journal.pone.0203256.

1048 IRENA, 2025. *Renewable Capacity Statistics 2025*. Technical Report. Inter-
1049 national Renewable Energy Agency. Abu Dhabi.

1050 Jägermeyr, J., Müller, C., Ruane, A.C., Elliott, J., Balkovic, J., Castillo, O.,
1051 Faye, B., Foster, I., Folberth, C., Franke, J.A., Fuchs, K., Guarin, J.R.,
1052 Heinke, J., Hoogenboom, G., Iizumi, T., Jain, A.K., Kelly, D., Khabarov,
1053 N., Lange, S., Lin, T.S., Liu, W., Mialyk, O., Minoli, S., Moyer, E.J.,
1054 Okada, M., Phillips, M., Porter, C., Rabin, S.S., Scheer, C., Schneider,

1055 J.M., Schyns, J.F., Skalsky, R., Smerald, A., Stella, T., Stephens, H.,
1056 Webber, H., Zabel, F., Rosenzweig, C., 2021. Climate impacts on global
1057 agriculture emerge earlier in new generation of climate and crop models.
1058 *Nature Food* 2, 873–885. doi:10.1038/s43016-021-00400-y.

1059 Karniadakis, G.E., Kevrekidis, I.G., Lu, L., Perdikaris, P., Wang, S., Yang,
1060 L., 2021. Physics-informed machine learning. *Nature Reviews Physics* 3,
1061 422–440. doi:10.1038/s42254-021-00314-5.

1062 Kingma, D.P., Ba, J., 2015. Adam: A method for stochastic optimization,
1063 in: *International Conference on Learning Representations*.

1064 Laub, M., Pataczek, L., Feuerbacher, A., Zikeli, S., Högy, P., 2022. Contrast-
1065 ing yield responses at varying levels of shade suggest different suitability
1066 of crops for dual land-use systems: a meta-analysis. *Agronomy for Sus-
1067 tainable Development* 42, 51. doi:10.1007/s13593-022-00783-7.

1068 Liu, B.Y., Jordan, R.C., 1960. The interrelationship and characteristic dis-
1069 tribution of direct, diffuse and total solar radiation. *Solar Energy* 4, 1–19.
1070 doi:10.1016/0038-092x(60)90062-1.

1071 Lorenzo, E., Narvarte, L., Muñoz, J., 2011. Tracking and back-tracking.
1072 *Progress in Photovoltaics: Research and Applications* 19, 747–753. doi:10.
1073 1002/pip.1085.

1074 Marrou, H., Dufour, L., Wery, J., 2013a. How does a shelter of solar panels
1075 influence water flows in a soil–crop system? *European Journal of Agron-
1076 omy* 50, 38–51. doi:10.1016/j.eja.2013.05.004.

1077 Marrou, H., Guilioni, L., Dufour, L., Dupraz, C., Wery, J., 2013b. Microcli-
1078 mate under agrivoltaic systems: Is crop growth rate affected in the partial
1079 shade of solar panels? *Agricultural and Forest Meteorology* 177, 117–132.
1080 doi:10.1016/j.agrformet.2013.04.012.

1081 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen,
1082 T., Lin, Z., Gimelshein, N., Antiga, L., et al., 2019. PyTorch: An impera-
1083 tive style, high-performance deep learning library, in: *Advances in Neural
1084 Information Processing Systems*, pp. 8024–8035.

1085 Ramos-Fuentes, I.A., Elamri, Y., Cheviron, B., Dejean, C., Belaud, G.,
1086 Fumey, D., 2023. Effects of shade and deficit irrigation on maize growth

1087 and development in fixed and dynamic agrivoltaic systems. *Agricultural*
1088 *Water Management* 280, 108187. doi:10.1016/j.agwat.2023.108187.

1089 Siebert, S., Ewert, F., Eyshi Rezaei, E., Kage, H., Graß, R., 2014. Impact
1090 of heat stress on crop yield—on the importance of considering canopy
1091 temperature. *Environmental Research Letters* 9, 044012. doi:10.1088/
1092 1748-9326/9/4/044012.

1093 Sparks, A., 2018. nasapower: A nasa power global meteorology, surface solar
1094 energy and climatology data client for r. *Journal of Open Source Software*
1095 3, 1035. doi:10.21105/joss.01035.

1096 Trommsdorff, M., Kang, J., Reise, C., Schindele, S., Bopp, G., Ehmann, A.,
1097 Weselek, A., Högy, P., Obergfell, T., 2021. Combining food and energy
1098 production: Design of an agrivoltaic system applied in arable and vegetable
1099 farming in germany. *Renewable and Sustainable Energy Reviews* 140,
1100 110694. doi:10.1016/j.rser.2020.110694.

1101 Valle, B., Simonneau, T., Sourd, F., Pechier, P., Hamard, P., Frisson, T.,
1102 Ryckewaert, M., Christophe, A., 2017. Increasing the total productivity of
1103 a land by combining mobile photovoltaic panels and food crops. *Applied*
1104 *Energy* 206, 1495–1507. doi:10.1016/j.apenergy.2017.09.113.

1105 Waha, K., Dietrich, J.P., Portmann, F.T., Siebert, S., Thornton, P.K., Bon-
1106 deau, A., Herrero, M., 2020. Multiple cropping systems of the world
1107 and the potential for increasing cropping intensity. *Global Environmental*
1108 *Change* 64, 102131. doi:10.1016/j.gloenvcha.2020.102131.

1109 Webber, H., White, J.W., Kimball, B.A., Ewert, F., Asseng, S., Eyshi Rezaei,
1110 E., Pinter, P.J., Hatfield, J.L., Reynolds, M.P., Ababaei, B., Bindi, M.,
1111 Doltra, J., Ferrise, R., Kage, H., Kassie, B.T., Kersebaum, K.C., Luig,
1112 A., Olesen, J.E., Semenov, M.A., Stratonovitch, P., Ratjen, A.M., LaM-
1113 orte, R.L., Leavitt, S.W., Hunsaker, D.J., Wall, G.W., Martre, P., 2018.
1114 Physical robustness of canopy temperature models for crop heat stress
1115 simulation across environments and production conditions. *Field Crops*
1116 *Research* 216, 75–88. doi:10.1016/j.fcr.2017.11.005.

1117 Weselek, A., Bauerle, A., Hartung, J., Zikeli, S., Lewandowski, I., Högy, P.,
1118 2021. Agrivoltaic system impacts on microclimate and yield of different

1119 crops within an organic crop rotation in a temperate climate. *Agronomy*
1120 for Sustainable Development 41, 59. doi:10.1007/s13593-021-00714-y.

1121 Weselek, A., Ehmann, A., Zikeli, S., Lewandowski, I., Schindele, S., Högy,
1122 P., 2019. Agrophotovoltaic systems: applications, challenges, and op-
1123 portunities. a review. *Agronomy for Sustainable Development* 39, 35.
1124 doi:10.1007/s13593-019-0581-3.

1125 Woli, P., Jones, J.W., Ingram, K.T., Fraisse, C.W., 2012. Agricultural
1126 reference index for drought (arid). *Agronomy Journal* 104, 287–300.
1127 doi:10.2134/agronj2011.0286.

1128 Yu, Q., You, L., Wood-Sichra, U., Ru, Y., Joglekar, A.K.B., Fritz, S., Xiong,
1129 W., Lu, M., Wu, W., Yang, P., 2020. A cultivated planet in 2010 – part 2:
1130 The global gridded agricultural-production maps. *Earth System Science*
1131 Data 12, 3545–3572. doi:10.5194/essd-12-3545-2020.

1132 Zhang, L., Geng, X., Ding, H., Cao, K., Wang, L., Chen, H., Deng, L., Wu,
1133 C., Xiao, M., Bao, E., 2026. Multi-objective optimization of photovoltaic
1134 array layouts on farmland via a hybrid algorithm framework for enhanced
1135 agricultural production and energy generation. *Artificial Intelligence in*
1136 *Agriculture* doi:10.1016/j.aiia.2026.09.002.

1137 Zhao, C., Liu, B., Piao, S., Wang, X., Lobell, D.B., Huang, Y., Huang, M.,
1138 Yao, Y., Bassu, S., Ciais, P., Durand, J.L., Elliott, J., Ewert, F., Janssens,
1139 I.A., Li, T., Lin, E., Liu, Q., Martre, P., Müller, C., Peng, S., Peñuelas,
1140 J., Ruane, A.C., Wallach, D., Wang, T., Wu, D., Liu, Z., Zhu, Y., Zhu,
1141 Z., Asseng, S., 2017. Temperature increase reduces global yields of major
1142 crops in four independent estimates. *Proceedings of the National Academy*
1143 *of Sciences* 114, 9326–9331. doi:10.1073/pnas.1701762114.

1144 Zhao, C., Liu, B., Xiao, L., Hoogenboom, G., Boote, K.J., Kassie, B.T.,
1145 Pavan, W., Shelia, V., Kim, K.S., Hernandez-Ochoa, I.M., Wallach, D.,
1146 Porter, C.H., Stockle, C.O., Zhu, Y., Asseng, S., 2019. A simple crop
1147 model. *European Journal of Agronomy* 104, 97–106. doi:10.1016/j.eja.
1148 2019.01.009.

Table 2: Held-out performance (92 test sites, seasons 2015–2019, canopy cooling $\kappa = 1^\circ\text{C}$). The density (GCR) of each site is the largest value on the grid 0.05, 0.075, ..., 0.5 at which the method’s mean retention in the site’s 2001–2014 seasons is ≥ 0.9 . E : electricity ($\text{MWh ha}^{-1} \text{ yr}^{-1}$); e : electricity relative to a conventional PV plant; r : mean yield retention; r_{prod} : production-weighted retention; C_s and C_{site} : share of seasons with $r \geq 0.9$ and share of sites whose test-season mean $r \geq 0.9$ (%); $\text{LER} = r + e$.

Method	GCR	E	e	r	r_{prod}	C_s	C_{site}	LER
<i>Baseline</i>								
AV static	0.06	217	0.16	0.917	0.914	93	96	1.07
Seasonal sharing	0.17	397	0.28	0.904	0.902	69	72	1.19
Phenology rule	0.17	465	0.33	0.904	0.902	69	75	1.24
Stress rule	0.13	387	0.28	0.907	0.907	76	82	1.18
SALS (ours)	0.17	510	0.36	0.905	0.902	70	74	1.27
<i>+1.5 °C</i>								
AV static	0.06	221	0.16	0.916	0.914	93	92	1.08
Seasonal sharing	0.16	393	0.28	0.905	0.903	72	74	1.19
Phenology rule	0.17	478	0.34	0.905	0.902	68	75	1.25
Stress rule	0.12	377	0.27	0.908	0.908	73	77	1.18
SALS (ours)	0.17	518	0.37	0.906	0.903	68	75	1.28
<i>+2 °C</i>								
AV static	0.06	223	0.16	0.916	0.914	92	93	1.08
Seasonal sharing	0.17	394	0.28	0.905	0.904	72	76	1.19
Phenology rule	0.17	483	0.35	0.905	0.903	69	74	1.25
Stress rule	0.12	370	0.26	0.909	0.909	75	78	1.17
SALS (ours)	0.17	521	0.38	0.907	0.904	69	75	1.28
<i>+3 °C</i>								
AV static	0.06	231	0.17	0.915	0.914	88	89	1.08
Seasonal sharing	0.17	397	0.29	0.906	0.904	69	75	1.19
Phenology rule	0.18	499	0.36	0.905	0.902	67	72	1.26
Stress rule	0.11	366	0.26	0.908	0.908	74	77	1.17
SALS (ours)	0.18	540	0.39	0.907	0.904	68	75	1.30

Table 3: LER of the controllers at different yield-retention floors ρ (main split; the historical design step is repeated for each floor).

Method	Baseline climate				+3 °C			
	0.80	0.85	0.90	0.95	0.80	0.85	0.90	0.95
AV static	1.159	1.115	1.072	1.011	1.174	1.128	1.083	1.013
Stress rule	1.380	1.282	1.182	1.077	1.364	1.269	1.172	1.071
Phenology rule	1.467	1.356	1.236	1.108	1.513	1.395	1.265	1.122
SALS (trained at $\rho = 0.9$)	1.505	1.392	1.270	1.141	1.553	1.432	1.297	1.154
SALS conditioned on ρ	1.504	1.391	1.269	1.141	1.553	1.431	1.296	1.154

Table 4: Rules and the learned controller pooled over the three spatial splits (76 distinct crop-site pairs, every third held-out pair; baseline climate; main shade response; history-based design on a GCR grid with step 0.05). The tuned phenology rule shares light while thermal time lies between 0.15 and 1.0 of the season requirement (window chosen on training-block seasons; the original rule uses 0.05 to 1.0).

Method	GCR	E/E_{PV}	Retention	Seasons ≥ 0.9 (%)	Pairs ≥ 0.9 (%)	LER
AV static	0.059	0.153	0.922	95	96	1.076
Seasonal sharing	0.151	0.249	0.915	84	91	1.164
Stress rule	0.118	0.255	0.918	83	88	1.173
Phenology rule	0.154	0.292	0.916	85	91	1.207
Tuned phenology rule	0.161	0.317	0.916	84	89	1.233
SALS	0.161	0.335	0.916	83	87	1.251

Table 5: Simulated relative yield at the shading of published field trials (static arrays, training-block sites). Observed values: rice, single trial; wheat and potato, the worse of the two seasons (the better season gave +2.7% and +11%); maize, mean of the three seasons of the dynamic system (DAV) and single values for the other two systems. The same trials enter the calibration of Section 2.6, so the comparison is not independent.

Trial	Shading	Observed	Conservative	Calibrated
Rice, Gonocruz et al. (2021)	0.25	0.84	0.77	0.79
Wheat, Weselek et al. (2021)	0.30	0.81	0.66	0.74
Potato, Weselek et al. (2021)	0.30	0.82	0.72	0.72
Maize, Ramos-Fuentes et al. (2023) (DAV)	0.35	0.83	0.70	0.70
Maize, Ramos-Fuentes et al. (2023) (half)	0.30	0.91	0.75	0.75
Maize, Ramos-Fuentes et al. (2023) (full)	0.50	0.70	0.54	0.54

Table 6: Sensitivity of the deployable light-sharing results to the crop shade response (rule controllers, baseline climate, floor $\rho = 0.9$, split 0, every fourth held-out site, 22 sites). Conservative: proportional response (main analysis); calibrated: radiation-use compensation fitted to Laub et al. (2022); field-fitted: fitted to the field trials, maize harvest index unaffected.

Shade response	Method	GCR	E/E_{PV}	Retention	LER
conservative	AV static	0.06	0.15	0.92	1.07
	Phenology rule	0.17	0.33	0.90	1.23
	Seasonal sharing	0.16	0.28	0.90	1.18
	Stress rule	0.11	0.24	0.91	1.15
calibrated	AV static	0.07	0.18	0.91	1.09
	Phenology rule	0.19	0.37	0.90	1.28
	Seasonal sharing	0.19	0.32	0.90	1.22
	Stress rule	0.13	0.29	0.91	1.20
field-fitted	AV static	0.16	0.41	0.91	1.32
	Phenology rule	0.37	0.70	0.92	1.62
	Seasonal sharing	0.37	0.61	0.92	1.53
	Stress rule	0.29	0.63	0.90	1.53

Table 7: Predicting relative yield under shade from 428 observations of 58 studies (Laub et al., 2022): leave-one-study-out cross-validation (every study held out once; observations with $s > 0$). Covariates: irradiation, climate zone, shade type, intercropping.

Model	RMSE	MAE	Median AE	Mean MAE per study
Proportional loss, $y = 1 - s$	0.430	0.305	0.210	0.224
One curve for all crops	0.372	0.265	0.193	0.208
Curve per crop type, no pooling	0.313	0.208	0.141	0.165
Curve per crop type, partial pooling	0.312	0.210	0.150	0.162
Partial pooling + covariates	0.311	0.206	0.148	0.161
Gradient-boosted trees	0.321	0.219	0.162	0.184

Table 8: Leave-one-study-out prediction error (RMSE of relative yield, observations with $s > 0$) by crop type: proportional loss, hierarchical curve with partial pooling, and gradient-boosted trees. The last columns give the difference hierarchical minus proportional with a 95% cluster-bootstrap interval over studies (negative: hierarchical better).

Group	Obs./studies	Proportional	Hierarchical	Boosted t
C3 cereals	30/10	0.214	0.187	0.205
Maize	40/10	0.225	0.235	0.248
Grain legumes	58/14	0.215	0.225	0.234
Tubers/root crops	10/2	0.250	0.254	0.257
Four crop types used by the simulator	138/30	0.220	0.222	0.234
Other types (berries, forages, fruits, vegetables)	201/29	0.528	0.361	0.369
All types	339/58	0.430	0.312	0.321

Table 9: Relative yield at 20% and 40% shading: data-driven model (median and 80% bootstrap interval over studies; crop-type curve, covariates at their reference values: mean irradiation, temperate zone, no solar-panel shade, no intercropping) and simulator under the two calibrated ensembles (mean of members; E1 fitted to hand-read curves, E2 fitted to the data-driven curves). The data-driven curve of rice is that of C3 cereals (no rice in the data set). Observations and studies: C3 cereals 41/10, maize 51/10, grain legumes 76/14, tubers 12/2.

Crop	Shading	Data-driven	Simulator, E1	Simulator, E2	Proportional
Wheat	20%	0.93 [0.88, 0.96]	0.88	0.92	0.80
Wheat	40%	0.72 [0.59, 0.78]	0.70	0.77	0.60
Rice	20%	0.93 [0.88, 0.96]	0.85	0.90	0.80
Rice	40%	0.72 [0.59, 0.78]	0.66	0.73	0.60
Maize	20%	0.73 [0.58, 0.86]	0.86	0.90	0.80
Maize	40%	0.42 [0.27, 0.59]	0.69	0.75	0.60
Soybean	20%	0.71 [0.65, 0.79]	0.77	0.74	0.80
Soybean	40%	0.40 [0.34, 0.53]	0.53	0.48	0.60
Potato	20%	0.93 [0.92, 0.97]	0.88	0.92	0.80
Potato	40%	0.71 [0.67, 0.83]	0.73	0.79	0.60

Table 10: Shade-response parameters of ensemble E2, calibrated by gradient descent through the simulator to the data-driven curves and the field trials (12 randomised fits; mean $\pm$ s.d.): radiation-use compensation c and multiplier h on the flowering-window harvest-index penalty (1 = assumed value).

Crop	c	h multiplier
Wheat	0.91 ± 0.31	0.74 ± 0.53
Rice	0.61 ± 0.25	0.83 ± 0.33
Maize	0.14 ± 0.11	0.16 ± 0.13
Soybean	0.05 ± 0.09	1.66 ± 0.38
Potato	0.62 ± 0.20	0.90 ± 0.50

Table 11: Land-equivalent ratio of rules and learned controllers under four true shade responses, with the array density designed under the true response (split 0, baseline climate, floor $\rho = 0.9$, every sixth held-out site, 21 sites). The proportional (conservative) and harsh responses lie outside the range of the calibrated ensemble.

Method	Proportional	Calibrated	Field-fitted	Harsh
AV static	1.052	1.059	1.254	1.037
Seasonal sharing	1.133	1.156	1.385	1.097
Stress rule	1.129	1.160	1.406	1.086
Feedback rule	1.139	1.164	1.392	1.089
Phenology rule	1.164	1.190	1.456	1.122
SALS (trained on proportional response)	1.178	1.209	1.489	1.132
SALS-robust (ensemble-trained)	1.179	1.204	1.491	1.127
SALS-adaptive (ensemble, biomass feedback)	1.176	1.202	1.495	1.124

Table 12: Array design under shade-response uncertainty, phenology rule, pooled over three spatial splits (76 distinct crop–site pairs; every third held-out pair). Truth: each of five members of the ensemble in turn; designs of the first group use the other four members (leave-one-out). Ensemble E2 is calibrated to the learned curves and field trials, E3 additionally carries the between-study spread (target errors 0.2). Pairs ≥ 0.9 : share whose mean retention reaches the floor. The last three rows keep the density rule of the main analysis but raise the floor used in the design.

Design	E2		E3	
	Pairs ≥ 0.9 (%)	LER	Pairs ≥ 0.9 (%)	LER
oracle (true response)	87	1.486	85	1.515
mean response	69	1.500	70	1.523
floor met in 50% of responses	65	1.510	67	1.541
floor met in 75% of responses	78	1.462	79	1.489
floor met in 90% of responses	89	1.412	88	1.409
floor 0.92 under mean response	89	1.434	85	1.468
floor 0.94 under mean response	98	1.351	93	1.407
floor 0.96 under mean response	99	1.258	98	1.313

Table 13: Assumed economic parameters (illustrative, not data). Ranges used in the sensitivity analysis are in parentheses.

Parameter	Value
Electricity price	60 USD MWh ⁻¹ (30–120)
Capex, AV single-axis tracker plant	1.0 USD Wp ⁻¹ (0.7–1.5)
Capex, conventional tracker plant	0.75 USD Wp ⁻¹
Discount rate, lifetime, O&M	6% (4–8%), 25 years, 1.5% of capex per year
Installed capacity	2.1 MWp per hectare of module area share g (0.21 kWp m ⁻²)
Crop price per tonne of dry matter	wheat 290, rice 400, maize 250, soybean 520, potato 1000 USD (0.

Table 14: Illustrative annual economics per hectare (76 crop–site pairs, three splits, main shade response, baseline climate; parameters in Table 13). Net value is relative to the open field (crop costs unchanged) in kUSD ha⁻¹ yr⁻¹ at three electricity prices; break-even is the price at which the net value is zero. The conventional plant replaces the crop.

Method	E (MWh)	Capex annuity	Crop loss	Net at 60	Net at 90	Net at 120
AV static	200	11.5	0.27	0.3	6.3	12.3
Seasonal sharing	326	29.6	0.31	-10.3	-0.5	9.2
Stress rule	336	23.2	0.30	-3.3	6.8	16.9
Phenology rule	382	30.1	0.30	-7.5	4.0	15.4
Tuned phenology rule	415	31.6	0.31	-7.0	5.5	17.9
SALS	438	31.6	0.31	-5.6	7.6	20.7
Conventional PV plant	1306	58.7	3.60	16.0	55.2	94.4

Table A.15: Paired comparison of SALS with each baseline on the 92 held-out sites of the main split (unit of analysis: site mean over test seasons). ΔLER : mean difference (SALS minus baseline) with cluster-bootstrap 95% CI; p : two-sided Wilcoxon signed-rank test; % better: share of sites where SALS has the higher LER; Δr : difference in mean yield retention with its Wilcoxon p -value.

Climate	Baseline	ΔLER [95% CI]	p	% better	Δr	p	Sites
Baseline	AV static	+0.205 [+0.176, +0.236]	$< 10^{-4}$	100	-0.011	$< 10^{-4}$	92
Baseline	Seasonal sharing	+0.088 [+0.068, +0.109]	$< 10^{-4}$	97	+0.001	0.922	92
Baseline	Phenology rule	+0.038 [+0.025, +0.053]	$< 10^{-4}$	84	+0.001	0.723	92
Baseline	Stress rule	+0.096 [+0.071, +0.124]	$< 10^{-4}$	93	-0.002	0.055	92
+1.5 °C	AV static	+0.208 [+0.180, +0.237]	$< 10^{-4}$	100	-0.010	$< 10^{-4}$	92
+1.5 °C	Seasonal sharing	+0.095 [+0.077, +0.114]	$< 10^{-4}$	98	+0.001	0.812	92
+1.5 °C	Phenology rule	+0.033 [+0.022, +0.045]	$< 10^{-4}$	86	+0.002	0.855	92
+1.5 °C	Stress rule	+0.108 [+0.083, +0.134]	$< 10^{-4}$	98	-0.001	0.041	92
+2 °C	AV static	+0.209 [+0.181, +0.238]	$< 10^{-4}$	100	-0.010	$< 10^{-4}$	92
+2 °C	Seasonal sharing	+0.097 [+0.079, +0.115]	$< 10^{-4}$	97	+0.001	0.876	92
+2 °C	Phenology rule	+0.032 [+0.022, +0.044]	$< 10^{-4}$	85	+0.002	0.496	92
+2 °C	Stress rule	+0.115 [+0.089, +0.142]	$< 10^{-4}$	98	-0.003	0.016	92
+3 °C	AV static	+0.219 [+0.190, +0.250]	$< 10^{-4}$	100	-0.009	$< 10^{-4}$	92
+3 °C	Seasonal sharing	+0.108 [+0.089, +0.129]	$< 10^{-4}$	96	+0.001	0.935	92
+3 °C	Phenology rule	+0.034 [+0.022, +0.046]	$< 10^{-4}$	86	+0.002	0.330	92
+3 °C	Stress rule	+0.131 [+0.102, +0.162]	$< 10^{-4}$	98	-0.002	0.142	92