

1 Light-sharing control for tracker-based agrivoltaics
2 under a yield floor: what learning adds and what
3 shade-response uncertainty costs

4 Quang-Vinh Dang*
5 *British University Vietnam, Hung Yen, Vietnam*

6 **Abstract**

Single-axis trackers let agrivoltaic arrays shift the daily rotation from maximum electricity towards letting light through (light sharing). We quantify its value under a yield-retention floor of 90% of open-field yield with a differentiable simulator of the whole system. The crop response to shade is the least certain part: a hierarchical model fitted to 428 observations from 58 field studies predicts a new study of the four staple crop types only to about 0.22 in relative yield, no better than proportional loss, so we turned the learned curves and four field trials into ensembles of plausible responses by gradient descent through the simulator. On the 92 held-out crop–site pairs of the main split, light-sharing rules allowed arrays 2.0–2.9 times as dense as static arrays at equal mean retention and raised the land-equivalent ratio (LER) from 1.07 to 1.18–1.24 (1.27 for a neural controller), with a phenology rule needing no forecast the best; the floor was met in 69–76% of seasons. Pooled over three splits (76 pairs), the neural controller exceeded the written rule by 0.044 LER and a rule tuned on the same seasons by 0.018 (95% interval 0.012–0.024); a CMA-ES-optimised schedule was 0.013 behind it, and training on ensembles or a farm-measurable biomass signal did not help. Design mattered more: a density chosen under the mean response left about 70% of pairs at the floor when another plausible response held (oracle 85–87%); a floor raised by 0.02 was as good as or slightly better than an ensemble-quantile rule, and no rule was safe outside the ensemble (7–33% of pairs under three fixed responses). Under illustrative costs, light sharing breaks even with the open field at 70–92 USD MWh^{−1} and static arrays at 59, and a conventional plant is more profitable above 48. For maize and soybean (53% of the upscaled 14 PWh yr^{−1}) the bias of the simulator is uncertain.

*Corresponding author
Preprint submitted to *Artificial Intelligence in Agriculture*
Email address: vinh.dq4@buv.edu.vn (Quang-Vinh Dang)

7 *Keywords:* Agrivoltaics, Light sharing, Solar tracking, Food–energy nexus,
8 Crop modelling, Differentiable simulation, Hierarchical shrinkage model

9 **1. Introduction**

10 Two of the defining challenges of this century compete for the same re-
11 source: land. The energy transition requires terawatts of new solar pho-
12 tovoltaic (PV) capacity, and ground-mounted solar farms need roughly one
13 to two hectares per megawatt. Food production must rise while crops are
14 exposed to more frequent heat waves and droughts. Warming alone is pro-
15 jected to reduce global yields of maize, wheat, rice and soybean by several
16 per cent per degree (Zhao et al., 2017; IPCC, 2021), and a new generation
17 of climate and crop models shows these impacts emerging earlier than pre-
18 viously expected (Jägermeyr et al., 2021). Croplands are also where solar
19 power potential is greatest (Adeh et al., 2019). They are flat, sunny, acces-
20 sible and close to rural electricity demand, so converting farmland into solar
21 farms is tempting and controversial at the same time.

22 Agrivoltaic (AV) systems combine crops and PV modules on the same
23 land and aim to resolve this conflict (Dupraz et al., 2011; Dinesh and Pearce,
24 2016; Weselek et al., 2019). By sharing sunlight between modules and plants,
25 a hectare can deliver more food plus energy than the two uses side by side,
26 with land-equivalent ratios (LER) well above one in field trials and mod-
27 elling studies (Dupraz et al., 2011; Amaducci et al., 2018; Valle et al., 2017;
28 Trommsdorff et al., 2021). The partial shade also changes the crop micro-
29 climate. It lowers canopy temperature and evaporative demand, conserves
30 soil moisture and improves water-use efficiency, especially in drylands (Mar-
31 rou et al., 2013b; Hassanpour Adeh et al., 2018; Barron-Gafford et al., 2019;
32 Elamri et al., 2018; Weselek et al., 2021). Agrivoltaics is therefore a can-
33 didate adaptation measure against heat and drought stress, as reviewed by
34 Mamun et al. (2022), in addition to a mitigation measure through renew-
35 able electricity. The same shade, however, also reduces the light available for
36 photosynthesis. Regulators increasingly require AV systems to retain most
37 of the reference crop yield. The German pre-standard DIN SPEC 91434, for
38 example, requires at least 66% of the reference yield (DIN, 2021), and other
39 schemes and many farmers set stricter floors. How much electricity can be
40 produced without sacrificing food security depends on the crop, the local
41 climate, the year and the growth stage.

Most AV designs fix the geometry of the array, i.e. the ground coverage ratio (GCR, module area per ground area) and the tilt or tracking rule, from generic rules of thumb. Single-axis trackers add a degree of freedom: the modules can be rotated away from the sun to let light through when the crop needs it, as demonstrated for lettuce by Valle et al. (2017), or towards the sun to maximise electricity. Choosing the rotation every day is a sequential decision problem. The value of light depends on the crop’s phenological stage, soil water and upcoming weather, and the effect of a decision appears in the harvest months later. Recent work in this journal has used evolutionary multi-objective optimisation to lay out PV arrays on farmland (Zhang et al., 2026). Machine learning for crop management has mainly used reinforcement learning with black-box simulators (Gautron et al., 2022), which is sample-inefficient and hard to scale to global, multi-decadal assessments.

Two methodological strands are relevant. Differentiable process models, which couple physical simulators with learning, allow parameters and policies to be fitted by gradient descent through the simulator (Reichstein et al., 2019; Shen et al., 2023), and we use them for both calibration and control. Decision making under deep uncertainty asks for designs that perform acceptably across plausible models and not only under the best estimate (Marchau et al., 2019), and scenario-based designs have finite-sample guarantees under convexity assumptions that we do not meet (Calafiore and Campi, 2006; Campi and Garatti, 2008); our design rules are heuristic counterparts and we compare them with a simple safety margin.

Two questions follow. First, how much of the electricity that an agri-voltaic system can produce under a food-security floor comes from the array density, and how much from dynamic light-sharing control of the trackers? Second, how robust is that answer to the weakest part of any such study, the crop response to shade, and does a controller learned from a simulator add anything to a well-designed hand-written rule? Answering them requires a simulator that couples the physics of light, PV, microclimate and crop growth, a protocol in which all controllers face the same design constraints, a check of the shade response against field data, and enough sites and seasons to separate real effects from noise.

This paper makes three contributions. The first is a quantification of light-sharing control under a yield floor. Light sharing is parameterised by one daily level between full backtracking and the rotation that minimises the shadow on the crop, and we compare static arrays, a seasonal schedule, a rule that follows crop phenology, a rule that reacts to heat and drought stress,

and a neural controller (SALS) trained through the simulator. The second is a deployable evaluation protocol. The array density of every site is chosen from its historical seasons only, with the same rule for every method, and performance is measured on held-out sites and years, over three training seeds and three spatial splits. The third is a differentiable agrivoltaic simulator that supports both, coupling hourly solar geometry, single-axis trackers with backtracking, row-to-row shading and PV electricity, the direct and diffuse light reaching the crop, canopy temperature and reference evapotranspiration under the modules, and the SIMPLE crop model (Zhao et al., 2019) extended with a flowering-stage shade penalty. All components are differentiable tensor operations (Karniadakis et al., 2021). The PV part is benchmarked against PVGIS. The crop response to shade is learned from the open data of a meta-analysis (Laub et al., 2022), checked against field trials on rice, wheat, potato and maize (Gonocruz et al., 2021; Weselek et al., 2021; Ramos-Fuentes et al., 2023), and propagated through the simulator as an ensemble.

We apply the protocol to 400 crop–site pairs that represent where maize, wheat, rice, soybean and potato are grown worldwide, sampled in proportion to their harvested area (Yu et al., 2020) with local crop calendars (Jägermeyr et al., 2021), and driven by NASA POWER daily weather for 2001–2019 (Sparks, 2018) under baseline and +1.5, +2 and +3 °C climates. The main findings are as follows.

1. Light sharing is the main lever. At the same mean yield retention of 90%, arrays with light-sharing control carry 2.0 to 2.9 times the module area of static arrays and produce 1.8–2.4 times as much electricity (Section 3.1).
2. The best hand-written rule follows crop phenology and needs no weather forecast. Rules that react to heat or drought stress are not better and depend on forecast quality (Section 3.3).
3. The crop response to shade can be learned from data. A hierarchical model fitted to 428 observations from 58 field studies predicts a held-out study of the four staple crop types to about 0.22 in relative yield, no better than proportional loss (it is better only for crop types the simulator does not represent), so between-study spread is large. The simulator is close to these curves for cereals and potato but milder than the meta-analysis for maize and soybean, whereas one maize trial is milder still; the gains for those two crops are therefore uncertain and possibly overestimated (Section 3.7).

- 117 4. A learned controller adds a small gain over the best rule: 0.044 LER
118 over the written phenology rule and 0.018 over a rule tuned on the
119 same training seasons. Training it on ensembles of shade responses, or
120 giving it a farm-measurable biomass signal, does not help, and with
121 the wider ensembles it did not retain its margin over the written rule
122 in our training budget (Section 3.8).
- 123 5. The bottleneck is the uncertainty of the crop response that enters the
124 design. A density chosen under the mean response leaves about 70% of
125 crop–site pairs at the floor when another plausible response holds (ora-
126 cle 85–87%). A rule built on the quantiles of an ensemble of responses
127 raises this by 9 points for 0.04 LER, but a floor raised by 0.02 under
128 the mean response is as good as or slightly better at equal compliance,
129 and no rule is safe outside the ensemble, so the ensemble is useful to
130 size a margin and not to replace it (Section 3.8).
- 131 6. Under illustrative costs, light sharing breaks even with the open field at
132 73 to 92 USD MWh⁻¹, learning is worth about 2 kUSD ha⁻¹ yr⁻¹ over a
133 tuned rule at 90 USD MWh⁻¹, and a conventional plant replacing the
134 crop is more profitable at every price (Section 3.9).
- 135 7. At the yield floor, agrivoltaics on 5% of the world’s staple cropland
136 could supply about 14 PWh yr⁻¹, but only with about 10 TWp of mod-
137 ules, several times the installed PV capacity of the world (Section 3.11).

138 2. Materials and methods

139 2.1. Global cropland sites, weather and crop calendars

140 We studied five staple crops that together cover about 670 Mha of har-
141 vested area and 580 Mha of physical cropland: wheat, rice, maize, soybean
142 and potato. Harvested areas (rainfed and irrigated) were taken from the
143 SPAM 2010 gridded production maps (Yu et al., 2020) and aggregated to 1°
144 cells. For each crop, 80 cells with at least 5000 ha of the crop were drawn
145 without replacement with probability proportional to harvested area. An un-
146 weighted average over the sample therefore approximates an area-weighted
147 average over global production. The 400 crop–site pairs lie in 381 distinct
148 cells in 50 countries (Fig. 2). Sowing and maturity dates came from the
149 GGCM Phase 3 crop calendar at 0.5° (Jägermeyr et al., 2021). For wheat
150 we used the winter-wheat calendar unless the coldest-month mean temper-
151 ature was below −8 °C, in which case the spring-wheat calendar was used.
152 Daily weather for 2001–2020 was retrieved from NASA POWER (Sparks,

2018): global and diffuse shortwave irradiance, maximum, minimum and dew-point temperature, precipitation and wind speed at 2 m. Each site contributes 19 seasons (sowing years 2001–2019), each simulated as the 365 days after sowing, which gives 7600 site-seasons.

Climate change was represented by the global warming levels +1.5, +2 and +3 °C relative to 2001–2020. Daily maximum, minimum and dew-point temperatures were shifted uniformly, which keeps relative humidity approximately constant and raises the vapour-pressure deficit. The atmospheric CO₂ concentration was set to 390, 450, 500 and 600 ppm for the baseline and the three warming levels. Precipitation and irradiance were left unchanged. This delta approach isolates the thermal signal of warming; its limitations are discussed in Section 4.5.

2.2. A differentiable agrivoltaic simulator

The simulator advances all site-seasons in parallel with a daily time step for the crop and an hourly time step for radiation and PV (Fig. 1). All equations are implemented as PyTorch tensor operations (Paszke et al., 2019), so that yields and electricity are differentiable with respect to the design and control variables.

Hourly irradiance. Daily global and diffuse irradiance were disaggregated into hourly values with the Collares-Pereira and Rabl (1979) and Liu and Jordan (1960) ratios and renormalised to the daily totals. Beam irradiance on the horizontal is the difference between the two. Solar position follows the standard declination and hour-angle equations.

Array geometry and PV electricity. We consider elevated horizontal single-axis trackers with north–south axes. The axis height is 1.25 module widths, e.g. 2.75 m for the common 2.2-m single-module-high (1P) tracker. The rotation ϕ is limited to $\pm 60^\circ$. The ground coverage ratio g is the module width divided by the row pitch and is restricted to $g \leq 0.5$. For 2.2-m modules this keeps a pitch of at least 4.4 m, and a free strip of at least 2.2 m between the rows when the modules are horizontal, which is the stow position during field operations. Let ψ be the solar zenith angle projected onto the east–west vertical plane. The energy-optimal rotation with backtracking (Lorenzo et al., 2011) is $\phi_{\text{bt}} = \psi - \text{sgn}(\psi) \arccos(\cos \psi / g)$ when $\cos \psi < g$, and $\phi_{\text{bt}} = \psi$ otherwise. The *light-sharing* rotation ϕ_{sh} is the feasible angle closest to edge-on

187 to the sun, which minimises the shadow on the crop. Given a light-sharing
 188 level $u \in [0, 1]$ chosen once per day, the rotation in every hour is

$$\phi = (1 - u) \phi_{\text{bt}} + u \phi_{\text{sh}}. \quad (1)$$

189 Plane-of-array irradiance combines three terms: beam irradiance on the mod-
 190 ule, reduced by the fraction of the module shaded by the neighbouring row,
 191 $\max\{0, 1 - \cos \psi / (g |\cos(\phi - \psi)|)\}$; isotropic sky diffuse; and ground-reflected
 192 light (albedo 0.2). The cell temperature follows the NOCT model (45 °C),
 193 with an hourly air temperature interpolated sinusoidally between the daily
 194 extremes. Electricity per unit land is plane-of-array irradiance $\times \eta \times g$, with
 195 $\eta = 0.21 [1 - 0.0035(T_{\text{cell}} - 25)] (1 - 0.14)$ for module efficiency, temperature
 196 losses and system losses. The installed capacity is $2.1 g \text{ MWp ha}^{-1}$. A con-
 197 ventional ground-mounted PV plant ($g = 0.4$, backtracking, no crop) is the
 198 energy reference E_{PV} .

199 *Validation of the PV model.* We compared the simulated specific yield of an
 200 isolated tracker row ($g = 0.05$, negligible row-to-row shading) with PVGIS 5.2
 201 (European Commission, Joint Research Centre, 2022), for a horizontal north-
 202 south single-axis tracker with 14% losses, at 30 randomly chosen sites (29
 203 returned data). Across sites, the simulated and PVGIS specific yields were
 204 correlated with $r = 0.971$. The simulator was on average 4.9% lower (mean
 205 absolute deviation 5.5%). This is consistent with its $\pm 60^\circ$ rotation limit and
 206 with the different irradiance databases (NASA POWER vs. SARA H2, ERA5
 207 and NSRDB). Electricity estimates are therefore, if anything, conservative.

208 *Light and microclimate at the crop.* Averaged over the row pitch, the fraction
 209 of ground reached by beam light is $1 - \min\{1, g |\cos(\phi - \psi)| / \cos \psi\}$. The
 210 fraction of sky diffuse reaching the ground is a two-dimensional isotropic
 211 sky-view factor. We pre-computed it for 25 values of g and 13 rotation
 212 angles by ray casting over 31 neighbouring rows, and interpolate it bilinearly
 213 (differentiably). The daily ratio s of the irradiance at the crop to that of the
 214 open field defines the crop's shading. The canopy maximum temperature is
 215 lowered by $\kappa(1 - s)$.

216 κ is uncertain. Field studies under PV modules report small reductions
 217 of daytime crop or canopy temperature and little change in air temperature
 218 (Marrou et al., 2013b; Barron-Gafford et al., 2019), and heat-stress thresholds
 219 calibrated on air temperature do not transfer directly to canopy temperature
 220 (Siebert et al., 2014; Webber et al., 2018). We therefore use a moderate

221 $\kappa = 1^\circ\text{C}$ as the main case and report $\kappa = 0$ (no cooling) and $\kappa = 2^\circ\text{C}$
 222 for all methods (Section 3.2). Nocturnal warming under the modules is not
 223 represented. The reference evapotranspiration under the modules, ET_0 , is re-
 224 computed with the FAO-56 Penman–Monteith equation (Allen et al., 1998)
 225 from the reduced radiation and canopy temperature. Soil water conserva-
 226 tion under the modules (Hassanpour Adeg et al., 2018; Elamri et al., 2018)
 227 therefore emerges from the model, but rain redistribution by the modules
 228 (Marrou et al., 2013a) is ignored.

229 *Crop growth.* Crop growth follows the SIMPLE model (Zhao et al., 2019)
 230 with the species parameters of its Table 1a (Table 1). Daily biomass growth
 231 is

$$\Delta B = R_{\text{crop}} \cdot \text{RUE} \cdot f_{\text{Solar}} \cdot f_{\text{CO}_2} \cdot f_{\text{Temp}} \cdot \min\{f_{\text{Heat}}, f_{\text{Water}}\}, \quad (2)$$

232 where $R_{\text{crop}} = sR$ is the radiation received by the crop. The terms are as
 233 follows.

- 234 • f_{Solar} is the canopy’s radiation interception. In SIMPLE it depends
 235 only on thermal time, with senescence accelerated by heat and drought
 236 stress. Shade therefore reduces the daily biomass increment but not
 237 leaf-area development.
- 238 • f_{Heat} decreases linearly between the heat threshold T_{heat} and the ex-
 239 treme temperature T_{ext} of the canopy maximum temperature.
- 240 • $f_{\text{Water}} = 1 - S_{\text{water}} \text{ARID}$, where the agricultural reference index for
 241 drought (Woli et al., 2012) is computed from a daily soil-water balance
 242 with curve-number runoff, drainage and root water uptake. Irrigated
 243 fields have $f_{\text{Water}} = 1$.

244 SIMPLE has a fixed harvest index and therefore ignores the high sen-
 245 sitivity of grain number to shading around flowering (Fischer, 1985; An-
 246 drade et al., 1999). It also ignores the differences in shade tolerance be-
 247 tween crop types. In the meta-analysis of Laub et al. (2022), maize and
 248 grain legumes lost yield strongly even at low shade levels, whereas C3 ce-
 249 reals and tuber crops initially lost less than proportionally; shading trials
 250 also show marked effects of shade on maize development (Ramos-Fuentes
 251 et al., 2023). We added a stage-dependent penalty. During a critical window
 252 ($0.45 \leq \text{TT}/T_{\text{sum}} \leq 0.65$), the harvest index is multiplied by $1 - h\bar{d}$, where $\bar{d}$
 253 is the mean relative radiation deficit $1 - s$ in that window. The sensitivity h

254 is highest for maize (0.6) and soybean (0.5), the crop types that Laub et al.
 255 (2022) found most shade-sensitive, lower for the C3 cereals (wheat 0.4; rice
 256 0.5, for which no specific curve was available) and lowest for potato (0.2).
 257 These values are assumptions that reproduce this ranking; they are not cal-
 258 ibrated estimates. Section 3.6 shows the resulting shade-response curves.
 259 Because RUE is linear, SIMPLE does not credit the higher light-use effi-
 260 ciency of shaded canopies under more diffuse light. The penalty, conversely,
 261 makes shading around flowering more damaging than proportional. The net
 262 effect is closer to the field evidence than either assumption alone.

263 The thermal-time requirement to maturity was calibrated for every site as
 264 the mean thermal time between the GGCM sowing and maturity dates under
 265 the 2001–2020 climate, and the canopy parameters were scaled accordingly.
 266 The cultivar is thus adapted to the present climate, and warming shortens the
 267 season. Growth stops smoothly once the thermal-time requirement is met.
 268 Yield is biomass multiplied by the (shade-adjusted) harvest index. Each site
 269 is represented by a rainfed and an irrigated sub-field in proportion to its
 270 SPAM irrigated fraction. Each sub-field carries its own array, so both yield
 271 and electricity are area-weighted over the two sub-fields.

272 2.3. SALS: simulation-trained adaptive light sharing

273 SALS learns a *control network* $u_t = \pi_\omega(\mathbf{x}_t; g)$ (Algorithm 1). It is a
 274 multilayer perceptron with two hidden layers of 64 units and a bounded-logit
 275 output, $u = \sigma(5 \tanh(z/5)) \in [0.007, 0.993]$, which keeps gradients non-zero
 276 at the extremes. It chooses the daily light-sharing level of Eq. (1). Its inputs
 277 are the information available on day t :

- 278 • the crop and soil state: relative thermal time TT/T_{sum} , radiation in-
 279 terception f_{Solar} , relative soil water and the previous day’s ARID;
- 280 • a *forecast* of the day’s maximum temperature (relative to T_{heat} and
 281 T_{opt}), irradiance and ET_0 ;
- 282 • static information: the irrigation status of the sub-field, the array’s
 283 GCR g , latitude, day of season and the crop.

284 Forecasts are perturbed with random errors, independently every day: Gaus-
 285 sian with a standard deviation of 2 °C for maximum temperature, and mul-
 286 tiplicative with standard deviations of 20% for irradiance and 15% for ET_0 .
 287 The physics always uses the true weather. The same forecast errors are

288 applied to the rule-based controllers. The controller is conditioned on g .
 289 During training, each season is simulated with a GCR close to the density
 290 that is sustainable at its training site. This is the largest value on a coarse
 291 grid (0.05, 0.10, ..., 0.50) at which seasonal sharing keeps the mean retention
 292 of the site’s training seasons at or above ρ , plus a uniform perturbation in
 293 $[-0.05, +0.10]$ truncated to $[0.05, 0.50]$. The network therefore learns to con-
 294 trol the densities that design selection will later pick. Outside the cropping
 295 season the trackers backtrack for maximum electricity.

296 For site-season i with open-field yield Y_i^{open} , let $r_i = Y_i/Y_i^{\text{open}}$ be the
 297 yield retention and $e_i = E_i/E_{\text{PV},i}$ the electricity relative to a conventional
 298 PV plant. Training minimises the penalised objective

$$\mathcal{L}(\omega) = \mathbb{E}_{i \sim \mathcal{B}} \left[-e_i + \mu_1 \max\{0, \rho - r_i\} + \mu_2 \max\{0, \rho - r_i\}^2 \right], \quad \rho = 0.9, \mu_1 = 5, \mu_2 = 30, \quad (3)$$

299 which combines an exact (hinge) and a quadratic penalty for the food-security
 300 floor ρ . Because r_i and e_i are outputs of the differentiable simulator, the
 301 exact gradient of Eq. (3) is obtained by back-propagation through the 365
 302 daily steps of every season (back-propagation through time). This is the key
 303 difference from reinforcement learning with black-box simulators (Gautron
 304 et al., 2022), which estimates gradients from sampled returns.

305 Each minibatch contains 512 site-seasons, each simulated as a rainfed and
 306 an irrigated sub-field, drawn from one of the four climates in turn. Seasons
 307 with a failed open-field crop ($Y^{\text{open}} < 0.2 \text{ t ha}^{-1}$) are excluded from training
 308 and evaluation. The network is fine-tuned through the simulator with Adam
 309 (Kingma and Ba, 2015) (learning rate 10^{-3} , cosine decay, gradient-norm
 310 clipping at 1) for 300 steps. Growth-stage logistic functions are evaluated in
 311 a numerically stable form, so that gradients remain finite even under extreme
 312 heat.

313 *Initialisation.* Trained from random weights, the controller converged to
 314 sharing light throughout the cropping window ($u > 0.99$ on all days), i.e. to
 315 the seasonal-sharing rule, because the gradient of electricity with respect to u
 316 is small and of the same sign for most days of a season. We therefore initialise
 317 the network by behaviour cloning of the phenology rule of Section 2.4: binary
 318 cross-entropy against the rule’s output (soft targets 0.95 and 0.05) on about
 319 1.7 million day-states visited by the rule in the training seasons, with 1500
 320 Adam steps at a learning rate of 3×10^{-3} . Fine-tuning then moves the
 321 policy away from this initialisation wherever doing so improves the objective.
 322 Consequently, the gain of SALS over the phenology rule measures what fine-

Algorithm 1 SALS training

Require: training site-seasons in four climates; floor ρ ; penalties μ_1, μ_2

- 1: initialise π_ω by behaviour cloning of the phenology rule
- 2: **for** step = $1, \dots, N$ **do**
- 3: sample a climate and a minibatch $\mathcal{B}$ of site-seasons (rainfed and irrigated sub-fields)
- 4: set g_i to the site’s sustainable density plus a perturbation; initialise crop and soil states
- 5: **for** day $t = 1, \dots, 365$ **do**
- 6: $u_{it} \leftarrow \pi_\omega(\mathbf{x}_{it}; g_i)$ from state and noisy forecast if the crop is in the field, else 0
- 7: hourly rotation (Eq. 1), PV electricity, crop light, canopy temperature, ET_0
- 8: advance the soil-water balance and the SIMPLE crop state (Eq. 2)
- 9: $r_i \leftarrow Y_i/Y_i^{\text{open}}, e_i \leftarrow E_i/E_{\text{PV},i}$ (area-weighted); loss $\mathcal{L}$ (Eq. 3)
- 10: back-propagate through the season; Adam update of ω

323 tuning through the simulator adds to that rule, not what can be learned
324 without it.

325 *2.4. Deployable design selection and baselines*

326 A real array has one fixed density, which must be chosen before it is built.
327 Every method, SALS included, therefore uses the same *history-based design*
328 *rule*. For a site, the GCR is the largest value on the grid $\{0.05, 0.075, \dots, 0.50\}$
329 for which the method’s mean yield retention over the site’s 2001–2014 sea-
330 sons is at least $\rho = 0.9$. If no value qualifies, no array is built and the site
331 remains open field. A stricter *chance-constrained* variant requires $r \geq 0.9$ in
332 at least 80% of the historical seasons. The historical seasons of the test sites
333 are never used for training, and the weather of the evaluation years is never
334 used for design. We compare four controllers:

- 335 • *AV static*: backtracking trackers without light sharing ($u = 0$);
- 336 • *Seasonal sharing*: the light-sharing rotation ($u = 1$) throughout the
337 cropping season;
- 338 • *Stress rule*: an expert rule that shares light ($u = 1$) during the main
339 growth phase ($0.1 < TT/T_{\text{sum}} < 0.9$) except on days with a heat-stress

340 forecast ($T_{\max} \geq T_{\text{heat}}$) or drought ($\text{ARID} > 0.5$), when the modules
 341 track the sun and shade the crop;

- 342 • *SALS*: the learned controller.

343 The open field and a conventional PV plant without crops ($\text{LER} = 1$ each)
 344 are the references. Because the design rule is identical for all controllers, dif-
 345 ferences between methods reflect the quality of control and the array density
 346 that each controller can sustain under the food-security floor. We also com-
 347 pare the controllers at identical GCR over the whole grid (matched-density
 348 frontiers).

349 2.5. A shade-response model learned from field data

350 The simulator needs a crop response to shade that is supported by ob-
 351 servations. We learned one from the open data of the meta-analysis of Laub
 352 et al. (2022) (Zenodo 10.5281/zenodo.5716091, CC BY 4.0): 428 observa-
 353 tions of yield relative to an unshaded control at relative shade reductions of
 354 0–90% from 58 studies, with crop and crop type, experiment type (artificial
 355 shading or intercropping), climate zone, irradiation and shade type. Rela-
 356 tive yield is modelled as $y(s) = \exp(-\beta s^\gamma)$ with Student- t errors ($\nu = 4$;
 357 maximum a posteriori estimation, so this is a hierarchical shrinkage regres-
 358 sion with bootstrap intervals and not a full Bayesian model). We compared
 359 six models: proportional loss $y = 1 - s$; one curve for all crops; one curve
 360 per crop type without pooling; one curve per crop type with partial pooling
 361 (Gaussian shrinkage of $\log \beta$ and $\log \gamma$ towards a common value, scale 0.5);
 362 partial pooling with covariates on $\log \beta$ (standardised irradiation, subtrop-
 363 ical zone, solar-panel shade and intercropping); and gradient-boosted trees
 364 on the same inputs. Fits use L-BFGS. Models were compared by leave-one-
 365 study-out cross-validation, in which every study is held out once, which tests
 366 the prediction of a new study and not of a new observation in a known study.
 367 Intervals for the crop-type curves and covariate effects come from a cluster
 368 bootstrap over studies (150 resamples).

369 2.6. Shade-response uncertainty: gradient calibration, robust controllers and 370 design

371 The crop’s response to shade is the least certain part of the simulator,
 372 so we treat it as an unknown with a distribution. Two parameters per crop
 373 enter the simulator: a radiation-use compensation c (growth is multiplied by

374 $1 + c(1 - s)$, s being the share of open-field irradiance that reaches the crop)
 375 and a multiplier h on the flowering-window harvest-index penalty. Because
 376 the simulator is differentiable, both can be fitted by gradient descent through
 377 whole simulated seasons to published yield-shading observations. The ob-
 378 servations are the relative yields at 20% and 40% shading of the data-driven
 379 curves of Section 2.5 for the five crops (median, with an error equal to the
 380 width of the 80% bootstrap interval divided by 2.56 and at least 0.04; rice
 381 takes the curve of C3 cereals, which contains no rice), the field trials on rice,
 382 wheat, potato and maize (Gonocruz et al., 2021; Weselek et al., 2021; Ramos-
 383 Fuentes et al., 2023) (each season of each trial as a separate observation, error
 384 0.06–0.08) and the unaffected harvest index of shaded maize (Ramos-Fuentes
 385 et al., 2023). We call the resulting ensemble E2. A wider ensemble, E3, uses
 386 the same procedure with a target error of 0.2 for the learned curves, which is
 387 the between-study spread of a new study (Section 3.7). An earlier ensemble,
 388 E1, was fitted in the same way to curves read by hand from the figures of
 389 Laub et al. (2022) and was used to train the robust and adaptive controllers
 390 below. Priors are $c \sim \mathcal{N}(0.4, 0.5^2)$ and $h \sim \mathcal{N}(1, 0.5^2)$, and the param-
 391 eters are clamped to $[0, 2.5]$ (c) and $[0, 2]$ (h) after every step, which is why
 392 a response milder than the data cannot be lowered further. We fitted 12
 393 ensemble members by randomised maximum a posteriori estimation: every
 394 member sees the observations perturbed by their errors and a prior draw
 395 perturbed by its width, and is optimised with Adam (40 steps, with the loss
 396 still falling slightly at the last step; 40 training-block seasons per crop and
 397 member; shading of the static arrays cached). The ensemble is a device for
 398 generating plausible responses and not a validated posterior: it has no cov-
 399 erage check, and for some parameters (the harvest-index multiplier of wheat
 400 and potato) its spread equals that of the prior. The design analyses use five
 401 fixed members (indices 0, 3, 5, 8 and 10 of 12), chosen in advance. The
 402 ensemble is used in three ways. All training and evaluation ran on a 4-core
 403 CPU; one training run takes about 20–25 minutes and the calibration of one
 404 ensemble about 7 minutes.

405 *Robust and adaptive controllers.* SALS-robust has the architecture and
 406 training of SALS (Section 2.3), but every training season is simulated with the
 407 response of a randomly drawn ensemble member that the controller does not
 408 observe, and the density jitter is widened to $[-0.05, +0.25]$ because feasible
 409 densities differ between members. SALS-adaptive additionally observes the
 410 realised biomass of the crop under the modules relative to an open reference
 411 plot, a quantity that can be measured on the farm by remote sensing or by

a reference plot. A feedback rule that shares light whenever this ratio falls below a target serves as the hand-written counterpart; its target (1.0) was chosen on training-block seasons under the ensemble.

Uncertainty-aware design. The design rule of Section 2.4 chooses the largest GCR whose historical retention meets the floor. Under shade-response uncertainty we choose the largest GCR that meets the floor in at least 75% of five ensemble members, and compare it with a design under the ensemble mean and with an oracle design under the true response. Evaluation of this rule is leave-one-out: the truth is the held-out member and the design uses the other four.

Stress tests. Controllers and rules are also evaluated under four fixed responses, the proportional response, the response calibrated to the meta-analysis, the response fitted to the field trials and a harsh response (proportional growth and a 1.5 times larger harvest-index penalty). The proportional and harsh responses lie outside the range of the ensemble.

2.7. Economic analysis

To judge whether the electricity gains matter economically, we computed an annual net value per hectare for each method relative to the open field: electricity revenue minus the annuity and operating cost of the installed capacity minus the loss of crop revenue (crop production costs are unchanged). The installed capacity is 2.1 g MWp per hectare (modules cover a share g of the land at 0.21 kWp m⁻²). Yield is dry matter, valued at assumed prices. A conventional ground-mounted plant (GCR 0.4) that replaces the crop is the comparator. All parameters (Table 14) are assumptions chosen to be representative, not data, and their ranges drive the conclusions; we therefore report break-even electricity prices and one-at-a-time sensitivities, and we do not monetise food security, which market crop prices do not capture.

2.8. Evaluation protocol, statistics and upscaling

Sites were split spatially: 10°×10° blocks were randomly assigned to training (70%) or test (30%), so that no test site has a training site in the same block. Training uses the 2001–2014 seasons of the training sites. Design selection uses the 2001–2014 seasons of the test sites. Evaluation uses the 2015–2019 seasons of the test sites. The same site-seasons are analysed in all four climates.

We report the following metrics:

- 447 • the mean electricity ($\text{MWh ha}^{-1} \text{yr}^{-1}$) and the relative electricity e ;
- 448 • the mean yield retention r and the production-weighted retention $\sum Y / \sum Y^{\text{open}}$;
- 449 • season-level compliance (share of seasons with $r \geq \rho$) and site-level
450 compliance (share of sites whose mean test-season retention is $\geq \rho$);
- 451 • the lower tail (10th percentile of r , share of seasons with $r < 0.8$);
- 452 • the land-equivalent ratio $\text{LER} = r + e$ (Dupraz et al., 2011).

453 Two aggregations are used and labelled in the text: tables report means over
454 all test seasons of a sample (season-pooled), whereas statistical comparisons
455 and intervals use the crop–site pair as the unit of analysis (mean of the pair’s
456 seasons, then mean over pairs), so the same contrast can differ slightly be-
457 tween the two (for example 0.034 and 0.038 for SALS minus the phenology
458 rule on the main split). Samples are also labelled: the main analysis uses
459 the 92 held-out pairs of one split and a GCR grid with step 0.025; analyses
460 pooled over three splits use the 76 distinct pairs of every third held-out pair
461 and a grid with step 0.05, and analyses of alternative shade responses on
462 one split use every sixth (21 pairs) or every fourth pair (22 pairs). Statis-
463 tical comparisons use site means as the unit of analysis. We report paired
464 Wilcoxon signed-rank tests and cluster (site) bootstrap 95% confidence in-
465 tervals of the mean difference. SALS was trained with three random seeds
466 on the main split, and with one seed on each of two further spatial splits.

467 For the global estimate, every site is evaluated in the split in which it is
468 held out. The mean electricity and yield per hectare of each crop are scaled
469 to the *physical* area of that crop in SPAM 2010, which avoids double counting
470 multiple-cropped land (Waha et al., 2020). We report the installed capacity
471 (TWp), the capacity factor and a supply curve.

472 3. Results

473 3.1. Electricity and food on held-out sites and years

474 Table 2 compares the four controllers on the 92 held-out crop–site pairs of
475 the main split (95 pairs are held out; pairs without a season whose open-field
476 yield is at least 0.2 t ha^{-1} in all four climates are excluded) (seasons sown
477 in 2015–2019), after each site’s density has been chosen from its 2001–2014
478 seasons with the same rule for all controllers. By construction all methods

Table 1: SIMPLE crop parameters (Zhao et al., 2019, Table 1a) and the harvest-index shade sensitivity h added in this study. T_{sum} , I_{50A} and I_{50B} are rescaled per site to the local crop calendar (Section 2.2). Units: temperatures in $^{\circ}\text{C}$, thermal times in $^{\circ}\text{C d}$, RUE in gMJ^{-1} ; S_{CO_2} is the relative RUE increase per 100 ppm above 350 ppm.

Crop (cultivar)	h	T_{sum}	HI	I_{50A}	I_{50B}	T_{base}	T_{opt}	RUE	$I_{50 \text{ max } H}$	$I_{50 \text{ max } W}$	T_{heat}	T_{ext}	S_{CO_2}	S_{water}
Wheat (Batten)	0.4	2150	0.34	280	50	0	15	1.24	100	25	34	45	0.08	0.4
Rice (IR72)	0.5	2300	0.47	850	200	9	26	1.24	100	10	34	50	0.08	1.0
Maize (McCurdy 84aa)	0.6	2050	0.50	500	50	8	28	2.10	100	12	34	50	0.01	1.2
Soybean (Williams 82)	0.5	2350	0.40	600	200	6	27	0.86	120	20	36	50	0.07	0.9
Potato (Sebago)	0.2	2400	0.85	500	350	4	22	1.30	50	30	34	45	0.10	0.4




Figure 1: Simulator and controller. The array density g of a site is fixed from its historical seasons. A control network maps the daily crop and soil state, the density g and a noisy forecast to the light-sharing level u_t . The differentiable simulator converts these decisions into PV electricity, light and microclimate at the crop, soil water and crop growth. The season’s electricity and yield retention enter a penalised objective whose exact gradient is back-propagated through the 365-day season.

479 keep the mean yield retention close to 90% (0.904–0.917), so differences in
 480 electricity are differences in the density that each controller can sustain under
 481 this floor.

482 **Static arrays are the weakest.** Without light sharing, the floor limits
 483 the ground coverage to 6% on average, and arrays produce $217 \text{ MWh ha}^{-1} \text{ yr}^{-1}$,
 484 16% of a conventional PV plant per hectare (LER 1.07). Rotating the
 485 modules to let light through during the cropping season (seasonal shar-
 486 ing) allows a coverage of 17% and $397 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ (LER 1.19). A rule
 487 that shares light only while the crop is growing (phenology rule) reaches
 488 $465 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ (LER 1.24). An expert rule that additionally shades the
 489 crop on forecast heat or drought days (stress rule) is worse than the phenol-
 490 ogy rule (LER 1.18): its density is limited to 13% because it reduces light at




Figure 2: The 400 crop–site pairs, sampled in proportion to the harvested area of each crop (SPAM 2010). Blue: training sites; orange: geographically held-out test sites of the main split.

491 times when the floor is most binding.

492 **The learned controller is best by a modest, consistent margin.**
493 SALS sustains a density of 17% (similar to the phenology rule), produces
494 $510 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ and reaches an LER of 1.270 at the same retention. Com-
495 pared with the phenology rule this is $+0.038$ in LER (95% CI over sites
496 $0.025\text{--}0.053$; $+10\%$ electricity) and it is higher at 84% of the 92 sites (Ta-
497 ble A.17). Compared with static arrays it is $+0.205$ ($0.176\text{--}0.236$), with 2.35
498 times the electricity at 2.9 times the density. The frontier in Fig. 3 shows
499 the same picture at matched densities: the learned policy lies on or slightly
500 above the phenology rule over the whole density range, both lie well above
501 seasonal sharing and the stress rule, and static arrays are far below.

502 **What the floor does and does not guarantee.** Enforcing the floor
503 on the historical mean does not guarantee it in every season. In the test
504 seasons, 70% of SALS seasons have $r \geq 0.9$ (phenology rule 69%, stress rule
505 76%, static arrays 93%), and 2.7% have $r < 0.8$ (1.3%, 0.9% and 0.9%).




Figure 3: Electricity-retention frontiers at identical array densities (labels: GCR) for the held-out sites. Each point is the mean over test seasons at a fixed GCR; the dashed line is the 90% floor.

506 The lower tail of SALS is therefore slightly heavier than that of the rules:
 507 its 10th percentile of r is 0.883, against 0.887, 0.886 and 0.903. Requiring
 508 $r \geq 0.9$ in at least 80% of a site’s historical seasons (chance-constrained
 509 design) raises the season-level compliance to 90–94% for the light-sharing
 510 controllers. It costs about 0.065 in LER for SALS (1.205), which remains the
 511 best controller (phenology rule 1.192; Section 3.2).

512 3.2. Robustness

513 **Seeds and splits.** Three training seeds of SALS give LERs of 1.270,
 514 1.268 and 1.269 under the baseline climate and 1.297, 1.298 and 1.297 at
 515 +3 °C. On two further spatial splits, in which different blocks of sites are
 516 held out, SALS exceeds the phenology rule by 0.047 and 0.041 in LER under
 517 the baseline climate and by 0.052 and 0.039 at +3 °C (main split, season-
 518 pooled: 0.034 and 0.032; the site-level paired contrasts of Table A.17 are
 519 0.038 and 0.034). In every split SALS ranks first, the phenology rule second
 520 and static arrays last.

521 **Canopy cooling.** The canopy-cooling coefficient κ is the most un-
 522 certain microclimate parameter. We repeated the whole protocol, includ-
 523 ing the historical design step, with $\kappa = 0, 1$ and 2 °C. Electricity changes
 524 with κ because cooler canopies tolerate denser arrays (SALS: 475, 510 and
 525 549 MWh ha⁻¹ yr⁻¹ under the baseline climate), but the retention is held at
 526 0.905 by the design rule and the ranking is unchanged. SALS exceeds the

phenology rule by 0.029, 0.034 and 0.042 in LER under the baseline climate (0.022, 0.032 and 0.051 at +3 °C).

Which ingredients matter. Two ablations test the design of SALS at the same training budget. Masking the stress inputs (soil water, ARID, and the maximum temperature, irradiance and ET_0 forecasts) leaves the LER unchanged (1.269 vs 1.270 under the baseline climate, 1.297 vs 1.297 at +3 °C). Training only on the baseline climate gives 1.268 and 1.291. The improvement over the rules therefore does not depend on stress-aware inputs or on training in warmed climates.

3.3. Floor conditioning, chance-constrained training and forecast errors

One network for all floors. Regulations and farmers set different floors. The network of the main experiments was trained at $\rho = 0.9$ only. We evaluated it at floors of 0.80, 0.85, 0.90 and 0.95, re-running the historical design step for each floor (Table 3). Its LER exceeds that of the phenology rule by 0.037, 0.036, 0.034 and 0.033 under the baseline climate (0.040, 0.037, 0.032 and 0.033 at +3 °C), so the benefit carries over to other floors without retraining. A network that received the floor as an additional input and was trained over all four floors did not improve on it (differences in LER ≤ 0.001). Because the design step already adapts the density to the floor, conditioning the controller on the floor is not needed.

Chance-constrained training. We also trained a network whose objective penalises a batch share of seasons below the floor that exceeds 20%. It did not improve the tail. With the mean-based design step its LER was 1.263 (baseline) and 1.290 (+3 °C) against 1.270 and 1.297, and with chance-constrained design it reached season-level compliance of 92.6% against 92.0% and an LER of 1.198 against 1.205. The season-level guarantee is obtained by the design step (Section 3.1), not by the training objective.

Forecast errors. We raised the forecast errors seen by the controllers to 0, 1, 2 and 3 times the errors used in training. The LER of SALS and of the phenology rule was identical at all four error levels to three decimals (densities 0.2 and 0.3: SALS 1.307 and 1.457, phenology rule 1.278 and 1.416). Neither controller uses the weather forecast: both rely on the crop’s thermal time, which is observable from on-farm temperature records. The stress rule, which depends on the forecast of maximum temperature, degrades with forecast error: at density 0.3 its retention falls from 0.778 to 0.751 and its LER from 1.432 to 1.418 as the errors triple.

563 **A tuned rule.** The phenology rule above was written, not tuned,
564 whereas SALS is trained. We therefore tuned the window of the rule on
565 the training-block seasons with the same objective (electricity at the density
566 where mean retention equals the floor); the best window starts at 0.15 in-
567 stead of 0.05 of the thermal-time requirement (Table 4). Pooled over three
568 splits and 76 distinct crop–site pairs, the tuned rule raises LER over the orig-
569 inal rule by 0.026 (95% cluster-bootstrap interval 0.016 to 0.037), and SALS
570 exceeds the tuned rule by 0.018 (0.012 to 0.024), or 5.7% more electricity at
571 equal retention, and the original rule by 0.044 (0.030 to 0.060). Tuning the
572 rule thus accounts for about 60% of the margin of SALS over the written
573 rule. A five-parameter ramp-plateau schedule optimised by CMA-ES on the
574 same training seasons performs like the tuned window (LER 1.238 against
575 1.233, difference +0.005, interval -0.001 to $+0.010$), and SALS exceeds it
576 by 0.013 (0.007 to 0.019); the optimum is an almost binary window (start
577 0.17, end 0.93, level 0.99), so more flexible schedules of this family do not
578 close the gap. The remaining margin of SALS is small but not zero; it is not
579 explained by the tuning of a window or by a smooth ramp, and we did not
580 test richer rule families (for example schedules that depend on soil water).
581 SALS meets the floor in fewer seasons than the rules (83% against 84–85%).

582 3.4. *What the learned policy does*

583 Averaged over the rainfed test seasons of the main split at a density of 0.3,
584 the light-sharing level of SALS is a smooth function of relative thermal time.
585 It is close to zero in the first 10% of the thermal time required for maturity,
586 rises through 0.55 at 10–20%, stays at 0.94–0.99 between 30% and 70%, and
587 falls through 0.58 at 80–90% to about zero at maturity. The plateau covers
588 the window in which our crop model penalises shading through the harvest
589 index. The ramps at both ends let the modules produce electricity when
590 additional light is worth least to the crop. Fig. 4 shows two example seasons.
591 The phenology rule, which is also the initialisation of the network, shares
592 light during the whole growing period. Fine-tuning through the simulator
593 found that sharing can be reduced at the edges of the season: SALS gains its
594 advantage by yielding light when it matters and collecting electricity when
595 it matters less.

596 The policy also shares less light on days with a forecast heat stress: the
597 mean level is 0.35 on such days against 0.60 (baseline) and 0.51 ($+3^\circ\text{C}$)
598 on other days, and 58–63% of sites show a lower level on heat days. This
599 cannot be read as a stress response, however. Heat days occur late in the




Figure 4: Learned daily control for two test seasons under $+3^{\circ}\text{C}$, at the density chosen for the site, simulated for the dominant sub-field of the site. Shading marks the cropping window. Orange: light-sharing level u_t ; green: fraction of open-field irradiance reaching the crop; grey: daily maximum temperature (right axis).

season, when the phenological ramp is already falling, and a policy without any stress input shows a similar contrast (0.37 vs 0.57 in a separate check). The ablation above confirms that stress information does not improve the result.

3.5. Crops and regions

The benefit of light sharing differs between crops (baseline climate, SALS). Soybean and maize tolerate the densest arrays (GCR 0.24 and 0.19; 691 and 577 $\text{MWh ha}^{-1} \text{yr}^{-1}$; LER 1.41 and 1.30), potato follows (0.19; 542; 1.31), and rice and wheat are the most restrictive (0.13 and 0.15; 395 and 428; 1.20 for both). The stress rule has a lower LER for every crop (1.26 for soybean, 1.14–1.21 for the others).

The floor is not met evenly. Production-weighted retention is 0.888 for soybean and 0.898 for maize, below 0.9, because high-yielding seasons lose more than low-yielding ones. Regionally, Latin America and Oceania (4 sites) combine the highest density (GCR 0.41) and electricity (1090 $\text{MWh ha}^{-1} \text{yr}^{-1}$) with a production-weighted retention of only 0.862. Africa and the Middle East (8 sites) reach 0.894, and only 63% of their sites keep a mean retention of at least 0.9. South Asia, the region with the most sites (39), reaches 0.902

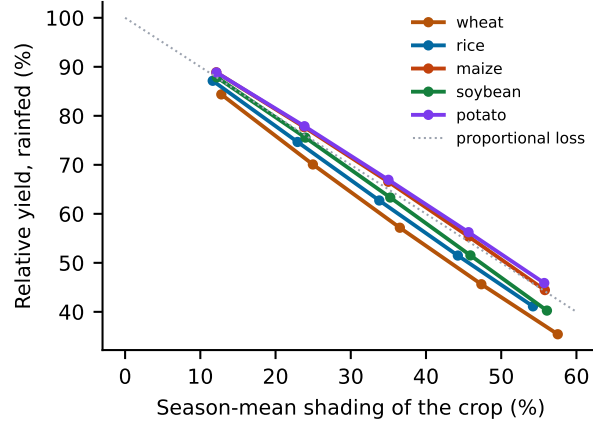


Figure 5: Simulated relative yield of rainfed crops under static arrays (GCR 0.1–0.5) against season-mean shading of the crop, baseline climate, held-out sites. The dotted line is proportional loss.

618 (67% of sites) at a density of 0.14. A regulatory floor enforced per site would
619 therefore need a region- and crop-specific design step, not a global one.

620 3.6. Model behaviour and validation

621 **PV output.** The simulated specific yield of an isolated tracker row agrees
622 with PVGIS 5.2 (European Commission, Joint Research Centre, 2022) at 29
623 sites with $r = 0.971$, a mean bias of -4.9% and a mean absolute deviation
624 of 5.5% (Section 2.2).

625 **Shade response.** Fig. 5 shows the simulated relative yield of rainfed
626 crops under static arrays as a function of the season-mean shading of the crop.
627 Yield declines close to proportionally with shading: 84–89% relative yield at
628 12% shading, 70–78% at 25%, 57–67% at 37% and 35–46% at 57%. Wheat
629 loses most and potato and maize least. This response is steeper than the *less*
630 *than proportional* losses that Laub et al. (2022) report for small shade levels
631 of C3 cereals and tubers, and it does not reproduce their finding that maize
632 and grain legumes are more shade sensitive than C3 cereals. In the simulator,
633 shading also relieves heat and drought in rainfed maize and soybean, which
634 offsets the higher harvest-index penalty we assigned to them. The absolute
635 effect of shading on each crop is therefore uncertain (Section 4.5).

636 **Comparison with field trials.** We fitted one radiation-use compen-
637 sation coefficient per crop to the yield–shading curves of Laub et al. (2022)

638 using training-block sites only (Methods, calibration script in the repository).
 639 The fit changed wheat and rice only (coefficients 0.4 and 0.1) because the
 640 proportional model already lies within about 0.02 of the C3-cereal and tu-
 641 ber curves, and it cannot lower the maize and legume responses below the
 642 proportional case (root-mean-square error 0.16 for maize and 0.08 for soy-
 643 bean). Against independent field trials (Table 5) the model underestimates
 644 retained yield: rice 0.79 against 0.84 (Gonocruz et al., 2021), wheat and
 645 potato 0.74 and 0.72 against 0.81–0.82 in the worse season of Weselek et al.
 646 (2021) (and above 1 in the better one), and maize 0.70–0.75 against 0.83–0.91
 647 at 30–35% shading (Ramos-Fuentes et al., 2023), where harvest index was
 648 not affected by shade. The model is therefore conservative for food reten-
 649 tion with respect to these trials, except for the stronger maize losses of the
 650 meta-analysis. Table 6 shows the consequence for the design results: with
 651 the calibrated response the allowed density and the LER of every method
 652 rise (phenology rule 1.28 against 1.23), and with a response fitted to the
 653 field trials the LER of the phenology rule is 1.62. The phenology rule is the
 654 best method in all three cases, and light sharing adds LER +0.17, +0.19 and
 655 +0.30 over static arrays (conservative, calibrated, field-fitted); only the order
 656 of the seasonal and stress rules changes. The absolute gains for rice, wheat
 657 and potato are thus conservative with respect to these trials; for maize and
 658 soybean the meta-analysis points the other way (Section 3.7).

659 *3.7. A shade-response model learned from field data*

660 Cross-validation by study shows what the data can and cannot support
 661 (Tables 7 and 8). Over all 339 shaded observations the root-mean-square
 662 error of a held-out study falls from 0.43 for proportional loss $y = 1 - s$ to 0.37
 663 for one curve for all crops and to 0.31 for crop-type curves, but this gain comes
 664 from crop types that the simulator does not represent (berries, forages, fruits
 665 and vegetables: 0.53 to 0.36). For the four crop types that the simulator uses,
 666 the learned curves are not better than proportional loss: the error is 0.222
 667 against 0.220 pooled (difference +0.002, 95% interval -0.023 to $+0.030$),
 668 0.187 against 0.214 for C3 cereals, and slightly worse for maize, grain legumes
 669 and tubers. Partial pooling is as good as separate curves (0.312 against 0.313)
 670 and the covariates add almost nothing (0.311). Gradient-boosted trees on the
 671 same inputs are worse (0.321) than the parametric model. The error of 0.22
 672 for the four crop types is large: yields of a new study differ from the learned
 673 curve by about 20 percentage points, which is the between-study spread
 674 that a design under uncertainty has to cover. Crop type is the information

675 that transfers between studies, and none of the covariate effects on $\log \beta$ is
676 clearly different from zero (irradiation -0.10 , 90% interval -0.32 to 0.09 ;
677 subtropical zone -0.37 , -0.73 to 0.07 ; intercropping 0.16 , -0.48 to 0.45).
678 The estimated effect of solar-panel shade is not interpretable, because all 46
679 panel observations are of leafy vegetables, mainly lettuce, and the effect is
680 confounded with crop type.

681 The learned curves show a clear ordering of crop types (Table 9). At
682 20% shading C3 cereals and tubers lose less than proportionally (relative
683 yield 0.93), whereas maize (0.73) and grain legumes (0.71) lose more; at 40%
684 shading the values are 0.72 and 0.71 against 0.42 and 0.40 . Intervals are wide
685 for maize (0.27 – 0.59 at 40%) and for tubers, which rest on two studies. The
686 simulator’s responses are close to the data for cereals and potato but, under
687 both calibrated ensembles, too mild for maize and soybean at 40% shading
688 (0.75 and 0.48 for E2 against 0.42 and 0.40 in the data). The calibration
689 cannot lower the maize response further, because the simulated shading also
690 relieves heat and drought stress and the compensation and harvest-index
691 parameters are bounded below at zero. The field trials of maize are also
692 milder than the meta-analysis: Ramos-Fuentes et al. (2023) report a harvest
693 index unaffected by shade, and Sekiyama and Nagashima (2019) a yield 5.6%
694 above the control under a low-density stilt-mounted array in one season at
695 one site. For maize and soybean the direction of the bias is uncertain, and
696 the gains reported in this paper are possibly overestimated.

697 3.8. Learning under shade-response uncertainty

698 **Calibration.** Gradient descent through the simulator fitted the five crops
699 to the data-driven curves and the field trials (ensemble E2, Table 10). The
700 data identify two features: the harvest index of maize is hardly affected by
701 shade (multiplier 0.16 ± 0.13 , $c \approx 0.14$), as observed by Ramos-Fuentes et al.
702 (2023), and soybean is more shade sensitive than the cereals (compensation
703 0.05 , harvest-index multiplier 1.66), as in Laub et al. (2022). Wheat has the
704 largest compensation (0.91 ± 0.31) and the largest uncertainty; the compen-
705 sation is an effective parameter that absorbs diffuse-light and light-saturation
706 effects that the simulator does not model explicitly, not a physical value.

707 **Learned controllers.** Controllers were evaluated under four fixed re-
708 sponses on 21 crop–site pairs of split 0 with one training seed, so the intervals
709 below, which are cluster bootstraps over pairs of per-pair mean LER differ-
710 ences, do not include seed or split variation (Table 11 gives season means,

711 which differ slightly). Under every response SALS, trained on the propor-
 712 tional response, is better than the written phenology rule: by +0.019 (0.007
 713 to 0.039), +0.024 (0.009 to 0.045), +0.038 (0.017 to 0.063) and +0.015 (0.000
 714 to 0.036) LER under the proportional, calibrated, field-fitted and harsh re-
 715 sponses. Training on the narrow ensemble E1 did not change this: SALS-
 716 robust differs from SALS by -0.005 to $+0.003$ (intervals contain zero except
 717 for the proportional response, 0.000 to 0.003), and SALS-adaptive, which
 718 sees the measured biomass ratio, by -0.010 to $+0.007$ (all intervals contain
 719 zero). Training on the wider ensembles did not help either: SALS-robust on
 720 E2 and E3 and SALS-adaptive on E3 are below SALS by 0.014, 0.017 and
 721 0.017 LER under the proportional response, 0.017, 0.021 and 0.020 under the
 722 calibrated response, 0.004, 0.005 and 0.011 under the field-fitted response (in-
 723 tervals contain zero) and 0.008, 0.014 and 0.017 under the harsh response;
 724 the intervals exclude zero in 6 of the 9 comparisons for the proportional, cal-
 725 ibrated and harsh responses. The E3 robust controller is not distinguishable
 726 from the written rule ($+0.001$ to $+0.003$, intervals contain zero) except un-
 727 der the field-fitted response ($+0.032$, 0.001 to 0.067). A tuned window and a
 728 CMA-ES ramp schedule, also evaluated under the four responses (Table 11),
 729 are 0.005 to 0.018 LER below SALS (intervals of the SALS contrast: $+0.004$
 730 to $+0.016$ against the tuned window under the proportional response, and
 731 containing zero for the ramp under the calibrated and field-fitted responses
 732 and for the window under the harsh response), and the ensemble-trained
 733 controllers are at or below the tuned window under all responses except the
 734 field-fitted one. With one seed and 21 pairs these controllers therefore did not
 735 retain the margin of SALS within our training budget; a controller that must
 736 perform across plausible responses is less specialised, but a larger budget or
 737 more seeds could change this. The feedback rule is worse than the written
 738 rule by 0.026–0.058 LER under every response (intervals exclude zero). A
 739 plausible reason is that the biomass ratio does not reveal the harvest-index
 740 loss from shading around flowering, but we did not test this explanation.

741 **Design under uncertainty.** The larger risk is in the design step (Ta-
 742 bles 12 and 13, Fig. 6). The designs below are evaluated by leaving one
 743 ensemble member out as the truth and using the other four, so that the
 744 truth never enters the design. Pooled over the three spatial splits (76 dis-
 745 tinct crop–site pairs), a density chosen under the mean response of E2 leaves
 746 70% of the pairs at the floor when the true response is another plausible one,
 747 against 87% for an oracle design (E3: 71% and 85%). Choosing the density
 748 so that the floor is met in at least 75% of the other members raises compli-




Figure 6: Share of crop-site pairs at the floor against LER for the phenology rule under design rules for shade-response uncertainty (leave-one-member-out, three splits, 76 pairs): a higher floor under the mean response (blue; the first point is the design under the mean response) and a floor met in a share of the other members (orange), with the oracle design (green).

749 ance by +7.9 points under E2 (95% cluster-bootstrap interval 4.5 to 11.6)
 750 and +7.9 under E3 (3.9 to 12.4), at a cost of 0.031 and 0.025 in LER. Raising
 751 the floor used in the design from 0.90 to 0.92 under the mean response does
 752 more: +18.0 points (13.4 to 22.6) at 0.068 LER under E2 and +14.2 points
 753 (9.7 to 19.2) at 0.056 under E3, and the pairs below the floor fall short by less
 754 (mean shortfall 0.0022 against 0.0045 for the quantile rule under E2). The
 755 two kinds of rule lie on one frontier (Fig. 6): at about 88–89% compliance the
 756 ensemble rule that requires all other members to meet the floor gives LER
 757 1.412 under E2 against 1.426 for a floor of 0.92, and under E3 1.409 against
 758 1.458, so a simple margin is as good as or better than the ensemble rule by
 759 0.01 to 0.03 LER. The ensemble is thus useful for sizing the margin (0.02 to
 760 0.04 above the floor for 85–97% of the pairs), not for replacing it. None of
 761 these rules is safe outside the ensemble (Table 13): under the proportional,
 762 calibrated and harsh responses the mean-response design leaves 7–17% of the
 763 pairs at the floor, a floor of 0.92 or the ensemble rules 12–33%, a floor of 0.94
 764 29–47% and a floor of 0.96 51–78%, at a cost of about 0.2 LER, whereas the
 765 field-fitted response, which lies inside the range spanned by the ensembles, is
 766 handled well by every rule (80–99%) except the mean-response design (68–
 767 79%). The guarantee therefore holds only within the range of responses that
 768 the ensemble spans, and a response harsher than the data supported by the
 769 ensemble breaks it.

770 3.9. Economics of light sharing

771 All figures in this section follow from the assumed parameters of Table 14
772 and should be read as a sensitivity analysis, not a forecast. The economics
773 are dominated by the capacity installed, not by the crop (Table 15): the
774 annuity of the modules is 11–32 kUSD per hectare and year, whereas the
775 crop revenue lost is about 0.3 kUSD, so a method is worth what it earns per
776 installed watt. With the densities designed for the floor, static arrays earn
777 +0.3 kUSD ha⁻¹ yr⁻¹ at 60 USD MWh⁻¹ and break even at 59, whereas all
778 light-sharing methods lose money at that price (–3 to –10 kUSD) and break
779 even at 70 (stress rule), 73 (SALS), 77 (tuned phenology rule), 80 (written
780 rule) and 92 (seasonal sharing) USD MWh⁻¹. At 90 USD MWh⁻¹ static
781 arrays (6.3 kUSD) are better than the written (4.0) and tuned (5.5) rules
782 and worse than SALS (7.6), which overtakes static arrays above about 85
783 USD MWh⁻¹. A conventional plant that replaces the crop breaks even at 48
784 USD MWh⁻¹ and is more profitable than every agrivoltaic option above that
785 price (16 kUSD at 60, 94 kUSD at 120, against –6 and 21 for SALS); below
786 it all options lose money and static arrays lose least (at 30 USD MWh⁻¹:
787 plant –23.2, SALS –18.7, static –5.7 kUSD). The break-even price of SALS
788 exceeds that of the plant by about 25 USD MWh⁻¹, which is the order of the
789 premium or the value of food production that agrivoltaics needs for parity
790 with the open field; parity with the plant needs more. The sensitivity is
791 large (Table 16): at 90 USD MWh⁻¹ the net value of SALS changes from
792 +17 to –8 kUSD between module costs of 0.7 and 1.5 USD Wp⁻¹, whereas
793 crop prices matter little.

794 Because the density was designed for the floor and not for profit, we also
795 let each site choose, among the densities that meet the floor in its history,
796 the one that maximises its historical net value (sites where none pays stay
797 open field; Fig. 7). Then no method pays at 30–40 USD MWh⁻¹; at 60 14–
798 59% of the pairs install an array (none for seasonal sharing) and the net
799 values are small (static arrays 0.9, SALS 0.3 kUSD); static arrays are the
800 best option up to about 75 USD MWh⁻¹, SALS leads from about 80 (4.9
801 against 4.3 at 80, 8.3 against 6.3 at 90, 20.7 against 12.3 at 120), and the
802 tuned rule overtakes static arrays only above about 105. The value of control
803 is therefore real but depends on a high price: at 90 USD MWh⁻¹ SALS is
804 worth 1.7 kUSD ha⁻¹ yr⁻¹ over the tuned rule with profit-optimal densities.

805 **Cost of the safety margin.** At 90 USD MWh⁻¹, moving from a density
806 designed under the mean response to a floor of 0.92 under that response costs
807 1.00 kUSD ha⁻¹ yr⁻¹ (17%) under E2 for 18 more points of compliance (70%




Figure 7: Net value per hectare relative to the open field against electricity price. Left: densities designed for the floor, with a conventional plant replacing the crop (dashed). Right: each site chooses the feasible density that maximises its historical net value. Illustrative parameters of Table 14.

to 88%) and 1.13 kUSD (20%) under E3 for 14 points (71% to 85%), about 56 and 79 USD ha⁻¹ yr⁻¹ per point. The ensemble rule that requires all other members to meet the floor costs 0.85 and 1.22 kUSD for 19 and 17 points (45 and 74 USD per point), so on this measure it is cheaper than the margin under E2 and similar under E3. At 60 USD MWh⁻¹ a margin has a negative cost, because every module beyond break-even loses money; the economics of the design margin therefore depend on the price level, whereas its effect on compliance does not.

3.10. Agrivoltaics under warming

Under warming, the open-field production of the held-out sites (Fig. 8) falls at +3 °C by 23% for maize and 18% for rice, is nearly unchanged for soybean, and rises for wheat (+3%) and potato (+10%). Agrivoltaic production with light-sharing control follows the open field at about 90% of its level, so it inherits the same climate impacts (maize -31%, rice -27% relative to the baseline open field at +3 °C). Agrivoltaics does not protect staple production against warming in this model. The LER of light-sharing systems increases slightly with warming (SALS 1.270 to 1.297; phenology rule 1.236 to 1.265), because the floor allows slightly denser arrays in warmer conditions (GCR 0.173 to 0.177 and electricity 510 to 540 MWh ha⁻¹ yr⁻¹). The margin of SALS over the rules does not grow with warming (Table A.17).




Figure 8: Production relative to the baseline open field (sum over held-out site-seasons) as a function of warming, for the open field, the stress rule and SALS.

3.11. Global potential and what it would take

Every site is evaluated when it is held out (219 of the 400 crop–site pairs, in at least one split; 222 are held out in at least one split and three have no valid season), and each crop’s mean electricity per hectare is scaled to its physical area in SPAM 2010 (581 Mha for the five crops, which avoids counting multiple-cropped land twice). Equipping 5% of this area (29.1 Mha) with SALS arrays at the floor would produce 14.4 PWh yr^{-1} under the baseline climate and 16.0 PWh yr^{-1} at $+3^\circ\text{C}$, with a production-weighted retention of 0.903 and 0.900. The phenology rule would give 13.0 and 14.5 PWh yr^{-1} , and static arrays 6.4 and 6.9 PWh yr^{-1} . For comparison, world electricity generation was about 30 PWh in 2023 (Ember, 2024).

These numbers require large amounts of hardware. SALS at 5% coverage needs 10.6 TWp of modules (11.3 TWp at $+3^\circ\text{C}$), at a capacity factor of 15.5% (16.2%). World solar PV capacity at the end of 2024 was about 1.9 TW (IRENA, 2025), so the requirement is about 5.7 times the installed PV fleet. The electricity is also variable, peaks at midday and would be curtailed or stored. The figures are technical potentials, not projections. Ordering sites by electricity per hectare gives a supply curve (Fig. 9): the best 2% of the staple cropland would supply 16 PWh yr^{-1} and the best 5% 40 PWh yr^{-1} with SALS. The best tenth of the sites are mostly rainfed maize and soybean sites in the Americas and in southern and eastern Africa (7% mean irrigated fraction against 29% overall), which are also among the regions where the floor is hardest to meet (Section 3.5). Siting by electricity alone would therefore concentrate agrivoltaics where food-security risk is highest.

Maize and soybean, for which the simulated shade response is probably too mild (Section 3.7), account for 53% of the electricity of the upscaling




Figure 9: Supply curve under the baseline climate: cumulative electricity when sites are equipped in order of decreasing electricity per hectare, against the share of staple cropland (physical area) equipped. Each controller is evaluated at the density that keeps its mean retention at 0.9.

854 (maize 30%, soybean 23%, wheat 29%, rice 15%, potato 3%). Counting only
855 wheat, rice and potato would give about 7 PWh yr^{-1} at the same coverage,
856 so the global figure should be read as an upper estimate for about half of its
857 content. By crop, the LER of SALS under the baseline climate is 1.20 for
858 wheat and rice, 1.31 for potato, 1.30 for maize and 1.41 for soybean.

859 4. Discussion

860 4.1. What the results establish

861 Four findings stand on consistent evidence. First, sharing light with the
862 crop is the main lever on the electricity that tracker-based agrivoltaic arrays
863 can produce under a yield floor. With static trackers the floor limits arrays
864 to about 6% ground coverage and an LER of 1.07. Any light-sharing control
865 allows 2.0 to 2.9 times the density, and the LER rises to 1.18–1.24 for hand-
866 written rules. Second, the best simple rule follows crop phenology. It shares
867 light across the middle of the season, around flowering and grain filling, and
868 produces electricity at the beginning and the end. It needs only the crop's
869 thermal time, which is observable on the farm, and it is insensitive to forecast
870 errors, whereas the rule that reacts to heat and drought stress is worse in
871 every climate and degrades with forecast error. Third, the qualitative result
872 does not depend on the weakest part of the model. The simulated shade




Figure 10: SALS electricity (a) and land-equivalent ratio (b) at the pooled held-out sites (baseline climate, mean of seasons 2015–2019).

873 response is conservative against the field trials we could use (Section 3.6);
 874 calibrating it raises the allowed density, and the phenology rule stays the
 875 best method, with a gain over static arrays of 0.17, 0.19 and 0.30 LER
 876 under the conservative, calibrated and field-fitted responses. However, the
 877 response learned from the meta-analysis shows steeper losses for maize and
 878 grain legumes than the simulator reproduces (Section 3.7), so the absolute
 879 gains are conservative for wheat, rice and potato and probably overestimates
 880 for maize and soybean (53% of the upscaled electricity). Fourth, a controller
 881 learned through the differentiable simulator is better than the best rule we
 882 could write, by 0.04 in LER against the written rule (0.018 against a rule
 883 whose window was tuned on the training seasons, 5.7% more electricity at
 884 equal retention), reproducibly over training seeds (range 0.002 in LER), three
 885 spatial splits, three canopy-cooling assumptions and four climates, and at 84–
 886 86% of the held-out sites. This gain is small compared with the gap between
 887 static and dynamic control.

888 4.2. What role does artificial intelligence play?

889 We tested three ways in which learning could matter, and the answers dif-
 890 fer. The first is learning the crop’s response to shade from data. The learned
 891 model is useful mainly as a measurement of how little can be learned: a hier-
 892 archical model fitted to 428 observations from 58 studies predicts a held-out
 893 study of the four staple crop types with an error of 0.22 in relative yield, no
 894 better than proportional loss, and the gain over proportional loss (0.43 to

0.31 over all types) comes from crop types that the simulator does not represent. Crop type is the information that transfers, whereas irradiation, climate zone and shade type add nothing measurable and a non-parametric model (gradient-boosted trees) is worse with this number of studies (Section 3.7). Two things follow. The shade response of the staple crops is a distribution with a spread of about 20 percentage points between studies and not a curve, and the process-based simulator does not reproduce the steep losses of maize and legumes seen in the data. Learning the response was limited by the data and not by the model: two studies support the tuber curve, no rice study exists, and 46 panel observations are all leafy vegetables.

The second is calibration through the differentiable simulator. Gradient descent through whole seasons fitted the response of each crop to the learned curves and to the field trials, and it produced an ensemble of plausible responses that reconciles them (Section 3.8). The ensemble informs the design decision: a density chosen under the mean response leaves about 70% of crop–site pairs at the floor if another plausible response holds, against 85–87% for an oracle, and a density chosen to meet the floor in 75% of plausible responses recovers about half of that gap at a cost of 0.04 LER. A simple margin, however, does as well: raising the floor from 0.90 to 0.92 under the mean response reaches 85–88% of the pairs, and at equal compliance it keeps 0.01 (E2) to 0.03 (E3) more LER than the ensemble rule that requires all other members to meet the floor, while the cost per point of compliance is similar. None of the rules is safe outside the ensemble (7–33% of the pairs at the floor under the proportional, calibrated and harsh responses for the mean design, the margin of 0.92 and the ensemble rules), and only a floor of 0.96 reaches 51–78% at a cost of about 0.2 LER. What the ensemble adds is the size of the margin, not a better rule. The quantile rule is a scenario-type design (Calafiore and Campi, 2006; Campi and Garatti, 2008); with five members it has none of the finite-sample guarantees of that theory, and the comparison with the margin shows that it is not needed in practice. The bottleneck of tracker-based agrivoltaics under a food-security floor is the uncertainty of the crop response that enters the design, and it is data on that response, not larger networks, that would improve the result; the learned curves show a spread of about 20 percentage points between studies. The guarantee holds only within the range of responses that the ensemble spans.

The third is control, and here learning contributes little. The learned controller adds 0.018 LER over a tuned phenology rule (and 0.044 over the written one), and 0.015–0.038 over the written rule under the four alternative

933 responses on one split (21 pairs); a reader who needs a simple, auditable
 934 controller can use that rule and obtain most of the electricity gain over static
 935 arrays. Neither of the two extensions that we expected to matter under
 936 uncertainty did: training on ensembles of responses (E1, E2 and E3) did not
 937 help and, for the wider ensembles, did not retain the margin over the written
 938 rule within our training budget (one seed, 21 pairs), and neither did giving
 939 the controller a farm-measurable signal, the biomass of the crop relative to
 940 a reference plot, nor a feedback rule built on that signal. The learned policy
 941 also has to be initialised from the phenology rule (Section 4.5). We therefore
 942 do not claim that learned control is what makes agrivoltaics work: in this
 943 model, the schedule that a learned policy converges to is a smooth function
 944 of crop phenology, which a rule can already express.

945 *4.3. What the results do not establish*

946 The early hypothesis that stress-aware control, that is, shading the crop
 947 on heat and drought days, would be a main source of benefit under warming
 948 was not supported. The stress rule performs worse than the phenology rule
 949 in every climate, and a learned controller without stress inputs performs as
 950 well as one with them. A plausible reason is that in our crop model the yield
 951 gain from shading on a hot or dry day is smaller than the loss of light on
 952 that day, so stress information does not help. Whether this holds for crop
 953 models with more detailed canopy-temperature and heat-stress responses is
 954 an open question. Agrivoltaics does not protect staple production against
 955 warming in our simulations: light-sharing systems produce about 90% of the
 956 open-field yield in every climate and follow the same decline in maize and
 957 rice (Fig. 8).

958 *4.4. Food security in practice*

959 Holding the historical mean retention at 90% does not hold it in each
 960 season or at each site. 70% of SALS seasons meet the floor, the lower tail is
 961 slightly heavier than for the rules, and production-weighted retention is below
 962 0.9 for soybean (0.888) and maize (0.898) and in several regions. Chance-
 963 constrained design raises season-level compliance to about 92% at a cost of
 964 0.065 in LER, and the ranking of the controllers is preserved. Whichever for-
 965 mulation is chosen, a per-site regulatory floor would need a region- and crop-
 966 specific design step. The upscaled potential (Section 3.11) shows the trade-off
 967 in magnitude. The 14 PWh yr⁻¹ that 5% of the staple cropland could supply
 968 need about 10 TWp of modules, several times the present global fleet, and

969 the best sites are rainfed maize and soybean areas where production-weighted
970 retention is lowest.

971 4.5. Limitations

- 972 • *Simulation only.* All results come from a simulator. PV output was
973 benchmarked against PVGIS (29 sites, mean bias -4.9%), but the crop
974 response to shade was learned from one meta-analysis data set and
975 compared with four field trials (Sections 3.6 and 3.7). The simulated
976 response is conservative with respect to the trials (rice, wheat, potato,
977 maize) but too mild for maize and soybean relative to the meta-analysis,
978 where the 40%-shading relative yields are 0.42 and 0.40 against 0.75
979 and 0.48 in the simulator. The harvest-index sensitivities h are assump-
980 tions, and the sensitivity analysis uses rule controllers on a subsample
981 of 22 sites.
- 982 • *Canopy cooling.* The effect of shade on canopy temperature was repre-
983 sented by one coefficient κ , applied to a heat-stress function that was
984 calibrated on air temperature (Siebert et al., 2014; Webber et al., 2018).
985 The ranking of the controllers is robust to κ between 0 and 2°C , but
986 the absolute electricity changes by about $\pm 8\%$. Nocturnal warming
987 under modules and rain redistribution are not modelled.
- 988 • *Crop model.* SIMPLE has a fixed harvest index and linear radiation-
989 use efficiency, and its CO_2 response is that of controlled experiments
990 that tend to overstate field responses (Long et al., 2006). More detailed
991 models with grain-number formation and canopy energy balance may
992 favour different schedules.
- 993 • *Layouts and models.* The baseline of static arrays is a backtracking
994 tracker held at the energy-optimal angle; fixed-tilt, stow-angle and ver-
995 tical bifacial layouts (Willockx et al., 2023; Campana et al., 2021) are
996 not modelled, and the ground coverage ratio is a coarse proxy of crop
997 productivity (Dupraz, 2023). The PV model is isotropic in the dif-
998 fuse component and has no bifacial gain; it was benchmarked against
999 PVGIS at low density only and has not been compared with a refer-
1000 ence implementation such as pvlib (F. Holmgren et al., 2018). Rotation
1001 limits of the trackers were not varied.

- 1002 • *Climate scenarios.* Warming was applied as uniform temperature shifts
1003 with fixed precipitation, radiation and relative humidity, and with un-
1004 changed cultivars and sowing dates. It is not a bias-adjusted climate
1005 projection.
- 1006 • *Initialisation and state information.* The network is initialised from the
1007 phenology rule, because training from random weights collapsed to full
1008 sharing. The gain of SALS is therefore the gain of fine-tuning that rule,
1009 and we cannot say whether a better optimiser or a richer policy class
1010 would find larger gains. The controller uses the simulated state of the
1011 crop and soil (thermal time, canopy interception, soil water) without
1012 measurement error, and the forecast errors we tested apply to weather
1013 only.
- 1014 • *Economics and deployment.* The economic analysis (Section 3.9) uses
1015 assumed prices and costs, no grid, curtailment, storage or subsidy
1016 model, and does not value food security; its break-even prices move
1017 with the capex and the discount rate (Fig. 7). The value of midday
1018 electricity, curtailment, storage, structural cost and farmer acceptance
1019 (Trommsdorff et al., 2021) will determine the realisable share. The av-
1020 erage row pitch at the chosen densities (about 13 m for 2.2-m modules)
1021 leaves room for machinery, but the densest designs (GCR up to 0.5)
1022 have a pitch of only 4.4 m.

1023 5. Conclusions

1024 Light sharing, the daily choice between maximum electricity and letting
1025 light through to the crop, is the main lever of tracker-based agrivoltaics under
1026 a food-security floor. In a differentiable simulator of trackers, PV output,
1027 crop light, microclimate and crop growth, evaluated on held-out crop-site
1028 pairs (92 on the main split, 76 distinct pairs pooled over three splits for the
1029 controller and design comparisons) with the array density of each site chosen
1030 from its own history, light-sharing rules allowed 2.0 to 2.9 times the module
1031 density of static arrays at a mean yield retention of 90%. They raised the
1032 land-equivalent ratio from 1.07 to 1.18–1.24, and a rule that follows crop
1033 phenology and needs no weather forecast was the best of them. The simu-
1034 lated shade response is conservative against the field trials we could use for
1035 rice, wheat and potato, and calibrating it raised the gains without chang-
1036 ing the ranking of methods; for maize and soybean, which supply 53% of

the upscaled electricity, the meta-analysis indicates larger losses than the simulator reproduces whereas one maize trial indicates smaller ones, so the gains for these crops are uncertain and possibly overestimated. The crop’s response to shade, the least certain input, was learned from 428 field observations: a hierarchical model predicted a held-out study of the staple crop types no better than proportional loss (error 0.22), which shows the size of the between-study spread, and gradient descent through the simulator turned it into an ensemble of plausible responses. A learned controller exceeded the written phenology rule by 0.044 LER and a rule with a window tuned on the training seasons by 0.018 (95% interval 0.012–0.024), by following a smooth phenology-dependent schedule that a rule can express; training it on ensembles of shade responses or giving it a farm-measurable biomass signal did not help, and with the wider ensembles it did not retain its margin over the written rule in our training budget (one seed, 21 pairs). The decision that mattered was the design. A density chosen under the mean response left about 70% of crop–site pairs at the floor when another plausible response held (oracle 85–87%); choosing the density so that the floor holds in 75% of the plausible responses recovered about half of the gap at a cost of 0.04 LER, but a floor raised by 0.02 under the mean response was as good as or slightly better at equal compliance, and no rule was safe outside the ensemble. The ensemble is therefore useful to size a safety margin, not to replace it. Under illustrative costs, light sharing breaks even with the open field at 73 (learned controller) to 92 USD MWh^{−1}; a conventional plant replacing the crop stays more profitable at every price, so agrivoltaics depends on a premium, land-use rules or a value for food production that markets do not give. The yield floor is met on average but not in every season or region, and the electricity potential of 5% of staple cropland (about 14 PWh yr^{−1}) requires about 10 TWp of modules. Future work should validate the crop response to shade under trackers with more field data, notably for maize and legumes where the meta-analysis and the trials disagree, extend the crop model to grain-number formation and canopy energy balance, add costs and grid constraints, and pursue controllers that guarantee the food-security floor per site and per season.

CRedit authorship contribution statement

Quang-Vinh Dang: Conceptualization, Methodology, Software, Formal analysis, Investigation, Data curation, Visualization, Writing – original draft,

1073 Writing – review & editing.

1074 Declaration of competing interest

1075 The author has no known competing financial interests or personal re-
1076 lationships that could have appeared to influence the work reported in this
1077 paper.

1078 Data and code availability

1079 All data used are public: SPAM 2010 v2r0 harvested areas (Yu et al.,
1080 2020), the GGCM Phase 3 crop calendar (Jägermeyr et al., 2021), NASA
1081 POWER daily weather (Sparks, 2018) and the SIMPLE crop parameters
1082 (Zhao et al., 2019). The code of the simulator, of SALS and of all experi-
1083 ments, together with the site list, is available from the author’s public reposi-
1084 tory (<https://github.com/vinhqdang/Artificial-Intelligence-in-Agriculture>).

1085 Funding

1086 This research did not receive any specific grant from funding agencies in
1087 the public, commercial, or not-for-profit sectors.

1088 Appendix A. Statistical comparison

1089 References

- 1090 Adeh, E.H., Good, S.P., Calaf, M., Higgins, C.W., 2019. Solar pv power
1091 potential is greatest over croplands. *Scientific Reports* 9, 11442. doi:10.
1092 1038/s41598-019-47803-3.
- 1093 Allen, R.G., Pereira, L.S., Raes, D., Smith, M., 1998. Crop evapotranspira-
1094 tion: Guidelines for computing crop water requirements. *FAO Irrigation*
1095 *and Drainage Paper* 56. Food and Agriculture Organization of the United
1096 Nations. Rome.
- 1097 Amaducci, S., Yin, X., Colauzzi, M., 2018. Agrivoltaic systems to optimise
1098 land use for electric energy production. *Applied Energy* 220, 545–561.
1099 doi:10.1016/j.apenergy.2018.03.081.

1100 Andrade, F.H., Vega, C., Uhart, S., Cirilo, A., Cantarero, M., Valentinuz, O.,
1101 1999. Kernel number determination in maize. *Crop Science* 39, 453–459.
1102 doi:10.2135/cropsci1999.0011183X0039000200026x.

1103 Barron-Gafford, G.A., Pavao-Zuckerman, M.A., Minor, R.L., Sutter, L.F.,
1104 Barnett-Moreno, I., Blackett, D.T., Thompson, M., Dimond, K., Gerlak,
1105 A.K., Nabhan, G.P., Macknick, J.E., 2019. Agrivoltaics provide mutual
1106 benefits across the food–energy–water nexus in drylands. *Nature Sustain-*
1107 *ability* 2, 848–855. doi:10.1038/s41893-019-0364-5.

1108 Calafiore, G.C., Campi, M.C., 2006. The scenario approach to robust control
1109 design. *IEEE Transactions on Automatic Control* 51, 742–753. doi:10.
1110 1109/TAC.2006.875041.

1111 Campana, P.E., Stridh, B., Amaducci, S., Colauzzi, M., 2021. Optimisation
1112 of vertically mounted agrivoltaic systems. *Journal of Cleaner Production*
1113 325, 129091. doi:10.1016/j.jclepro.2021.129091.

1114 Campi, M.C., Garatti, S., 2008. The exact feasibility of randomized solutions
1115 of uncertain convex programs. *SIAM Journal on Optimization* 19, 1211–
1116 1230. doi:10.1137/07069821X.

1117 Collares-Pereira, M., Rabl, A., 1979. The average distribution of solar
1118 radiation-correlations between diffuse and hemispherical and between daily
1119 and hourly insolation values. *Solar Energy* 22, 155–164. doi:10.1016/
1120 0038-092x(79)90100-2.

1121 DIN, 2021. DIN SPEC 91434:2021-05: Agri-photovoltaic systems – require-
1122 ments for primary agricultural use. Beuth Verlag, Berlin.

1123 Dinesh, H., Pearce, J.M., 2016. The potential of agrivoltaic systems. *Renew-*
1124 *able and Sustainable Energy Reviews* 54, 299–308. doi:10.1016/j.rser.
1125 2015.10.024.

1126 Dupraz, C., 2023. Assessment of the ground coverage ratio of agrivoltaic
1127 systems as a proxy for potential crop productivity. *Agroforestry Systems*
1128 98, 2679–2696. doi:10.1007/s10457-023-00906-3.

1129 Dupraz, C., Marrou, H., Talbot, G., Dufour, L., Nogier, A., Ferard, Y., 2011.
1130 Combining solar photovoltaic panels and food crops for optimising land

1131 use: Towards new agrivoltaic schemes. *Renewable Energy* 36, 2725–2732.
1132 doi:10.1016/j.renene.2011.03.005.

1133 Elamri, Y., Cheviron, B., Lopez, J.M., Dejean, C., Belaud, G., 2018. Water
1134 budget and crop modelling for agrivoltaic systems: Application to irrigated
1135 lettuces. *Agricultural Water Management* 208, 440–453. doi:10.1016/j.
1136 agwat.2018.07.001.

1137 Ember, 2024. *Global Electricity Review 2024*. Technical Report. Ember.
1138 London.

1139 European Commission, Joint Research Centre, 2022. Photovoltaic geo-
1140 graphical information system (PVGIS), version 5.2. [https://re.jrc.ec.
1141 europa.eu/pvg_tools/](https://re.jrc.ec.europa.eu/pvg_tools/).

1142 F. Holmgren, W., W. Hansen, C., A. Mikofski, M., 2018. pvlib python: a
1143 python package for modeling solar energy systems. *Journal of Open Source
1144 Software* 3, 884. doi:10.21105/joss.00884.

1145 Fischer, R.A., 1985. Number of kernels in wheat crops and the influence of
1146 solar radiation and temperature. *The Journal of Agricultural Science* 105,
1147 447–461. doi:10.1017/S0021859600056495.

1148 Gautron, R., Maillard, O.A., Preux, P., Corbeels, M., Sabbadin, R., 2022.
1149 Reinforcement learning for crop management support: Review, prospects
1150 and challenges. *Computers and Electronics in Agriculture* 200, 107182.
1151 doi:10.1016/j.compag.2022.107182.

1152 Gonocruz, R.A., Nakamura, R., Yoshino, K., Homma, M., Doi, T., Yoshida,
1153 Y., Tani, A., 2021. Analysis of the rice yield under an agrivoltaic
1154 system: A case study in japan. *Environments* 8, 65. doi:10.3390/
1155 environments8070065.

1156 Hassanpour Adeg, E., Selker, J.S., Higgins, C.W., 2018. Remarkable agri-
1157 voltaic influence on soil moisture, micrometeorology and water-use effi-
1158 ciency. *PLOS ONE* 13, e0203256. doi:10.1371/journal.pone.0203256.

1159 IPCC, 2021. *Climate Change 2021: The Physical Science Basis*. Contribution
1160 of Working Group I to the Sixth Assessment Report of the Intergovern-
1161 mental Panel on Climate Change. Cambridge University Press, Cambridge.
1162 doi:10.1017/9781009157896.

1163 IRENA, 2025. Renewable Capacity Statistics 2025. Technical Report. Inter-
1164 national Renewable Energy Agency. Abu Dhabi.

1165 Jägermeyr, J., Müller, C., Ruane, A.C., Elliott, J., Balkovic, J., Castillo, O.,
1166 Faye, B., Foster, I., Folberth, C., Franke, J.A., Fuchs, K., Guarin, J.R.,
1167 Heinke, J., Hoogenboom, G., Iizumi, T., Jain, A.K., Kelly, D., Khabarov,
1168 N., Lange, S., Lin, T.S., Liu, W., Mialyk, O., Minoli, S., Moyer, E.J.,
1169 Okada, M., Phillips, M., Porter, C., Rabin, S.S., Scheer, C., Schneider,
1170 J.M., Schyns, J.F., Skalsky, R., Smerald, A., Stella, T., Stephens, H.,
1171 Webber, H., Zabel, F., Rosenzweig, C., 2021. Climate impacts on global
1172 agriculture emerge earlier in new generation of climate and crop models.
1173 *Nature Food* 2, 873–885. doi:10.1038/s43016-021-00400-y.

1174 Karniadakis, G.E., Kevrekidis, I.G., Lu, L., Perdikaris, P., Wang, S., Yang,
1175 L., 2021. Physics-informed machine learning. *Nature Reviews Physics* 3,
1176 422–440. doi:10.1038/s42254-021-00314-5.

1177 Kingma, D.P., Ba, J., 2015. Adam: A method for stochastic optimization,
1178 in: *International Conference on Learning Representations*.

1179 Laub, M., Pataczek, L., Feuerbacher, A., Zikeli, S., Högy, P., 2022. Contrast-
1180 ing yield responses at varying levels of shade suggest different suitability
1181 of crops for dual land-use systems: a meta-analysis. *Agronomy for Sus-
1182 tainable Development* 42, 51. doi:10.1007/s13593-022-00783-7.

1183 Liu, B.Y., Jordan, R.C., 1960. The interrelationship and characteristic dis-
1184 tribution of direct, diffuse and total solar radiation. *Solar Energy* 4, 1–19.
1185 doi:10.1016/0038-092x(60)90062-1.

1186 Long, S.P., Ainsworth, E.A., Leakey, A.D.B., Nösberger, J., Ort, D.R., 2006.
1187 Food for thought: Lower-than-expected crop yield stimulation with ris-
1188 ing co₂ concentrations. *Science* 312, 1918–1921. doi:10.1126/science.
1189 1114722.

1190 Lorenzo, E., Narvarte, L., Muñoz, J., 2011. Tracking and back-tracking.
1191 *Progress in Photovoltaics: Research and Applications* 19, 747–753. doi:10.
1192 1002/pip.1085.

1193 Mamun, M.A.A., Dargusch, P., Wadley, D., Zulkarnain, N.A., Aziz, A.A.,
1194 2022. A review of research on agrivoltaic systems. *Renewable and Sustain-
1195 able Energy Reviews* 161, 112351. doi:10.1016/j.rser.2022.112351.

- 1196 Marchau, V.A.W.J., Walker, W.E., Bloemen, P.J.T.M., Popper, S.W. (Eds.),
1197 2019. Decision Making under Deep Uncertainty: From Theory to Practice.
1198 Springer. doi:10.1007/978-3-030-05252-2.
- 1199 Marrou, H., Dufour, L., Wery, J., 2013a. How does a shelter of solar panels
1200 influence water flows in a soil-crop system? *European Journal of Agron-*
1201 *omy* 50, 38–51. doi:10.1016/j.eja.2013.05.004.
- 1202 Marrou, H., Guilioni, L., Dufour, L., Dupraz, C., Wery, J., 2013b. Microcli-
1203 mate under agrivoltaic systems: Is crop growth rate affected in the partial
1204 shade of solar panels? *Agricultural and Forest Meteorology* 177, 117–132.
1205 doi:10.1016/j.agrformet.2013.04.012.
- 1206 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen,
1207 T., Lin, Z., Gimelshein, N., Antiga, L., et al., 2019. PyTorch: An impera-
1208 tive style, high-performance deep learning library, in: *Advances in Neural*
1209 *Information Processing Systems*, pp. 8024–8035.
- 1210 Ramos-Fuentes, I.A., Elamri, Y., Cheviron, B., Dejean, C., Belaud, G.,
1211 Fumey, D., 2023. Effects of shade and deficit irrigation on maize growth
1212 and development in fixed and dynamic agrivoltaic systems. *Agricultural*
1213 *Water Management* 280, 108187. doi:10.1016/j.agwat.2023.108187.
- 1214 Reichstein, M., Camps-Valls, G., Stevens, B., Jung, M., Denzler, J., Car-
1215 valhais, N., Prabhat, 2019. Deep learning and process understanding for
1216 data-driven Earth system science. *Nature* 566, 195–204. doi:10.1038/
1217 s41586-019-0912-1.
- 1218 Sekiyama, T., Nagashima, A., 2019. Solar sharing for both food and
1219 clean energy production: Performance of agrivoltaic systems for corn,
1220 a typical shade-intolerant crop. *Environments* 6, 65. doi:10.3390/
1221 environments6060065.
- 1222 Shen, C., Appling, A.P., Gentine, P., Bandai, T., Gupta, H., Tartakovsky, A.,
1223 Baity-Jesi, M., Fenicia, F., Kifer, D., Li, L., Liu, X., Ren, W., Zheng, Y.,
1224 Harman, C.J., Clark, M., Farthing, M., Feng, D., Kumar, P., Aboelyazeed,
1225 D., Rahmani, F., Song, Y., Beck, H.E., Bindas, T., Dwivedi, D., Fang,
1226 K., Höge, M., Rackauckas, C., Mohanty, B., Roy, T., Xu, C., Lawson,
1227 K., 2023. Differentiable modelling to unify machine learning and physical

1228 models for geosciences. *Nature Reviews Earth & Environment* 4, 552–567.
1229 doi:10.1038/s43017-023-00450-9.

1230 Siebert, S., Ewert, F., Eyshi Rezaei, E., Kage, H., Graß, R., 2014. Impact
1231 of heat stress on crop yield—on the importance of considering canopy
1232 temperature. *Environmental Research Letters* 9, 044012. doi:10.1088/
1233 1748-9326/9/4/044012.

1234 Sparks, A., 2018. nasapower: A nasa power global meteorology, surface solar
1235 energy and climatology data client for r. *Journal of Open Source Software*
1236 3, 1035. doi:10.21105/joss.01035.

1237 Trommsdorff, M., Kang, J., Reise, C., Schindele, S., Bopp, G., Ehmann, A.,
1238 Weselek, A., Högy, P., Obergfell, T., 2021. Combining food and energy
1239 production: Design of an agrivoltaic system applied in arable and vegetable
1240 farming in germany. *Renewable and Sustainable Energy Reviews* 140,
1241 110694. doi:10.1016/j.rser.2020.110694.

1242 Valle, B., Simonneau, T., Sourd, F., Pechier, P., Hamard, P., Frisson, T.,
1243 Ryckewaert, M., Christophe, A., 2017. Increasing the total productivity of
1244 a land by combining mobile photovoltaic panels and food crops. *Applied*
1245 *Energy* 206, 1495–1507. doi:10.1016/j.apenergy.2017.09.113.

1246 Waha, K., Dietrich, J.P., Portmann, F.T., Siebert, S., Thornton, P.K., Bon-
1247 deau, A., Herrero, M., 2020. Multiple cropping systems of the world
1248 and the potential for increasing cropping intensity. *Global Environmental*
1249 *Change* 64, 102131. doi:10.1016/j.gloenvcha.2020.102131.

1250 Webber, H., White, J.W., Kimball, B.A., Ewert, F., Asseng, S., Eyshi Rezaei,
1251 E., Pinter, P.J., Hatfield, J.L., Reynolds, M.P., Ababaei, B., Bindi, M.,
1252 Doltra, J., Ferrise, R., Kage, H., Kassie, B.T., Kersebaum, K.C., Luig,
1253 A., Olesen, J.E., Semenov, M.A., Stratonovitch, P., Ratjen, A.M., LaM-
1254 orte, R.L., Leavitt, S.W., Hunsaker, D.J., Wall, G.W., Martre, P., 2018.
1255 Physical robustness of canopy temperature models for crop heat stress
1256 simulation across environments and production conditions. *Field Crops*
1257 *Research* 216, 75–88. doi:10.1016/j.fcr.2017.11.005.

1258 Weselek, A., Bauerle, A., Hartung, J., Zikeli, S., Lewandowski, I., Högy, P.,
1259 2021. Agrivoltaic system impacts on microclimate and yield of different

1260 crops within an organic crop rotation in a temperate climate. *Agronomy*
1261 for Sustainable Development 41, 59. doi:10.1007/s13593-021-00714-y.

1262 Weselek, A., Ehmann, A., Zikeli, S., Lewandowski, I., Schindele, S., Högy,
1263 P., 2019. Agrophotovoltaic systems: applications, challenges, and op-
1264 portunities. a review. *Agronomy for Sustainable Development* 39, 35.
1265 doi:10.1007/s13593-019-0581-3.

1266 Willockx, B., Lavaert, C., Cappelle, J., 2023. Performance evaluation of
1267 vertical bifacial and single-axis tracked agrivoltaic systems on arable land.
1268 *Renewable Energy* 217, 119181. doi:10.1016/j.renene.2023.119181.

1269 Woli, P., Jones, J.W., Ingram, K.T., Fraisse, C.W., 2012. Agricultural
1270 reference index for drought (arid). *Agronomy Journal* 104, 287–300.
1271 doi:10.2134/agronj2011.0286.

1272 Yu, Q., You, L., Wood-Sichra, U., Ru, Y., Joglekar, A.K.B., Fritz, S., Xiong,
1273 W., Lu, M., Wu, W., Yang, P., 2020. A cultivated planet in 2010 – part 2:
1274 The global gridded agricultural-production maps. *Earth System Science*
1275 *Data* 12, 3545–3572. doi:10.5194/essd-12-3545-2020.

1276 Zhang, L., Geng, X., Ding, H., Cao, K., Wang, L., Chen, H., Deng, L., Wu,
1277 C., Xiao, M., Bao, E., 2026. Multi-objective optimization of photovoltaic
1278 array layouts on farmland via a hybrid algorithm framework for enhanced
1279 agricultural production and energy generation. *Artificial Intelligence in*
1280 *Agriculture* doi:10.1016/j.aiaa.2026.09.002.

1281 Zhao, C., Liu, B., Piao, S., Wang, X., Lobell, D.B., Huang, Y., Huang, M.,
1282 Yao, Y., Bassu, S., Ciais, P., Durand, J.L., Elliott, J., Ewert, F., Janssens,
1283 I.A., Li, T., Lin, E., Liu, Q., Martre, P., Müller, C., Peng, S., Peñuelas,
1284 J., Ruane, A.C., Wallach, D., Wang, T., Wu, D., Liu, Z., Zhu, Y., Zhu,
1285 Z., Asseng, S., 2017. Temperature increase reduces global yields of major
1286 crops in four independent estimates. *Proceedings of the National Academy*
1287 *of Sciences* 114, 9326–9331. doi:10.1073/pnas.1701762114.

1288 Zhao, C., Liu, B., Xiao, L., Hoogenboom, G., Boote, K.J., Kassie, B.T.,
1289 Pavan, W., Shelia, V., Kim, K.S., Hernandez-Ochoa, I.M., Wallach, D.,
1290 Porter, C.H., Stockle, C.O., Zhu, Y., Asseng, S., 2019. A simple crop
1291 model. *European Journal of Agronomy* 104, 97–106. doi:10.1016/j.eja.
1292 2019.01.009.

Table 2: Held-out performance (92 test sites, seasons 2015–2019, canopy cooling $\kappa = 1^\circ\text{C}$). The density (GCR) of each site is the largest value on the grid 0.05, 0.075, ..., 0.5 at which the method’s mean retention in the site’s 2001–2014 seasons is ≥ 0.9 . E : electricity ($\text{MWh ha}^{-1} \text{ yr}^{-1}$); e : electricity relative to a conventional PV plant; r : mean yield retention; r_{prod} : production-weighted retention; C_s and C_{site} : share of seasons with $r \geq 0.9$ and share of sites whose test-season mean $r \geq 0.9$ (%); $\text{LER} = r + e$.

Method	GCR	E	e	r	r_{prod}	C_s	C_{site}	LER
<i>Baseline</i>								
AV static	0.06	217	0.16	0.917	0.914	93	96	1.07
Seasonal sharing	0.17	397	0.28	0.904	0.902	69	72	1.19
Phenology rule	0.17	465	0.33	0.904	0.902	69	75	1.24
Stress rule	0.13	387	0.28	0.907	0.907	76	82	1.18
SALS (ours)	0.17	510	0.36	0.905	0.902	70	74	1.27
<i>+1.5 °C</i>								
AV static	0.06	221	0.16	0.916	0.914	93	92	1.08
Seasonal sharing	0.16	393	0.28	0.905	0.903	72	74	1.19
Phenology rule	0.17	478	0.34	0.905	0.902	68	75	1.25
Stress rule	0.12	377	0.27	0.908	0.908	73	77	1.18
SALS (ours)	0.17	518	0.37	0.906	0.903	68	75	1.28
<i>+2 °C</i>								
AV static	0.06	223	0.16	0.916	0.914	92	93	1.08
Seasonal sharing	0.17	394	0.28	0.905	0.904	72	76	1.19
Phenology rule	0.17	483	0.35	0.905	0.903	69	74	1.25
Stress rule	0.12	370	0.26	0.909	0.909	75	78	1.17
SALS (ours)	0.17	521	0.38	0.907	0.904	69	75	1.28
<i>+3 °C</i>								
AV static	0.06	231	0.17	0.915	0.914	88	89	1.08
Seasonal sharing	0.17	397	0.29	0.906	0.904	69	75	1.19
Phenology rule	0.18	499	0.36	0.905	0.902	67	72	1.26
Stress rule	0.11	366	0.26	0.908	0.908	74	77	1.17
SALS (ours)	0.18	540	0.39	0.907	0.904	68	75	1.30

Table 3: LER of the controllers at different yield-retention floors ρ (main split; the historical design step is repeated for each floor).

Method	Baseline climate				+3 °C			
	0.80	0.85	0.90	0.95	0.80	0.85	0.90	0.95
AV static	1.159	1.115	1.072	1.011	1.174	1.128	1.083	1.013
Stress rule	1.380	1.282	1.182	1.077	1.364	1.269	1.172	1.071
Phenology rule	1.467	1.356	1.236	1.108	1.513	1.395	1.265	1.122
SALS (trained at $\rho = 0.9$)	1.505	1.392	1.270	1.141	1.553	1.432	1.297	1.154
SALS conditioned on ρ	1.504	1.391	1.269	1.141	1.553	1.431	1.296	1.154

Table 4: Rules and the learned controller pooled over the three spatial splits (76 distinct crop-site pairs, every third held-out pair; baseline climate; main shade response; history-based design on a GCR grid with step 0.05, coarser than the 0.025 grid of the main analysis). The tuned phenology rule shares light while thermal time lies between 0.15 and 1.0 of the season requirement (window chosen on training-block seasons of split 0; the written rule uses 0.05 to 1.0). The optimised ramp schedule has five parameters (start 0.17, end 0.93, ramp widths 0.02, level 0.99) found by CMA-ES on the same training seasons.

Method	GCR	E/E_{PV}	Retention	Seasons ≥ 0.9 (%)	Pairs ≥ 0.9 (%)	LER
AV static	0.059	0.153	0.922	95	96	1.076
Seasonal sharing	0.151	0.249	0.915	84	91	1.164
Stress rule	0.118	0.255	0.918	83	88	1.173
Phenology rule	0.154	0.292	0.916	85	91	1.207
Tuned phenology rule	0.161	0.317	0.916	84	89	1.233
Optimised ramp schedule	0.155	0.322	0.916	83	87	1.238
SALS	0.161	0.335	0.916	83	87	1.251

Table 5: Simulated relative yield at the shading of published field trials (static arrays, training-block sites). Observed values: rice, single trial; wheat and potato, the worse of the two seasons (the better season gave +2.7% and +11%); maize, mean of the three seasons of the dynamic system (DAV) and single values for the other two systems. The same trials enter the calibration of Section 2.6, so the comparison is not independent.

Trial	Shading	Observed	Conservative	Calibrated
Rice, Gonocruz et al. (2021)	0.25	0.84	0.77	0.79
Wheat, Weselek et al. (2021)	0.30	0.81	0.66	0.74
Potato, Weselek et al. (2021)	0.30	0.82	0.72	0.72
Maize, Ramos-Fuentes et al. (2023) (DAV)	0.35	0.83	0.70	0.70
Maize, Ramos-Fuentes et al. (2023) (half)	0.30	0.91	0.75	0.75
Maize, Ramos-Fuentes et al. (2023) (full)	0.50	0.70	0.54	0.54

Table 6: Sensitivity of the deployable light-sharing results to the crop shade response (rule controllers, baseline climate, floor $\rho = 0.9$, split 0, every fourth held-out site, 22 sites). Conservative: proportional response (main analysis); calibrated: radiation-use compensation fitted to Laub et al. (2022); field-fitted: fitted to the field trials, maize harvest index unaffected.

Shade response	Method	GCR	E/E_{PV}	Retention	LER
conservative	AV static	0.06	0.15	0.92	1.07
	Phenology rule	0.17	0.33	0.90	1.23
	Seasonal sharing	0.16	0.28	0.90	1.18
	Stress rule	0.11	0.24	0.91	1.15
calibrated	AV static	0.07	0.18	0.91	1.09
	Phenology rule	0.19	0.37	0.90	1.28
	Seasonal sharing	0.19	0.32	0.90	1.22
	Stress rule	0.13	0.29	0.91	1.20
field-fitted	AV static	0.16	0.41	0.91	1.32
	Phenology rule	0.37	0.70	0.92	1.62
	Seasonal sharing	0.37	0.61	0.92	1.53
	Stress rule	0.29	0.63	0.90	1.53

Table 7: Predicting relative yield under shade from 428 observations of 58 studies (Laub et al., 2022): leave-one-study-out cross-validation (every study held out once; observations with $s > 0$). Covariates: irradiation, climate zone, shade type, intercropping.

Model	RMSE	MAE	Median AE	Mean MAE per study
Proportional loss, $y = 1 - s$	0.430	0.305	0.210	0.224
One curve for all crops	0.372	0.265	0.193	0.208
Curve per crop type, no pooling	0.313	0.208	0.141	0.165
Curve per crop type, partial pooling	0.312	0.210	0.150	0.162
Partial pooling + covariates	0.311	0.206	0.148	0.161
Gradient-boosted trees	0.321	0.219	0.162	0.184

Table 8: Leave-one-study-out prediction error (RMSE of relative yield, observations with $s > 0$) by crop type: proportional loss, hierarchical curve with partial pooling, and gradient-boosted trees. The last columns give the difference hierarchical minus proportional with a 95% cluster-bootstrap interval over studies (negative: hierarchical better).

Group	Obs./studies	Proportional	Hierarchical	Boosted t
C3 cereals	30/10	0.214	0.187	0.205
Maize	40/10	0.225	0.235	0.248
Grain legumes	58/14	0.215	0.225	0.234
Tubers/root crops	10/2	0.250	0.254	0.257
Four crop types used by the simulator	138/30	0.220	0.222	0.234
Other types (berries, forages, fruits, vegetables)	201/29	0.528	0.361	0.369
All types	339/58	0.430	0.312	0.321

Table 9: Relative yield at 20% and 40% shading: data-driven model (median and 80% bootstrap interval over studies; crop-type curve, covariates at their reference values: mean irradiation, temperate zone, no solar-panel shade, no intercropping) and simulator under the two calibrated ensembles (mean of members; E1 fitted to hand-read curves, E2 fitted to the data-driven curves). The data-driven curve of rice is that of C3 cereals (no rice in the data set). Observations and studies: C3 cereals 41/10, maize 51/10, grain legumes 76/14, tubers 12/2.

Crop	Shading	Data-driven	Simulator, E1	Simulator, E2	Proportional
Wheat	20%	0.93 [0.88, 0.96]	0.88	0.92	0.80
Wheat	40%	0.72 [0.59, 0.78]	0.70	0.77	0.60
Rice	20%	0.93 [0.88, 0.96]	0.85	0.90	0.80
Rice	40%	0.72 [0.59, 0.78]	0.66	0.73	0.60
Maize	20%	0.73 [0.58, 0.86]	0.86	0.90	0.80
Maize	40%	0.42 [0.27, 0.59]	0.69	0.75	0.60
Soybean	20%	0.71 [0.65, 0.79]	0.77	0.74	0.80
Soybean	40%	0.40 [0.34, 0.53]	0.53	0.48	0.60
Potato	20%	0.93 [0.92, 0.97]	0.88	0.92	0.80
Potato	40%	0.71 [0.67, 0.83]	0.73	0.79	0.60

Table 10: Shade-response parameters of ensemble E2, calibrated by gradient descent through the simulator to the data-driven curves and the field trials (12 randomised fits; mean $\pm$ s.d.): radiation-use compensation c and multiplier h on the flowering-window harvest-index penalty (1 = assumed value).

Crop	c	h multiplier
Wheat	0.91 ± 0.31	0.74 ± 0.53
Rice	0.61 ± 0.25	0.83 ± 0.33
Maize	0.14 ± 0.11	0.16 ± 0.13
Soybean	0.05 ± 0.09	1.66 ± 0.38
Potato	0.62 ± 0.20	0.90 ± 0.50

Table 11: Land-equivalent ratio of rules and learned controllers under four true shade responses, with the array density designed under the true response (split 0, baseline climate, floor $\rho = 0.9$, every sixth held-out pair, 21 pairs, coarse GCR grid, one training seed; season means, so differences differ slightly from the pair-level intervals in the text). SALS is trained on the proportional response; robust controllers are trained on random members of ensemble E1, E2 or E3 (adaptive: with the measured biomass ratio). The proportional and harsh responses lie outside the ensembles.

Method	Proportional	Calibrated	Field-fitted	Harsh
AV static	1.052	1.059	1.254	1.037
Seasonal sharing	1.133	1.156	1.385	1.097
Stress rule	1.129	1.160	1.406	1.086
Feedback rule	1.139	1.164	1.392	1.089
Phenology rule	1.164	1.190	1.456	1.122
Tuned phenology rule	1.170	1.199	1.472	1.129
Optimised ramp schedule	1.170	1.199	1.477	1.122
SALS	1.178	1.209	1.489	1.132
SALS-robust, E1	1.179	1.204	1.491	1.127
SALS-adaptive, E1	1.176	1.202	1.495	1.124
SALS-robust, E2	1.165	1.192	1.483	1.126
SALS-robust, E3	1.163	1.191	1.481	1.121
SALS-adaptive, E3	1.164	1.191	1.475	1.118

Table 12: Array design under shade-response uncertainty, phenology rule, three spatial splits pooled (76 distinct crop-site pairs, every third held-out pair). Truth: each of five members of the ensemble in turn; the mean response and the member-based rules use the other four members (leave-one-out, so the truth never enters the design). E2: ensemble calibrated to the learned curves and field trials; E3: same procedure with target errors of 0.2 on the learned curves. Pairs: share with mean retention ≥ 0.9 ; shortfall: mean of $\max(0, 0.9 - \text{retention})$ per pair.

Design	E2		LER	E3		
	Pairs ≥ 0.9 (%)	Shortfall		Pairs ≥ 0.9 (%)	Shortfall	
oracle (true response)	87	0.0025	1.486	85	0.0029	1
mean response	70	0.0062	1.493	71	0.0093	1
floor met in 50% of other members	65	0.0084	1.510	67	0.0125	1
floor met in 75% of other members	78	0.0045	1.462	79	0.0083	1
floor met in all other members	89	0.0019	1.412	88	0.0032	1
floor 0.92 under mean response	88	0.0022	1.426	85	0.0051	1
floor 0.94 under mean response	97	0.0007	1.345	92	0.0027	1
floor 0.96 under mean response	99	0.0003	1.251	96	0.0015	1

Table 13: Design rules outside the ensemble: share of crop–site pairs (%) with mean retention ≥ 0.9 under four fixed responses (P proportional, C calibrated to the meta-analysis, F field-fitted, H harsh), phenology rule, three splits pooled (76 distinct pairs). Designs use the whole ensemble E2 or E3 (five members and their mean); the oracle uses the true response. LER in the last column is the mean over the four responses.

Design	E2				E3				LER	
	P	C	F	H	P	C	F	H	E2	E3
oracle (true response)	91	89	87	84	91	89	87	84	1.283	1.283
mean response	13	13	68	12	13	17	79	7	1.435	1.448
floor met in 75% of members	16	16	80	12	22	25	93	13	1.384	1.397
floor met in all members	17	17	80	12	29	32	93	16	1.348	1.336
floor 0.92 under mean response	17	20	86	12	29	33	91	20	1.376	1.399
floor 0.94 under mean response	45	47	97	34	45	45	97	29	1.303	1.345
floor 0.96 under mean response	63	78	99	54	58	63	99	51	1.220	1.264

Table 14: Assumed economic parameters (illustrative, not data). Ranges used in the sensitivity analysis are in parentheses.

Parameter	Value
Electricity price	60 USD MWh ⁻¹ (30–120)
Capex, AV single-axis tracker plant	1.0 USD Wp ⁻¹ (0.7–1.5)
Capex, conventional tracker plant	0.75 USD Wp ⁻¹
Discount rate, lifetime, O&M	6% (4–8%), 25 years, 1.5% of capex per year
Installed capacity	2.1 MWp per hectare of module area share g (0.21 kWp m ⁻²)
Crop price per tonne of dry matter	wheat 290, rice 400, maize 250, soybean 520, potato 1000 USD (0.

Table 15: Illustrative annual economics per hectare (76 crop–site pairs, three splits, main shade response, baseline climate; parameters in Table 14). Net value is relative to the open field (crop costs unchanged) in kUSD ha⁻¹ yr⁻¹ at three electricity prices; break-even is the price at which the net value is zero. The conventional plant replaces the crop.

Method	E (MWh)	Capex annuity	Crop loss	Net at 60	Net at 90	Net at 120
AV static	200	11.5	0.27	0.3	6.3	12.3
Seasonal sharing	326	29.6	0.31	-10.3	-0.5	9.2
Stress rule	336	23.2	0.30	-3.3	6.8	16.9
Phenology rule	382	30.1	0.30	-7.5	4.0	15.4
Tuned phenology rule	415	31.6	0.31	-7.0	5.5	17.9
SALS	438	31.6	0.31	-5.6	7.6	20.7
Conventional PV plant	1306	58.7	3.60	16.0	55.2	94.4

Table 16: One-at-a-time sensitivity of the net value (kUSD ha⁻¹ yr⁻¹, relative to the open field) at 90 USD MWh⁻¹, densities designed for the floor as in the main analysis. The capex of the conventional plant is 0.75/0.70 of that of the agrivoltaic plant in the capex rows.

Parameter	Value	Static arrays	Tuned rule	SALS	Conventional plant
Base		6.3	5.5	7.6	55.2
AV capex (USD/Wp)	0.7	9.7	14.9	17.1	72.8
AV capex (USD/Wp)	1.5	0.6	-10.3	-8.2	25.8
Discount rate	4%	8.0	10.3	12.4	64.1
Discount rate	8%	4.4	0.2	2.4	45.4
Crop prices	0.5×	6.4	5.6	7.7	57.0
Crop prices	2.0×	6.0	5.2	7.3	51.6
Electricity price	60	0.3	-7.0	-5.6	16.0
Electricity price	120	12.3	17.9	20.7	94.4

Table A.17: Paired comparison of SALS with each baseline on the 92 held-out sites of the main split (unit of analysis: site mean over test seasons). Δ LER: mean difference (SALS minus baseline) with cluster-bootstrap 95% CI; p : two-sided Wilcoxon signed-rank test; % better: share of sites where SALS has the higher LER; Δr : difference in mean yield retention with its Wilcoxon p -value.

Climate	Baseline	Δ LER [95% CI]	p	% better	Δr	p	Sites
Baseline	AV static	+0.205 [+0.176, +0.236]	$< 10^{-4}$	100	-0.011	$< 10^{-4}$	92
Baseline	Seasonal sharing	+0.088 [+0.068, +0.109]	$< 10^{-4}$	97	+0.001	0.922	92
Baseline	Phenology rule	+0.038 [+0.025, +0.053]	$< 10^{-4}$	84	+0.001	0.723	92
Baseline	Stress rule	+0.096 [+0.071, +0.124]	$< 10^{-4}$	93	-0.002	0.055	92
+1.5 °C	AV static	+0.208 [+0.180, +0.237]	$< 10^{-4}$	100	-0.010	$< 10^{-4}$	92
+1.5 °C	Seasonal sharing	+0.095 [+0.077, +0.114]	$< 10^{-4}$	98	+0.001	0.812	92
+1.5 °C	Phenology rule	+0.033 [+0.022, +0.045]	$< 10^{-4}$	86	+0.002	0.855	92
+1.5 °C	Stress rule	+0.108 [+0.083, +0.134]	$< 10^{-4}$	98	-0.001	0.041	92
+2 °C	AV static	+0.209 [+0.181, +0.238]	$< 10^{-4}$	100	-0.010	$< 10^{-4}$	92
+2 °C	Seasonal sharing	+0.097 [+0.079, +0.115]	$< 10^{-4}$	97	+0.001	0.876	92
+2 °C	Phenology rule	+0.032 [+0.022, +0.044]	$< 10^{-4}$	85	+0.002	0.496	92
+2 °C	Stress rule	+0.115 [+0.089, +0.142]	$< 10^{-4}$	98	-0.003	0.016	92
+3 °C	AV static	+0.219 [+0.190, +0.250]	$< 10^{-4}$	100	-0.009	$< 10^{-4}$	92
+3 °C	Seasonal sharing	+0.108 [+0.089, +0.129]	$< 10^{-4}$	96	+0.001	0.935	92
+3 °C	Phenology rule	+0.034 [+0.022, +0.046]	$< 10^{-4}$	86	+0.002	0.330	92
+3 °C	Stress rule	+0.131 [+0.102, +0.162]	$< 10^{-4}$	98	-0.002	0.142	92