

1 Light-sharing control for tracker-based agrivoltaics
2 under a yield floor: a simulation study of what
3 learning adds and what shade-response uncertainty
4 costs

5 Quang-Vinh Dang*

6 *British University Vietnam, Hung Yen, Vietnam*

7 **Abstract**

Single-axis trackers can shift their daily rotation from maximum electricity towards letting light through (light sharing). We quantify its value under a 90% yield-retention floor with a differentiable simulator. On 92 held-out crop–site pairs and the proportional-growth shade response, light-sharing rules allowed arrays 2.1–2.9 times as dense as static arrays at equal mean retention and raised the land-equivalent ratio (LER) from 1.07 to 1.18–1.24, with the floor met in 69–76% of seasons (static arrays: 93%). The crop response to shade is the least certain input: a hierarchical model fitted to 428 observations from 58 studies predicts a new study of the four staple crop types (138 observations, 30 studies) only to about 0.22 in relative yield, no better than proportional loss, so we built ensembles of plausible responses from the learned curves and field trials. Within these ensembles, a density chosen under the mean response left about 70% of pairs at the floor when another plausible response held, against 85–87% for an oracle design with the true response; a floor raised by 0.02 and an ensemble-quantile rule were not distinguishable at equal compliance, and neither was safe outside the ensemble, whose spread is narrower than the between-study spread. Pooled over three splits (76 pairs), a neural controller initialised from the written phenology rule beat it by 0.044 LER but a CMA-ES-optimised window, re-tuned on each split, by only 0.007 (95% interval 0.002–0.014; 0.003 at matched retention, interval including zero); ensemble or biomass-signal training did not help within our budget. Under illustrative costs, light sharing breaks even with the open field at floor-designed densities at 70–92 USD MWh^{−1}

*Corresponding author.
Preprint submitted to Artificial Intelligence in Agriculture
Email address: vinh.dq@buv.edu.vn (Quang-Vinh Dang)

(static arrays at 59) and a conventional plant is more profitable above 48. For maize and soybean (53% of the upscaled 14 PWh yr⁻¹) the simulator's bias is uncertain.

8 *Keywords:* Agrivoltaics, Light sharing, Solar tracking, Food–energy nexus,
9 Crop modelling, Differentiable simulation, Decision under uncertainty

10 1. Introduction

11 Two of the defining challenges of this century compete for the same re-
12 source: land. The energy transition requires terawatts of new solar pho-
13 tovoltaic (PV) capacity, and ground-mounted solar farms need roughly one
14 to two hectares per megawatt. Food production must rise while crops are
15 exposed to more frequent heat waves and droughts. Warming alone is pro-
16 jected to reduce global yields of maize, wheat, rice and soybean by several
17 per cent per degree (Zhao et al., 2017; IPCC, 2021), and a new generation
18 of climate and crop models shows these impacts emerging earlier than pre-
19 viously expected (Jägermeyr et al., 2021). Croplands are also where solar
20 power potential is greatest (Adeh et al., 2019). They are flat, sunny, acces-
21 sible and close to rural electricity demand, so converting farmland into solar
22 farms is tempting and controversial at the same time.

23 Agrivoltaic (AV) systems combine crops and PV modules on the same
24 land and aim to resolve this conflict (Dupraz et al., 2011; Dinesh and Pearce,
25 2016; Weselek et al., 2019). By sharing sunlight between modules and plants,
26 a hectare can deliver more food plus energy than the two uses side by side,
27 with land-equivalent ratios (LER) well above one in field trials and mod-
28 elling studies (Dupraz et al., 2011; Amaducci et al., 2018; Valle et al., 2017;
29 Trommsdorff et al., 2021). The partial shade also changes the crop micro-
30 climate. It lowers canopy temperature and evaporative demand, conserves
31 soil moisture and improves water-use efficiency, especially in drylands (Mar-
32 rou et al., 2013b; Hassanpour Adeh et al., 2018; Barron-Gafford et al., 2019;
33 Elamri et al., 2018; Weselek et al., 2021). Agrivoltaics is therefore a can-
34 didate adaptation measure against heat and drought stress, as reviewed by
35 Mamun et al. (2022), in addition to a mitigation measure through renew-
36 able electricity. The same shade, however, also reduces the light available for
37 photosynthesis. Regulators increasingly require AV systems to retain most
38 of the reference crop yield. The German pre-standard DIN SPEC 91434, for
39 example, requires at least 66% of the reference yield (DIN, 2021), and other

schemes and many farmers set stricter floors. How much electricity can be produced without sacrificing food security depends on the crop, the local climate, the year and the growth stage.

Most AV designs fix the geometry of the array, i.e. the ground coverage ratio (GCR, module area per ground area) and the tilt or tracking rule, from generic rules of thumb. Single-axis trackers add a degree of freedom: the modules can be rotated away from the sun to let light through when the crop needs it, as demonstrated for lettuce by Valle et al. (2017), or towards the sun to maximise electricity. Choosing the rotation every day is a sequential decision problem. The value of light depends on the crop’s phenological stage, soil water and upcoming weather, and the effect of a decision appears in the harvest months later. Recent work in this journal has used evolutionary multi-objective optimisation to lay out PV arrays on farmland (Zhang et al., 2026). Machine learning for crop management has mainly used reinforcement learning with black-box simulators (Gautron et al., 2022), which would need very many simulated seasons for global, multi-decadal assessments.

Two methodological strands are relevant. Differentiable process models, which couple physical simulators with learning, allow parameters and policies to be fitted by gradient descent through the simulator (Reichstein et al., 2019; Shen et al., 2023), and we use them for both calibration and control. Decision making under deep uncertainty asks for designs that perform acceptably across plausible models and not only under the best estimate (Marchau et al., 2019), and scenario-based designs have finite-sample guarantees under convexity assumptions that we do not meet (Calafiore and Campi, 2006; Campi and Garatti, 2008); our design rules are heuristic counterparts and we compare them with a simple safety margin.

Two questions follow. First, how much of the electricity that an agri-voltaic system can produce under a yield-retention floor comes from the array density, and how much from dynamic light-sharing control of the trackers? Second, how robust is that answer to the weakest part of any such study, the crop response to shade, and does a controller learned from a simulator add anything to a well-designed hand-written rule? Answering them requires a simulator that couples the physics of light, PV, microclimate and crop growth, a protocol in which all controllers face the same design constraints, a check of the shade response against field data, and enough sites and seasons to separate real effects from noise.

This paper makes three contributions. The first is a quantification of light-sharing control under a yield floor. Light sharing is parameterised by

one daily level between full backtracking and the rotation that minimises the shadow on the crop, and we compare static arrays, a seasonal schedule, a rule that follows crop phenology, a rule that reacts to heat and drought stress, and a neural controller (SALS) trained through the simulator. The second is a deployable evaluation protocol. The array density of every site is chosen from its historical seasons only, with the same rule for every method, and performance is measured on held-out sites and years, over three spatial splits and three training seeds per split. The third is a differentiable agrivoltaic simulator that supports both, coupling hourly solar geometry, single-axis trackers with backtracking, row-to-row shading and PV electricity, the direct and diffuse light reaching the crop, canopy temperature and reference evapotranspiration under the modules, and the SIMPLE crop model (Zhao et al., 2019) extended with a flowering-stage shade penalty. All components are differentiable tensor operations (Karniadakis et al., 2021). The PV part is benchmarked against PVGIS. The crop response to shade is learned from the open data of a meta-analysis (Laub et al., 2022), checked against field trials on rice, wheat, potato and maize (Gonocruz et al., 2021; Weselek et al., 2021; Ramos-Fuentes et al., 2023), and propagated through the simulator as an ensemble.

We apply the protocol to 400 crop–site pairs that represent where maize, wheat, rice, soybean and potato are grown worldwide, sampled in proportion to their harvested area (Yu et al., 2020) with local crop calendars (Jägermeyr et al., 2021), and driven by NASA POWER daily weather for 2001–2019 (Sparks, 2018) under baseline and +1.5, +2 and +3 °C climates. The main findings are as follows.

1. Light sharing is the main lever on electricity at a given floor, although economically it pays only at high prices (item 6). At the same mean yield retention of 90%, arrays with light-sharing control carry 2.1 to 2.9 times the module area of static arrays and produce 1.8–2.4 times as much electricity (Section 3.1).
2. The best hand-written rule follows crop phenology and needs no weather forecast. Rules that react to heat or drought stress are not better and depend on forecast quality (Section 3.3).
3. The crop response to shade can be learned from data. A hierarchical model fitted to 428 observations from 58 field studies predicts a held-out study of the four staple crop types to about 0.22 in relative yield, no better than proportional loss (it is better only for crop types the

- 115 simulator does not represent), so between-study spread is large. The
 116 simulator is close to these curves for cereals and potato but milder
 117 than the meta-analysis for maize and soybean, whereas one maize trial
 118 is milder still; the gains for those two crops are therefore uncertain and
 119 possibly overestimated (Section 3.7).
- 120 4. A learned controller adds a small gain over the best rule: 0.044 LER
 121 over the written phenology rule but only 0.007 over a window optimised
 122 on the training seasons of each split. Training it on ensembles of shade
 123 responses, or giving it a farm-measurable biomass signal, does not help,
 124 and with the wider ensembles it did not retain its margin over the
 125 written rule in our training budget (Section 3.8).
 - 126 5. The bottleneck is the uncertainty of the crop response that enters the
 127 design. A density chosen under the mean response leaves about 70% of
 128 crop–site pairs at the floor when another plausible response holds (ora-
 129 cle 85–87%). A rule built on the quantiles of an ensemble of responses
 130 raises this by 8 points for 0.03 LER, and a floor raised by 0.02 under
 131 the mean response gives no detectable difference at equal compliance
 132 (the intervals are too wide to establish equivalence); no rule is safe out-
 133 side the ensemble (whose spread is narrower than the between-study
 134 spread), so the ensemble is useful to size a margin (Section 3.8).
 - 135 6. Under illustrative costs, light sharing breaks even with the open field,
 136 at floor-designed densities, at 70 to 92 USD MWh⁻¹ (static arrays at
 137 59), learning is worth at most about 0.3–0.4 kUSD ha⁻¹ yr⁻¹ over an
 138 optimised window at 90 USD MWh⁻¹ (the 95% intervals reach from
 139 about -0.1 to 0.8, and sensing and compute costs are excluded), and
 140 a conventional plant replacing the crop is more profitable above its
 141 break-even of 48 USD MWh⁻¹ (Section 3.9).
 - 142 7. At the yield floor, agrivoltaics on 5% of the world’s staple cropland
 143 could supply about 14 PWh yr⁻¹, but only with about 10 TWp of mod-
 144 ules, several times the installed PV capacity of the world (Section 3.11).

145 2. Materials and methods

146 2.1. Global cropland sites, weather and crop calendars

147 We studied five staple crops that together cover about 670 Mha of har-
 148 vested area and 580 Mha of physical cropland: wheat, rice, maize, soybean
 149 and potato. Harvested areas (rainfed and irrigated) were taken from the

150 SPAM 2010 gridded production maps (Yu et al., 2020) and aggregated to 1°
 151 cells. For each crop, 80 cells with at least 5000 ha of the crop were drawn
 152 without replacement with probability proportional to harvested area. An un-
 153 weighted average over the sample therefore approximates an area-weighted
 154 average over global production. The 400 crop-site pairs lie in 381 distinct
 155 cells in 50 countries (Fig. 2). Sowing and maturity dates came from the
 156 GGCM Phase 3 crop calendar at 0.5° (Jägermeyr et al., 2021). For wheat
 157 we used the winter-wheat calendar unless the coldest-month mean temper-
 158 ature was below -8°C , in which case the spring-wheat calendar was used.
 159 Daily weather for 2001–2020 was retrieved from NASA POWER (Sparks,
 160 2018): global and diffuse shortwave irradiance, maximum, minimum and
 161 dew-point temperature, precipitation and wind speed at 2 m. Each site con-
 162 tributes 19 seasons (sowing years 2001–2019), each simulated as the 365 days
 163 after sowing, which gives 7600 site-seasons.

164 Climate change was represented by the global warming levels +1.5, +2
 165 and +3 °C relative to 2001–2020. Daily maximum, minimum and dew-point
 166 temperatures were shifted uniformly, which keeps relative humidity approx-
 167 imately constant and raises the vapour-pressure deficit. The atmospheric
 168 CO₂ concentration was set to 390, 450, 500 and 600 ppm for the baseline and
 169 the three warming levels. Precipitation and irradiance were left unchanged.
 170 This delta approach isolates the thermal signal of warming; its limitations
 171 are discussed in Section 4.5.

172 *2.2. A differentiable agrivoltaic simulator*

173 The simulator advances all site-seasons in parallel with a daily time step
 174 for the crop and an hourly time step for radiation and PV (Fig. 1). All equa-
 175 tions are implemented as PyTorch tensor operations (Paszke et al., 2019), so
 176 that yields and electricity are differentiable with respect to the design and
 177 control variables.

178 *Hourly irradiance.* Daily global and diffuse irradiance were disaggregated
 179 into hourly values with the Collares-Pereira and Rabl (1979) and Liu and
 180 Jordan (1960) ratios and renormalised to the daily totals. Beam irradiance
 181 on the horizontal is the difference between the two. Solar position follows
 182 the standard declination and hour-angle equations.

183 *Array geometry and PV electricity.* We consider elevated horizontal single-
 184 axis trackers with north–south axes. The axis height is 1.25 module widths,

185 e.g. 2.75 m for the common 2.2-m single-module-high (1P) tracker. The ro-
 186 tation ϕ is limited to $\pm 60^\circ$. The ground coverage ratio g is the module width
 187 divided by the row pitch and is restricted to $g \leq 0.5$. For 2.2-m modules this
 188 keeps a pitch of at least 4.4 m, and a free strip of at least 2.2 m between the
 189 rows when the modules are horizontal, which is the stow position during field
 190 operations. Let ψ be the solar zenith angle projected onto the east–west ver-
 191 tical plane. The energy-optimal rotation with backtracking (Lorenzo et al.,
 192 2011) is $\phi_{\text{bt}} = \psi - \text{sgn}(\psi) \arccos(\cos \psi / g)$ when $\cos \psi < g$, and $\phi_{\text{bt}} = \psi$ oth-
 193 erwise. The *light-sharing* rotation ϕ_{sh} is the feasible angle closest to edge-on
 194 to the sun, which minimises the shadow on the crop. Given a light-sharing
 195 level $u \in [0, 1]$ chosen once per day, the rotation in every hour is

$$\phi = (1 - u) \phi_{\text{bt}} + u \phi_{\text{sh}}. \quad (1)$$

196 Plane-of-array irradiance combines three terms: beam irradiance on the mod-
 197 ule, reduced by the fraction of the module shaded by the neighbouring row,
 198 $\max\{0, 1 - \cos \psi / (g |\cos(\phi - \psi)|)\}$; isotropic sky diffuse; and ground-reflected
 199 light (albedo 0.2). The cell temperature follows the NOCT model (45°C),
 200 with an hourly air temperature interpolated sinusoidally between the daily
 201 extremes. Electricity per unit land is plane-of-array irradiance $\times \eta \times g$, with
 202 $\eta = 0.21 [1 - 0.0035(T_{\text{cell}} - 25)] (1 - 0.14)$ for module efficiency, temperature
 203 losses and system losses. The installed capacity is $2.1 g \text{ MWp ha}^{-1}$. A con-
 204 ventional ground-mounted PV plant ($g = 0.4$, backtracking, no crop) is the
 205 energy reference E_{PV} .

206 *Validation of the PV model.* We compared the simulated specific yield of an
 207 isolated tracker row ($g = 0.05$, negligible row-to-row shading) with PVGIS 5.2
 208 (European Commission, Joint Research Centre, 2022), for a horizontal north–
 209 south single-axis tracker with 14% losses, at 30 randomly chosen sites (29
 210 returned data). Across sites, the simulated and PVGIS specific yields were
 211 correlated with $r = 0.971$. The simulator was on average 4.9% lower (mean
 212 absolute deviation 5.5%). This is consistent with its $\pm 60^\circ$ rotation limit
 213 and with the different irradiance databases (NASA POWER vs. SARA2,
 214 ERA5 and NSRDB). The benchmark covers an isolated row at low density
 215 only; electricity at high density, with row-to-row shading and diffuse-light
 216 effects, was not validated.

217 *Light and microclimate at the crop.* Averaged over the row pitch, the fraction
 218 of ground reached by beam light is $1 - \min\{1, g |\cos(\phi - \psi)| / \cos \psi\}$. The

fraction of sky diffuse reaching the ground is a two-dimensional isotropic sky-view factor. We pre-computed it for 25 values of g and 13 rotation angles by ray casting over 31 neighbouring rows, and interpolate it bilinearly (differentiably). The daily ratio s of the irradiance at the crop to that of the open field defines the crop’s shading. The canopy maximum temperature is lowered by $\kappa(1 - s)$.

κ is uncertain. Field studies under PV modules report small reductions of daytime crop or canopy temperature and little change in air temperature (Marrou et al., 2013b; Barron-Gafford et al., 2019), and heat-stress thresholds calibrated on air temperature do not transfer directly to canopy temperature (Siebert et al., 2014; Webber et al., 2018). We therefore use a moderate $\kappa = 1^\circ\text{C}$ as the main case and report $\kappa = 0$ (no cooling) and $\kappa = 2^\circ\text{C}$ for all methods (Section 3.2). Nocturnal warming under the modules is not represented. The reference evapotranspiration under the modules, ET_0 , is re-computed with the FAO-56 Penman–Monteith equation (Allen et al., 1998) from the reduced radiation and canopy temperature. Soil water conservation under the modules (Hassanpour Adeh et al., 2018; Elamri et al., 2018) therefore emerges from the model, but rain redistribution by the modules (Marrou et al., 2013a) is ignored.

Crop growth. Crop growth follows the SIMPLE model (Zhao et al., 2019) with the species parameters of its Table 1a (Table 1). Daily biomass growth is

$$\Delta B = R_{\text{crop}} \cdot \text{RUE} \cdot f_{\text{Solar}} \cdot f_{\text{CO}_2} \cdot f_{\text{Temp}} \cdot \min\{f_{\text{Heat}}, f_{\text{Water}}\}, \quad (2)$$

where $R_{\text{crop}} = sR$ is the radiation received by the crop. The terms are as follows.

- f_{Solar} is the canopy’s radiation interception. In SIMPLE it depends only on thermal time, with senescence accelerated by heat and drought stress. Shade therefore reduces the daily biomass increment but not leaf-area development.
- f_{Heat} decreases linearly between the heat threshold T_{heat} and the extreme temperature T_{ext} of the canopy maximum temperature.
- $f_{\text{Water}} = 1 - S_{\text{water}} \text{ARID}$, where the agricultural reference index for drought (Woli et al., 2012) is computed from a daily soil-water balance with curve-number runoff, drainage and root water uptake. Irrigated fields have $f_{\text{Water}} = 1$.

253 SIMPLE has a fixed harvest index and therefore ignores the high sen-
 254 sitivity of grain number to shading around flowering (Fischer, 1985; An-
 255 drade et al., 1999). It also ignores the differences in shade tolerance be-
 256 tween crop types. In the meta-analysis of Laub et al. (2022), maize and
 257 grain legumes lost yield strongly even at low shade levels, whereas C3 ce-
 258 reals and tuber crops initially lost less than proportionally; shading trials
 259 also show marked effects of shade on maize development (Ramos-Fuentes
 260 et al., 2023). We added a stage-dependent penalty. During a critical window
 261 ($0.45 \leq \text{TT}/T_{\text{sum}} \leq 0.65$), the harvest index is multiplied by $1 - h\bar{d}$, where $\bar{d}$
 262 is the mean relative radiation deficit $1 - s$ in that window. The sensitivity h
 263 is highest for maize (0.6) and soybean (0.5), the crop types that Laub et al.
 264 (2022) found most shade-sensitive, lower for the C3 cereals (wheat 0.4; rice
 265 0.5, for which no specific curve was available) and lowest for potato (0.2).
 266 These values are assumptions that reproduce this ranking; they are not cal-
 267 ibrated estimates. Section 3.6 shows the resulting shade-response curves.
 268 Because RUE is linear, SIMPLE does not credit the higher light-use effi-
 269 ciency of shaded canopies under more diffuse light. The penalty, conversely,
 270 makes shading around flowering more damaging than proportional. The net
 271 effect is closer to the field evidence than either assumption alone.

272 The thermal-time requirement to maturity was calibrated for every site as
 273 the mean thermal time between the GGCM sowing and maturity dates under
 274 the 2001–2020 climate, and the canopy parameters were scaled accordingly.
 275 The cultivar is thus adapted to the present climate, and warming shortens the
 276 season. Growth stops smoothly once the thermal-time requirement is met.
 277 Yield is biomass multiplied by the (shade-adjusted) harvest index. Each site
 278 is represented by a rainfed and an irrigated sub-field in proportion to its
 279 SPAM irrigated fraction. Each sub-field carries its own array, so both yield
 280 *and* electricity are area-weighted over the two sub-fields.

281 2.3. SALS: simulation-trained adaptive light sharing

282 SALS learns a *control network* $u_t = \pi_{\omega}(\mathbf{x}_t; g)$ (Algorithm 1). It is a
 283 multilayer perceptron with two hidden layers of 64 units and a bounded-logit
 284 output, $u = \sigma(5 \tanh(z/5)) \in [0.007, 0.993]$, which keeps gradients non-zero
 285 at the extremes. It chooses the daily light-sharing level of Eq. (1). Its inputs
 286 are the information available on day t :

- 287 • the crop and soil state: relative thermal time TT/T_{sum} , radiation in-
 288 terception f_{Solar} , relative soil water and the previous day’s ARID;

- a *forecast* of the day’s maximum temperature (relative to T_{heat} and T_{opt}), irradiance and ET_0 ;
- static information: the irrigation status of the sub-field, the array’s GCR g , latitude, day of season and the crop.

Forecasts are perturbed with random errors, independently every day: Gaussian with a standard deviation of 2°C for maximum temperature, and multiplicative with standard deviations of 20% for irradiance and 15% for ET_0 . The physics always uses the true weather. The same forecast errors are applied to the rule-based controllers. The controller is conditioned on g . During training, each season is simulated with a GCR close to the density that is sustainable at its training site. This is the largest value on a coarse grid $(0.05, 0.10, \dots, 0.50)$ at which seasonal sharing keeps the mean retention of the site’s training seasons at or above ρ , plus a uniform perturbation in $[-0.05, +0.10]$ truncated to $[0.05, 0.50]$. The network therefore learns to control the densities that design selection will later pick. Outside the cropping season the trackers backtrack for maximum electricity.

For site-season i with open-field yield Y_i^{open} , let $r_i = Y_i/Y_i^{\text{open}}$ be the yield retention and $e_i = E_i/E_{\text{PV},i}$ the electricity relative to a conventional PV plant. Training minimises the penalised objective

$$\mathcal{L}(\omega) = \mathbb{E}_{i \sim \mathcal{B}} \left[-e_i + \mu_1 \max\{0, \rho - r_i\} + \mu_2 \max\{0, \rho - r_i\}^2 \right], \quad \rho = 0.9, \mu_1 = 5, \mu_2 = 30, \quad (3)$$

which combines an exact (hinge) and a quadratic penalty for the yield-retention floor ρ . Because r_i and e_i are outputs of the differentiable simulator, the exact gradient of Eq. (3) is obtained by back-propagation through the 365 daily steps of every season (back-propagation through time). This is the key difference from reinforcement learning with black-box simulators (Gautron et al., 2022), which estimates gradients from sampled returns.

Each minibatch contains 512 site-seasons, each simulated as a rainfed and an irrigated sub-field, drawn from one of the four climates in turn. Seasons with a failed open-field crop ($Y^{\text{open}} < 0.2 \text{ t ha}^{-1}$) are excluded from training and evaluation. The network is fine-tuned through the simulator with Adam (Kingma and Ba, 2015) (learning rate 10^{-3} , cosine decay, gradient-norm clipping at 1) for 300 steps. Growth-stage logistic functions are evaluated in a numerically stable form, so that gradients remain finite even under extreme heat.

Algorithm 1 SALS training

Require: training site-seasons in four climates; floor ρ ; penalties μ_1, μ_2

- 1: initialise π_ω by behaviour cloning of the phenology rule
- 2: **for** step = $1, \dots, N$ **do**
- 3: sample a climate and a minibatch $\mathcal{B}$ of site-seasons (rainfed and irrigated sub-fields)
- 4: set g_i to the site’s sustainable density plus a perturbation; initialise crop and soil states
- 5: **for** day $t = 1, \dots, 365$ **do**
- 6: $u_{it} \leftarrow \pi_\omega(\mathbf{x}_{it}; g_i)$ from state and noisy forecast if the crop is in the field, else 0
- 7: hourly rotation (Eq. 1), PV electricity, crop light, canopy temperature, ET_0
- 8: advance the soil-water balance and the SIMPLE crop state (Eq. 2)
- 9: $r_i \leftarrow Y_i/Y_i^{\text{open}}, e_i \leftarrow E_i/E_{\text{PV},i}$ (area-weighted); loss $\mathcal{L}$ (Eq. 3)
- 10: back-propagate through the season; Adam update of ω

322 *Initialisation.* Trained from random weights, the controller converged to
323 sharing light throughout the cropping window ($u > 0.99$ on all days), i.e. to
324 the seasonal-sharing rule, because the gradient of electricity with respect to u
325 is small and of the same sign for most days of a season. We therefore initialise
326 the network by behaviour cloning of the phenology rule of Section 2.4: binary
327 cross-entropy against the rule’s output (soft targets 0.95 and 0.05) on about
328 1.7 million day-states visited by the rule in the training seasons, with 1500
329 Adam steps at a learning rate of 3×10^{-3} . Fine-tuning then moves the
330 policy away from this initialisation wherever doing so improves the objective.
331 Consequently, the gain of SALS over the phenology rule measures what fine-
332 tuning through the simulator adds to that rule, not what can be learned
333 without it.

334 *2.4. Deployable design selection and baselines*

335 A real array has one fixed density, which must be chosen before it is built.
336 Every method, SALS included, therefore uses the same *history-based design*
337 *rule*. For a site, the GCR is the largest value on the grid $\{0.05, 0.075, \dots, 0.50\}$
338 for which the method’s mean yield retention over the site’s 2001–2014 sea-
339 sons is at least $\rho = 0.9$. If no value qualifies, no array is built and the site
340 remains open field. A stricter *chance-constrained* variant requires $r \geq 0.9$ in

at least 80% of the historical seasons. The historical seasons of the test sites are never used for training, and the weather of the evaluation years is never used for design. We compare four controllers:

- *AV static*: backtracking trackers without light sharing ($u = 0$);
- *Seasonal sharing*: the light-sharing rotation ($u = 1$) throughout the cropping season;
- *Stress rule*: an expert rule that shares light ($u = 1$) during the main growth phase ($0.1 < TT/T_{\text{sum}} < 0.9$) except on days with a heat-stress forecast ($T_{\text{max}} \geq T_{\text{heat}}$) or drought ($\text{ARID} > 0.5$), when the modules track the sun and shade the crop;
- *SALS*: the learned controller.

The open field and a conventional PV plant without crops ($\text{LER} = 1$ each) are the references. Because the design rule is identical for all controllers, differences between methods reflect the quality of control and the array density that each controller can sustain under the yield-retention floor. We also compare the controllers at identical GCR over the whole grid (matched-density frontiers).

2.5. A shade-response model learned from field data

The simulator needs a crop response to shade that is supported by observations. We learned one from the open data of the meta-analysis of Laub et al. (2022) (Zenodo 10.5281/zenodo.5716091, CC BY 4.0): 428 observations of yield relative to an unshaded control at relative shade reductions of 0–90% from 58 studies, with crop and crop type, experiment type (artificial shading or intercropping), climate zone, irradiation and shade type. Relative yield is modelled as $y(s) = \exp(-\beta s^\gamma)$ with Student- t errors ($\nu = 4$; maximum a posteriori estimation, so this is a hierarchical shrinkage regression with bootstrap intervals and not a full Bayesian model). We compared six models: proportional loss $y = 1 - s$; one curve for all crops; one curve per crop type without pooling; one curve per crop type with partial pooling (Gaussian shrinkage of $\log \beta$ and $\log \gamma$ towards a common value, scale 0.5); partial pooling with covariates on $\log \beta$ (standardised irradiation, subtropical zone, solar-panel shade and intercropping); and gradient-boosted trees on the same inputs. Fits use L-BFGS. Models were compared by leave-one-study-out cross-validation, in which every study is held out once, which tests

375 the prediction of a new study and not of a new observation in a known study.
 376 Intervals for the crop-type curves and covariate effects come from a cluster
 377 bootstrap over studies (150 resamples).

378 2.6. *Shade-response uncertainty: gradient calibration, robust controllers and* 379 *design*

380 The crop’s response to shade is the least certain part of the simulator,
 381 so we treat it as an unknown with a distribution. Two parameters per crop
 382 enter the simulator: a radiation-use compensation c (growth is multiplied by
 383 $1 + c(1 - s)$, s being the share of open-field irradiance that reaches the crop)
 384 and a multiplier h on the flowering-window harvest-index penalty. Because
 385 the simulator is differentiable, both can be fitted by gradient descent through
 386 whole simulated seasons to published yield–shading observations. The ob-
 387 servations are the relative yields at 20% and 40% shading of the data-driven
 388 curves of Section 2.5 for the five crops (median, with an error equal to the
 389 width of the 80% bootstrap interval divided by 2.56 and at least 0.04; rice
 390 takes the curve of C3 cereals, which contains no rice), the field trials on rice,
 391 wheat, potato and maize (Gonocruz et al., 2021; Weselek et al., 2021; Ramos-
 392 Fuentes et al., 2023) (each season of each trial as a separate observation, error
 393 0.06–0.08) and the unaffected harvest index of shaded maize (Ramos-Fuentes
 394 et al., 2023). We call the resulting ensemble E2. A wider ensemble, E3, uses
 395 the same procedure with a target error of 0.2 for the learned curves, which
 396 is the between-study spread of a new study (Section 3.7); the spread of
 397 the resulting members is smaller (standard deviation of relative yield across
 398 members 0.04–0.09 at 20% shading and 0.05–0.12 at 40% shading, against
 399 0.02–0.04 and 0.03–0.06 for E2), because the fit also weights the field trials
 400 and the priors. An earlier ensemble, E1, was fitted in the same way to curves
 401 read by hand from the figures of Laub et al. (2022) and was used to train
 402 the robust and adaptive controllers below. Priors are $c \sim \mathcal{N}(0.4, 0.5^2)$ and
 403 $h \sim \mathcal{N}(1, 0.5^2)$, and the parameters are clamped to $[0, 2.5]$ (c) and $[0, 2]$ (h)
 404 after every step, which is why a response milder than the data cannot be
 405 lowered further. We fitted 12 ensemble members by randomised maximum a
 406 posteriori estimation: every member sees the observations perturbed by their
 407 errors and a prior draw perturbed by its width, and is optimised with Adam
 408 (40 steps, with the loss still falling slightly at the last step; 40 training-block
 409 seasons per crop and member; shading of the static arrays cached). The
 410 ensemble is a device for generating plausible responses and not a validated
 411 posterior: it has no coverage check, and for some parameters (the harvest-

index multiplier of wheat and potato) its spread equals that of the prior. The design analyses use five fixed members (indices 0, 3, 5, 8 and 10 of 12), chosen in advance. The ensemble is used in three ways. All training and evaluation ran on a 4-core CPU; one training run takes about 20–25 minutes and the calibration of one ensemble about 7 minutes.

Tuned rules. The written phenology rule shares light while thermal time lies between 0.05 and 1.0 of the season requirement. To avoid comparing a trained controller with an untuned rule, we tuned this window on the training-block seasons of each split (500 random seasons, baseline climate, main shade response) with the objective used for all rules: the relative electricity at the density where the mean retention equals the floor, found by interpolation over five densities (0.1 to 0.5). A 4×3 grid of start (0, 0.05, 0.15, 0.3) and end (0.6, 0.8, 1.0) gave the grid-tuned window. A five-parameter schedule (start in $[0, 0.6]$, end in $[0.5, 1.2]$, two ramp widths and a level in $[0, 1]$, all through logistic or exponential transforms) was then optimised by CMA-ES (Hansen and Ostermeier, 2001) with 300 seasons, population 8 and 25 generations (200 evaluations) per run; run 1 started at the grid-tuned window and runs 2 and 3 at random points, and the run with the best training objective was evaluated on that split (the CMA-ES schedule and the grid window were tuned separately on each split’s training block; on split 0 run 2 was best, on split 1 run 2 and on split 2 run 1). The grid window was 0.15 to 1.0 on all three splits, and the schedules of the best runs had starts of 0.20, 0.20 and 0.22, ends of 0.93, 0.92 and 0.93, a level of 1.0 and ramp widths below 0.01 of the requirement. On split 0 run 3 stopped in a worse optimum (objective 0.345 against 0.377) and on split 1 as well (0.349 against 0.379), so three runs give no guarantee of global optimality.

Robust and adaptive controllers. SALS-robust has the architecture and training of SALS (Section 2.3), but every training season is simulated with the response of a randomly drawn ensemble member that the controller does not observe, and the density jitter is widened to $[-0.05, +0.25]$ because feasible densities differ between members. SALS-adaptive additionally observes the realised biomass of the crop under the modules relative to an open reference plot, a quantity that can be measured on the farm by remote sensing or by a reference plot. A feedback rule that shares light whenever this ratio falls below a target serves as the hand-written counterpart; its target (1.0) was chosen on training-block seasons under the ensemble.

Uncertainty-aware design. The design rule of Section 2.4 chooses the largest GCR whose historical retention meets the floor. Under shade-response

450 uncertainty we choose the largest GCR that meets the floor in at least 75%
 451 of the ensemble members (three of four in the leave-one-out analysis, four
 452 of five otherwise), and compare it with a design under the ensemble mean
 453 and with an oracle design under the true response. Evaluation of this rule
 454 is leave-one-out: the truth is the held-out member and the design uses the
 455 other four.

456 *Design rules compared.* With the ensemble members indexed k and one
 457 member held out as the truth, the density of a pair is chosen from its 2001–
 458 2014 seasons as (i) the *oracle*: the largest GCR whose mean retention under
 459 the truth meets the floor; (ii) the *mean-response design*: the same rule under
 460 the mean response of the other four members (the mean of their parameters);
 461 (iii) a *margin*: the same rule under that mean response with the floor raised
 462 to 0.92, 0.94 or 0.96; and (iv) the *ensemble rule*: the largest GCR at which
 463 the mean retention under at least a share q of the other four members meets
 464 the floor ($q = 0.5$: two of four; 0.75 : three of four; all: four of four). Out
 465 of the ensemble, the five members and their mean (the mean of all twelve
 466 members for the mean-response design) are used in the same way and the
 467 truth is one of the four fixed responses below. The margin grid was chosen
 468 after a first analysis of one ensemble and is not pre-specified.

469 *Stress tests.* Controllers and rules are also evaluated under four fixed
 470 responses, the proportional response, the response calibrated to the meta-
 471 analysis, the response fitted to the field trials and a harsh response (propor-
 472 tional growth and a 1.5 times larger harvest-index penalty). The proportional
 473 and harsh responses lie outside the range of the ensemble.

474 2.7. Economic analysis

475 To judge whether the electricity gains matter economically, we computed
 476 an annual net value per hectare for each method relative to the open field:
 477 electricity revenue minus the annualised capital cost (annuity) and operating
 478 cost of the installed capacity minus the loss of crop revenue (crop produc-
 479 tion costs are unchanged). The installed capacity is 2.1 g MWp per hectare
 480 (modules cover a share g of the land at 0.21 kWp m^{-2}). Yield is dry matter,
 481 valued at assumed prices. A conventional ground-mounted plant (GCR 0.4)
 482 that replaces the crop is the comparator. All parameters (Table 15) are as-
 483 sumptions chosen to be representative, not data, and their ranges drive the
 484 conclusions; we therefore report break-even electricity prices and one-at-a-
 485 time sensitivities, and we do not monetise food security, which market crop
 486 prices do not capture. The analysis omits land rent, module degradation,

the extra structure cost of sparse arrays, curtailment, subsidies and the uncertainty of the parameters themselves, which the bootstrap intervals do not include.

2.8. Evaluation protocol, statistics and upscaling

Sites were split spatially: $10^\circ \times 10^\circ$ blocks were randomly assigned to training (70%) or test (30%), so that no test site has a training site in the same block. Training uses the 2001–2014 seasons of the training sites. Design selection uses the 2001–2014 seasons of the test sites. Evaluation uses the 2015–2019 seasons of the test sites. The same site-seasons are analysed in all four climates.

We report the following metrics:

- the mean electricity ($\text{MWh ha}^{-1} \text{yr}^{-1}$) and the relative electricity e ;
- the mean yield retention r and the production-weighted retention $\sum Y / \sum Y^{\text{open}}$;
- season-level compliance (share of seasons with $r \geq \rho$) and site-level compliance (share of sites whose mean test-season retention is $\geq \rho$);
- the lower tail (10th percentile of r , share of seasons with $r < 0.8$);
- the land-equivalent ratio $\text{LER} = r + e$ (Dupraz et al., 2011), which gives equal weight to the retained yield and to the electricity relative to a conventional plant, and is relative to the open field of the same climate (a higher LER under warming reflects the lower reference electricity share and not a benefit of agrivoltaics).

Two aggregations are used and labelled in the text: tables report means over all test seasons of a sample (season-pooled), whereas statistical comparisons and intervals use the crop–site pair as the unit of analysis (mean of the pair’s seasons, then mean over pairs), so the same contrast can differ slightly between the two (for example 0.034 and 0.038 for SALS minus the phenology rule on the main split). Samples are also labelled: the main analysis uses the 92 held-out pairs of one split and a GCR grid with step 0.025; analyses pooled over three splits use the 76 distinct pairs of every third held-out pair and a grid with step 0.05, and analyses of alternative shade responses on one split use every sixth (21 pairs) or every fourth pair (22 pairs). Statistical comparisons use site means as the unit of analysis. We report paired

Table 1: SIMPLE crop parameters (Zhao et al., 2019, Table 1a) and the harvest-index shade sensitivity h added in this study. T_{sum} , I_{50A} and I_{50B} are rescaled per site to the local crop calendar (Section 2.2). Units: temperatures in $^{\circ}\text{C}$, thermal times in $^{\circ}\text{C d}$, RUE in gMJ^{-1} ; S_{CO_2} is the relative RUE increase per 100 ppm above 350 ppm.

Crop (cultivar)	h	T_{sum}	HI	I_{50A}	I_{50B}	T_{base}	T_{opt}	RUE	$I_{50 \text{ max } H}$	$I_{50 \text{ max } W}$	T_{heat}	T_{ext}	S_{CO_2}	S_{water}
Wheat (Batten)	0.4	2150	0.34	280	50	0	15	1.24	100	25	34	45	0.08	0.4
Rice (IR72)	0.5	2300	0.47	850	200	9	26	1.24	100	10	34	50	0.08	1.0
Maize (McCurdy 84aa)	0.6	2050	0.50	500	50	8	28	2.10	100	12	34	50	0.01	1.2
Soybean (Williams 82)	0.5	2350	0.40	600	200	6	27	0.86	120	20	36	50	0.07	0.9
Potato (Sebago)	0.2	2400	0.85	500	350	4	22	1.30	50	30	34	45	0.10	0.4

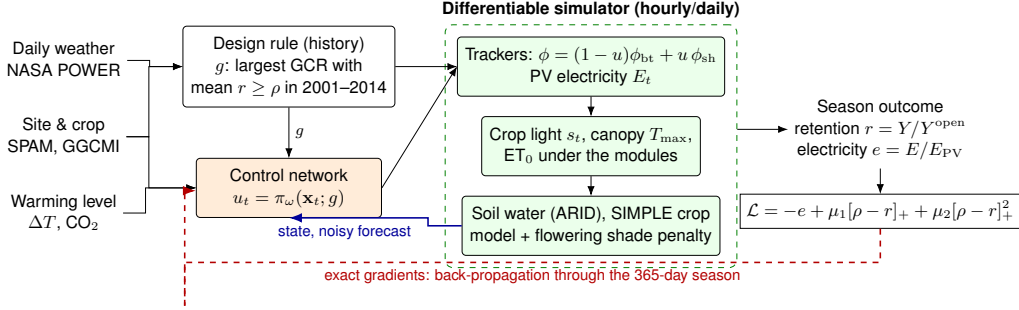


Figure 1: Simulator and controller. The array density g of a site is fixed from its historical seasons. A control network maps the daily crop and soil state, the density g and a noisy forecast to the light-sharing level u_t . The differentiable simulator converts these decisions into PV electricity, light and microclimate at the crop, soil water and crop growth. The season’s electricity and yield retention enter a penalised objective whose exact gradient is back-propagated through the 365-day season.

519 Wilcoxon signed-rank tests and cluster (site) bootstrap 95% confidence in-
520 tervals of the mean difference. SALS was trained with three random seeds on
521 each of the three spatial splits; the controller comparisons pooled over splits
522 use the mean of the three seeds, and the economic analysis uses seed 0.

523 For the global estimate, every site is evaluated in the split in which it is
524 held out. The mean electricity and yield per hectare of each crop are scaled
525 to the *physical* area of that crop in SPAM 2010, which avoids double counting
526 multiple-cropped land (Waha et al., 2020). We report the installed capacity
527 (TWp), the capacity factor and a supply curve.




Figure 2: The 400 crop–site pairs, sampled in proportion to the harvested area of each crop (SPAM 2010). Blue: training sites; orange: geographically held-out test sites of the main split.

528 3. Results

529 3.1. Electricity and food on held-out sites and years

530 Table 2 compares the four controllers on the 92 held-out crop–site pairs of
 531 the main split (95 pairs are held out; pairs without a season whose open-field
 532 yield is at least 0.2 t ha^{-1} in all four climates are excluded) (seasons sown
 533 in 2015–2019), after each site’s density has been chosen from its 2001–2014
 534 seasons with the same rule for all controllers. By construction all methods
 535 keep the mean yield retention close to 90% (0.904–0.917), so differences in
 536 electricity are differences in the density that each controller can sustain under
 537 this floor.

538 **Static arrays are the weakest.** Without light sharing, the floor limits
 539 the ground coverage to 6% on average, and arrays produce $217 \text{ MWh ha}^{-1} \text{ yr}^{-1}$,
 540 16% of a conventional PV plant per hectare (LER 1.07). Rotating the
 541 modules to let light through during the cropping season (seasonal shar-
 542 ing) allows a coverage of 17% and $397 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ (LER 1.19). A rule




Figure 3: Electricity-retention frontiers at identical array densities (labels: GCR) for the held-out sites. Each point is the mean over test seasons at a fixed GCR; the dashed line is the 90% floor.

that shares light only while the crop is growing (phenology rule) reaches 465 MWh ha⁻¹ yr⁻¹ (LER 1.24). An expert rule that additionally shades the crop on forecast heat or drought days (stress rule) is worse than the phenology rule (LER 1.18): its density is limited to 13% because it reduces light at times when the floor is most binding.

The learned controller is best by a modest, consistent margin. SALS sustains a density of 17% (similar to the phenology rule), produces 510 MWh ha⁻¹ yr⁻¹ and reaches an LER of 1.270 at the same retention. Compared with the phenology rule this is +0.038 in LER (95% CI over sites 0.025–0.053; +10% electricity) and it is higher at 84% of the 92 sites (Table A.18). Compared with static arrays it is +0.205 (0.176–0.236), with 2.35 times the electricity at 2.9 times the density. The frontier in Fig. 3 shows the same picture at matched densities: the learned policy lies on or slightly above the phenology rule over the whole density range, both lie well above seasonal sharing and the stress rule, and static arrays are far below.

What the floor does and does not guarantee. Enforcing the floor on the historical mean does not guarantee it in every season. In the test seasons, 70% of SALS seasons have $r \geq 0.9$ (phenology rule 69%, stress rule 76%, static arrays 93%), and 2.7% have $r < 0.8$ (1.3%, 0.9% and 0.9%). The lower tail of SALS is therefore slightly heavier than that of the rules: its 10th percentile of r is 0.883, against 0.887, 0.886 and 0.903. Requiring $r \geq 0.9$ in at least 80% of a site’s historical seasons (chance-constrained design) raises the season-level compliance to 90–94% for the light-sharing

566 controllers. It costs about 0.065 in LER for SALS (1.205), which remains the
 567 best controller (phenology rule 1.192; Section 3.2).

568 *3.2. Robustness*

569 **Seeds and splits.** Three training seeds of SALS give LERs of 1.270,
 570 1.268 and 1.269 under the baseline climate and 1.297, 1.298 and 1.297 at
 571 $+3^\circ\text{C}$. On two further spatial splits, in which different blocks of sites are
 572 held out (training seed 0; the pooled analysis of Table 5 uses three seeds per
 573 split), SALS exceeds the phenology rule by 0.047 and 0.041 in LER under the
 574 baseline climate and by 0.052 and 0.039 at $+3^\circ\text{C}$ (main split, season-pooled:
 575 0.034 and 0.032; the site-level paired contrasts of Table A.18 are 0.038 and
 576 0.034). In every split SALS ranks first, the phenology rule second and static
 577 arrays last.

578 **Canopy cooling.** The canopy-cooling coefficient κ is the most un-
 579 certain microclimate parameter. We repeated the whole protocol, includ-
 580 ing the historical design step, with $\kappa = 0, 1$ and 2°C . Electricity changes
 581 with κ because cooler canopies tolerate denser arrays (SALS: 475, 510 and
 582 $549 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ under the baseline climate), but the retention is held at
 583 0.905 by the design rule and the ranking is unchanged. SALS exceeds the
 584 phenology rule by 0.029, 0.034 and 0.042 in LER under the baseline climate
 585 (0.022, 0.032 and 0.051 at $+3^\circ\text{C}$).

586 **Which ingredients matter.** Two ablations test the design of SALS
 587 at the same training budget. Masking the stress inputs (soil water, ARID,
 588 and the maximum temperature, irradiance and ET_0 forecasts) leaves the
 589 LER unchanged (1.269 vs 1.270 under the baseline climate, 1.297 vs 1.297
 590 at $+3^\circ\text{C}$). Training only on the baseline climate gives 1.268 and 1.291. The
 591 improvement over the rules therefore does not depend on stress-aware inputs
 592 or on training in warmed climates.

593 *3.3. Floor conditioning, chance-constrained training and forecast errors*

594 **One network for all floors.** Regulations and farmers set different
 595 floors. The network of the main experiments was trained at $\rho = 0.9$ only.
 596 We evaluated it at floors of 0.80, 0.85, 0.90 and 0.95, re-running the historical
 597 design step for each floor (Table 3). Its LER exceeds that of the phenology
 598 rule by 0.037, 0.036, 0.034 and 0.033 under the baseline climate (0.040, 0.037,
 599 0.032 and 0.033 at $+3^\circ\text{C}$), so the benefit carries over to other floors without
 600 retraining. A network that received the floor as an additional input and was
 601 trained over all four floors did not improve on it (differences in LER ≤ 0.001).

602 Because the design step already adapts the density to the floor, conditioning
 603 the controller on the floor is not needed.

604 **Layout, rotation limit and harvest-index assumptions.** We checked
 605 three assumptions that could drive the ranking (Table 4; 21 pairs of the main
 606 split, one training seed). The rotation limit of the trackers sets how much
 607 light sharing can deliver: with a limit of 45° the LER of the phenology rule
 608 is 1.138, with 60° 1.170 and with 80° 1.206, whereas static trackers stay at
 609 1.05–1.06, so the benefit of light sharing grows with the limit and the ranking
 610 is unchanged. Fixed non-tracking arrays held flat or rotated by 30° about
 611 the north–south axis (facing east, not the equator; equator-facing tilts are
 612 not modelled) do no better than static trackers (LER 1.048 and 1.043 against
 613 1.053; differences of 0.005 are within the noise of 21 pairs and one training
 614 seed), so static trackers are no worse than the fixed arrays we tested; verti-
 615 cal bifacial layouts are not modelled. The ranking of the rules also does not
 616 depend on the assumed harvest-index penalty. The window test is weakly in-
 617 formative about the timing of the phenology rule, which shares light through
 618 almost the whole growing period, so a shifted window cannot move it much;
 619 the informative case is the removal of the penalty. Moving the critical window
 620 to 0.30–0.50 or 0.60–0.80 of the thermal-time requirement changes the LER
 621 of the phenology rule by at most 0.001, and removing the penalty altogether
 622 raises all light-sharing LERs (phenology rule 1.328, SALS 1.353) without
 623 changing the order: the phenology rule exceeds seasonal sharing by 0.03 to
 624 0.05 in every case (0.035 with the early window, 0.034 with the late window,
 625 0.053 without penalty), because it produces electricity before canopy devel-
 626 opment and after maturity, when the crop does not use the light, and not
 627 because of the assumed window. The margin of SALS over the written rule
 628 changes with these assumptions (0.005 to 0.025 LER), and its margin over
 629 the tuned rule is 0.000 to 0.009 (at a rotation limit of 45° SALS equals the
 630 tuned rule).

631 **A floor of 66%.** The German pre-standard DIN SPEC 91434 requires
 632 66% of the reference yield (DIN, 2021). Repeating the design step with
 633 $\rho = 0.66$ (baseline climate, main split, same network trained at 0.9) the
 634 ground coverage cap of 0.5 binds for the sharing methods (mean GCR 0.48
 635 for the phenology rule and 0.48 for SALS, 0.44 for the stress rule), static
 636 arrays reach a coverage of 0.24 and LER 1.28, the phenology rule 1.66 and
 637 SALS 1.71. At this floor the density is limited by the cap and by the mean
 638 retention of 0.73 reached at the cap, not by the floor, so the comparison of
 639 controllers says little about control; we do not interpret it further.

640 **Chance-constrained training.** We also trained a network whose ob-
641 jective penalises a batch share of seasons below the floor that exceeds 20%.
642 It did not improve the tail. With the mean-based design step its LER was
643 1.263 (baseline) and 1.290 (+3 °C) against 1.270 and 1.297, and with chance-
644 constrained design it reached season-level compliance of 92.6% against 92.0%
645 and an LER of 1.198 against 1.205. The season-level guarantee is obtained
646 by the design step (Section 3.1), not by the training objective.

647 **Forecast errors.** We raised the forecast errors seen by the controllers
648 to 0, 1, 2 and 3 times the errors used in training. The LER of SALS and
649 of the phenology rule was identical at all four error levels to three decimals
650 (densities 0.2 and 0.3: SALS 1.307 and 1.457, phenology rule 1.278 and
651 1.416). Neither controller uses the weather forecast: both rely on the crop’s
652 thermal time, which is observable from on-farm temperature records. The
653 stress rule, which depends on the forecast of maximum temperature, degrades
654 with forecast error: at density 0.3 its retention falls from 0.778 to 0.751 and
655 its LER from 1.432 to 1.418 as the errors triple.

656 **Tuned rules.** The phenology rule above was written, not tuned, whereas
657 SALS is trained. We therefore tuned its window on the training-block seasons
658 with the same objective as the other rules (electricity at the density where
659 mean retention equals the floor), first on a coarse 4×3 grid and then, as a
660 derivative-free baseline, by CMA-ES over a five-parameter schedule (start,
661 end, two ramp widths and a level; three runs of 200 evaluations, Methods).
662 The grid-tuned window starts at 0.15 instead of 0.05 of the thermal-time
663 requirement (Table 5). Two of the three CMA-ES runs, one started at the
664 grid-tuned window and one at a random point, reached the same almost
665 binary window (all three runs stopped at the generation cap) from 0.20 to
666 0.93 of the requirement at level 1.00; the third stopped in a worse optimum.
667 Pooled over three splits and 76 distinct crop–site pairs, the grid-tuned win-
668 dow raises LER over the written rule by 0.026 (95% cluster-bootstrap interval
669 0.016 to 0.037), the CMA-ES window by 0.036 (0.023 to 0.052) and SALS by
670 0.044 (0.030 to 0.060), so the optimisation of a window accounts for about
671 84% of the margin of SALS over the written rule. SALS exceeds the CMA-
672 ES window by 0.007 LER (0.002 to 0.014), which is 2.4% more electricity at
673 equal retention, and the grid-tuned window by 0.018 (0.011 to 0.025). The
674 remaining margin is small (it covers the pairs only: SALS is the mean of three
675 training seeds per split, whose pooled margins over the CMA-ES window are
676 0.007, 0.008 and 0.007, and the CMA-ES schedule and the grid window were
677 re-tuned on each split’s own training block); it is not explained by the tuning

678 of a window or a ramp, and it cannot come from soil water, forecasts or
 679 aridity, because masking them leaves SALS unchanged (Section 3.2); canopy
 680 development and density remain untested. SALS and the CMA-ES window
 681 meet the floor in 83% of seasons, the written, grid-tuned and seasonal rules
 682 in 84–85%. Per split, SALS exceeds the CMA-ES window by 0.008 (95%
 683 interval 0.001 to 0.016; 32 pairs), 0.010 (0.001 to 0.022; 26 pairs) and 0.005
 684 (−0.002 to 0.012; 35 pairs), with per-pair medians of 0.001 to 0.002, so the
 685 margin is positive on every split but concentrated in a few pairs. The design
 686 step quantises the density to a grid of 0.05 (about 30% of the mean density),
 687 which is much coarser than the margin. We therefore also compared the
 688 controllers on the test seasons at matched mean retention, reading the LER
 689 at a retention of 0.90 off the retention–LER curve of each pair (interpolated
 690 over the grid): SALS exceeds the CMA-ES window by 0.0033 (95% inter-
 691 val −0.0004 to 0.0071; 73 pairs), the grid-tuned window by 0.011 (0.007 to
 692 0.016) and the written rule by 0.036 (0.023 to 0.051). At fixed densities of
 693 0.10 to 0.30 the difference to the CMA-ES window is 0.001 to 0.003 LER,
 694 with intervals that include zero. At matched retention the margin over the
 695 optimised window is therefore not distinguishable from zero.

696 3.4. *What the learned policy does*

697 Averaged over the rainfed test seasons of the main split at a density of
 698 0.3, the light-sharing level of SALS is a smooth function of relative thermal
 699 time. It is close to zero in the first 10% of the thermal time required for
 700 maturity, rises through 0.55 at 10–20%, stays at 0.94–0.99 between 30% and
 701 70%, and falls through 0.58 at 80–90% to about zero at maturity. The plateau
 702 overlaps the window in which our crop model penalises shading through the
 703 harvest index; we did not test whether this alignment is learned or inherited
 704 from the initialisation. The ramps at both ends let the modules produce
 705 electricity when additional light is worth least to the crop. Fig. 4 shows two
 706 example seasons. The phenology rule, which is also the initialisation of the
 707 network, shares light during the whole growing period. Fine-tuning through
 708 the simulator found that sharing can be reduced at the edges of the season:
 709 this pattern describes the average policy at a density of 0.3 (the designed
 710 mean is 0.17) and we did not test that it causes the gain; the CMA-ES
 711 window is almost binary, so a smooth ramp is not needed to reproduce most
 712 of it.

713 The policy also shares less light on days with a forecast heat stress: the
 714 mean level is 0.35 on such days against 0.60 (baseline) and 0.51 (+3 °C)




Figure 4: Learned daily control for two test seasons under $+3^{\circ}\text{C}$, at the density chosen for the site, simulated for the dominant sub-field of the site. Shading marks the cropping window. Orange: light-sharing level u_t ; green: fraction of open-field irradiance reaching the crop; grey: daily maximum temperature (right axis).

on other days, and 58–63% of sites show a lower level on heat days. This cannot be read as a stress response, however. Heat days occur late in the season, when the phenological ramp is already falling, and a policy without any stress input shows a similar contrast (0.37 vs 0.57 in a separate check). The ablation above confirms that stress information does not improve the result.

3.5. Crops and regions

The benefit of light sharing differs between crops (baseline climate, SALS). Soybean and maize tolerate the densest arrays (GCR 0.24 and 0.19; 691 and 577 $\text{MWh ha}^{-1}\text{yr}^{-1}$; LER 1.41 and 1.30), potato follows (0.19; 542; 1.31), and rice and wheat are the most restrictive (0.13 and 0.15; 395 and 428; 1.20 for both). The stress rule has a lower LER for every crop (1.26 for soybean, 1.14–1.21 for the others).

The floor is not met evenly. Production-weighted retention is 0.888 for soybean and 0.898 for maize, below 0.9, because high-yielding seasons lose more than low-yielding ones. Regionally, Latin America and Oceania (4 sites) combine the highest density (GCR 0.41) and electricity (1090 $\text{MWh ha}^{-1}\text{yr}^{-1}$) with a production-weighted retention of only 0.862. Africa and the Middle

East (8 sites) reach 0.894, and only 63% of their sites keep a mean retention of at least 0.9. South Asia, the region with the most sites (39), reaches 0.902 (67% of sites) at a density of 0.14. The regional figures rest on 4 to 39 pairs without intervals and should be read as indications. A regulatory floor enforced per site would therefore need a region- and crop-specific design step, not a global one.

3.6. Model behaviour and validation

PV output. The simulated specific yield of an isolated tracker row agrees with PVGIS 5.2 (European Commission, Joint Research Centre, 2022) at 29 sites with $r = 0.971$, a mean bias of -4.9% and a mean absolute deviation of 5.5% (Section 2.2).

Shade response. Fig. 5 shows the simulated relative yield of rainfed crops under static arrays as a function of the season-mean shading of the crop. Yield declines close to proportionally with shading: 84–89% relative yield at 12% shading, 70–78% at 25%, 57–67% at 37% and 35–46% at 57%. Wheat loses most and potato and maize least. This response is steeper than the *less than proportional* losses that Laub et al. (2022) report for small shade levels of C3 cereals and tubers, and it does not reproduce their finding that maize and grain legumes are more shade sensitive than C3 cereals. In the simulator, shading also relieves heat and drought in rainfed maize and soybean, which offsets the higher harvest-index penalty we assigned to them. The absolute effect of shading on each crop is therefore uncertain (Section 4.5).

Comparison with field trials. We fitted one radiation-use compensation coefficient per crop to the yield–shading curves of Laub et al. (2022) using training-block sites only (Methods, calibration script in the repository). The fit changed wheat and rice only (coefficients 0.4 and 0.1) because the proportional model already lies within about 0.02 of the C3-cereal and tuber curves, and it cannot lower the maize and legume responses below the proportional case (root-mean-square error 0.16 for maize and 0.08 for soybean). Against field trials (not independent of the choice of ensemble, Table 6) (Table 6) the model underestimates retained yield: rice 0.79 against 0.84 (Gonocruz et al., 2021), wheat and potato 0.74 and 0.72 against 0.81–0.82 in the worse season of Weselek et al. (2021) (and above 1 in the better one), and maize 0.70–0.75 against 0.83–0.91 at 30–35% shading (Ramos-Fuentes et al., 2023), where harvest index was not affected by shade. The model is therefore conservative for food retention with respect to these trials, except for the stronger maize

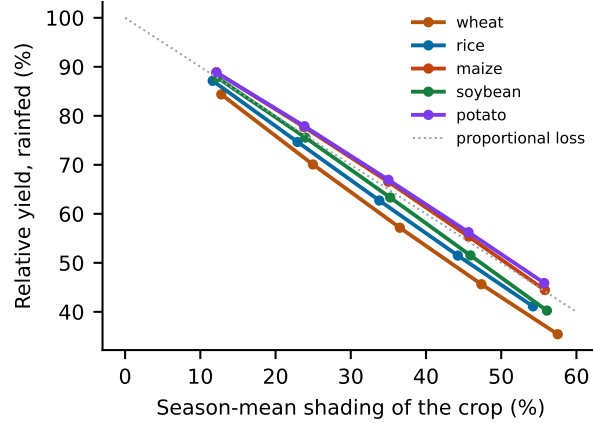


Figure 5: Simulated relative yield of rainfed crops under static arrays (GCR 0.1–0.5) against season-mean shading of the crop, baseline climate, held-out sites. The dotted line is proportional loss.

losses of the meta-analysis. Table 7 shows the consequence for the design results: with the calibrated response the allowed density and the LER of every method rise (phenology rule 1.28 against 1.23), and with a response fitted to the field trials the LER of the phenology rule is 1.62. The phenology rule is the best method in all three cases, and light sharing adds LER +0.17, +0.19 and +0.30 over static arrays (conservative, calibrated, field-fitted); only the order of the seasonal and stress rules changes. The absolute gains for rice, wheat and potato are thus conservative with respect to these trials; for maize and soybean the meta-analysis points the other way (Section 3.7).

3.7. A shade-response model learned from field data

Cross-validation by study shows what the data can and cannot support (Tables 8 and 9). Over all 339 shaded observations the root-mean-square error of a held-out study falls from 0.43 for proportional loss $y = 1 - s$ to 0.37 for one curve for all crops and to 0.31 for crop-type curves, but this gain comes from crop types that the simulator does not represent (berries, forages, fruits and vegetables: 0.53 to 0.36). For the four crop types that the simulator uses, the learned curves are not better than proportional loss: the error is 0.222 against 0.220 pooled (difference +0.002, 95% interval -0.023 to $+0.030$), 0.187 against 0.214 for C3 cereals, and slightly worse for maize, grain legumes and tubers. Partial pooling is as good as separate curves (0.312 against 0.313)

789 and the covariates add almost nothing (0.311). Gradient-boosted trees on the
790 same inputs are worse (0.321) than the parametric model. The error of 0.22
791 for the four crop types is large: yields of a new study differ from the learned
792 curve by about 20 percentage points, which is the between-study spread
793 that a design under uncertainty has to cover. Crop type is the information
794 that transfers between studies, and none of the covariate effects on $\log \beta$ is
795 clearly different from zero (irradiation -0.10 , 90% interval -0.32 to 0.09 ;
796 subtropical zone -0.37 , -0.73 to 0.07 ; intercropping 0.16 , -0.48 to 0.45).
797 The estimated effect of solar-panel shade is not interpretable, because all 46
798 panel observations are of leafy vegetables, mainly lettuce, and the effect is
799 confounded with crop type.

800 The learned curves show a clear ordering of crop types (Table 10). At
801 20% shading C3 cereals and tubers lose less than proportionally (relative
802 yield 0.93), whereas maize (0.73) and grain legumes (0.71) lose more; at 40%
803 shading the values are 0.72 and 0.71 against 0.42 and 0.40. Intervals are
804 wide for maize (0.27–0.59 at 40%) and for tubers, which rest on two studies.
805 The simulator’s responses are close to the data for cereals and potato but,
806 under both calibrated ensembles, too mild for maize at 40% shading (0.75 for
807 E2 against 0.42 in the data); for soybean the E2 value (0.48) lies within the
808 interval of the data (0.40, 80% interval 0.34–0.53). The calibration cannot
809 lower the maize response further, because the simulated shading also relieves
810 heat and drought stress and the compensation and harvest-index parameters
811 are bounded below at zero. The field trials of maize are also milder than the
812 meta-analysis: Ramos-Fuentes et al. (2023) report a harvest index unaffected
813 by shade, and Sekiyama and Nagashima (2019) a yield 5.6% above the control
814 under a low-density stilt-mounted array in one season at one site. For maize
815 and soybean the direction of the bias is uncertain, and the gains reported in
816 this paper are possibly overestimated.

817 3.8. *Learning under shade-response uncertainty*

818 **Calibration.** Gradient descent through the simulator fitted the five crops
819 to the data-driven curves and the field trials (ensemble E2, Table 11). The
820 data identify two features: the harvest index of maize is hardly affected by
821 shade (multiplier 0.16 ± 0.13 , $c \approx 0.14$), as observed by Ramos-Fuentes et al.
822 (2023), and soybean is more shade sensitive than the cereals (compensation
823 0.05, harvest-index multiplier 1.66), as in Laub et al. (2022). Wheat has the
824 largest compensation (0.91 ± 0.31) and the largest uncertainty; the compen-
825 sation is an effective parameter that absorbs diffuse-light and light-saturation

826 effects that the simulator does not model explicitly, not a physical value.

827 **Learned controllers.** Controllers were evaluated under four fixed re-
828 sponses on 21 crop–site pairs of split 0 with one training seed, so the intervals
829 below, which are cluster bootstraps over pairs of per-pair mean LER differ-
830 ences, do not include seed or split variation (Table 12 gives season means,
831 which differ slightly). Under every response SALS, trained on the propor-
832 tional response, is better than the written phenology rule: by +0.019 (+0.007
833 to +0.039), +0.024 (+0.009 to +0.046), +0.038 (+0.018 to +0.062) and
834 +0.015 (+0.001 to +0.037) LER under the proportional, calibrated, field-
835 fitted and harsh responses. It is also above the grid-tuned window and the
836 CMA-ES optimised window, but by less: 0.005 to 0.018 and 0.005 to 0.010
837 LER, with intervals that exclude zero for the grid-tuned window under 3 of
838 the four responses and for the CMA-ES window under none. Training on
839 the narrow ensemble E1 did not change this: SALS-robust differs from SALS
840 by -0.005 to $+0.003$ (intervals contain zero except for one response) and
841 SALS-adaptive, which sees the measured biomass ratio, by -0.010 to $+0.007$.
842 Training on the wider ensembles did not help either (the robust controllers
843 fall below SALS even under responses inside their training ensemble, which
844 may point to under-training with our budget of 300 steps): SALS-robust
845 on E2 and E3 and SALS-adaptive on E3 are below SALS by 0.014, 0.017,
846 0.017 LER under the proportional response, 0.017, 0.021, 0.020 under the
847 calibrated response, 0.004, 0.005, 0.011 under the field-fitted response (in-
848 tervals contain zero) and 0.008, 0.014, 0.017 under the harsh response; the
849 intervals exclude zero in 7 of the 9 comparisons for the proportional, cali-
850 brated and harsh responses. The E3 robust controller is not distinguishable
851 from the written rule ($+0.001$ to $+0.032$) under any response (all intervals
852 contain zero). With one seed and 21 pairs these controllers therefore did not
853 retain the margin of SALS within our training budget; a controller that must
854 perform across plausible responses is less specialised, but a larger budget or
855 more seeds could change this. The feedback rule is worse than the written
856 rule by 0.026–0.058 LER under every response (intervals exclude zero). A
857 plausible reason is that the biomass ratio does not reveal the harvest-index
858 loss from shading around flowering, but we did not test this explanation.

859 **Design under uncertainty.** The larger risk is in the design step (Ta-
860 bles 13 and 14, Fig. 6). The designs below are evaluated by leaving one
861 ensemble member out as the truth and using the other four, so that the
862 truth never enters the design. Pooled over the three spatial splits (76 dis-
863 tinct crop–site pairs), a density chosen under the mean response of E2 leaves

70% of the pairs at the floor when the true response is another plausible one,
 against 87% for an oracle design (E3: 71% and 85%). Choosing the density
 so that the floor is met in at least 75% of the other members raises compli-
 ance by +7.9 points under E2 (95% cluster-bootstrap interval 4.5 to 11.6)
 and +7.9 under E3 (3.9 to 12.1), at a cost of 0.031 and 0.025 in LER. Raising
 the floor used in the design from 0.90 to 0.92 under the mean response does
 more: +18.0 points (13.4 to 22.6) at 0.068 LER under E2 and +14.3 points
 (9.5 to 19.2) at 0.056 under E3, and the shortfall is smaller (mean over pairs
 of the maximum of zero and 0.9 minus retention: 0.0022 against 0.0045 for
 the quantile rule under E2). The two kinds of rule lie on one frontier (Fig. 6).
 Interpolating the margin frontier (floors 0.90 to 0.96) at the compliance of
 the ensemble rules, the LER of the margin minus that of the rule is +0.001
 (-0.009 to +0.012) for the rule with 75% of the other members and +0.003
 (-0.016 to +0.021) for the rule with all of them under E2, and -0.006 (-0.019
 to +0.005) and +0.024 (-0.011 to +0.058) under E3 (95% intervals). We find
 no detectable difference, but the intervals (down to -0.02 and up to +0.06
 LER) are too wide to establish equivalence. The ensemble is thus useful for
 sizing the margin (0.02 to 0.04 above the floor for 85–97% of the pairs), not
 for replacing it. None of these rules is safe outside the ensemble (Table 14):
 under the proportional, calibrated and harsh responses the mean-response
 design leaves 7–17% of the pairs at the floor, a floor of 0.92 or the ensemble
 rules 12–33%, a floor of 0.94 29–47% and a floor of 0.96 51–78%, at a cost
 of about 0.2 LER, whereas the field-fitted response, which lies within the
 range spanned by the ensembles, is handled well by every rule (80–99%) ex-
 cept the mean-response design (68–79%). Outside the ensemble the shortfall
 of the pairs below the floor is also large: under E2 the mean shortfall per
 pair (maximum of zero and 0.9 minus retention) is 0.075, 0.060 and 0.104
 for the mean-response design under the proportional, calibrated and harsh
 responses, against 0.053, 0.038 and 0.078 for a floor of 0.92 and 0.012, 0.007
 and 0.024 for a floor of 0.96. Across the five members of an ensemble, the
 standard deviation of the simulated relative yield of a crop (static arrays,
 baseline climate, held-out sites of split 0) is 0.02 to 0.04 at 20% shading and
 0.02 to 0.06 at 40% under E2, and 0.02 to 0.09 and 0.03 to 0.15 under E3
 (largest for soybean and potato; ranges between the extreme members up to
 0.36), against a leave-one-study-out error of 0.22 for new studies. The two
 quantities differ (a spread of curves against an error over observations), but
 the ensembles are narrower than the between-study spread for all crops ex-
 cept, possibly, soybean under E3. The guarantee therefore holds only within




Figure 6: Share of crop-site pairs at the floor against LER for the phenology rule under design rules for shade-response uncertainty (leave-one-member-out, three splits, 76 pairs): a higher floor under the mean response (blue; the first point is the design under the mean response) and a floor met in a share of the other members (orange), with the oracle design (green).

the range of responses that the ensemble spans; we did not test an ensemble that spans the proportional and harsh responses, which would protect against them at a higher cost in LER.

3.9. Economics of light sharing

All figures in this section follow from the assumed parameters of Table 15 and should be read as a sensitivity analysis, not a forecast. The economics are dominated by the capacity installed, not by the crop (Table 16): the annuity and operating cost of the modules is 11–32 kUSD per hectare and year, whereas the crop revenue lost is about 0.3 kUSD, so a method is worth what it earns per installed watt. With the densities designed for the floor, static arrays earn $+0.3 \text{ kUSD ha}^{-1} \text{ yr}^{-1}$ at 60 USD MWh^{-1} and break even at 59, whereas all light-sharing methods lose money at that price (−3 to −10 kUSD) and break even at 70 (stress rule), 73 (SALS), 73 (CMA-ES window), 77 (grid-tuned window), 80 (written rule) and 92 (seasonal sharing) USD MWh^{-1} . The 95% cluster-bootstrap intervals over pairs are 56–61 for static arrays, 69–76 for SALS, 70–77 for the CMA-ES window and 75–84 for the written rule, and 46–50 for the plant (break-even prices of SALS and the CMA-ES window differ by −0.4, −1.1 to 0.4); they cover the sampling of pairs and not the cost assumptions. At 90 USD MWh^{-1} static arrays (6.3 kUSD) are better than the written (4.0) and grid-tuned (5.5) rules, close

922 to the CMA-ES window (7.3) and worse than SALS (7.6), which overtakes
 923 static arrays above about 85 USD MWh⁻¹; the difference between SALS and
 924 static arrays at 90 USD MWh⁻¹ is +1.3 kUSD (95% interval -0.04 to 2.7, 76
 925 pairs) and not distinguishable from zero, whereas SALS exceeds the CMA-
 926 ES window by 0.3 (-0.1 to 0.8, not distinguishable from zero), the grid-
 927 tuned window by 2.1 (1.6 to 2.7) and the written rule by 3.6 (2.8 to 4.5). A
 928 conventional plant that replaces the crop breaks even at 48 USD MWh⁻¹ and
 929 is more profitable than every agrivoltaic option above that price (16 kUSD
 930 at 60, 94 kUSD at 120, against -6 and 21 for SALS); below it all options
 931 lose money and static arrays lose least (at 30 USD MWh⁻¹: plant -23.2,
 932 SALS -18.7, static -5.7 kUSD). Per megawatt-hour, the electricity of SALS
 933 costs about 73 USD including the crop lost, against 48 USD for the plant, a
 934 difference of 25 USD MWh⁻¹. A premium of that size would equalise the unit
 935 cost but not the profit per hectare, because a hectare of plant produces three
 936 times as much electricity (1306 against 438 MWh): for the same electricity
 937 as one hectare of plant, SALS needs about three hectares of array, which
 938 costs about 25 USD MWh⁻¹ more including the crop lost but keeps about
 939 90% of the yield of that cropland; relative to the plant, SALS would need a
 940 premium of about 49 USD MWh⁻¹ at a market price of 60 and about 109 at
 941 90 to match its profit per hectare, and only about 13 at 60 (none above 73) to
 942 match the open field. The sensitivity is large (Table 17), and the ranking of
 943 static arrays and light sharing flips with the module cost (at 1.5 USD Wp⁻¹
 944 static arrays are better than SALS, at 0.7 SALS leads by 7 kUSD): at 90
 945 USD MWh⁻¹ the net value of SALS changes from +17 to -8 kUSD between
 946 module costs of 0.7 and 1.5 USD Wp⁻¹, whereas crop prices matter little.

947 Because the density was designed for the floor and not for profit, we also
 948 let each site choose, among the densities that meet the floor in its history, the
 949 one that maximises its historical net value (sites where none pays stay open
 950 field; Fig. 7). Then no method pays at 30-40 USD MWh⁻¹; at 60 14-59%
 951 of the pairs install an array (none for seasonal sharing) and the net values are
 952 small (static arrays 0.9, SALS 0.3 kUSD); static arrays are the best option up
 953 to about 75 USD MWh⁻¹, SALS overtakes them at about 75 (4.9 against 4.3
 954 at 80, 8.3 against 6.3 at 90, 20.7 against 12.3 at 120), the CMA-ES window
 955 at about 78, the stress rule at about 81, the grid-tuned window at about
 956 87 and the written rule at about 97. The value of control is therefore real
 957 but depends on a high price: at 90 USD MWh⁻¹ and with profit-optimal
 958 densities SALS exceeds the CMA-ES window by 0.4 kUSD ha⁻¹ yr⁻¹ (95%
 959 cluster-bootstrap interval 0.01 to 0.8), the grid-tuned window by 1.7 (1.2 to




Figure 7: Net value per hectare relative to the open field against electricity price. Left: densities designed for the floor, with a conventional plant replacing the crop (dashed). Right: each site chooses the feasible density that maximises its historical net value. Illustrative parameters of Table 15.

2.2), the written rule by 2.8 (2.0 to 3.6) and static arrays by 2.0 (0.8 to 3.3); the CMA-ES window exceeds static arrays by about 1.6 (0.4 to 2.8), so most of the value of control is obtained with a well-optimised rule.

Cost of the safety margin. At 90 USD MWh⁻¹, moving from a density designed under the mean response to a floor of 0.92 under that response costs 1.00 kUSD ha⁻¹ yr⁻¹ (17%) under E2 for 18 more points of compliance (70% to 88%) and 1.13 kUSD (20%) under E3 for 14 points (71% to 85%), about 56 and 79 USD ha⁻¹ yr⁻¹ per point. The ensemble rule that requires all other members to meet the floor costs 0.85 and 1.22 kUSD for 19 and 17 points (45 and 74 USD per point), so on this measure it is cheaper than the margin under E2 and similar under E3. At 60 USD MWh⁻¹ a margin has a negative cost, because every module beyond break-even loses money; the economics of the design margin therefore depend on the price level, whereas its effect on compliance does not.

3.10. Agrivoltaics under warming

Under warming, the open-field production of the held-out sites (Fig. 8) falls at +3 °C by 23% for maize and 18% for rice, is nearly unchanged for soybean, and rises for wheat (+3%) and potato (+10%). Agrivoltaic production with light-sharing control follows the open field at about 90% of its level, so it inherits the same climate impacts (maize -31%, rice -27% relative to the baseline open field at +3 °C). Agrivoltaics does not protect staple production




Figure 8: Production relative to the baseline open field (sum over held-out site-seasons) as a function of warming, for the open field, the stress rule and SALS.

981 against warming in this model. The LER of light-sharing systems increases
 982 slightly with warming (SALS 1.270 to 1.297; phenology rule 1.236 to 1.265),
 983 because the floor allows slightly denser arrays in warmer conditions (GCR
 984 0.173 to 0.177 and electricity 510 to 540 MWh ha⁻¹ yr⁻¹). The margin of
 985 SALS over the rules does not grow with warming (Table A.18).

986 3.11. Global potential and what it would take

987 Every site is evaluated when it is held out (222 of the 400 crop–site pairs
 988 are held out in at least one split, and three of them have no valid season,
 989 leaving 219), and each crop’s mean electricity per hectare is scaled to its
 990 physical area in SPAM 2010 (581 Mha for the five crops, which avoids count-
 991 ing multiple-cropped land twice). Equipping 5% of this area (29.1 Mha) with
 992 SALS arrays at the floor would produce 14.4 PWh yr⁻¹ under the baseline
 993 climate and 16.0 PWh yr⁻¹ at +3 °C, with a production-weighted retention
 994 of 0.903 and 0.900. The phenology rule would give 13.0 and 14.5 PWh yr⁻¹,
 995 and static arrays 6.4 and 6.9 PWh yr⁻¹. For comparison, world electricity
 996 generation was about 30 PWh in 2023 (Ember, 2024).

997 These numbers require large amounts of hardware. SALS at 5% cover-
 998 age needs 10.6 TWp of modules (11.3 TWp at +3 °C), at a capacity factor
 999 of 15.5% (16.2%). World solar PV capacity at the end of 2024 was about
 1000 1.9 TW (IRENA, 2025), so the requirement is about 5.6 times the installed
 1001 PV fleet. The electricity is also variable, peaks at midday and would be
 1002 curtailed or stored. The figures are technical potentials, not projections.
 1003 Ordering sites by electricity per hectare gives a supply curve (Fig. 9): the
 1004 best 2% of the staple cropland would supply about 21 PWh yr⁻¹ and the best
 1005 4.7% about 40 PWh yr⁻¹ with SALS. The best tenth of the sites are mostly




Figure 9: Supply curve under the baseline climate: cumulative electricity when sites are equipped in order of decreasing electricity per hectare, against the share of staple cropland (physical area) equipped. Each controller is evaluated at the density that keeps its mean retention at 0.9.

rainfed maize and soybean sites in the Americas and in southern and eastern Africa (7% mean irrigated fraction against 29% overall), which are also among the regions where the floor is hardest to meet (Section 3.5). Siting by electricity alone would therefore concentrate agrivoltaics where food-security risk is highest.

Maize and soybean, for which the simulated shade response is milder than the meta-analysis (Section 3.7), account for 53% of the electricity of the upscaling (a scaling convention of 5% of the cropland, reported without an interval) (maize 30%, soybean 23%, wheat 29%, rice 15%, potato 3%). Counting only wheat, rice and potato would give about 7 PWh yr⁻¹ at the same coverage, so the global figure should be read as an upper estimate for about half of its content. By crop, the LER of SALS under the baseline climate is 1.20 for wheat and rice, 1.31 for potato, 1.30 for maize and 1.41 for soybean.

4. Discussion

4.1. What the results establish

Four findings are discussed; the first three stand on consistent evidence and the fourth is not established. First, sharing light with the crop is the main lever on the electricity that tracker-based agrivoltaic arrays can produce




Figure 10: SALS electricity (a) and land-equivalent ratio (b) at the pooled held-out sites (baseline climate, mean of seasons 2015–2019).

under a yield floor. With static trackers the floor limits arrays to about 6% ground coverage and an LER of 1.07. Any light-sharing control allows 2.1 to 2.9 times the density, and the LER rises to 1.18–1.24 for hand-written rules. Second, the best simple rule follows crop phenology. It shares light from the onset of canopy development (5% of the thermal-time requirement) until maturity and produces electricity before and after, when the crop does not use the light. It needs only the crop’s thermal time, which is observable on the farm, and it is insensitive to forecast errors, whereas the rule that reacts to heat and drought stress is worse in every climate and degrades with forecast error. Third, the qualitative result does not depend on the weakest part of the model. The simulated shade response is conservative against the field trials we could use (Section 3.6); calibrating it raises the allowed density, and the phenology rule stays the best method, with a gain over static arrays of 0.17, 0.19 and 0.30 LER under the conservative, calibrated and field-fitted responses. However, the response learned from the meta-analysis shows steeper losses for maize and grain legumes than the simulator reproduces (Section 3.7), so the absolute gains are conservative for wheat, rice and potato and uncertain, possibly overestimated, for maize and soybean (53% of the upscaled electricity). Fourth, a controller learned through the differentiable simulator is better than the best rule we could write, by 0.044 in LER against the written rule but by only 0.007 (0.002 to 0.014; 2.4% more electricity at equal retention, and not distinguishable from zero at matched retention) against a window optimised on the training seasons of each split.

The margin over the written rule is consistent over training seeds (main split, range 0.002 in LER), the three spatial splits, three canopy-cooling assumptions and the baseline and +3°C climates, whereas the margin over the optimised window is small and varies between splits (0.008, 0.010 and 0.005 LER on splits 0, 1 and 2, mean of three training seeds; per-pair medians 0.001 to 0.002), so it is not established as a reproducible gain. This gain is small compared with the gap between static and dynamic control.

4.2. What role does artificial intelligence play?

We tested three ways in which learning could matter, and the answers differ. The first is learning the crop’s response to shade from data. The learned model is useful mainly as a measurement of how little can be learned: a hierarchical model fitted to 428 observations from 58 studies predicts a held-out study of the four staple crop types with an error of 0.22 in relative yield, no better than proportional loss, and the gain over proportional loss (0.43 to 0.31 over all types) comes from crop types that the simulator does not represent. Crop type is the information that transfers, whereas irradiation, climate zone and shade type add nothing measurable and a non-parametric model (gradient-boosted trees) is worse with this number of studies (Section 3.7). Two things follow. The shade response of the staple crops is a distribution with a spread of about 20 percentage points between studies and not a curve, and the process-based simulator does not reproduce the steep losses of maize and legumes seen in the data. Learning the response was limited by the data and not by the model: two studies support the tuber curve, no rice study exists, and 46 panel observations are all leafy vegetables.

The second is calibration through the differentiable simulator. Gradient descent through whole seasons fitted the response of each crop to the learned curves and to the field trials, and it produced an ensemble of plausible responses that fits the trials and the curves for cereals and potato but not the maize and soybean curves (Section 3.8). The ensemble informs the design decision: a density chosen under the mean response leaves about 70% of crop–site pairs at the floor if another plausible response holds, against 85–87% for an oracle, and a density chosen to meet the floor in 75% of plausible responses recovers about half of that gap at a cost of 0.03 LER. A simple margin reaches similar compliance: raising the floor from 0.90 to 0.92 under the mean response reaches 85–88% of the pairs, and interpolating the margin frontier at the compliance of the ensemble rules gives LER differences (margin minus rule) of +0.001 (−0.009 to 0.012) and +0.003 (−0.016 to 0.021)

under E2 and -0.006 (-0.019 to 0.005) and $+0.024$ (-0.011 to 0.058) under E3, so the two are not distinguishable on this sample (equivalence is not established, since the intervals reach ± 0.02 to 0.06 LER). None of the rules is safe outside the ensemble (7–33% of the pairs at the floor under the proportional, calibrated and harsh responses for the mean design, the margin of 0.92 and the ensemble rules), and only a floor of 0.96 reaches 51–78% at a cost of about 0.2 LER. What the ensemble adds is the size of the margin, not a better rule. The quantile rule is a scenario-type design (Calafiore and Campi, 2006; Campi and Garatti, 2008); with five members it has none of the finite-sample guarantees of that theory, and since the comparison with the margin shows no detectable difference (with intervals too wide to establish equivalence), we see no evidence that it is needed. The bottleneck of tracker-based agrivoltaics under a yield-retention floor is the uncertainty of the crop response that enters the design, and it is data on that response, not larger networks, that would improve the result; the learned curves show a spread of about 20 percentage points between studies. The guarantee holds only within the range of responses that the ensemble spans.

The third is control, and here learning contributes little. Differentiability is what lets a network of 5441 weights be trained in 300 gradient steps through whole seasons; a derivative-free search such as CMA-ES handled a five-parameter schedule with 200 simulator evaluations per run and found an almost binary window that is only 0.007 LER behind SALS, but would not scale to the network. The learned controller adds 0.007 LER over an optimised window (and 0.017 over a grid-tuned one, 0.044 over the written rule), and 0.005–0.010 over the optimised window under the four alternative responses on one split (21 pairs, intervals containing zero); a reader who needs a simple, auditable controller can use that window and obtain about 84% of the LER gain of SALS over the written rule. Neither of the two extensions that we expected to matter under uncertainty did: training on ensembles of responses (E1, E2 and E3) did not help and, for the wider ensembles, did not retain the margin over the written rule within our training budget (one seed, 21 pairs), and neither did giving the controller a farm-measurable signal, the biomass of the crop relative to a reference plot, nor a feedback rule built on that signal. The learned policy also has to be initialised from the phenology rule (Section 4.5). We therefore do not claim that learned control is what makes agrivoltaics work: in this model, the schedule that a learned policy converges to is a smooth function of crop phenology that an optimised window already approximates to within 0.007 LER, and the source of the

1123 remaining margin is not identified.

1124 *4.3. What the results do not establish*

1125 The early hypothesis that stress-aware control, that is, shading the crop
1126 on heat and drought days, would be a main source of benefit under warming
1127 was not supported. The stress rule performs worse than the phenology rule
1128 in every climate, and a learned controller without stress inputs performs as
1129 well as one with them. A plausible reason is that in our crop model the yield
1130 gain from shading on a hot or dry day is smaller than the loss of light on
1131 that day, so stress information does not help. Whether this holds for crop
1132 models with more detailed canopy-temperature and heat-stress responses is
1133 an open question. Agrivoltaics does not protect staple production against
1134 warming in our simulations: light-sharing systems produce about 90% of the
1135 open-field yield in every climate and follow the same decline in maize and
1136 rice (Fig. 8).

1137 *4.4. Food security in practice*

1138 Holding the historical mean retention at 90% does not hold it in each
1139 season or at each site. 70% of SALS seasons meet the floor, the lower tail is
1140 slightly heavier than for the rules, and production-weighted retention is below
1141 0.9 for soybean (0.888) and maize (0.898) and in several regions. Chance-
1142 constrained design raises season-level compliance to about 92% at a cost of
1143 0.065 in LER, and the ranking of the controllers is preserved. Whichever for-
1144 mulation is chosen, a per-site regulatory floor would need a region- and crop-
1145 specific design step. The upscaled potential (Section 3.11) shows the trade-off
1146 in magnitude. The 14 PWh yr⁻¹ that 5% of the staple cropland could supply
1147 need about 10 TWp of modules, several times the present global fleet, and
1148 the best sites are rainfed maize and soybean areas where production-weighted
1149 retention is lowest.

1150 **Who bears the risk.** A floor defined on the multi-year mean leaves
1151 the farmer exposed in the seasons below it: with the density designed for
1152 the floor, about 30% of seasons fall short on the main split (70% compliance
1153 for SALS), the shortfall is mostly small (2.7% of SALS seasons are below
1154 0.8 of the open-field yield and the 10th percentile of retention is 0.883), and
1155 a per-season guarantee through chance-constrained design costs 0.065 LER.
1156 Siting by electricity favours maize and soybean areas, where compliance is
1157 lowest (Section 3.11), so the risk concentrates where the energy gain is largest.
1158 The crop revenue at stake (about 0.3 kUSD ha⁻¹ yr⁻¹) is small against 20–30

kUSD of electricity revenue, so the exposure is mainly one of compliance with the floor and of the contract, and it coincides with the largest uncertainty of the shade response, that of maize and soybean. Who carries it, through contracts, insurance or a premium, is a design question that the model does not address, and food is valued only through crop revenue at market prices.

4.5. Limitations

- Simulation only.* All results come from a simulator. PV output was benchmarked against PVGIS (29 sites, mean bias -4.9%), but the crop response to shade was learned from one meta-analysis data set and compared with four field trials (Sections 3.6 and 3.7). The simulated response is conservative with respect to the field trials of rice, wheat and potato; for maize it is conservative with respect to the trials but milder than the meta-analysis, and for soybean the simulator (0.48 at 40% shading) lies within the interval of the data (0.40, 80% interval 0.34–0.53), while for maize the data give 0.42 against 0.75 in the simulator, so for these crops the direction of the bias depends on the source. The harvest-index sensitivities h are assumptions, and the sensitivity analyses use subsamples of crop–site pairs (22 for the rule comparison under alternative responses, 21 for the controller comparison, one training seed).
- Canopy cooling.* The effect of shade on canopy temperature was represented by one coefficient κ , applied to a heat-stress function that was calibrated on air temperature (Siebert et al., 2014; Webber et al., 2018). The ranking of the controllers is robust to κ between 0 and 2°C , but the absolute electricity changes by about $\pm 8\%$. Nocturnal warming under modules and rain redistribution are not modelled.
- Crop model.* SIMPLE has a fixed harvest index and linear radiation-use efficiency, and its CO_2 response is that of controlled experiments that tend to overstate field responses (Long et al., 2006). More detailed models with grain-number formation and canopy energy balance may favour different schedules.
- Layouts and models.* The baseline of static arrays is a backtracking tracker held at the energy-optimal angle; flat and 30° -rotated (east-facing, not equator-facing) fixed arrays did no better in a check on 21 pairs, but vertical bifacial layouts (Willockx et al., 2023; Campana

et al., 2021) are not modelled, and the ground coverage ratio is a coarse proxy of crop productivity (Dupraz, 2023). The PV model is isotropic in the diffuse component and has no bifacial gain; it was benchmarked against PVGIS at low density only and has not been compared with a reference implementation such as pvlib (F. Holmgren et al., 2018). The rotation limit was varied between 45 and 80 degrees and changes the size of the benefit but not the ranking; the harvest-index window is assumed to be common to all crops, and shifting it or removing the penalty did not change the ranking of the rules.

- *Climate scenarios.* Warming was applied as uniform temperature shifts with fixed precipitation, radiation and relative humidity, and with unchanged cultivars and sowing dates. It is not a bias-adjusted climate projection.
- *Initialisation and state information.* The network is initialised from the phenology rule, because training from random weights collapsed to full sharing. The gain of SALS is therefore the gain of fine-tuning that rule, and we cannot say whether a better optimiser or a richer policy class would find larger gains. The controller uses the simulated state of the crop and soil (thermal time, canopy interception, soil water) without measurement error, and the forecast errors we tested apply to weather only.
- *Economics and deployment.* The economic analysis (Section 3.9) uses assumed prices and costs, no grid, curtailment, storage or subsidy model, and does not value food security; its break-even prices move with the capex and the discount rate (Fig. 7). The value of midday electricity, curtailment, storage, structural cost and farmer acceptance (Trommsdorff et al., 2021) will determine the realisable share. The average row pitch at the chosen densities (about 13 m for 2.2-m modules) leaves room for machinery, but the densest designs (GCR up to 0.5) have a pitch of only 4.4 m.

5. Conclusions

Light sharing, the daily choice between maximum electricity and letting light through to the crop, is the main lever on the electricity of tracker-based agrivoltaics at a given yield-retention floor. In a differentiable simulator,

1228 evaluated on held-out crop–site pairs (92 on the main split, 76 distinct pairs
1229 pooled over three splits for the controller and design comparisons) with the
1230 density of each site chosen from its own history, light-sharing rules allowed
1231 2.1 to 2.9 times the module density of static arrays at a mean yield retention
1232 of 90% and raised the land-equivalent ratio from 1.07 to 1.18–1.24 under
1233 the proportional-growth shade response. A rule that follows crop phenology
1234 and needs no weather forecast was the best written rule. The lever pays
1235 only at high electricity prices: with illustrative costs it breaks even with
1236 the open field, at floor-designed densities, at 70 to 92 USD MWh^{−1} (73 for
1237 the learned controller, 59 for static arrays), static arrays earn more below
1238 about 75 to 85 USD MWh^{−1}, and a conventional plant replacing the crop
1239 is more profitable above its break-even of 48 USD MWh^{−1}. In this model
1240 agrivoltaics therefore pays only with a premium, land-use rules or a value for
1241 food production that the model does not include.

1242 The crop’s response to shade, the least certain input, was learned from 428
1243 field observations. For the staple crop types a hierarchical model predicted a
1244 held-out study no better than proportional loss (error 0.22), which measures
1245 the between-study spread, and gradient descent through the simulator turned
1246 the learned curves and four field trials into ensembles of plausible responses.
1247 The simulator is conservative against the trials for rice, wheat and potato; for
1248 maize and soybean, which supply 53% of the upscaled electricity, the meta-
1249 analysis indicates larger losses than simulated and one maize trial smaller
1250 ones, so the gains for these crops are uncertain and possibly overestimated.
1251 A learned controller exceeded the written phenology rule by 0.044 LER but a
1252 CMA-ES-optimised window by only 0.007 (95% interval 0.002–0.014, mean of
1253 three training seeds per split; at matched retention 0.003, interval −0.0004
1254 to 0.007). The source of this margin is not identified: it is not explained
1255 by the tuning of a window or a ramp, and masking soil water, aridity and
1256 forecasts leaves the controller unchanged. Training on ensembles of shade
1257 responses, or giving the controller a farm-measurable biomass signal, did not
1258 help within our budget (one seed, 21 pairs).

1259 Compliance depends more on the response assumed in the design than
1260 on the controller. A density chosen under the mean response left about 70%
1261 of crop–site pairs at the floor when another plausible response held, against
1262 85–87% for an oracle design; this gap of 15 points (a model-misspecification
1263 effect) is larger than the 4 points between light-sharing controllers at a fixed
1264 design rule under a known response, but the two are not like-for-like. De-
1265 signing for the floor in 75% of the plausible responses recovered about half

1266 of the gap at a cost of 0.03 LER, and a floor raised by 0.02 under the mean
1267 response gave no detectable difference at equal compliance (the intervals in-
1268 clude zero and are too wide to establish equivalence); no rule was safe outside
1269 the ensemble, whose members differ less from one another than a new study
1270 differs from the learned curve. The ensemble is therefore useful to size a
1271 safety margin, not to replace it. The yield floor is met on average but not in
1272 every season or region, and the electricity potential of 5% of staple cropland
1273 (about 14 PWh yr⁻¹, a scaling convention without an interval) requires about
1274 10 TWp of modules. Future work should validate the crop response to shade
1275 under trackers with more field data, notably for maize and legumes where the
1276 meta-analysis and the trials disagree, extend the crop model to grain-number
1277 formation and canopy energy balance, add costs and grid constraints, and
1278 pursue controllers that guarantee the yield-retention floor per site and per
1279 season.

1280 **CRedit authorship contribution statement**

1281 **Quang-Vinh Dang:** Conceptualization, Methodology, Software, Formal
1282 analysis, Investigation, Data curation, Visualization, Writing – original draft,
1283 Writing – review & editing.

1284 **Declaration of competing interest**

1285 The author has no known competing financial interests or personal re-
1286 lationships that could have appeared to influence the work reported in this
1287 paper.

1288 **Data and code availability**

1289 All data used are public: SPAM 2010 v2r0 harvested areas (Yu et al.,
1290 2020), the GGCM Phase 3 crop calendar (Jägermeyr et al., 2021), NASA
1291 POWER daily weather (Sparks, 2018) the SIMPLE crop parameters (Zhao
1292 et al., 2019) and the shade-response meta-analysis data of Laub et al. (2022)
1293 (Zenodo 10.5281/zenodo.5716091, CC BY 4.0, copied into the repository).
1294 The code of the simulator, of SALS and of all experiments, together with the
1295 site list, is available from the author’s public repository (<https://github.com/vinhqdang/Artificial-Intelligence-in-Agriculture>); a versioned
1296 archive with a DOI will be deposited at acceptance.

1298 **Funding**

1299 This research did not receive any specific grant from funding agencies in
1300 the public, commercial, or not-for-profit sectors.

1301 **Appendix A. Statistical comparison**

1302 **References**

- 1303 Adeh, E.H., Good, S.P., Calaf, M., Higgins, C.W., 2019. Solar pv power
1304 potential is greatest over croplands. *Scientific Reports* 9, 11442. doi:10.
1305 1038/s41598-019-47803-3.
- 1306 Allen, R.G., Pereira, L.S., Raes, D., Smith, M., 1998. Crop evapotranspira-
1307 tion: Guidelines for computing crop water requirements. *FAO Irrigation*
1308 *and Drainage Paper* 56. Food and Agriculture Organization of the United
1309 Nations. Rome.
- 1310 Amaducci, S., Yin, X., Colauzzi, M., 2018. Agrivoltaic systems to optimise
1311 land use for electric energy production. *Applied Energy* 220, 545–561.
1312 doi:10.1016/j.apenergy.2018.03.081.
- 1313 Andrade, F.H., Vega, C., Uhart, S., Cirilo, A., Cantarero, M., Valentinuz, O.,
1314 1999. Kernel number determination in maize. *Crop Science* 39, 453–459.
1315 doi:10.2135/cropsci1999.0011183X0039000200026x.
- 1316 Barron-Gafford, G.A., Pavao-Zuckerman, M.A., Minor, R.L., Sutter, L.F.,
1317 Barnett-Moreno, I., Blackett, D.T., Thompson, M., Dimond, K., Gerlak,
1318 A.K., Nabhan, G.P., Macknick, J.E., 2019. Agrivoltaics provide mutual
1319 benefits across the food–energy–water nexus in drylands. *Nature Sustain-*
1320 *ability* 2, 848–855. doi:10.1038/s41893-019-0364-5.
- 1321 Calafiore, G.C., Campi, M.C., 2006. The scenario approach to robust control
1322 design. *IEEE Transactions on Automatic Control* 51, 742–753. doi:10.
1323 1109/TAC.2006.875041.
- 1324 Campana, P.E., Stridh, B., Amaducci, S., Colauzzi, M., 2021. Optimisation
1325 of vertically mounted agrivoltaic systems. *Journal of Cleaner Production*
1326 325, 129091. doi:10.1016/j.jclepro.2021.129091.

1327 Campi, M.C., Garatti, S., 2008. The exact feasibility of randomized solutions
 1328 of uncertain convex programs. *SIAM Journal on Optimization* 19, 1211–
 1329 1230. doi:10.1137/07069821X.

1330 Collares-Pereira, M., Rabl, A., 1979. The average distribution of solar
 1331 radiation-correlations between diffuse and hemispherical and between daily
 1332 and hourly insolation values. *Solar Energy* 22, 155–164. doi:10.1016/
 1333 0038-092x(79)90100-2.

1334 DIN, 2021. DIN SPEC 91434:2021-05: Agri-photovoltaic systems – require-
 1335 ments for primary agricultural use. Beuth Verlag, Berlin.

1336 Dinesh, H., Pearce, J.M., 2016. The potential of agrivoltaic systems. *Renew-
 1337 able and Sustainable Energy Reviews* 54, 299–308. doi:10.1016/j.rser.
 1338 2015.10.024.

1339 Dupraz, C., 2023. Assessment of the ground coverage ratio of agrivoltaic
 1340 systems as a proxy for potential crop productivity. *Agroforestry Systems*
 1341 98, 2679–2696. doi:10.1007/s10457-023-00906-3.

1342 Dupraz, C., Marrou, H., Talbot, G., Dufour, L., Nogier, A., Ferard, Y., 2011.
 1343 Combining solar photovoltaic panels and food crops for optimising land
 1344 use: Towards new agrivoltaic schemes. *Renewable Energy* 36, 2725–2732.
 1345 doi:10.1016/j.renene.2011.03.005.

1346 Elamri, Y., Cheviron, B., Lopez, J.M., Dejean, C., Belaud, G., 2018. Water
 1347 budget and crop modelling for agrivoltaic systems: Application to irrigated
 1348 lettuces. *Agricultural Water Management* 208, 440–453. doi:10.1016/j.
 1349 agwat.2018.07.001.

1350 Ember, 2024. Global Electricity Review 2024. Technical Report. Ember.
 1351 London.

1352 European Commission, Joint Research Centre, 2022. Photovoltaic geo-
 1353 graphical information system (PVGIS), version 5.2. [https://re.jrc.ec.
 1354 europa.eu/pvg_tools/](https://re.jrc.ec.europa.eu/pvg_tools/).

1355 F. Holmgren, W., W. Hansen, C., A. Mikofski, M., 2018. pvlib python: a
 1356 python package for modeling solar energy systems. *Journal of Open Source
 1357 Software* 3, 884. doi:10.21105/joss.00884.

1358 Fischer, R.A., 1985. Number of kernels in wheat crops and the influence of
1359 solar radiation and temperature. *The Journal of Agricultural Science* 105,
1360 447–461. doi:10.1017/S0021859600056495.

1361 Gautron, R., Maillard, O.A., Preux, P., Corbeels, M., Sabbadin, R., 2022.
1362 Reinforcement learning for crop management support: Review, prospects
1363 and challenges. *Computers and Electronics in Agriculture* 200, 107182.
1364 doi:10.1016/j.compag.2022.107182.

1365 Gonocruz, R.A., Nakamura, R., Yoshino, K., Homma, M., Doi, T., Yoshida,
1366 Y., Tani, A., 2021. Analysis of the rice yield under an agrivoltaic
1367 system: A case study in japan. *Environments* 8, 65. doi:10.3390/
1368 environments8070065.

1369 Hansen, N., Ostermeier, A., 2001. Completely derandomized self-adaptation
1370 in evolution strategies. *Evolutionary Computation* 9, 159–195. doi:10.
1371 1162/106365601750190398.

1372 Hassanpour Adeg, E., Selker, J.S., Higgins, C.W., 2018. Remarkable agri-
1373 voltaic influence on soil moisture, micrometeorology and water-use effi-
1374 ciency. *PLOS ONE* 13, e0203256. doi:10.1371/journal.pone.0203256.

1375 IPCC, 2021. *Climate Change 2021: The Physical Science Basis. Contribution*
1376 *of Working Group I to the Sixth Assessment Report of the Intergovern-*
1377 *mental Panel on Climate Change.* Cambridge University Press, Cambridge.
1378 doi:10.1017/9781009157896.

1379 IRENA, 2025. *Renewable Capacity Statistics 2025.* Technical Report. Inter-
1380 national Renewable Energy Agency. Abu Dhabi.

1381 Jägermeyr, J., Müller, C., Ruane, A.C., Elliott, J., Balkovic, J., Castillo, O.,
1382 Faye, B., Foster, I., Folberth, C., Franke, J.A., Fuchs, K., Guarin, J.R.,
1383 Heinke, J., Hoogenboom, G., Iizumi, T., Jain, A.K., Kelly, D., Khabarov,
1384 N., Lange, S., Lin, T.S., Liu, W., Mialyk, O., Minoli, S., Moyer, E.J.,
1385 Okada, M., Phillips, M., Porter, C., Rabin, S.S., Scheer, C., Schneider,
1386 J.M., Schyns, J.F., Skalsky, R., Smerald, A., Stella, T., Stephens, H.,
1387 Webber, H., Zabel, F., Rosenzweig, C., 2021. Climate impacts on global
1388 agriculture emerge earlier in new generation of climate and crop models.
1389 *Nature Food* 2, 873–885. doi:10.1038/s43016-021-00400-y.

- 1390 Karniadakis, G.E., Kevrekidis, I.G., Lu, L., Perdikaris, P., Wang, S., Yang,
1391 L., 2021. Physics-informed machine learning. *Nature Reviews Physics* 3,
1392 422–440. doi:10.1038/s42254-021-00314-5.
- 1393 Kingma, D.P., Ba, J., 2015. Adam: A method for stochastic optimization,
1394 in: *International Conference on Learning Representations*.
- 1395 Laub, M., Pataczek, L., Feuerbacher, A., Zikeli, S., Högy, P., 2022. Contrast-
1396 ing yield responses at varying levels of shade suggest different suitability
1397 of crops for dual land-use systems: a meta-analysis. *Agronomy for Sus-
1398 tainable Development* 42, 51. doi:10.1007/s13593-022-00783-7.
- 1399 Liu, B.Y., Jordan, R.C., 1960. The interrelationship and characteristic dis-
1400 tribution of direct, diffuse and total solar radiation. *Solar Energy* 4, 1–19.
1401 doi:10.1016/0038-092x(60)90062-1.
- 1402 Long, S.P., Ainsworth, E.A., Leakey, A.D.B., Nösberger, J., Ort, D.R., 2006.
1403 Food for thought: Lower-than-expected crop yield stimulation with ris-
1404 ing CO_2 concentrations. *Science* 312, 1918–1921. doi:10.1126/science.
1405 1114722.
- 1406 Lorenzo, E., Narvarte, L., Muñoz, J., 2011. Tracking and back-tracking.
1407 *Progress in Photovoltaics: Research and Applications* 19, 747–753. doi:10.
1408 1002/pip.1085.
- 1409 Mamun, M.A.A., Dargusch, P., Wadley, D., Zulkarnain, N.A., Aziz, A.A.,
1410 2022. A review of research on agrivoltaic systems. *Renewable and Sustain-
1411 able Energy Reviews* 161, 112351. doi:10.1016/j.rser.2022.112351.
- 1412 Marchau, V.A.W.J., Walker, W.E., Bloemen, P.J.T.M., Popper, S.W. (Eds.),
1413 2019. *Decision Making under Deep Uncertainty: From Theory to Practice*.
1414 Springer. doi:10.1007/978-3-030-05252-2.
- 1415 Marrou, H., Dufour, L., Wery, J., 2013a. How does a shelter of solar panels
1416 influence water flows in a soil–crop system? *European Journal of Agron-
1417 omy* 50, 38–51. doi:10.1016/j.eja.2013.05.004.
- 1418 Marrou, H., Guillioni, L., Dufour, L., Dupraz, C., Wery, J., 2013b. Microcli-
1419 mate under agrivoltaic systems: Is crop growth rate affected in the partial
1420 shade of solar panels? *Agricultural and Forest Meteorology* 177, 117–132.
1421 doi:10.1016/j.agrformet.2013.04.012.

- 1422 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen,
1423 T., Lin, Z., Gimelshein, N., Antiga, L., et al., 2019. PyTorch: An impera-
1424 tive style, high-performance deep learning library, in: Advances in Neural
1425 Information Processing Systems, pp. 8024–8035.
- 1426 Ramos-Fuentes, I.A., Elamri, Y., Cheviron, B., Dejean, C., Belaud, G.,
1427 Fumey, D., 2023. Effects of shade and deficit irrigation on maize growth
1428 and development in fixed and dynamic agrivoltaic systems. *Agricultural*
1429 *Water Management* 280, 108187. doi:10.1016/j.agwat.2023.108187.
- 1430 Reichstein, M., Camps-Valls, G., Stevens, B., Jung, M., Denzler, J., Car-
1431 valhais, N., Prabhat, 2019. Deep learning and process understanding for
1432 data-driven Earth system science. *Nature* 566, 195–204. doi:10.1038/
1433 s41586-019-0912-1.
- 1434 Sekiyama, T., Nagashima, A., 2019. Solar sharing for both food and
1435 clean energy production: Performance of agrivoltaic systems for corn,
1436 a typical shade-intolerant crop. *Environments* 6, 65. doi:10.3390/
1437 environments6060065.
- 1438 Shen, C., Appling, A.P., Gentine, P., Bandai, T., Gupta, H., Tartakovsky, A.,
1439 Baity-Jesi, M., Fenicia, F., Kifer, D., Li, L., Liu, X., Ren, W., Zheng, Y.,
1440 Harman, C.J., Clark, M., Farthing, M., Feng, D., Kumar, P., Aboelyazeed,
1441 D., Rahmani, F., Song, Y., Beck, H.E., Bindas, T., Dwivedi, D., Fang,
1442 K., Höge, M., Rackauckas, C., Mohanty, B., Roy, T., Xu, C., Lawson,
1443 K., 2023. Differentiable modelling to unify machine learning and physical
1444 models for geosciences. *Nature Reviews Earth & Environment* 4, 552–567.
1445 doi:10.1038/s43017-023-00450-9.
- 1446 Siebert, S., Ewert, F., Eyshi Rezaei, E., Kage, H., Graß, R., 2014. Impact
1447 of heat stress on crop yield—on the importance of considering canopy
1448 temperature. *Environmental Research Letters* 9, 044012. doi:10.1088/
1449 1748-9326/9/4/044012.
- 1450 Sparks, A., 2018. nasapower: A nasa power global meteorology, surface solar
1451 energy and climatology data client for r. *Journal of Open Source Software*
1452 3, 1035. doi:10.21105/joss.01035.
- 1453 Trommsdorff, M., Kang, J., Reise, C., Schindele, S., Bopp, G., Ehmann, A.,
1454 Weselek, A., Högy, P., Obergfell, T., 2021. Combining food and energy

1455 production: Design of an agrivoltaic system applied in arable and vegetable
 1456 farming in germany. *Renewable and Sustainable Energy Reviews* 140,
 1457 110694. doi:10.1016/j.rser.2020.110694.

1458 Valle, B., Simonneau, T., Sourd, F., Pechier, P., Hamard, P., Frisson, T.,
 1459 Ryckewaert, M., Christophe, A., 2017. Increasing the total productivity of
 1460 a land by combining mobile photovoltaic panels and food crops. *Applied*
 1461 *Energy* 206, 1495–1507. doi:10.1016/j.apenergy.2017.09.113.

1462 Waha, K., Dietrich, J.P., Portmann, F.T., Siebert, S., Thornton, P.K., Bon-
 1463 deau, A., Herrero, M., 2020. Multiple cropping systems of the world
 1464 and the potential for increasing cropping intensity. *Global Environmental*
 1465 *Change* 64, 102131. doi:10.1016/j.gloenvcha.2020.102131.

1466 Webber, H., White, J.W., Kimball, B.A., Ewert, F., Asseng, S., Eyshi Rezaei,
 1467 E., Pinter, P.J., Hatfield, J.L., Reynolds, M.P., Ababaei, B., Bindi, M.,
 1468 Doltra, J., Ferrise, R., Kage, H., Kassie, B.T., Kersebaum, K.C., Luig,
 1469 A., Olesen, J.E., Semenov, M.A., Stratonovitch, P., Ratjen, A.M., LaM-
 1470 orte, R.L., Leavitt, S.W., Hunsaker, D.J., Wall, G.W., Martre, P., 2018.
 1471 Physical robustness of canopy temperature models for crop heat stress
 1472 simulation across environments and production conditions. *Field Crops*
 1473 *Research* 216, 75–88. doi:10.1016/j.fcr.2017.11.005.

1474 Weselek, A., Bauerle, A., Hartung, J., Zikeli, S., Lewandowski, I., Högy, P.,
 1475 2021. Agrivoltaic system impacts on microclimate and yield of different
 1476 crops within an organic crop rotation in a temperate climate. *Agronomy*
 1477 *for Sustainable Development* 41, 59. doi:10.1007/s13593-021-00714-y.

1478 Weselek, A., Ehmann, A., Zikeli, S., Lewandowski, I., Schindele, S., Högy,
 1479 P., 2019. Agrophotovoltaic systems: applications, challenges, and op-
 1480 portunities. a review. *Agronomy for Sustainable Development* 39, 35.
 1481 doi:10.1007/s13593-019-0581-3.

1482 Willockx, B., Lavaert, C., Cappelle, J., 2023. Performance evaluation of
 1483 vertical bifacial and single-axis tracked agrivoltaic systems on arable land.
 1484 *Renewable Energy* 217, 119181. doi:10.1016/j.renene.2023.119181.

1485 Woli, P., Jones, J.W., Ingram, K.T., Fraisse, C.W., 2012. Agricultural
 1486 reference index for drought (arid). *Agronomy Journal* 104, 287–300.
 1487 doi:10.2134/agronj2011.0286.

- 1488 Yu, Q., You, L., Wood-Sichra, U., Ru, Y., Joglekar, A.K.B., Fritz, S., Xiong,
1489 W., Lu, M., Wu, W., Yang, P., 2020. A cultivated planet in 2010 – part 2:
1490 The global gridded agricultural-production maps. *Earth System Science*
1491 *Data* 12, 3545–3572. doi:10.5194/essd-12-3545-2020.
- 1492 Zhang, L., Geng, X., Ding, H., Cao, K., Wang, L., Chen, H., Deng, L., Wu,
1493 C., Xiao, M., Bao, E., 2026. Multi-objective optimization of photovoltaic
1494 array layouts on farmland via a hybrid algorithm framework for enhanced
1495 agricultural production and energy generation. *Artificial Intelligence in*
1496 *Agriculture* doi:10.1016/j.aiia.2026.09.002.
- 1497 Zhao, C., Liu, B., Piao, S., Wang, X., Lobell, D.B., Huang, Y., Huang, M.,
1498 Yao, Y., Bassu, S., Ciais, P., Durand, J.L., Elliott, J., Ewert, F., Janssens,
1499 I.A., Li, T., Lin, E., Liu, Q., Martre, P., Müller, C., Peng, S., Peñuelas,
1500 J., Ruane, A.C., Wallach, D., Wang, T., Wu, D., Liu, Z., Zhu, Y., Zhu,
1501 Z., Asseng, S., 2017. Temperature increase reduces global yields of major
1502 crops in four independent estimates. *Proceedings of the National Academy*
1503 *of Sciences* 114, 9326–9331. doi:10.1073/pnas.1701762114.
- 1504 Zhao, C., Liu, B., Xiao, L., Hoogenboom, G., Boote, K.J., Kassie, B.T.,
1505 Pavan, W., Shelia, V., Kim, K.S., Hernandez-Ochoa, I.M., Wallach, D.,
1506 Porter, C.H., Stockle, C.O., Zhu, Y., Asseng, S., 2019. A simple crop
1507 model. *European Journal of Agronomy* 104, 97–106. doi:10.1016/j.eja.
1508 2019.01.009.

Table 2: Held-out performance (92 test sites, seasons 2015–2019, canopy cooling $\kappa = 1^\circ\text{C}$). The density (GCR) of each site is the largest value on the grid 0.05, 0.075, ..., 0.5 at which the method’s mean retention in the site’s 2001–2014 seasons is ≥ 0.9 . E : electricity ($\text{MWh ha}^{-1} \text{ yr}^{-1}$); e : electricity relative to a conventional PV plant; r : mean yield retention; r_{prod} : production-weighted retention; C_s and C_{site} : share of seasons with $r \geq 0.9$ and share of sites whose test-season mean $r \geq 0.9$ (%); $\text{LER} = r + e$.

Method	GCR	E	e	r	r_{prod}	C_s	C_{site}	LER
<i>Baseline</i>								
AV static	0.06	217	0.16	0.917	0.914	93	96	1.07
Seasonal sharing	0.17	397	0.28	0.904	0.902	69	72	1.19
Phenology rule	0.17	465	0.33	0.904	0.902	69	75	1.24
Stress rule	0.13	387	0.28	0.907	0.907	76	82	1.18
SALS	0.17	510	0.36	0.905	0.902	70	74	1.27
<i>+1.5 °C</i>								
AV static	0.06	221	0.16	0.916	0.914	93	92	1.08
Seasonal sharing	0.16	393	0.28	0.905	0.903	72	74	1.19
Phenology rule	0.17	478	0.34	0.905	0.902	68	75	1.25
Stress rule	0.12	377	0.27	0.908	0.908	73	77	1.18
SALS	0.17	518	0.37	0.906	0.903	68	75	1.28
<i>+2 °C</i>								
AV static	0.06	223	0.16	0.916	0.914	92	93	1.08
Seasonal sharing	0.17	394	0.28	0.905	0.904	72	76	1.19
Phenology rule	0.17	483	0.35	0.905	0.903	69	74	1.25
Stress rule	0.12	370	0.26	0.909	0.909	75	78	1.17
SALS	0.17	521	0.38	0.907	0.904	69	75	1.28
<i>+3 °C</i>								
AV static	0.06	231	0.17	0.915	0.914	88	89	1.08
Seasonal sharing	0.17	397	0.29	0.906	0.904	69	75	1.19
Phenology rule	0.18	499	0.36	0.905	0.902	67	72	1.26
Stress rule	0.11	366	0.26	0.908	0.908	74	77	1.17
SALS	0.18	540	0.39	0.907	0.904	68	75	1.30

Table 3: LER of the controllers at different yield-retention floors ρ (main split; the historical design step is repeated for each floor).

Method	Baseline climate				+3 °C			
	0.80	0.85	0.90	0.95	0.80	0.85	0.90	0.95
AV static	1.159	1.115	1.072	1.011	1.174	1.128	1.083	1.013
Stress rule	1.380	1.282	1.182	1.077	1.364	1.269	1.172	1.071
Phenology rule	1.467	1.356	1.236	1.108	1.513	1.395	1.265	1.122
SALS (trained at $\rho = 0.9$)	1.505	1.392	1.270	1.141	1.553	1.432	1.297	1.154
SALS conditioned on ρ	1.504	1.391	1.269	1.141	1.553	1.431	1.296	1.154

Table 4: Sensitivity of the controller comparison to the rotation limit of the trackers and to the harvest-index penalty (LER; main split, every sixth held-out pair, 21 pairs, main shade response, history-based design on a GCR grid with step 0.05). SALS was trained for the main setting. Fixed arrays (no tracking; flat and tilted by 30° from the horizontal about the north–south axis) reach LER 1.048 and 1.043 in the same sample, against 1.053 for static trackers.

Setting	Static tracker	Phenology rule	Grid-tuned window	SALS
Main (limit 60°, window 0.45–0.65)	1.053	1.170	1.180	1.189
Rotation limit 45°	1.057	1.138	1.144	1.144
Rotation limit 80°	1.053	1.206	1.219	1.224
Window 0.30–0.50	1.053	1.169	1.176	1.185
Window 0.60–0.80	1.053	1.171	1.181	1.182
No harvest-index penalty	1.109	1.328	1.347	1.353

Table 5: Rules and the learned controller (SALS: mean of three training seeds per split) pooled over the three spatial splits (76 distinct crop–site pairs, every third held-out pair; baseline climate; main shade response; history-based design on a GCR grid with step 0.05, coarser than the 0.025 grid of the main analysis). The grid-tuned window shares light while thermal time lies between 0.15 and 1.0 of the season requirement (best of a 4×3 grid on the training-block seasons of each split, the same on all three; the written rule uses 0.05 to 1.0). The CMA-ES schedule has five parameters (start, end, two ramp widths, level) was tuned on each split’s training block and stopped, at the generation cap, at an almost binary window from 0.20 to 0.93 of the requirement at level 1.00 (see Methods).

Method	GCR	E/E_{PV}	Retention	Seasons ≥ 0.9 (%)	Pairs ≥ 0.9 (%)	LER
AV static	0.059	0.153	0.922	95	96	1.076
Seasonal sharing	0.151	0.249	0.915	84	91	1.164
Stress rule	0.118	0.255	0.918	83	88	1.173
Phenology rule (written)	0.154	0.292	0.916	85	91	1.207
Grid-tuned window	0.161	0.317	0.916	84	89	1.233
CMA-ES optimised window	0.158	0.327	0.916	83	87	1.244
SALS	0.161	0.335	0.916	83	87	1.251

Table 6: Simulated relative yield at the shading of published field trials (static arrays, training-block sites). Observed values: rice, single trial; wheat and potato, the worse of the two seasons (the better season gave +2.7% and +11%); maize, mean of the three seasons of the dynamic system (DAV) and single values for the other two systems. The same trials enter the calibration of Section 2.6, so the comparison is not independent.

Trial	Shading	Observed	Conservative	Calibrated
Rice, Gonocruz et al. (2021)	0.25	0.84	0.77	0.79
Wheat, Weselek et al. (2021)	0.30	0.81	0.66	0.74
Potato, Weselek et al. (2021)	0.30	0.82	0.72	0.72
Maize, Ramos-Fuentes et al. (2023) (DAV)	0.35	0.83	0.70	0.70
Maize, Ramos-Fuentes et al. (2023) (half)	0.30	0.91	0.75	0.75
Maize, Ramos-Fuentes et al. (2023) (full)	0.50	0.70	0.54	0.54

Table 7: Sensitivity of the deployable light-sharing results to the crop shade response (rule controllers, baseline climate, floor $\rho = 0.9$, split 0, every fourth held-out site, 22 sites). Conservative: proportional response (main analysis); calibrated: radiation-use compensation fitted to Laub et al. (2022); field-fitted: fitted to the field trials, maize harvest index unaffected.

Shade response	Method	GCR	E/E_{PV}	Retention	LER
conservative	AV static	0.06	0.15	0.92	1.07
	Phenology rule	0.17	0.33	0.90	1.23
	Seasonal sharing	0.16	0.28	0.90	1.18
	Stress rule	0.11	0.24	0.91	1.15
calibrated	AV static	0.07	0.18	0.91	1.09
	Phenology rule	0.19	0.37	0.90	1.28
	Seasonal sharing	0.19	0.32	0.90	1.22
	Stress rule	0.13	0.29	0.91	1.20
field-fitted	AV static	0.16	0.41	0.91	1.32
	Phenology rule	0.37	0.70	0.92	1.62
	Seasonal sharing	0.37	0.61	0.92	1.53
	Stress rule	0.29	0.63	0.90	1.53

Table 8: Predicting relative yield under shade from 428 observations of 58 studies (Laub et al., 2022): leave-one-study-out cross-validation (every study held out once; observations with $s > 0$). Covariates: irradiation, climate zone, shade type, intercropping.

Model	RMSE	MAE	Median AE	Mean MAE per study
Proportional loss, $y = 1 - s$	0.430	0.305	0.210	0.224
One curve for all crops	0.372	0.265	0.193	0.208
Curve per crop type, no pooling	0.313	0.208	0.141	0.165
Curve per crop type, partial pooling	0.312	0.210	0.150	0.162
Partial pooling + covariates	0.311	0.206	0.148	0.161
Gradient-boosted trees	0.321	0.219	0.162	0.184

Table 9: Leave-one-study-out prediction error (RMSE of relative yield, observations with $s > 0$) by crop type: proportional loss, hierarchical curve with partial pooling, and gradient-boosted trees. The last columns give the difference hierarchical minus proportional with a 95% cluster-bootstrap interval over studies (negative: hierarchical better).

Group	Obs./studies	Proportional	Hierarchical	Boosted t
C3 cereals	30/10	0.214	0.187	0.205
Maize	40/10	0.225	0.235	0.248
Grain legumes	58/14	0.215	0.225	0.234
Tubers/root crops	10/2	0.250	0.254	0.257
Four crop types used by the simulator	138/30	0.220	0.222	0.234
Other types (berries, forages, fruits, vegetables)	201/29	0.528	0.361	0.369
All types	339/58	0.430	0.312	0.321

Table 10: Relative yield at 20% and 40% shading: data-driven model (median and 80% bootstrap interval over studies; crop-type curve, covariates at their reference values: mean irradiation, temperate zone, no solar-panel shade, no intercropping) and simulator under the two calibrated ensembles (mean of members; E1 fitted to hand-read curves, E2 fitted to the data-driven curves). The data-driven curve of rice is that of C3 cereals (no rice in the data set). Observations and studies: C3 cereals 41/10, maize 51/10, grain legumes 76/14, tubers 12/2.

Crop	Shading	Data-driven	Simulator, E1	Simulator, E2	Proportional
Wheat	20%	0.93 [0.88, 0.96]	0.88	0.92	0.80
Wheat	40%	0.72 [0.59, 0.78]	0.70	0.77	0.60
Rice	20%	0.93 [0.88, 0.96]	0.85	0.90	0.80
Rice	40%	0.72 [0.59, 0.78]	0.66	0.73	0.60
Maize	20%	0.73 [0.58, 0.86]	0.86	0.90	0.80
Maize	40%	0.42 [0.27, 0.59]	0.69	0.75	0.60
Soybean	20%	0.71 [0.65, 0.79]	0.77	0.74	0.80
Soybean	40%	0.40 [0.34, 0.53]	0.53	0.48	0.60
Potato	20%	0.93 [0.92, 0.97]	0.88	0.92	0.80
Potato	40%	0.71 [0.67, 0.83]	0.73	0.79	0.60

Table 11: Shade-response parameters of ensemble E2, calibrated by gradient descent through the simulator to the data-driven curves and the field trials (12 randomised fits; mean $\pm$ s.d.): radiation-use compensation c and multiplier h on the flowering-window harvest-index penalty (1 = assumed value).

Crop	c	h multiplier
Wheat	0.91 ± 0.31	0.74 ± 0.53
Rice	0.61 ± 0.25	0.83 ± 0.33
Maize	0.14 ± 0.11	0.16 ± 0.13
Soybean	0.05 ± 0.09	1.66 ± 0.38
Potato	0.62 ± 0.20	0.90 ± 0.50

Table 12: Land-equivalent ratio of rules and learned controllers under four true shade responses, with the array density designed under the true response (split 0, baseline climate, floor $\rho = 0.9$, every sixth held-out pair, 21 pairs, coarse GCR grid, one training seed; season means, so differences differ slightly from the pair-level intervals in the text). SALS is trained on the proportional response; robust controllers are trained on random members of ensemble E1, E2 or E3 (adaptive: with the measured biomass ratio). The proportional and harsh responses lie outside the ensembles.

Method	Proportional	Calibrated	Field-fitted	Harsh
AV static	1.052	1.059	1.254	1.037
Seasonal sharing	1.133	1.156	1.385	1.097
Stress rule	1.129	1.160	1.406	1.086
Feedback rule	1.139	1.164	1.392	1.089
Phenology rule (written)	1.164	1.190	1.456	1.122
Grid-tuned window	1.170	1.199	1.472	1.129
CMA-ES optimised window	1.173	1.198	1.479	1.125
SALS	1.178	1.209	1.489	1.132
SALS-robust, E1	1.179	1.204	1.491	1.127
SALS-adaptive, E1	1.176	1.202	1.495	1.124
SALS-robust, E2	1.165	1.192	1.483	1.126
SALS-robust, E3	1.163	1.191	1.481	1.121
SALS-adaptive, E3	1.164	1.191	1.475	1.118

Table 13: Array design under shade-response uncertainty, phenology rule, three spatial splits pooled (76 distinct crop–site pairs, every third held-out pair). Truth: each of five members of the ensemble in turn; the mean response and the member-based rules use the other four members (leave-one-out, so the truth never enters the design). E2: ensemble calibrated to the learned curves and field trials; E3: same procedure with target errors of 0.2 on the learned curves. Pairs: share with mean retention ≥ 0.9 ; shortfall: mean of $\max(0, 0.9 - \text{retention})$ per pair.

Design	E2		LER	E3		
	Pairs ≥ 0.9 (%)	Shortfall		Pairs ≥ 0.9 (%)	Shortfall	
oracle (true response)	87	0.0025	1.486	85	0.0029	1
mean response	70	0.0062	1.493	71	0.0093	1
floor met in 50% of other members	65	0.0084	1.510	67	0.0125	1
floor met in 75% of other members	78	0.0045	1.462	79	0.0083	1
floor met in all other members	89	0.0019	1.412	88	0.0032	1
floor 0.92 under mean response	88	0.0022	1.426	85	0.0051	1
floor 0.94 under mean response	97	0.0007	1.345	92	0.0027	1
floor 0.96 under mean response	99	0.0003	1.251	96	0.0015	1

Table 14: Design rules outside the ensemble: share of crop–site pairs (%) with mean retention ≥ 0.9 under four fixed responses (P proportional, C calibrated to the meta-analysis, F field-fitted, H harsh), phenology rule, three splits pooled (76 distinct pairs). The member-based rules use the five members of E2 or E3, the mean-response design and the margins use the mean of all twelve members, and the oracle uses the true response. LER in the last two columns is the mean over the four responses; the oracle rows are identical for E2 and E3 because the oracle uses the true response.

Design	E2				E3				LER	
	P	C	F	H	P	C	F	H	E2	E3
oracle (true response)	91	89	87	84	91	89	87	84	1.283	1.283
mean response	13	13	68	12	13	17	79	7	1.435	1.448
floor met in 75% of members	16	16	80	12	22	25	93	13	1.384	1.397
floor met in all members	17	17	80	12	29	32	93	16	1.348	1.336
floor 0.92 under mean response	17	20	86	12	29	33	91	20	1.376	1.399
floor 0.94 under mean response	45	47	97	34	45	45	97	29	1.303	1.345
floor 0.96 under mean response	63	78	99	54	58	63	99	51	1.220	1.264

Table 15: Assumed economic parameters (illustrative, not data). Ranges used in the sensitivity analysis are in parentheses.

Parameter	Value
Electricity price	60 USD MWh ⁻¹ (30–120)
Capex, AV single-axis tracker plant	1.0 USD Wp ⁻¹ (0.7–1.5)
Capex, conventional tracker plant	0.75 USD Wp ⁻¹
Discount rate, lifetime, O&M	6% (4–8%), 25 years, 1.5% of capex per year
Installed capacity	2.1 MWp per hectare of module area share g (0.21 kWp m ⁻²)
Crop price per tonne of dry matter	wheat 290, rice 400, maize 250, soybean 520, potato 1000 USD (0.

Table 16: Illustrative annual economics per hectare (76 crop–site pairs, three splits, main shade response, baseline climate; parameters in Table 15). Net value is relative to the open field (crop costs unchanged) in kUSD ha⁻¹ yr⁻¹ at three electricity prices; break-even is the price at which the net value is zero. The conventional plant replaces the crop.

Method	E (MWh)	Capex annuity + O&M	Crop loss	Net at 60	Net at 90
AV static	200	11.5	0.27	0.3	6.3
Seasonal sharing	326	29.6	0.31	-10.3	-0.5
Stress rule	336	23.2	0.30	-3.3	6.8
Phenology rule (written)	382	30.1	0.30	-7.5	4.0
Grid-tuned window	415	31.6	0.31	-7.0	5.5
CMA-ES optimised window	428	31.0	0.30	-5.6	7.3
SALS	438	31.6	0.31	-5.6	7.6
Conventional PV plant	1306	58.7	3.60	16.0	55.2

Table 17: One-at-a-time sensitivity of the net value (kUSD ha⁻¹ yr⁻¹, relative to the open field) at 90 USD MWh⁻¹, densities designed for the floor as in the main analysis. In the capex rows the capex of the conventional plant is set to 0.75 times that of the agrivoltaic plant.

Parameter	Value	Static arrays	Optimised window	SALS	Conventional plant
Base		6.3	7.3	7.6	55.2
AV capex (USD/Wp)	0.7	9.7	16.6	17.1	72.8
AV capex (USD/Wp)	1.5	0.6	-8.2	-8.2	25.8
Discount rate	4%	8.0	12.0	12.4	64.1
Discount rate	8%	4.4	2.1	2.4	45.4
Crop prices	0.5×	6.4	7.4	7.7	57.0
Crop prices	2.0×	6.0	7.0	7.3	51.6
Electricity price	60	0.3	-5.6	-5.6	16.0
Electricity price	120	12.3	20.1	20.7	94.4

Table A.18: Paired comparison of SALS with each baseline on the 92 held-out sites of the main split (unit of analysis: site mean over test seasons). ΔLER : mean difference (SALS minus baseline) with cluster-bootstrap 95% CI; p : two-sided Wilcoxon signed-rank test; % better: share of sites where SALS has the higher LER; Δr : difference in mean yield retention with its Wilcoxon p -value.

Climate	Baseline	ΔLER [95% CI]	p	% better	Δr	p	Sites
Baseline	AV static	+0.205 [+0.176, +0.236]	$< 10^{-4}$	100	-0.011	$< 10^{-4}$	92
Baseline	Seasonal sharing	+0.088 [+0.068, +0.109]	$< 10^{-4}$	97	+0.001	0.922	92
Baseline	Phenology rule	+0.038 [+0.025, +0.053]	$< 10^{-4}$	84	+0.001	0.723	92
Baseline	Stress rule	+0.096 [+0.071, +0.124]	$< 10^{-4}$	93	-0.002	0.055	92
+1.5 °C	AV static	+0.208 [+0.180, +0.237]	$< 10^{-4}$	100	-0.010	$< 10^{-4}$	92
+1.5 °C	Seasonal sharing	+0.095 [+0.077, +0.114]	$< 10^{-4}$	98	+0.001	0.812	92
+1.5 °C	Phenology rule	+0.033 [+0.022, +0.045]	$< 10^{-4}$	86	+0.002	0.855	92
+1.5 °C	Stress rule	+0.108 [+0.083, +0.134]	$< 10^{-4}$	98	-0.001	0.041	92
+2 °C	AV static	+0.209 [+0.181, +0.238]	$< 10^{-4}$	100	-0.010	$< 10^{-4}$	92
+2 °C	Seasonal sharing	+0.097 [+0.079, +0.115]	$< 10^{-4}$	97	+0.001	0.876	92
+2 °C	Phenology rule	+0.032 [+0.022, +0.044]	$< 10^{-4}$	85	+0.002	0.496	92
+2 °C	Stress rule	+0.115 [+0.089, +0.142]	$< 10^{-4}$	98	-0.003	0.016	92
+3 °C	AV static	+0.219 [+0.190, +0.250]	$< 10^{-4}$	100	-0.009	$< 10^{-4}$	92
+3 °C	Seasonal sharing	+0.108 [+0.089, +0.129]	$< 10^{-4}$	96	+0.001	0.935	92
+3 °C	Phenology rule	+0.034 [+0.022, +0.046]	$< 10^{-4}$	86	+0.002	0.330	92
+3 °C	Stress rule	+0.131 [+0.102, +0.162]	$< 10^{-4}$	98	-0.002	0.142	92