

1 Light-sharing control for tracker-based agrivoltaics
2 under a yield floor: a simulation study of what
3 learning adds and what shade-response uncertainty
4 costs

5 Quang-Vinh Dang*

6 *British University Vietnam, Hung Yen, Vietnam*

7 **Abstract**

Single-axis trackers let agrivoltaic arrays shift their daily rotation from maximum electricity towards letting light through (light sharing). We quantify its value under a yield-retention floor of 90% with a differentiable simulator of the whole system. The crop response to shade is the least certain part: a hierarchical model fitted to 428 observations from 58 field studies predicts a new study of the four staple crop types only to about 0.22 in relative yield, no better than proportional loss, so we turned the learned curves and four field trials into ensembles of plausible responses by gradient descent through the simulator. With the proportional-growth shade response, on 92 held-out crop–site pairs, light-sharing rules allowed arrays 2.1–2.9 times as dense as static arrays at equal retention and raised the land-equivalent ratio (LER) from 1.07 to 1.18–1.24 (1.27 for a neural controller); the floor was met in 69–76% of seasons. Pooled over three splits, the neural controller beat the written phenology rule by 0.044 LER, a tuned rule by 0.018 (95% interval 0.012–0.024) and a CMA-ES-optimised schedule by 0.013; training on ensembles or a biomass signal did not help within our budget (one seed, 21 pairs). The design step carried more compliance risk than control (pairs at the floor: 70% to 88% across design rules, 87% to 91% across light-sharing controllers): a density chosen under the mean response left about 70% when another plausible response held (oracle 85–87%); a floor raised by 0.02 and an ensemble-quantile rule gave no detectable difference at equal compliance; neither was safe outside the ensemble. Under illustrative costs, light sharing breaks even with the open field at 70–92 USD MWh^{−1} (static arrays at 59)

*Corresponding author.
Preprint submitted to *Artificial Intelligence in Agriculture*
Email address: vinh.dang@buv.edu.vn (Quang-Vinh Dang)

and a conventional plant is more profitable above 48. The simulator’s bias is uncertain for maize and soybean (53% of the upscaled 14 PWh yr⁻¹).

8 *Keywords:* Agrivoltaics, Light sharing, Solar tracking, Food–energy nexus,
9 Crop modelling, Differentiable simulation, Hierarchical shrinkage model

10 1. Introduction

11 Two of the defining challenges of this century compete for the same re-
12 source: land. The energy transition requires terawatts of new solar pho-
13 tovoltaic (PV) capacity, and ground-mounted solar farms need roughly one
14 to two hectares per megawatt. Food production must rise while crops are
15 exposed to more frequent heat waves and droughts. Warming alone is pro-
16 jected to reduce global yields of maize, wheat, rice and soybean by several
17 per cent per degree (Zhao et al., 2017; IPCC, 2021), and a new generation
18 of climate and crop models shows these impacts emerging earlier than pre-
19 viously expected (Jägermeyr et al., 2021). Croplands are also where solar
20 power potential is greatest (Adeh et al., 2019). They are flat, sunny, acces-
21 sible and close to rural electricity demand, so converting farmland into solar
22 farms is tempting and controversial at the same time.

23 Agrivoltaic (AV) systems combine crops and PV modules on the same
24 land and aim to resolve this conflict (Dupraz et al., 2011; Dinesh and Pearce,
25 2016; Weselek et al., 2019). By sharing sunlight between modules and plants,
26 a hectare can deliver more food plus energy than the two uses side by side,
27 with land-equivalent ratios (LER) well above one in field trials and mod-
28 elling studies (Dupraz et al., 2011; Amaducci et al., 2018; Valle et al., 2017;
29 Trommsdorff et al., 2021). The partial shade also changes the crop micro-
30 climate. It lowers canopy temperature and evaporative demand, conserves
31 soil moisture and improves water-use efficiency, especially in drylands (Mar-
32 rou et al., 2013b; Hassanpour Adeh et al., 2018; Barron-Gafford et al., 2019;
33 Elamri et al., 2018; Weselek et al., 2021). Agrivoltaics is therefore a can-
34 didate adaptation measure against heat and drought stress, as reviewed by
35 Mamun et al. (2022), in addition to a mitigation measure through renew-
36 able electricity. The same shade, however, also reduces the light available for
37 photosynthesis. Regulators increasingly require AV systems to retain most
38 of the reference crop yield. The German pre-standard DIN SPEC 91434, for
39 example, requires at least 66% of the reference yield (DIN, 2021), and other
40 schemes and many farmers set stricter floors. How much electricity can be

41 produced without sacrificing food security depends on the crop, the local
42 climate, the year and the growth stage.

43 Most AV designs fix the geometry of the array, i.e. the ground coverage
44 ratio (GCR, module area per ground area) and the tilt or tracking rule, from
45 generic rules of thumb. Single-axis trackers add a degree of freedom: the
46 modules can be rotated away from the sun to let light through when the crop
47 needs it, as demonstrated for lettuce by Valle et al. (2017), or towards the
48 sun to maximise electricity. Choosing the rotation every day is a sequential
49 decision problem. The value of light depends on the crop’s phenological stage,
50 soil water and upcoming weather, and the effect of a decision appears in the
51 harvest months later. Recent work in this journal has used evolutionary
52 multi-objective optimisation to lay out PV arrays on farmland (Zhang et al.,
53 2026). Machine learning for crop management has mainly used reinforcement
54 learning with black-box simulators (Gautron et al., 2022), which would need
55 very many simulated seasons for global, multi-decadal assessments.

56 Two methodological strands are relevant. Differentiable process models,
57 which couple physical simulators with learning, allow parameters and poli-
58 cies to be fitted by gradient descent through the simulator (Reichstein et al.,
59 2019; Shen et al., 2023), and we use them for both calibration and control.
60 Decision making under deep uncertainty asks for designs that perform accept-
61 ably across plausible models and not only under the best estimate (Marchau
62 et al., 2019), and scenario-based designs have finite-sample guarantees un-
63 der convexity assumptions that we do not meet (Calafiore and Campi, 2006;
64 Campi and Garatti, 2008); our design rules are heuristic counterparts and
65 we compare them with a simple safety margin.

66 Two questions follow. First, how much of the electricity that an agri-
67 voltaic system can produce under a yield-retention floor comes from the array
68 density, and how much from dynamic light-sharing control of the trackers?
69 Second, how robust is that answer to the weakest part of any such study, the
70 crop response to shade, and does a controller learned from a simulator add
71 anything to a well-designed hand-written rule? Answering them requires
72 a simulator that couples the physics of light, PV, microclimate and crop
73 growth, a protocol in which all controllers face the same design constraints,
74 a check of the shade response against field data, and enough sites and seasons
75 to separate real effects from noise.

76 This paper makes three contributions. The first is a quantification of
77 light-sharing control under a yield floor. Light sharing is parameterised by
78 one daily level between full backtracking and the rotation that minimises the

79 shadow on the crop, and we compare static arrays, a seasonal schedule, a rule
80 that follows crop phenology, a rule that reacts to heat and drought stress,
81 and a neural controller (SALS) trained through the simulator. The second
82 is a deployable evaluation protocol. The array density of every site is chosen
83 from its historical seasons only, with the same rule for every method, and
84 performance is measured on held-out sites and years, over three training seeds
85 and three spatial splits. The third is a differentiable agrivoltaic simulator
86 that supports both, coupling hourly solar geometry, single-axis trackers with
87 backtracking, row-to-row shading and PV electricity, the direct and diffuse
88 light reaching the crop, canopy temperature and reference evapotranspiration
89 under the modules, and the SIMPLE crop model (Zhao et al., 2019) extended
90 with a flowering-stage shade penalty. All components are differentiable tensor
91 operations (Karniadakis et al., 2021). The PV part is benchmarked against
92 PVGIS. The crop response to shade is learned from the open data of a meta-
93 analysis (Laub et al., 2022), checked against field trials on rice, wheat, potato
94 and maize (Gonocruz et al., 2021; Weselek et al., 2021; Ramos-Fuentes et al.,
95 2023), and propagated through the simulator as an ensemble.

96 We apply the protocol to 400 crop–site pairs that represent where maize,
97 wheat, rice, soybean and potato are grown worldwide, sampled in proportion
98 to their harvested area (Yu et al., 2020) with local crop calendars (Jägermeyr
99 et al., 2021), and driven by NASA POWER daily weather for 2001–2019
100 (Sparks, 2018) under baseline and +1.5, +2 and +3 °C climates. The main
101 findings are as follows.

- 102 1. Light sharing is the main lever on electricity at a given floor, although
103 economically it pays only at high prices (item 6). At the same mean
104 yield retention of 90%, arrays with light-sharing control carry 2.1 to
105 2.9 times the module area of static arrays and produce 1.8–2.4 times
106 as much electricity (Section 3.1).
- 107 2. The best hand-written rule follows crop phenology and needs no weather
108 forecast. Rules that react to heat or drought stress are not better and
109 depend on forecast quality (Section 3.3).
- 110 3. The crop response to shade can be learned from data. A hierarchical
111 model fitted to 428 observations from 58 field studies predicts a held-
112 out study of the four staple crop types to about 0.22 in relative yield,
113 no better than proportional loss (it is better only for crop types the
114 simulator does not represent), so between-study spread is large. The
115 simulator is close to these curves for cereals and potato but milder

- 116 than the meta-analysis for maize and soybean, whereas one maize trial
 117 is milder still; the gains for those two crops are therefore uncertain and
 118 possibly overestimated (Section 3.7).
- 119 4. A learned controller adds a small gain over the best rule: 0.044 LER
 120 over the written phenology rule and 0.018 over a rule tuned on the
 121 same training seasons. Training it on ensembles of shade responses, or
 122 giving it a farm-measurable biomass signal, does not help, and with
 123 the wider ensembles it did not retain its margin over the written rule
 124 in our training budget (Section 3.8).
 - 125 5. The bottleneck is the uncertainty of the crop response that enters the
 126 design. A density chosen under the mean response leaves about 70% of
 127 crop–site pairs at the floor when another plausible response holds (ora-
 128 cle 85–87%). A rule built on the quantiles of an ensemble of responses
 129 raises this by 8 points for 0.03 LER, and a floor raised by 0.02 under the
 130 mean response gives no detectable difference at equal compliance (the
 131 intervals are too wide to establish equivalence); no rule is safe outside
 132 the ensemble, so the ensemble is useful to size a margin (Section 3.8).
 - 133 6. Under illustrative costs, light sharing breaks even with the open field
 134 at 70 to 92 USD MWh⁻¹ (static arrays at 59), learning is worth about
 135 2kUSD ha⁻¹ yr⁻¹ over a tuned rule at 90 USD MWh⁻¹, and a conven-
 136 tional plant replacing the crop is more profitable above its break-even
 137 of 48 USD MWh⁻¹ (Section 3.9).
 - 138 7. At the yield floor, agrivoltaics on 5% of the world’s staple cropland
 139 could supply about 14 PWh yr⁻¹, but only with about 10 TWp of mod-
 140 ules, several times the installed PV capacity of the world (Section 3.11).

141 2. Materials and methods

142 2.1. Global cropland sites, weather and crop calendars

143 We studied five staple crops that together cover about 670 Mha of har-
 144 vested area and 580 Mha of physical cropland: wheat, rice, maize, soybean
 145 and potato. Harvested areas (rainfed and irrigated) were taken from the
 146 SPAM 2010 gridded production maps (Yu et al., 2020) and aggregated to 1°
 147 cells. For each crop, 80 cells with at least 5000 ha of the crop were drawn
 148 without replacement with probability proportional to harvested area. An un-
 149 weighted average over the sample therefore approximates an area-weighted
 150 average over global production. The 400 crop–site pairs lie in 381 distinct

cells in 50 countries (Fig. 2). Sowing and maturity dates came from the GGCM Phase 3 crop calendar at 0.5° (Jägermeyr et al., 2021). For wheat we used the winter-wheat calendar unless the coldest-month mean temperature was below -8°C , in which case the spring-wheat calendar was used. Daily weather for 2001–2020 was retrieved from NASA POWER (Sparks, 2018): global and diffuse shortwave irradiance, maximum, minimum and dew-point temperature, precipitation and wind speed at 2 m. Each site contributes 19 seasons (sowing years 2001–2019), each simulated as the 365 days after sowing, which gives 7600 site-seasons.

Climate change was represented by the global warming levels $+1.5$, $+2$ and $+3^\circ\text{C}$ relative to 2001–2020. Daily maximum, minimum and dew-point temperatures were shifted uniformly, which keeps relative humidity approximately constant and raises the vapour-pressure deficit. The atmospheric CO_2 concentration was set to 390, 450, 500 and 600 ppm for the baseline and the three warming levels. Precipitation and irradiance were left unchanged. This delta approach isolates the thermal signal of warming; its limitations are discussed in Section 4.5.

2.2. A differentiable agrivoltaic simulator

The simulator advances all site-seasons in parallel with a daily time step for the crop and an hourly time step for radiation and PV (Fig. 1). All equations are implemented as PyTorch tensor operations (Paszke et al., 2019), so that yields and electricity are differentiable with respect to the design and control variables.

Hourly irradiance. Daily global and diffuse irradiance were disaggregated into hourly values with the Collares-Pereira and Rabl (1979) and Liu and Jordan (1960) ratios and renormalised to the daily totals. Beam irradiance on the horizontal is the difference between the two. Solar position follows the standard declination and hour-angle equations.

Array geometry and PV electricity. We consider elevated horizontal single-axis trackers with north–south axes. The axis height is 1.25 module widths, e.g. 2.75 m for the common 2.2-m single-module-high (1P) tracker. The rotation ϕ is limited to $\pm 60^\circ$. The ground coverage ratio g is the module width divided by the row pitch and is restricted to $g \leq 0.5$. For 2.2-m modules this keeps a pitch of at least 4.4 m, and a free strip of at least 2.2 m between the rows when the modules are horizontal, which is the stow position during field

186 operations. Let ψ be the solar zenith angle projected onto the east–west ver-
 187 tical plane. The energy-optimal rotation with backtracking (Lorenzo et al.,
 188 2011) is $\phi_{\text{bt}} = \psi - \text{sgn}(\psi) \arccos(\cos \psi / g)$ when $\cos \psi < g$, and $\phi_{\text{bt}} = \psi$ oth-
 189 erwise. The *light-sharing* rotation ϕ_{sh} is the feasible angle closest to edge-on
 190 to the sun, which minimises the shadow on the crop. Given a light-sharing
 191 level $u \in [0, 1]$ chosen once per day, the rotation in every hour is

$$\phi = (1 - u) \phi_{\text{bt}} + u \phi_{\text{sh}}. \quad (1)$$

192 Plane-of-array irradiance combines three terms: beam irradiance on the mod-
 193 ule, reduced by the fraction of the module shaded by the neighbouring row,
 194 $\max\{0, 1 - \cos \psi / (g |\cos(\phi - \psi)|)\}$; isotropic sky diffuse; and ground-reflected
 195 light (albedo 0.2). The cell temperature follows the NOCT model (45 °C),
 196 with an hourly air temperature interpolated sinusoidally between the daily
 197 extremes. Electricity per unit land is plane-of-array irradiance $\times \eta \times g$, with
 198 $\eta = 0.21 [1 - 0.0035(T_{\text{cell}} - 25)] (1 - 0.14)$ for module efficiency, temperature
 199 losses and system losses. The installed capacity is $2.1 g \text{ MWp ha}^{-1}$. A con-
 200 ventional ground-mounted PV plant ($g = 0.4$, backtracking, no crop) is the
 201 energy reference E_{PV} .

202 *Validation of the PV model.* We compared the simulated specific yield of an
 203 isolated tracker row ($g = 0.05$, negligible row-to-row shading) with PVGIS 5.2
 204 (European Commission, Joint Research Centre, 2022), for a horizontal north–
 205 south single-axis tracker with 14% losses, at 30 randomly chosen sites (29
 206 returned data). Across sites, the simulated and PVGIS specific yields were
 207 correlated with $r = 0.971$. The simulator was on average 4.9% lower (mean
 208 absolute deviation 5.5%). This is consistent with its $\pm 60^\circ$ rotation limit
 209 and with the different irradiance databases (NASA POWER vs. SARA2,
 210 ERA5 and NSRDB). The benchmark covers an isolated row at low density
 211 only; electricity at high density, with row-to-row shading and diffuse-light
 212 effects, was not validated.

213 *Light and microclimate at the crop.* Averaged over the row pitch, the fraction
 214 of ground reached by beam light is $1 - \min\{1, g |\cos(\phi - \psi)| / \cos \psi\}$. The
 215 fraction of sky diffuse reaching the ground is a two-dimensional isotropic
 216 sky-view factor. We pre-computed it for 25 values of g and 13 rotation
 217 angles by ray casting over 31 neighbouring rows, and interpolate it bilinearly
 218 (differentiably). The daily ratio s of the irradiance at the crop to that of the

open field defines the crop’s shading. The canopy maximum temperature is lowered by $\kappa(1 - s)$.

κ is uncertain. Field studies under PV modules report small reductions of daytime crop or canopy temperature and little change in air temperature (Marrou et al., 2013b; Barron-Gafford et al., 2019), and heat-stress thresholds calibrated on air temperature do not transfer directly to canopy temperature (Siebert et al., 2014; Webber et al., 2018). We therefore use a moderate $\kappa = 1^\circ\text{C}$ as the main case and report $\kappa = 0$ (no cooling) and $\kappa = 2^\circ\text{C}$ for all methods (Section 3.2). Nocturnal warming under the modules is not represented. The reference evapotranspiration under the modules, ET_0 , is re-computed with the FAO-56 Penman–Monteith equation (Allen et al., 1998) from the reduced radiation and canopy temperature. Soil water conservation under the modules (Hassanpour Adeh et al., 2018; Elamri et al., 2018) therefore emerges from the model, but rain redistribution by the modules (Marrou et al., 2013a) is ignored.

Crop growth. Crop growth follows the SIMPLE model (Zhao et al., 2019) with the species parameters of its Table 1a (Table 1). Daily biomass growth is

$$\Delta B = R_{\text{crop}} \cdot \text{RUE} \cdot f_{\text{Solar}} \cdot f_{\text{CO}_2} \cdot f_{\text{Temp}} \cdot \min\{f_{\text{Heat}}, f_{\text{Water}}\}, \quad (2)$$

where $R_{\text{crop}} = sR$ is the radiation received by the crop. The terms are as follows.

- f_{Solar} is the canopy’s radiation interception. In SIMPLE it depends only on thermal time, with senescence accelerated by heat and drought stress. Shade therefore reduces the daily biomass increment but not leaf-area development.
- f_{Heat} decreases linearly between the heat threshold T_{heat} and the extreme temperature T_{ext} of the canopy maximum temperature.
- $f_{\text{Water}} = 1 - S_{\text{water}} \text{ ARID}$, where the agricultural reference index for drought (Woli et al., 2012) is computed from a daily soil-water balance with curve-number runoff, drainage and root water uptake. Irrigated fields have $f_{\text{Water}} = 1$.

SIMPLE has a fixed harvest index and therefore ignores the high sensitivity of grain number to shading around flowering (Fischer, 1985; Andrade et al., 1999). It also ignores the differences in shade tolerance between crop types. In the meta-analysis of Laub et al. (2022), maize and

253 grain legumes lost yield strongly even at low shade levels, whereas C3 ce-
 254 reals and tuber crops initially lost less than proportionally; shading trials
 255 also show marked effects of shade on maize development (Ramos-Fuentes
 256 et al., 2023). We added a stage-dependent penalty. During a critical window
 257 ($0.45 \leq TT/T_{\text{sum}} \leq 0.65$), the harvest index is multiplied by $1 - h\bar{d}$, where $\bar{d}$
 258 is the mean relative radiation deficit $1 - s$ in that window. The sensitivity h
 259 is highest for maize (0.6) and soybean (0.5), the crop types that Laub et al.
 260 (2022) found most shade-sensitive, lower for the C3 cereals (wheat 0.4; rice
 261 0.5, for which no specific curve was available) and lowest for potato (0.2).
 262 These values are assumptions that reproduce this ranking; they are not cal-
 263 ibrated estimates. Section 3.6 shows the resulting shade-response curves.
 264 Because RUE is linear, SIMPLE does not credit the higher light-use effi-
 265 ciency of shaded canopies under more diffuse light. The penalty, conversely,
 266 makes shading around flowering more damaging than proportional. The net
 267 effect is closer to the field evidence than either assumption alone.

268 The thermal-time requirement to maturity was calibrated for every site as
 269 the mean thermal time between the GGCMI sowing and maturity dates under
 270 the 2001–2020 climate, and the canopy parameters were scaled accordingly.
 271 The cultivar is thus adapted to the present climate, and warming shortens the
 272 season. Growth stops smoothly once the thermal-time requirement is met.
 273 Yield is biomass multiplied by the (shade-adjusted) harvest index. Each site
 274 is represented by a rainfed and an irrigated sub-field in proportion to its
 275 SPAM irrigated fraction. Each sub-field carries its own array, so both yield
 276 *and* electricity are area-weighted over the two sub-fields.

277 2.3. SALS: simulation-trained adaptive light sharing

278 SALS learns a *control network* $u_t = \pi_{\omega}(\mathbf{x}_t; g)$ (Algorithm 1). It is a
 279 multilayer perceptron with two hidden layers of 64 units and a bounded-logit
 280 output, $u = \sigma(5 \tanh(z/5)) \in [0.007, 0.993]$, which keeps gradients non-zero
 281 at the extremes. It chooses the daily light-sharing level of Eq. (1). Its inputs
 282 are the information available on day t :

- 283 • the crop and soil state: relative thermal time TT/T_{sum} , radiation in-
 284 terception f_{Solar} , relative soil water and the previous day’s ARID;
- 285 • a *forecast* of the day’s maximum temperature (relative to T_{heat} and
 286 T_{opt}), irradiance and ET_0 ;

- static information: the irrigation status of the sub-field, the array’s GCR g , latitude, day of season and the crop.

Forecasts are perturbed with random errors, independently every day: Gaussian with a standard deviation of 2 °C for maximum temperature, and multiplicative with standard deviations of 20% for irradiance and 15% for ET_0 . The physics always uses the true weather. The same forecast errors are applied to the rule-based controllers. The controller is conditioned on g . During training, each season is simulated with a GCR close to the density that is sustainable at its training site. This is the largest value on a coarse grid (0.05, 0.10, ..., 0.50) at which seasonal sharing keeps the mean retention of the site’s training seasons at or above ρ , plus a uniform perturbation in $[-0.05, +0.10]$ truncated to $[0.05, 0.50]$. The network therefore learns to control the densities that design selection will later pick. Outside the cropping season the trackers backtrack for maximum electricity.

For site-season i with open-field yield Y_i^{open} , let $r_i = Y_i/Y_i^{\text{open}}$ be the yield retention and $e_i = E_i/E_{\text{PV},i}$ the electricity relative to a conventional PV plant. Training minimises the penalised objective

$$\mathcal{L}(\omega) = \mathbb{E}_{i \sim \mathcal{B}} \left[-e_i + \mu_1 \max\{0, \rho - r_i\} + \mu_2 \max\{0, \rho - r_i\}^2 \right], \quad \rho = 0.9, \mu_1 = 5, \mu_2 = 30, \quad (3)$$

which combines an exact (hinge) and a quadratic penalty for the food-security floor ρ . Because r_i and e_i are outputs of the differentiable simulator, the exact gradient of Eq. (3) is obtained by back-propagation through the 365 daily steps of every season (back-propagation through time). This is the key difference from reinforcement learning with black-box simulators (Gautron et al., 2022), which estimates gradients from sampled returns.

Each minibatch contains 512 site-seasons, each simulated as a rainfed and an irrigated sub-field, drawn from one of the four climates in turn. Seasons with a failed open-field crop ($Y^{\text{open}} < 0.2 \text{ t ha}^{-1}$) are excluded from training and evaluation. The network is fine-tuned through the simulator with Adam (Kingma and Ba, 2015) (learning rate 10^{-3} , cosine decay, gradient-norm clipping at 1) for 300 steps. Growth-stage logistic functions are evaluated in a numerically stable form, so that gradients remain finite even under extreme heat.

Initialisation. Trained from random weights, the controller converged to sharing light throughout the cropping window ($u > 0.99$ on all days), i.e. to the seasonal-sharing rule, because the gradient of electricity with respect to u

Algorithm 1 SALS training

Require: training site-seasons in four climates; floor ρ ; penalties μ_1, μ_2

- 1: initialise π_ω by behaviour cloning of the phenology rule
- 2: **for** step = $1, \dots, N$ **do**
- 3: sample a climate and a minibatch $\mathcal{B}$ of site-seasons (rainfed and irrigated sub-fields)
- 4: set g_i to the site’s sustainable density plus a perturbation; initialise crop and soil states
- 5: **for** day $t = 1, \dots, 365$ **do**
- 6: $u_{it} \leftarrow \pi_\omega(\mathbf{x}_{it}; g_i)$ from state and noisy forecast if the crop is in the field, else 0
- 7: hourly rotation (Eq. 1), PV electricity, crop light, canopy temperature, ET_0
- 8: advance the soil-water balance and the SIMPLE crop state (Eq. 2)
- 9: $r_i \leftarrow Y_i/Y_i^{\text{open}}, e_i \leftarrow E_i/E_{PV,i}$ (area-weighted); loss $\mathcal{L}$ (Eq. 3)
- 10: back-propagate through the season; Adam update of ω

321 is small and of the same sign for most days of a season. We therefore initialise
322 the network by behaviour cloning of the phenology rule of Section 2.4: binary
323 cross-entropy against the rule’s output (soft targets 0.95 and 0.05) on about
324 1.7 million day-states visited by the rule in the training seasons, with 1500
325 Adam steps at a learning rate of 3×10^{-3} . Fine-tuning then moves the
326 policy away from this initialisation wherever doing so improves the objective.
327 Consequently, the gain of SALS over the phenology rule measures what fine-
328 tuning through the simulator adds to that rule, not what can be learned
329 without it.

330 *2.4. Deployable design selection and baselines*

331 A real array has one fixed density, which must be chosen before it is built.
332 Every method, SALS included, therefore uses the same *history-based design*
333 *rule*. For a site, the GCR is the largest value on the grid $\{0.05, 0.075, \dots, 0.50\}$
334 for which the method’s mean yield retention over the site’s 2001–2014 sea-
335 sons is at least $\rho = 0.9$. If no value qualifies, no array is built and the site
336 remains open field. A stricter *chance-constrained* variant requires $r \geq 0.9$ in
337 at least 80% of the historical seasons. The historical seasons of the test sites
338 are never used for training, and the weather of the evaluation years is never
339 used for design. We compare four controllers:

- 340 • *AV static*: backtracking trackers without light sharing ($u = 0$);
- 341 • *Seasonal sharing*: the light-sharing rotation ($u = 1$) throughout the
- 342 cropping season;
- 343 • *Stress rule*: an expert rule that shares light ($u = 1$) during the main
- 344 growth phase ($0.1 < TT/T_{\text{sum}} < 0.9$) except on days with a heat-stress
- 345 forecast ($T_{\text{max}} \geq T_{\text{heat}}$) or drought ($\text{ARID} > 0.5$), when the modules
- 346 track the sun and shade the crop;
- 347 • *SALS*: the learned controller.

348 The open field and a conventional PV plant without crops ($\text{LER} = 1$ each)
 349 are the references. Because the design rule is identical for all controllers, dif-
 350 ferences between methods reflect the quality of control and the array density
 351 that each controller can sustain under the food-security floor. We also com-
 352 pare the controllers at identical GCR over the whole grid (matched-density
 353 frontiers).

354 2.5. A shade-response model learned from field data

355 The simulator needs a crop response to shade that is supported by ob-
 356 servations. We learned one from the open data of the meta-analysis of Laub
 357 et al. (2022) (Zenodo 10.5281/zenodo.5716091, CC BY 4.0): 428 observa-
 358 tions of yield relative to an unshaded control at relative shade reductions of
 359 0–90% from 58 studies, with crop and crop type, experiment type (artificial
 360 shading or intercropping), climate zone, irradiation and shade type. Rela-
 361 tive yield is modelled as $y(s) = \exp(-\beta s^\gamma)$ with Student- t errors ($\nu = 4$;
 362 maximum a posteriori estimation, so this is a hierarchical shrinkage regres-
 363 sion with bootstrap intervals and not a full Bayesian model). We compared
 364 six models: proportional loss $y = 1 - s$; one curve for all crops; one curve
 365 per crop type without pooling; one curve per crop type with partial pooling
 366 (Gaussian shrinkage of $\log \beta$ and $\log \gamma$ towards a common value, scale 0.5);
 367 partial pooling with covariates on $\log \beta$ (standardised irradiation, subtrop-
 368 ical zone, solar-panel shade and intercropping); and gradient-boosted trees
 369 on the same inputs. Fits use L-BFGS. Models were compared by leave-one-
 370 study-out cross-validation, in which every study is held out once, which tests
 371 the prediction of a new study and not of a new observation in a known study.
 372 Intervals for the crop-type curves and covariate effects come from a cluster
 373 bootstrap over studies (150 resamples).

374 2.6. Shade-response uncertainty: gradient calibration, robust controllers and 375 design

376 The crop’s response to shade is the least certain part of the simulator,
377 so we treat it as an unknown with a distribution. Two parameters per crop
378 enter the simulator: a radiation-use compensation c (growth is multiplied by
379 $1 + c(1 - s)$, s being the share of open-field irradiance that reaches the crop)
380 and a multiplier h on the flowering-window harvest-index penalty. Because
381 the simulator is differentiable, both can be fitted by gradient descent through
382 whole simulated seasons to published yield–shading observations. The ob-
383 servations are the relative yields at 20% and 40% shading of the data-driven
384 curves of Section 2.5 for the five crops (median, with an error equal to the
385 width of the 80% bootstrap interval divided by 2.56 and at least 0.04; rice
386 takes the curve of C3 cereals, which contains no rice), the field trials on rice,
387 wheat, potato and maize (Gonocruz et al., 2021; Weselek et al., 2021; Ramos-
388 Fuentes et al., 2023) (each season of each trial as a separate observation, error
389 0.06–0.08) and the unaffected harvest index of shaded maize (Ramos-Fuentes
390 et al., 2023). We call the resulting ensemble E2. A wider ensemble, E3, uses
391 the same procedure with a target error of 0.2 for the learned curves, which
392 is the between-study spread of a new study (Section 3.7); the spread of
393 the resulting members is smaller (standard deviation of relative yield across
394 members 0.04–0.09 at 20% shading and 0.05–0.12 at 40% shading, against
395 0.02–0.04 and 0.03–0.06 for E2), because the fit also weights the field trials
396 and the priors. An earlier ensemble, E1, was fitted in the same way to curves
397 read by hand from the figures of Laub et al. (2022) and was used to train
398 the robust and adaptive controllers below. Priors are $c \sim \mathcal{N}(0.4, 0.5^2)$ and
399 $h \sim \mathcal{N}(1, 0.5^2)$, and the parameters are clamped to $[0, 2.5]$ (c) and $[0, 2]$ (h)
400 after every step, which is why a response milder than the data cannot be
401 lowered further. We fitted 12 ensemble members by randomised maximum a
402 posteriori estimation: every member sees the observations perturbed by their
403 errors and a prior draw perturbed by its width, and is optimised with Adam
404 (40 steps, with the loss still falling slightly at the last step; 40 training-block
405 seasons per crop and member; shading of the static arrays cached). The
406 ensemble is a device for generating plausible responses and not a validated
407 posterior: it has no coverage check, and for some parameters (the harvest-
408 index multiplier of wheat and potato) its spread equals that of the prior.
409 The design analyses use five fixed members (indices 0, 3, 5, 8 and 10 of 12),
410 chosen in advance. The ensemble is used in three ways. All training and
411 evaluation ran on a 4-core CPU; one training run takes about 20–25 minutes

412 and the calibration of one ensemble about 7 minutes.

413 *Robust and adaptive controllers.* SALS-robust has the architecture and
414 training of SALS (Section 2.3), but every training season is simulated with the
415 response of a randomly drawn ensemble member that the controller does not
416 observe, and the density jitter is widened to $[-0.05, +0.25]$ because feasible
417 densities differ between members. SALS-adaptive additionally observes the
418 realised biomass of the crop under the modules relative to an open reference
419 plot, a quantity that can be measured on the farm by remote sensing or by
420 a reference plot. A feedback rule that shares light whenever this ratio falls
421 below a target serves as the hand-written counterpart; its target (1.0) was
422 chosen on training-block seasons under the ensemble.

423 *Uncertainty-aware design.* The design rule of Section 2.4 chooses the
424 largest GCR whose historical retention meets the floor. Under shade-response
425 uncertainty we choose the largest GCR that meets the floor in at least 75%
426 of five ensemble members, and compare it with a design under the ensemble
427 mean and with an oracle design under the true response. Evaluation of this
428 rule is leave-one-out: the truth is the held-out member and the design uses
429 the other four.

430 *Design rules compared.* With the ensemble members indexed k and one
431 member held out as the truth, the density of a pair is chosen from its 2001–
432 2014 seasons as (i) the *oracle*: the largest GCR whose mean retention under
433 the truth meets the floor; (ii) the *mean-response design*: the same rule under
434 the mean response of the other four members (the mean of their parameters);
435 (iii) a *margin*: the same rule under that mean response with the floor raised
436 to 0.92, 0.94 or 0.96; and (iv) the *ensemble rule*: the largest GCR at which
437 the mean retention under at least a share q of the other four members meets
438 the floor ($q = 0.5$: two of four; 0.75 : three of four; all: four of four). Out
439 of the ensemble, the five members and their mean (the mean of all twelve
440 members for the mean-response design) are used in the same way and the
441 truth is one of the four fixed responses below. The margin grid was chosen
442 after a first analysis of one ensemble and is not pre-specified.

443 *Stress tests.* Controllers and rules are also evaluated under four fixed
444 responses, the proportional response, the response calibrated to the meta-
445 analysis, the response fitted to the field trials and a harsh response (propor-
446 tional growth and a 1.5 times larger harvest-index penalty). The proportional
447 and harsh responses lie outside the range of the ensemble.

448 2.7. *Economic analysis*

449 To judge whether the electricity gains matter economically, we computed
 450 an annual net value per hectare for each method relative to the open field:
 451 electricity revenue minus the annuity and operating cost of the installed ca-
 452 pacity minus the loss of crop revenue (crop production costs are unchanged).
 453 The installed capacity is $2.1 g$ MWp per hectare (modules cover a share g of
 454 the land at 0.21 kWp m^{-2}). Yield is dry matter, valued at assumed prices.
 455 A conventional ground-mounted plant (GCR 0.4) that replaces the crop is
 456 the comparator. All parameters (Table 15) are assumptions chosen to be
 457 representative, not data, and their ranges drive the conclusions; we therefore
 458 report break-even electricity prices and one-at-a-time sensitivities, and we do
 459 not monetise food security, which market crop prices do not capture.

460 2.8. *Evaluation protocol, statistics and upscaling*

461 Sites were split spatially: $10^\circ \times 10^\circ$ blocks were randomly assigned to train-
 462 ing (70%) or test (30%), so that no test site has a training site in the same
 463 block. Training uses the 2001–2014 seasons of the training sites. Design
 464 selection uses the 2001–2014 seasons of the test sites. Evaluation uses the
 465 2015–2019 seasons of the test sites. The same site-seasons are analysed in all
 466 four climates.

467 We report the following metrics:

- 468 • the mean electricity ($\text{MWh ha}^{-1} \text{ yr}^{-1}$) and the relative electricity e ;
- 469 • the mean yield retention r and the production-weighted retention $\sum Y / \sum Y^{\text{open}}$;
- 470 • season-level compliance (share of seasons with $r \geq \rho$) and site-level
 471 compliance (share of sites whose mean test-season retention is $\geq \rho$);
- 472 • the lower tail (10th percentile of r , share of seasons with $r < 0.8$);
- 473 • the land-equivalent ratio $\text{LER} = r + e$ (Dupraz et al., 2011).

474 Two aggregations are used and labelled in the text: tables report means over
 475 all test seasons of a sample (season-pooled), whereas statistical comparisons
 476 and intervals use the crop–site pair as the unit of analysis (mean of the pair’s
 477 seasons, then mean over pairs), so the same contrast can differ slightly be-
 478 tween the two (for example 0.034 and 0.038 for SALS minus the phenology
 479 rule on the main split). Samples are also labelled: the main analysis uses
 480 the 92 held-out pairs of one split and a GCR grid with step 0.025; analyses

Table 1: SIMPLE crop parameters (Zhao et al., 2019, Table 1a) and the harvest-index shade sensitivity h added in this study. T_{sum} , I_{50A} and I_{50B} are rescaled per site to the local crop calendar (Section 2.2). Units: temperatures in $^{\circ}\text{C}$, thermal times in $^{\circ}\text{C d}$, RUE in gMJ^{-1} ; S_{CO_2} is the relative RUE increase per 100 ppm above 350 ppm.

Crop (cultivar)	h	T_{sum}	HI	I_{50A}	I_{50B}	T_{base}	T_{opt}	RUE	$I_{50 \text{ max } H}$	$I_{50 \text{ max } W}$	T_{heat}	T_{ext}	S_{CO_2}	S_{water}
Wheat (Batten)	0.4	2150	0.34	280	50	0	15	1.24	100	25	34	45	0.08	0.4
Rice (IR72)	0.5	2300	0.47	850	200	9	26	1.24	100	10	34	50	0.08	1.0
Maize (McCurdy 84aa)	0.6	2050	0.50	500	50	8	28	2.10	100	12	34	50	0.01	1.2
Soybean (Williams 82)	0.5	2350	0.40	600	200	6	27	0.86	120	20	36	50	0.07	0.9
Potato (Sebago)	0.2	2400	0.85	500	350	4	22	1.30	50	30	34	45	0.10	0.4

pooled over three splits use the 76 distinct pairs of every third held-out pair and a grid with step 0.05, and analyses of alternative shade responses on one split use every sixth (21 pairs) or every fourth pair (22 pairs). Statistical comparisons use site means as the unit of analysis. We report paired Wilcoxon signed-rank tests and cluster (site) bootstrap 95% confidence intervals of the mean difference. SALS was trained with three random seeds on the main split, and with one seed on each of two further spatial splits.

For the global estimate, every site is evaluated in the split in which it is held out. The mean electricity and yield per hectare of each crop are scaled to the *physical* area of that crop in SPAM 2010, which avoids double counting multiple-cropped land (Waha et al., 2020). We report the installed capacity (TWp), the capacity factor and a supply curve.

3. Results

3.1. Electricity and food on held-out sites and years

Table 2 compares the four controllers on the 92 held-out crop–site pairs of the main split (95 pairs are held out; pairs without a season whose open-field yield is at least 0.2 t ha^{-1} in all four climates are excluded) (seasons sown in 2015–2019), after each site’s density has been chosen from its 2001–2014 seasons with the same rule for all controllers. By construction all methods keep the mean yield retention close to 90% (0.904–0.917), so differences in electricity are differences in the density that each controller can sustain under this floor.

Static arrays are the weakest. Without light sharing, the floor limits the ground coverage to 6% on average, and arrays produce $217 \text{ MWh ha}^{-1} \text{ yr}^{-1}$, 16% of a conventional PV plant per hectare (LER 1.07). Rotating the




Figure 1: Simulator and controller. The array density g of a site is fixed from its historical seasons. A control network maps the daily crop and soil state, the density g and a noisy forecast to the light-sharing level u_t . The differentiable simulator converts these decisions into PV electricity, light and microclimate at the crop, soil water and crop growth. The season’s electricity and yield retention enter a penalised objective whose exact gradient is back-propagated through the 365-day season.

modules to let light through during the cropping season (seasonal sharing) allows a coverage of 17% and $397 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ (LER 1.19). A rule that shares light only while the crop is growing (phenology rule) reaches $465 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ (LER 1.24). An expert rule that additionally shades the crop on forecast heat or drought days (stress rule) is worse than the phenology rule (LER 1.18): its density is limited to 13% because it reduces light at times when the floor is most binding.

The learned controller is best by a modest, consistent margin. SALS sustains a density of 17% (similar to the phenology rule), produces $510 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ and reaches an LER of 1.270 at the same retention. Compared with the phenology rule this is +0.038 in LER (95% CI over sites 0.025–0.053; +10% electricity) and it is higher at 84% of the 92 sites (Table A.18). Compared with static arrays it is +0.205 (0.176–0.236), with 2.35 times the electricity at 2.9 times the density. The frontier in Fig. 3 shows the same picture at matched densities: the learned policy lies on or slightly above the phenology rule over the whole density range, both lie well above seasonal sharing and the stress rule, and static arrays are far below.

What the floor does and does not guarantee. Enforcing the floor on the historical mean does not guarantee it in every season. In the test seasons, 70% of SALS seasons have $r \geq 0.9$ (phenology rule 69%, stress rule 76%, static arrays 93%), and 2.7% have $r < 0.8$ (1.3%, 0.9% and 0.9%). The lower tail of SALS is therefore slightly heavier than that of the rules:

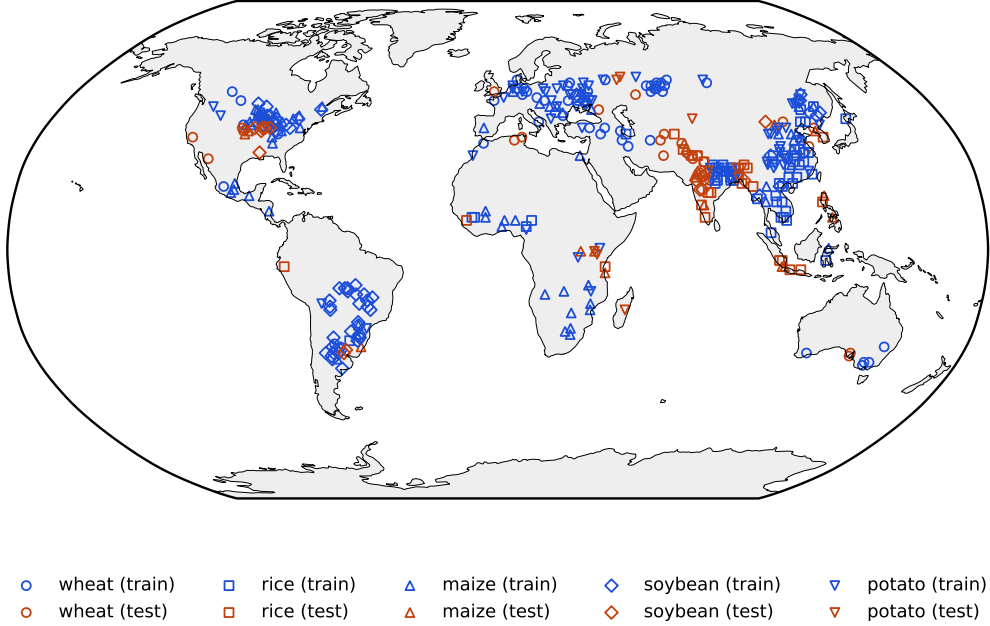


Figure 2: The 400 crop–site pairs, sampled in proportion to the harvested area of each crop (SPAM 2010). Blue: training sites; orange: geographically held-out test sites of the main split.

its 10th percentile of r is 0.883, against 0.887, 0.886 and 0.903. Requiring $r \geq 0.9$ in at least 80% of a site’s historical seasons (chance-constrained design) raises the season-level compliance to 90–94% for the light-sharing controllers. It costs about 0.065 in LER for SALS (1.205), which remains the best controller (phenology rule 1.192; Section 3.2).

3.2. Robustness

Seeds and splits. Three training seeds of SALS give LERs of 1.270, 1.268 and 1.269 under the baseline climate and 1.297, 1.298 and 1.297 at +3 °C. On two further spatial splits, in which different blocks of sites are held out, SALS exceeds the phenology rule by 0.047 and 0.041 in LER under the baseline climate and by 0.052 and 0.039 at +3 °C (main split, season-pooled: 0.034 and 0.032; the site-level paired contrasts of Table A.18 are 0.038 and 0.034). In every split SALS ranks first, the phenology rule second and static arrays last.




Figure 3: Electricity-retention frontiers at identical array densities (labels: GCR) for the held-out sites. Each point is the mean over test seasons at a fixed GCR; the dashed line is the 90% floor.

542 **Canopy cooling.** The canopy-cooling coefficient κ is the most un-
543 certain microclimate parameter. We repeated the whole protocol, includ-
544 ing the historical design step, with $\kappa = 0, 1$ and 2°C . Electricity changes
545 with κ because cooler canopies tolerate denser arrays (SALS: 475, 510 and
546 $549 \text{ MWh ha}^{-1} \text{ yr}^{-1}$ under the baseline climate), but the retention is held at
547 0.905 by the design rule and the ranking is unchanged. SALS exceeds the
548 phenology rule by 0.029, 0.034 and 0.042 in LER under the baseline climate
549 (0.022, 0.032 and 0.051 at $+3^\circ\text{C}$).

550 **Which ingredients matter.** Two ablations test the design of SALS
551 at the same training budget. Masking the stress inputs (soil water, ARID,
552 and the maximum temperature, irradiance and ET_0 forecasts) leaves the
553 LER unchanged (1.269 vs 1.270 under the baseline climate, 1.297 vs 1.297
554 at $+3^\circ\text{C}$). Training only on the baseline climate gives 1.268 and 1.291. The
555 improvement over the rules therefore does not depend on stress-aware inputs
556 or on training in warmed climates.

557 3.3. Floor conditioning, chance-constrained training and forecast errors

558 **One network for all floors.** Regulations and farmers set different
559 floors. The network of the main experiments was trained at $\rho = 0.9$ only.
560 We evaluated it at floors of 0.80, 0.85, 0.90 and 0.95, re-running the historical
561 design step for each floor (Table 3). Its LER exceeds that of the phenology
562 rule by 0.037, 0.036, 0.034 and 0.033 under the baseline climate (0.040, 0.037,
563 0.032 and 0.033 at $+3^\circ\text{C}$), so the benefit carries over to other floors without

retraining. A network that received the floor as an additional input and was trained over all four floors did not improve on it (differences in LER ≤ 0.001). Because the design step already adapts the density to the floor, conditioning the controller on the floor is not needed.

Layout, rotation limit and harvest-index assumptions. We checked three assumptions that could drive the ranking (Table 4; 21 pairs of the main split, one training seed). The rotation limit of the trackers sets how much light sharing can deliver: with a limit of 45° the LER of the phenology rule is 1.138, with 60° 1.170 and with 80° 1.206, whereas static trackers stay at 1.05–1.06, so the benefit of light sharing grows with the limit and the ranking is unchanged. Fixed non-tracking arrays held flat or rotated by 30° about the north–south axis (facing east, not the equator; equator-facing tilts are not modelled) do no better than static trackers (LER 1.048 and 1.043 against 1.053; differences of 0.005 are within the noise of 21 pairs and one training seed), so the static tracker is not a weak baseline relative to them; vertical bifacial layouts are not modelled. The ranking of the rules also does not depend on the assumed harvest-index penalty. The window test is weakly informative about the timing of the phenology rule, which shares light through almost the whole growing period, so a shifted window cannot move it much; the informative case is the removal of the penalty. Moving the critical window to 0.30–0.50 or 0.60–0.80 of the thermal-time requirement changes the LER of the phenology rule by at most 0.001, and removing the penalty altogether raises all light-sharing LERs (phenology rule 1.328, SALS 1.353) without changing the order: the phenology rule exceeds seasonal sharing by 0.03 to 0.05 in every case (0.035 with the early window, 0.034 with the late window, 0.053 without penalty), because it produces electricity before canopy development and after maturity, when the crop does not use the light, and not because of the assumed window. The margin of SALS over the written rule changes with these assumptions (0.005 to 0.025 LER), and its margin over the tuned rule is 0.000 to 0.009 (at a rotation limit of 45° SALS equals the tuned rule).

A floor of 66%. The German pre-standard DIN SPEC 91434 requires 66% of the reference yield (DIN, 2021). Repeating the design step with $\rho = 0.66$ (baseline climate, main split, same network trained at 0.9) the ground coverage cap of 0.5 binds for the sharing methods (mean GCR 0.48 for the phenology rule and 0.48 for SALS, 0.44 for the stress rule), static arrays reach a coverage of 0.24 and LER 1.28, the phenology rule 1.66 and SALS 1.71. At this floor the density is limited by the cap and by the mean

602 retention of 0.73 reached at the cap, not by the floor, so the comparison of
 603 controllers says little about control; we do not interpret it further.

604 **Chance-constrained training.** We also trained a network whose ob-
 605 jective penalises a batch share of seasons below the floor that exceeds 20%.
 606 It did not improve the tail. With the mean-based design step its LER was
 607 1.263 (baseline) and 1.290 (+3 °C) against 1.270 and 1.297, and with chance-
 608 constrained design it reached season-level compliance of 92.6% against 92.0%
 609 and an LER of 1.198 against 1.205. The season-level guarantee is obtained
 610 by the design step (Section 3.1), not by the training objective.

611 **Forecast errors.** We raised the forecast errors seen by the controllers
 612 to 0, 1, 2 and 3 times the errors used in training. The LER of SALS and
 613 of the phenology rule was identical at all four error levels to three decimals
 614 (densities 0.2 and 0.3: SALS 1.307 and 1.457, phenology rule 1.278 and
 615 1.416). Neither controller uses the weather forecast: both rely on the crop’s
 616 thermal time, which is observable from on-farm temperature records. The
 617 stress rule, which depends on the forecast of maximum temperature, degrades
 618 with forecast error: at density 0.3 its retention falls from 0.778 to 0.751 and
 619 its LER from 1.432 to 1.418 as the errors triple.

620 **A tuned rule.** The phenology rule above was written, not tuned,
 621 whereas SALS is trained. We therefore tuned the window of the rule on
 622 the training-block seasons with the same objective (electricity at the density
 623 where mean retention equals the floor); the best window starts at 0.15 in-
 624 stead of 0.05 of the thermal-time requirement (Table 5). Pooled over three
 625 splits and 76 distinct crop–site pairs, the tuned rule raises LER over the orig-
 626 inal rule by 0.026 (95% cluster-bootstrap interval 0.016 to 0.037), and SALS
 627 exceeds the tuned rule by 0.018 (0.012 to 0.024), or 5.7% more electricity at
 628 equal retention, and the original rule by 0.044 (0.030 to 0.060). Tuning the
 629 rule thus accounts for about 60% of the margin of SALS over the written
 630 rule. A five-parameter ramp-plateau schedule optimised by CMA-ES on the
 631 same training seasons performs like the tuned window (LER 1.238 against
 632 1.233, difference +0.005, interval −0.001 to +0.010), and SALS exceeds it
 633 by 0.013 (0.007 to 0.019); the optimum is an almost binary window (start
 634 0.17, end 0.93, level 0.99), so more flexible schedules of this family do not
 635 close the gap. The remaining margin of SALS is small but not zero; it is not
 636 explained by the tuning of a window or by a smooth ramp, and we did not
 637 test richer rule families (for example schedules that depend on soil water).
 638 SALS meets the floor in fewer seasons than the rules (83% against 84–85%).




Figure 4: Learned daily control for two test seasons under $+3^{\circ}\text{C}$, at the density chosen for the site, simulated for the dominant sub-field of the site. Shading marks the cropping window. Orange: light-sharing level u_t ; green: fraction of open-field irradiance reaching the crop; grey: daily maximum temperature (right axis).

3.4. What the learned policy does

Averaged over the rainfed test seasons of the main split at a density of 0.3, the light-sharing level of SALS is a smooth function of relative thermal time. It is close to zero in the first 10% of the thermal time required for maturity, rises through 0.55 at 10–20%, stays at 0.94–0.99 between 30% and 70%, and falls through 0.58 at 80–90% to about zero at maturity. The plateau covers the window in which our crop model penalises shading through the harvest index. The ramps at both ends let the modules produce electricity when additional light is worth least to the crop. Fig. 4 shows two example seasons. The phenology rule, which is also the initialisation of the network, shares light during the whole growing period. Fine-tuning through the simulator found that sharing can be reduced at the edges of the season: SALS gains its advantage by yielding light when it matters and collecting electricity when it matters less.

The policy also shares less light on days with a forecast heat stress: the mean level is 0.35 on such days against 0.60 (baseline) and 0.51 ($+3^{\circ}\text{C}$) on other days, and 58–63% of sites show a lower level on heat days. This cannot be read as a stress response, however. Heat days occur late in the season, when the phenological ramp is already falling, and a policy without

any stress input shows a similar contrast (0.37 vs 0.57 in a separate check). The ablation above confirms that stress information does not improve the result.

3.5. Crops and regions

The benefit of light sharing differs between crops (baseline climate, SALS). Soybean and maize tolerate the densest arrays (GCR 0.24 and 0.19; 691 and 577 MWh ha⁻¹ yr⁻¹; LER 1.41 and 1.30), potato follows (0.19; 542; 1.31), and rice and wheat are the most restrictive (0.13 and 0.15; 395 and 428; 1.20 for both). The stress rule has a lower LER for every crop (1.26 for soybean, 1.14–1.21 for the others).

The floor is not met evenly. Production-weighted retention is 0.888 for soybean and 0.898 for maize, below 0.9, because high-yielding seasons lose more than low-yielding ones. Regionally, Latin America and Oceania (4 sites) combine the highest density (GCR 0.41) and electricity (1090 MWh ha⁻¹ yr⁻¹) with a production-weighted retention of only 0.862. Africa and the Middle East (8 sites) reach 0.894, and only 63% of their sites keep a mean retention of at least 0.9. South Asia, the region with the most sites (39), reaches 0.902 (67% of sites) at a density of 0.14. The regional figures rest on 4 to 39 pairs without intervals and should be read as indications. A regulatory floor enforced per site would therefore need a region- and crop-specific design step, not a global one.

3.6. Model behaviour and validation

PV output. The simulated specific yield of an isolated tracker row agrees with PVGIS 5.2 (European Commission, Joint Research Centre, 2022) at 29 sites with $r = 0.971$, a mean bias of -4.9% and a mean absolute deviation of 5.5% (Section 2.2).

Shade response. Fig. 5 shows the simulated relative yield of rainfed crops under static arrays as a function of the season-mean shading of the crop. Yield declines close to proportionally with shading: 84–89% relative yield at 12% shading, 70–78% at 25%, 57–67% at 37% and 35–46% at 57%. Wheat loses most and potato and maize least. This response is steeper than the *less than proportional* losses that Laub et al. (2022) report for small shade levels of C3 cereals and tubers, and it does not reproduce their finding that maize and grain legumes are more shade sensitive than C3 cereals. In the simulator, shading also relieves heat and drought in rainfed maize and soybean, which




Figure 5: Simulated relative yield of rainfed crops under static arrays (GCR 0.1–0.5) against season-mean shading of the crop, baseline climate, held-out sites. The dotted line is proportional loss.

offsets the higher harvest-index penalty we assigned to them. The absolute effect of shading on each crop is therefore uncertain (Section 4.5).

Comparison with field trials. We fitted one radiation-use compensation coefficient per crop to the yield–shading curves of Laub et al. (2022) using training-block sites only (Methods, calibration script in the repository). The fit changed wheat and rice only (coefficients 0.4 and 0.1) because the proportional model already lies within about 0.02 of the C3-cereal and tuber curves, and it cannot lower the maize and legume responses below the proportional case (root-mean-square error 0.16 for maize and 0.08 for soybean). Against independent field trials (Table 6) the model underestimates retained yield: rice 0.79 against 0.84 (Gonocruz et al., 2021), wheat and potato 0.74 and 0.72 against 0.81–0.82 in the worse season of Weselek et al. (2021) (and above 1 in the better one), and maize 0.70–0.75 against 0.83–0.91 at 30–35% shading (Ramos-Fuentes et al., 2023), where harvest index was not affected by shade. The model is therefore conservative for food retention with respect to these trials, except for the stronger maize losses of the meta-analysis. Table 7 shows the consequence for the design results: with the calibrated response the allowed density and the LER of every method rise (phenology rule 1.28 against 1.23), and with a response fitted to the field trials the LER of the phenology rule is 1.62. The phenology rule is the best method in all three cases, and light sharing adds LER +0.17, +0.19 and

+0.30 over static arrays (conservative, calibrated, field-fitted); only the order of the seasonal and stress rules changes. The absolute gains for rice, wheat and potato are thus conservative with respect to these trials; for maize and soybean the meta-analysis points the other way (Section 3.7).

3.7. *A shade-response model learned from field data*

Cross-validation by study shows what the data can and cannot support (Tables 8 and 9). Over all 339 shaded observations the root-mean-square error of a held-out study falls from 0.43 for proportional loss $y = 1 - s$ to 0.37 for one curve for all crops and to 0.31 for crop-type curves, but this gain comes from crop types that the simulator does not represent (berries, forages, fruits and vegetables: 0.53 to 0.36). For the four crop types that the simulator uses, the learned curves are not better than proportional loss: the error is 0.222 against 0.220 pooled (difference +0.002, 95% interval -0.023 to $+0.030$), 0.187 against 0.214 for C3 cereals, and slightly worse for maize, grain legumes and tubers. Partial pooling is as good as separate curves (0.312 against 0.313) and the covariates add almost nothing (0.311). Gradient-boosted trees on the same inputs are worse (0.321) than the parametric model. The error of 0.22 for the four crop types is large: yields of a new study differ from the learned curve by about 20 percentage points, which is the between-study spread that a design under uncertainty has to cover. Crop type is the information that transfers between studies, and none of the covariate effects on $\log \beta$ is clearly different from zero (irradiation -0.10 , 90% interval -0.32 to 0.09 ; subtropical zone -0.37 , -0.73 to 0.07 ; intercropping 0.16 , -0.48 to 0.45). The estimated effect of solar-panel shade is not interpretable, because all 46 panel observations are of leafy vegetables, mainly lettuce, and the effect is confounded with crop type.

The learned curves show a clear ordering of crop types (Table 10). At 20% shading C3 cereals and tubers lose less than proportionally (relative yield 0.93), whereas maize (0.73) and grain legumes (0.71) lose more; at 40% shading the values are 0.72 and 0.71 against 0.42 and 0.40. Intervals are wide for maize (0.27–0.59 at 40%) and for tubers, which rest on two studies. The simulator’s responses are close to the data for cereals and potato but, under both calibrated ensembles, too mild for maize and soybean at 40% shading (0.75 and 0.48 for E2 against 0.42 and 0.40 in the data). The calibration cannot lower the maize response further, because the simulated shading also relieves heat and drought stress and the compensation and harvest-index parameters are bounded below at zero. The field trials of maize are also

milder than the meta-analysis: Ramos-Fuentes et al. (2023) report a harvest index unaffected by shade, and Sekiyama and Nagashima (2019) a yield 5.6% above the control under a low-density stilt-mounted array in one season at one site. For maize and soybean the direction of the bias is uncertain, and the gains reported in this paper are possibly overestimated.

3.8. *Learning under shade-response uncertainty*

Calibration. Gradient descent through the simulator fitted the five crops to the data-driven curves and the field trials (ensemble E2, Table 11). The data identify two features: the harvest index of maize is hardly affected by shade (multiplier 0.16 ± 0.13 , $c \approx 0.14$), as observed by Ramos-Fuentes et al. (2023), and soybean is more shade sensitive than the cereals (compensation 0.05, harvest-index multiplier 1.66), as in Laub et al. (2022). Wheat has the largest compensation (0.91 ± 0.31) and the largest uncertainty; the compensation is an effective parameter that absorbs diffuse-light and light-saturation effects that the simulator does not model explicitly, not a physical value.

Learned controllers. Controllers were evaluated under four fixed responses on 21 crop–site pairs of split 0 with one training seed, so the intervals below, which are cluster bootstraps over pairs of per-pair mean LER differences, do not include seed or split variation (Table 12 gives season means, which differ slightly). Under every response SALS, trained on the proportional response, is better than the written phenology rule: by +0.019 (0.007 to 0.039), +0.024 (0.009 to 0.045), +0.038 (0.017 to 0.063) and +0.015 (0.000 to 0.036) LER under the proportional, calibrated, field-fitted and harsh responses. Training on the narrow ensemble E1 did not change this: SALS-robust differs from SALS by -0.005 to $+0.003$ (intervals contain zero except for the proportional response, 0.000 to 0.003), and SALS-adaptive, which sees the measured biomass ratio, by -0.010 to $+0.007$ (all intervals contain zero). Training on the wider ensembles did not help either (the robust controllers fall below SALS even under responses inside their training ensemble, which may point to under-training with our budget of 300 steps): SALS-robust on E2 and E3 and SALS-adaptive on E3 are below SALS by 0.014, 0.017 and 0.017 LER under the proportional response, 0.017, 0.021 and 0.020 under the calibrated response, 0.004, 0.005 and 0.011 under the field-fitted response (intervals contain zero) and 0.008, 0.014 and 0.017 under the harsh response; the intervals exclude zero in 6 of the 9 comparisons for the proportional, calibrated and harsh responses. The E3 robust controller is not distinguishable

787 from the written rule (+0.001 to +0.003, intervals contain zero) except un-
788 der the field-fitted response (+0.032, 0.001 to 0.067). A tuned window and a
789 CMA-ES ramp schedule, also evaluated under the four responses (Table 12),
790 are 0.005 to 0.018 LER below SALS (intervals of the SALS contrast: +0.004
791 to +0.016 against the tuned window under the proportional response, and
792 containing zero for the ramp under the calibrated and field-fitted responses
793 and for the window under the harsh response), and the ensemble-trained
794 controllers are at or below the tuned window under all responses except the
795 field-fitted one. With one seed and 21 pairs these controllers therefore did not
796 retain the margin of SALS within our training budget; a controller that must
797 perform across plausible responses is less specialised, but a larger budget or
798 more seeds could change this. The feedback rule is worse than the written
799 rule by 0.026–0.058 LER under every response (intervals exclude zero). A
800 plausible reason is that the biomass ratio does not reveal the harvest-index
801 loss from shading around flowering, but we did not test this explanation.

802 **Design under uncertainty.** The larger risk is in the design step (Ta-
803 bles 13 and 14, Fig. 6). The designs below are evaluated by leaving one
804 ensemble member out as the truth and using the other four, so that the
805 truth never enters the design. Pooled over the three spatial splits (76 dis-
806 tinct crop–site pairs), a density chosen under the mean response of E2 leaves
807 70% of the pairs at the floor when the true response is another plausible
808 one, against 87% for an oracle design (E3: 71% and 85%). Choosing the
809 density so that the floor is met in at least 75% of the other members raises
810 compliance by +7.9 points under E2 (95% cluster-bootstrap interval 4.5 to
811 11.6) and +7.9 under E3 (3.9 to 12.1), at a cost of 0.031 and 0.025 in LER.
812 Raising the floor used in the design from 0.90 to 0.92 under the mean re-
813 sponse does more: +18.0 points (13.4 to 22.6) at 0.068 LER under E2 and
814 +14.3 points (9.5 to 19.2) at 0.056 under E3, and the shortfall is smaller
815 (mean over pairs of the maximum of zero and 0.9 minus retention: 0.0022
816 against 0.0045 for the quantile rule under E2). The two kinds of rule lie on
817 one frontier (Fig. 6). Interpolating the margin frontier (floors 0.90 to 0.96)
818 at the compliance of the ensemble rules, the LER of the margin minus that
819 of the rule is +0.001 (-0.009 to +0.012) for the rule with 75% of the other
820 members and +0.003 (-0.016 to +0.021) for the rule with all of them under
821 E2, and -0.006 (-0.019 to +0.005) and +0.024 (-0.011 to +0.058) under E3
822 (95% intervals). We find no detectable difference, but the intervals (down
823 to -0.02 and up to +0.06 LER) are too wide to establish equivalence. The
824 ensemble is thus useful for sizing the margin (0.02 to 0.04 above the floor for




Figure 6: Share of crop-site pairs at the floor against LER for the phenology rule under design rules for shade-response uncertainty (leave-one-member-out, three splits, 76 pairs): a higher floor under the mean response (blue; the first point is the design under the mean response) and a floor met in a share of the other members (orange), with the oracle design (green).

85–97% of the pairs), not for replacing it. None of these rules is safe outside the ensemble (Table 14): under the proportional, calibrated and harsh responses the mean-response design leaves 7–17% of the pairs at the floor, a floor of 0.92 or the ensemble rules 12–33%, a floor of 0.94 29–47% and a floor of 0.96 51–78%, at a cost of about 0.2 LER, whereas the field-fitted response, which lies within the range spanned by the ensembles, is handled well by every rule (80–99%) except the mean-response design (68–79%). Outside the ensemble the shortfall of the pairs below the floor is also large: under E2 the mean shortfall per pair (maximum of zero and 0.9 minus retention) is 0.075, 0.060 and 0.104 for the mean-response design under the proportional, calibrated and harsh responses, against 0.053, 0.038 and 0.078 for a floor of 0.92 and 0.012, 0.007 and 0.024 for a floor of 0.96. The guarantee therefore holds only within the range of responses that the ensemble spans; we did not test an ensemble that spans the proportional and harsh responses, which would protect against them at a higher cost in LER.

3.9. Economics of light sharing

All figures in this section follow from the assumed parameters of Table 15 and should be read as a sensitivity analysis, not a forecast. The economics are dominated by the capacity installed, not by the crop (Table 16): the annuity of the modules is 11–32 kUSD per hectare and year, whereas the crop

845 revenue lost is about 0.3 kUSD, so a method is worth what it earns per in-
 846 stalled watt. With the densities designed for the floor, static arrays earn +0.3
 847 kUSD ha⁻¹ yr⁻¹ at 60 USD MWh⁻¹ and break even at 59, whereas all light-
 848 sharing methods lose money at that price (−3 to −10 kUSD) and break even
 849 at 70 (stress rule), 73 (SALS), 77 (tuned phenology rule), 80 (written rule)
 850 and 92 (seasonal sharing) USD MWh⁻¹. At 90 USD MWh⁻¹ static arrays
 851 (6.3 kUSD) are better than the written (4.0) and tuned (5.5) rules and worse
 852 than SALS (7.6), which overtakes static arrays above about 85 USD MWh⁻¹;
 853 the difference between SALS and static arrays at 90 USD MWh⁻¹ is +1.3
 854 kUSD (95% interval −0.04 to 2.7, 76 pairs) and not distinguishable from
 855 zero, whereas SALS exceeds the tuned rule by 2.1 (1.6 to 2.7) and the writ-
 856 ten rule by 3.6 (2.8 to 4.5). A conventional plant that replaces the crop
 857 breaks even at 48 USD MWh⁻¹ and is more profitable than every agrivoltaic
 858 option above that price (16 kUSD at 60, 94 kUSD at 120, against −6 and 21
 859 for SALS); below it all options lose money and static arrays lose least (at 30
 860 USD MWh⁻¹: plant −23.2, SALS −18.7, static −5.7 kUSD). Per megawatt-
 861 hour, the electricity of SALS costs about 73 USD including the crop lost,
 862 against 48 USD for the plant, a difference of 25 USD MWh⁻¹. A premium of
 863 that size would equalise the unit cost but not the profit per hectare, because
 864 a hectare of plant produces three times as much electricity (1306 against
 865 438 MWh): relative to the plant, SALS would need a premium of about 49
 866 USD MWh⁻¹ at a market price of 60 and about 109 at 90 to match its profit
 867 per hectare, and only about 13 at 60 (none above 73) to match the open
 868 field. The sensitivity is large (Table 17): at 90 USD MWh⁻¹ the net value of
 869 SALS changes from +17 to −8 kUSD between module costs of 0.7 and 1.5
 870 USD Wp⁻¹, whereas crop prices matter little.

871 Because the density was designed for the floor and not for profit, we also
 872 let each site choose, among the densities that meet the floor in its history,
 873 the one that maximises its historical net value (sites where none pays stay
 874 open field; Fig. 7). Then no method pays at 30–40 USD MWh⁻¹; at 60 14–
 875 59% of the pairs install an array (none for seasonal sharing) and the net
 876 values are small (static arrays 0.9, SALS 0.3 kUSD); static arrays are the
 877 best option up to about 75 USD MWh⁻¹, SALS overtakes them at about 75
 878 (4.9 against 4.3 at 80, 8.3 against 6.3 at 90, 20.7 against 12.3 at 120), the
 879 stress rule at about 81, the tuned rule at about 87 and the written rule at
 880 about 97. The value of control is therefore real but depends on a high price:
 881 at 90 USD MWh⁻¹ and with profit-optimal densities SALS exceeds the tuned
 882 rule by 1.7 kUSD ha⁻¹ yr⁻¹ (95% cluster-bootstrap interval 1.2 to 2.3), the




Figure 7: Net value per hectare relative to the open field against electricity price. Left: densities designed for the floor, with a conventional plant replacing the crop (dashed). Right: each site chooses the feasible density that maximises its historical net value. Illustrative parameters of Table 15.

883 written rule by 2.8 (2.0 to 3.6) and static arrays by 2.0 (0.8 to 3.3).

884 **Cost of the safety margin.** At 90 USD MWh⁻¹, moving from a density
885 designed under the mean response to a floor of 0.92 under that response costs
886 1.00 kUSD ha⁻¹ yr⁻¹ (17%) under E2 for 18 more points of compliance (70%
887 to 88%) and 1.13 kUSD (20%) under E3 for 14 points (71% to 85%), about
888 56 and 79 USD ha⁻¹ yr⁻¹ per point. The ensemble rule that requires all other
889 members to meet the floor costs 0.85 and 1.22 kUSD for 19 and 17 points
890 (45 and 74 USD per point), so on this measure it is cheaper than the margin
891 under E2 and similar under E3. At 60 USD MWh⁻¹ a margin has a negative
892 cost, because every module beyond break-even loses money; the economics
893 of the design margin therefore depend on the price level, whereas its effect
894 on compliance does not.

895 3.10. Agrivoltaics under warming

896 Under warming, the open-field production of the held-out sites (Fig. 8)
897 falls at +3 °C by 23% for maize and 18% for rice, is nearly unchanged for soy-
898 bean, and rises for wheat (+3%) and potato (+10%). Agrivoltaic production
899 with light-sharing control follows the open field at about 90% of its level, so
900 it inherits the same climate impacts (maize -31%, rice -27% relative to the
901 baseline open field at +3 °C). Agrivoltaics does not protect staple production
902 against warming in this model. The LER of light-sharing systems increases
903 slightly with warming (SALS 1.270 to 1.297; phenology rule 1.236 to 1.265),




Figure 8: Production relative to the baseline open field (sum over held-out site-seasons) as a function of warming, for the open field, the stress rule and SALS.

904 because the floor allows slightly denser arrays in warmer conditions (GCR
 905 0.173 to 0.177 and electricity 510 to 540 MWh ha⁻¹ yr⁻¹). The margin of
 906 SALS over the rules does not grow with warming (Table A.18).

907 3.11. Global potential and what it would take

908 Every site is evaluated when it is held out (219 of the 400 crop–site pairs,
 909 in at least one split; 222 are held out in at least one split and three have
 910 no valid season), and each crop’s mean electricity per hectare is scaled to its
 911 physical area in SPAM 2010 (581 Mha for the five crops, which avoids count-
 912 ing multiple-cropped land twice). Equipping 5% of this area (29.1 Mha) with
 913 SALS arrays at the floor would produce 14.4 PWh yr⁻¹ under the baseline
 914 climate and 16.0 PWh yr⁻¹ at +3 °C, with a production-weighted retention
 915 of 0.903 and 0.900. The phenology rule would give 13.0 and 14.5 PWh yr⁻¹,
 916 and static arrays 6.4 and 6.9 PWh yr⁻¹. For comparison, world electricity
 917 generation was about 30 PWh in 2023 (Ember, 2024).

918 These numbers require large amounts of hardware. SALS at 5% cover-
 919 age needs 10.6 TWp of modules (11.3 TWp at +3 °C), at a capacity factor of
 920 15.5% (16.2%). World solar PV capacity at the end of 2024 was about 1.9 TW
 921 (IRENA, 2025), so the requirement is about 5.7 times the installed PV fleet.
 922 The electricity is also variable, peaks at midday and would be curtailed or
 923 stored. The figures are technical potentials, not projections. Ordering sites
 924 by electricity per hectare gives a supply curve (Fig. 9): the best 2% of the sta-
 925 ple cropland would supply 16 PWh yr⁻¹ and the best 5% 40 PWh yr⁻¹ with
 926 SALS. The best tenth of the sites are mostly rainfed maize and soybean sites
 927 in the Americas and in southern and eastern Africa (7% mean irrigated frac-
 928 tion against 29% overall), which are also among the regions where the floor




Figure 9: Supply curve under the baseline climate: cumulative electricity when sites are equipped in order of decreasing electricity per hectare, against the share of staple cropland (physical area) equipped. Each controller is evaluated at the density that keeps its mean retention at 0.9.

929 is hardest to meet (Section 3.5). Siting by electricity alone would therefore
 930 concentrate agrivoltaics where food-security risk is highest.

931 Maize and soybean, for which the simulated shade response is milder
 932 than the meta-analysis (Section 3.7), account for 53% of the electricity of
 933 the upscaling (maize 30%, soybean 23%, wheat 29%, rice 15%, potato 3%).
 934 Counting only wheat, rice and potato would give about 7 PWh yr⁻¹ at the
 935 same coverage, so the global figure should be read as an upper estimate for
 936 about half of its content. By crop, the LER of SALS under the baseline
 937 climate is 1.20 for wheat and rice, 1.31 for potato, 1.30 for maize and 1.41
 938 for soybean.

939 4. Discussion

940 4.1. What the results establish

941 Four findings stand on consistent evidence. First, sharing light with the
 942 crop is the main lever on the electricity that tracker-based agrivoltaic arrays
 943 can produce under a yield floor. With static trackers the floor limits arrays
 944 to about 6% ground coverage and an LER of 1.07. Any light-sharing con-
 945 trol allows 2.1 to 2.9 times the density, and the LER rises to 1.18–1.24 for
 946 hand-written rules. Second, the best simple rule follows crop phenology. It
 947 shares light from the onset of canopy development (5% of the thermal-time




Figure 10: SALS electricity (a) and land-equivalent ratio (b) at the pooled held-out sites (baseline climate, mean of seasons 2015–2019).

requirement) until maturity and produces electricity before and after, when
 the crop does not use the light. It needs only the crop's thermal time, which
 is observable on the farm, and it is insensitive to forecast errors, whereas
 the rule that reacts to heat and drought stress is worse in every climate and
 degrades with forecast error. Third, the qualitative result does not depend
 on the weakest part of the model. The simulated shade response is conser-
 vative against the field trials we could use (Section 3.6); calibrating it raises
 the allowed density, and the phenology rule stays the best method, with a
 gain over static arrays of 0.17, 0.19 and 0.30 LER under the conservative,
 calibrated and field-fitted responses. However, the response learned from
 the meta-analysis shows steeper losses for maize and grain legumes than the
 simulator reproduces (Section 3.7), so the absolute gains are conservative
 for wheat, rice and potato and uncertain, possibly overestimated, for maize
 and soybean (53% of the upscaled electricity). Fourth, a controller learned
 through the differentiable simulator is better than the best rule we could
 write, by 0.04 in LER against the written rule (0.018 against a rule whose
 window was tuned on the training seasons, 5.7% more electricity at equal re-
 tention), reproducibly over training seeds (range 0.002 in LER), three spatial
 splits, three canopy-cooling assumptions and four climates, and at 84–86% of
 the held-out sites. This gain is small compared with the gap between static
 and dynamic control.

969 *4.2. What role does artificial intelligence play?*

970 We tested three ways in which learning could matter, and the answers dif-
971 fer. The first is learning the crop’s response to shade from data. The learned
972 model is useful mainly as a measurement of how little can be learned: a hier-
973 archical model fitted to 428 observations from 58 studies predicts a held-out
974 study of the four staple crop types with an error of 0.22 in relative yield, no
975 better than proportional loss, and the gain over proportional loss (0.43 to
976 0.31 over all types) comes from crop types that the simulator does not repre-
977 sent. Crop type is the information that transfers, whereas irradiation, climate
978 zone and shade type add nothing measurable and a non-parametric model
979 (gradient-boosted trees) is worse with this number of studies (Section 3.7).
980 Two things follow. The shade response of the staple crops is a distribution
981 with a spread of about 20 percentage points between studies and not a curve,
982 and the process-based simulator does not reproduce the steep losses of maize
983 and legumes seen in the data. Learning the response was limited by the data
984 and not by the model: two studies support the tuber curve, no rice study
985 exists, and 46 panel observations are all leafy vegetables.

986 The second is calibration through the differentiable simulator. Gradient
987 descent through whole seasons fitted the response of each crop to the learned
988 curves and to the field trials, and it produced an ensemble of plausible re-
989 sponses that fits the trials and the curves for cereals and potato but not the
990 maize and soybean curves (Section 3.8). The ensemble informs the design
991 decision: a density chosen under the mean response leaves about 70% of
992 crop–site pairs at the floor if another plausible response holds, against 85–
993 87% for an oracle, and a density chosen to meet the floor in 75% of plausible
994 responses recovers about half of that gap at a cost of 0.03 LER. A simple mar-
995 gin, however, does as well: raising the floor from 0.90 to 0.92 under the mean
996 response reaches 85–88% of the pairs, and interpolating the margin frontier
997 at the compliance of the ensemble rules gives LER differences (margin minus
998 rule) of +0.001 (−0.009 to 0.012) and +0.003 (−0.016 to 0.021) under E2
999 and −0.006 (−0.019 to 0.005) and +0.024 (−0.011 to 0.058) under E3, so the
1000 two are not distinguishable on this sample, and the cost per point of com-
1001 pliance is similar. None of the rules is safe outside the ensemble (7–33% of
1002 the pairs at the floor under the proportional, calibrated and harsh responses
1003 for the mean design, the margin of 0.92 and the ensemble rules), and only a
1004 floor of 0.96 reaches 51–78% at a cost of about 0.2 LER. What the ensem-
1005 ble adds is the size of the margin, not a better rule. The quantile rule is a
1006 scenario-type design (Calafiore and Campi, 2006; Campi and Garatti, 2008);

with five members it has none of the finite-sample guarantees of that theory, and the comparison with the margin shows that it is not needed in practice. The bottleneck of tracker-based agrivoltaics under a yield-retention floor is the uncertainty of the crop response that enters the design, and it is data on that response, not larger networks, that would improve the result; the learned curves show a spread of about 20 percentage points between studies. The guarantee holds only within the range of responses that the ensemble spans.

The third is control, and here learning contributes little. Differentiability is what lets a network of 5441 weights be trained in 300 gradient steps through whole seasons; a derivative-free search such as CMA-ES handled a five-parameter schedule with a few hundred simulator evaluations and reached the performance of the tuned window, but would not scale to the network, and the network’s remaining margin over it is small (0.013 LER). The learned controller adds 0.018 LER over a tuned phenology rule (and 0.044 over the written one), and 0.015–0.038 over the written rule under the four alternative responses on one split (21 pairs); a reader who needs a simple, auditable controller can use that rule and obtain most of the electricity gain over static arrays. Neither of the two extensions that we expected to matter under uncertainty did: training on ensembles of responses (E1, E2 and E3) did not help and, for the wider ensembles, did not retain the margin over the written rule within our training budget (one seed, 21 pairs), and neither did giving the controller a farm-measurable signal, the biomass of the crop relative to a reference plot, nor a feedback rule built on that signal. The learned policy also has to be initialised from the phenology rule (Section 4.5). We therefore do not claim that learned control is what makes agrivoltaics work: in this model, the schedule that a learned policy converges to is a smooth function of crop phenology, which a rule can already express.

4.3. What the results do not establish

The early hypothesis that stress-aware control, that is, shading the crop on heat and drought days, would be a main source of benefit under warming was not supported. The stress rule performs worse than the phenology rule in every climate, and a learned controller without stress inputs performs as well as one with them. A plausible reason is that in our crop model the yield gain from shading on a hot or dry day is smaller than the loss of light on that day, so stress information does not help. Whether this holds for crop models with more detailed canopy-temperature and heat-stress responses is

an open question. Agrivoltaics does not protect staple production against warming in our simulations: light-sharing systems produce about 90% of the open-field yield in every climate and follow the same decline in maize and rice (Fig. 8).

4.4. Food security in practice

Holding the historical mean retention at 90% does not hold it in each season or at each site. 70% of SALS seasons meet the floor, the lower tail is slightly heavier than for the rules, and production-weighted retention is below 0.9 for soybean (0.888) and maize (0.898) and in several regions. Chance-constrained design raises season-level compliance to about 92% at a cost of 0.065 in LER, and the ranking of the controllers is preserved. Whichever formulation is chosen, a per-site regulatory floor would need a region- and crop-specific design step. The upscaled potential (Section 3.11) shows the trade-off in magnitude. The 14 PWh yr⁻¹ that 5% of the staple cropland could supply need about 10 TWp of modules, several times the present global fleet, and the best sites are rainfed maize and soybean areas where production-weighted retention is lowest.

Who bears the risk. A floor defined on the multi-year mean leaves the farmer exposed in the seasons below it: with the density designed for the floor, about 30% of seasons fall short on the main split (70% compliance for SALS), and a per-season guarantee through chance-constrained design costs 0.065 LER. Siting by electricity favours maize and soybean areas, where compliance is lowest (Section 3.11), so the risk concentrates where the energy gain is largest. Who carries it, through contracts, insurance or a premium, is a design question that the model does not address, and food is valued only through crop revenue at market prices.

4.5. Limitations

- *Simulation only.* All results come from a simulator. PV output was benchmarked against PVGIS (29 sites, mean bias -4.9%), but the crop response to shade was learned from one meta-analysis data set and compared with four field trials (Sections 3.6 and 3.7). The simulated response is conservative with respect to the trials (rice, wheat, potato, maize) but too mild for maize and soybean relative to the meta-analysis, where the 40%-shading relative yields are 0.42 and 0.40 against 0.75 and 0.48 in the simulator. The harvest-index sensitivities h are assumptions, and the sensitivity analyses use subsamples of crop-site

1080 pairs (22 for the rule comparison under alternative responses, 21 for
1081 the controller comparison, one training seed).

- 1082 • *Canopy cooling.* The effect of shade on canopy temperature was repre-
1083 sented by one coefficient κ , applied to a heat-stress function that was
1084 calibrated on air temperature (Siebert et al., 2014; Webber et al., 2018).
1085 The ranking of the controllers is robust to κ between 0 and 2 °C, but
1086 the absolute electricity changes by about $\pm 8\%$. Nocturnal warming
1087 under modules and rain redistribution are not modelled.
- 1088 • *Crop model.* SIMPLE has a fixed harvest index and linear radiation-
1089 use efficiency, and its CO₂ response is that of controlled experiments
1090 that tend to overstate field responses (Long et al., 2006). More detailed
1091 models with grain-number formation and canopy energy balance may
1092 favour different schedules.
- 1093 • *Layouts and models.* The baseline of static arrays is a backtracking
1094 tracker held at the energy-optimal angle; flat and 30°-tilted fixed ar-
1095 rays did no better in a check on 21 pairs, but vertical bifacial layouts
1096 (Willockx et al., 2023; Campana et al., 2021) are not modelled, and the
1097 ground coverage ratio is a coarse proxy of crop productivity (Dupraz,
1098 2023). The PV model is isotropic in the diffuse component and has no
1099 bifacial gain; it was benchmarked against PVGIS at low density only
1100 and has not been compared with a reference implementation such as
1101 pvlib (F. Holmgren et al., 2018). The rotation limit was varied between
1102 45 and 80 degrees and changes the size of the benefit but not the rank-
1103 ing; the harvest-index window is assumed to be common to all crops,
1104 and shifting it or removing the penalty did not change the ranking of
1105 the rules.
- 1106 • *Climate scenarios.* Warming was applied as uniform temperature shifts
1107 with fixed precipitation, radiation and relative humidity, and with un-
1108 changed cultivars and sowing dates. It is not a bias-adjusted climate
1109 projection.
- 1110 • *Initialisation and state information.* The network is initialised from the
1111 phenology rule, because training from random weights collapsed to full
1112 sharing. The gain of SALS is therefore the gain of fine-tuning that rule,
1113 and we cannot say whether a better optimiser or a richer policy class

would find larger gains. The controller uses the simulated state of the crop and soil (thermal time, canopy interception, soil water) without measurement error, and the forecast errors we tested apply to weather only.

- *Economics and deployment.* The economic analysis (Section 3.9) uses assumed prices and costs, no grid, curtailment, storage or subsidy model, and does not value food security; its break-even prices move with the capex and the discount rate (Fig. 7). The value of midday electricity, curtailment, storage, structural cost and farmer acceptance (Trommsdorff et al., 2021) will determine the realisable share. The average row pitch at the chosen densities (about 13 m for 2.2-m modules) leaves room for machinery, but the densest designs (GCR up to 0.5) have a pitch of only 4.4 m.

5. Conclusions

Light sharing, the daily choice between maximum electricity and letting light through to the crop, is the main lever on the electricity of tracker-based agrivoltaics at a given yield-retention floor, although it pays only at high electricity prices: below about 80 USD MWh⁻¹ static arrays earn more. In a differentiable simulator of trackers, PV output, crop light, microclimate and crop growth, evaluated on held-out crop-site pairs (92 on the main split, 76 distinct pairs pooled over three splits for the controller and design comparisons) with the array density of each site chosen from its own history, light-sharing rules allowed 2.1 to 2.9 times the module density of static arrays at a mean yield retention of 90%. They raised the land-equivalent ratio from 1.07 to 1.18–1.24, and a rule that follows crop phenology and needs no weather forecast was the best of them. The simulated shade response is conservative against the field trials we could use for rice, wheat and potato, and calibrating it raised the gains without changing the ranking of methods; for maize and soybean, which supply 53% of the upscaled electricity, the meta-analysis indicates larger losses than the simulator reproduces whereas one maize trial indicates smaller ones, so the gains for these crops are uncertain and possibly overestimated. The crop’s response to shade, the least certain input, was learned from 428 field observations: a hierarchical model predicted a held-out study of the staple crop types no better than proportional loss (error 0.22), which shows the size of the between-study spread, and gradient descent through the simulator turned it into an ensemble of plausible responses.

1150 A learned controller exceeded the written phenology rule by 0.044 LER and
 1151 a rule with a window tuned on the training seasons by 0.018 (95% interval
 1152 0.012–0.024), by following a smooth phenology-dependent schedule that a
 1153 window or a ramp rule approximates within 0.013 LER (the rest may come
 1154 from its state inputs such as soil water and density, which we did not test);
 1155 training it on ensembles of shade responses or giving it a farm-measurable
 1156 biomass signal did not help, and with the wider ensembles it did not retain
 1157 its margin over the written rule in our training budget (one seed, 21 pairs).
 1158 The share of pairs at the floor varied by up to 18 points between design
 1159 rules (70% to 88%) and by 4 points between controllers (87% to 91% among
 1160 light-sharing controllers), so the design step carried more compliance risk
 1161 than the control. A density chosen under the mean response left about 70%
 1162 of crop–site pairs at the floor when another plausible response held (oracle
 1163 85–87%); choosing the density so that the floor holds in 75% of the plausible
 1164 responses recovered about half of the gap at a cost of 0.03 LER, and a floor
 1165 raised by 0.02 under the mean response gave no detectable difference at equal
 1166 compliance (differences of -0.006 to $+0.024$ LER with intervals containing
 1167 zero, too wide to establish equivalence); no rule was safe outside the ensemble.
 1168 The ensemble is therefore useful to size a safety margin, not to replace
 1169 it. Under illustrative costs, light sharing breaks even with the open field at
 1170 70 to 92 USD MWh⁻¹ (73 for the learned controller, 59 for static arrays); a
 1171 conventional plant replacing the crop is more profitable above its break-even
 1172 of 48 USD MWh⁻¹, so agrivoltaics depends on a premium, land-use rules or
 1173 a value for food production that markets do not give. The yield floor is met
 1174 on average but not in every season or region, and the electricity potential of
 1175 5% of staple cropland (about 14 PWh yr⁻¹) requires about 10 TWp of mod-
 1176 ules. Future work should validate the crop response to shade under trackers
 1177 with more field data, notably for maize and legumes where the meta-analysis
 1178 and the trials disagree, extend the crop model to grain-number formation and
 1179 canopy energy balance, add costs and grid constraints, and pursue controllers
 1180 that guarantee the yield-retention floor per site and per season.

1181 **CRedit authorship contribution statement**

1182 **Quang-Vinh Dang:** Conceptualization, Methodology, Software, Formal
 1183 analysis, Investigation, Data curation, Visualization, Writing – original draft,
 1184 Writing – review & editing.

1185 Declaration of competing interest

1186 The author has no known competing financial interests or personal re-
1187 lationships that could have appeared to influence the work reported in this
1188 paper.

1189 Data and code availability

1190 All data used are public: SPAM 2010 v2r0 harvested areas (Yu et al.,
1191 2020), the GGCM Phase 3 crop calendar (Jägermeyr et al., 2021), NASA
1192 POWER daily weather (Sparks, 2018) the SIMPLE crop parameters (Zhao
1193 et al., 2019) and the shade-response meta-analysis data of Laub et al. (2022)
1194 (Zenodo 10.5281/zenodo.5716091, CC BY 4.0, copied into the repository).
1195 The code of the simulator, of SALS and of all experiments, together with the
1196 site list, is available from the author’s public repository (<https://github.com/vinhq dang/Artificial-Intelligence-in-Agriculture>).
1197

1198 Funding

1199 This research did not receive any specific grant from funding agencies in
1200 the public, commercial, or not-for-profit sectors.

1201 Appendix A. Statistical comparison

1202 References

- 1203 Adeh, E.H., Good, S.P., Calaf, M., Higgins, C.W., 2019. Solar pv power
1204 potential is greatest over croplands. Scientific Reports 9, 11442. doi:10.
1205 1038/s41598-019-47803-3.
- 1206 Allen, R.G., Pereira, L.S., Raes, D., Smith, M., 1998. Crop evapotranspira-
1207 tion: Guidelines for computing crop water requirements. FAO Irrigation
1208 and Drainage Paper 56. Food and Agriculture Organization of the United
1209 Nations. Rome.
- 1210 Amaducci, S., Yin, X., Colauzzi, M., 2018. Agrivoltaic systems to optimise
1211 land use for electric energy production. Applied Energy 220, 545–561.
1212 doi:10.1016/j.apenergy.2018.03.081.

1213 Andrade, F.H., Vega, C., Uhart, S., Cirilo, A., Cantarero, M., Valentinuz, O.,
1214 1999. Kernel number determination in maize. *Crop Science* 39, 453–459.
1215 doi:10.2135/cropsci1999.0011183X0039000200026x.

1216 Barron-Gafford, G.A., Pavao-Zuckerman, M.A., Minor, R.L., Sutter, L.F.,
1217 Barnett-Moreno, I., Blackett, D.T., Thompson, M., Dimond, K., Gerlak,
1218 A.K., Nabhan, G.P., Macknick, J.E., 2019. Agrivoltaics provide mutual
1219 benefits across the food–energy–water nexus in drylands. *Nature Sustain-*
1220 *ability* 2, 848–855. doi:10.1038/s41893-019-0364-5.

1221 Calafiore, G.C., Campi, M.C., 2006. The scenario approach to robust control
1222 design. *IEEE Transactions on Automatic Control* 51, 742–753. doi:10.
1223 1109/TAC.2006.875041.

1224 Campana, P.E., Stridh, B., Amaducci, S., Colauzzi, M., 2021. Optimisation
1225 of vertically mounted agrivoltaic systems. *Journal of Cleaner Production*
1226 325, 129091. doi:10.1016/j.jclepro.2021.129091.

1227 Campi, M.C., Garatti, S., 2008. The exact feasibility of randomized solutions
1228 of uncertain convex programs. *SIAM Journal on Optimization* 19, 1211–
1229 1230. doi:10.1137/07069821X.

1230 Collares-Pereira, M., Rabl, A., 1979. The average distribution of solar
1231 radiation-correlations between diffuse and hemispherical and between daily
1232 and hourly insolation values. *Solar Energy* 22, 155–164. doi:10.1016/
1233 0038-092x(79)90100-2.

1234 DIN, 2021. DIN SPEC 91434:2021-05: Agri-photovoltaic systems – require-
1235 ments for primary agricultural use. Beuth Verlag, Berlin.

1236 Dinesh, H., Pearce, J.M., 2016. The potential of agrivoltaic systems. *Renew-*
1237 *able and Sustainable Energy Reviews* 54, 299–308. doi:10.1016/j.rser.
1238 2015.10.024.

1239 Dupraz, C., 2023. Assessment of the ground coverage ratio of agrivoltaic
1240 systems as a proxy for potential crop productivity. *Agroforestry Systems*
1241 98, 2679–2696. doi:10.1007/s10457-023-00906-3.

1242 Dupraz, C., Marrou, H., Talbot, G., Dufour, L., Nogier, A., Ferard, Y., 2011.
1243 Combining solar photovoltaic panels and food crops for optimising land

1244 use: Towards new agrivoltaic schemes. *Renewable Energy* 36, 2725–2732.
1245 doi:10.1016/j.renene.2011.03.005.

1246 Elamri, Y., Cheviron, B., Lopez, J.M., Dejean, C., Belaud, G., 2018. Water
1247 budget and crop modelling for agrivoltaic systems: Application to irrigated
1248 lettuces. *Agricultural Water Management* 208, 440–453. doi:10.1016/j.
1249 agwat.2018.07.001.

1250 Ember, 2024. *Global Electricity Review 2024*. Technical Report. Ember.
1251 London.

1252 European Commission, Joint Research Centre, 2022. Photovoltaic geo-
1253 graphical information system (PVGIS), version 5.2. [https://re.jrc.ec.](https://re.jrc.ec.europa.eu/pvg_tools/)
1254 [europa.eu/pvg_tools/](https://re.jrc.ec.europa.eu/pvg_tools/).

1255 F. Holmgren, W., W. Hansen, C., A. Mikofski, M., 2018. pvlib python: a
1256 python package for modeling solar energy systems. *Journal of Open Source*
1257 *Software* 3, 884. doi:10.21105/joss.00884.

1258 Fischer, R.A., 1985. Number of kernels in wheat crops and the influence of
1259 solar radiation and temperature. *The Journal of Agricultural Science* 105,
1260 447–461. doi:10.1017/S0021859600056495.

1261 Gautron, R., Maillard, O.A., Preux, P., Corbeels, M., Sabbadin, R., 2022.
1262 Reinforcement learning for crop management support: Review, prospects
1263 and challenges. *Computers and Electronics in Agriculture* 200, 107182.
1264 doi:10.1016/j.compag.2022.107182.

1265 Gonocruz, R.A., Nakamura, R., Yoshino, K., Homma, M., Doi, T., Yoshida,
1266 Y., Tani, A., 2021. Analysis of the rice yield under an agrivoltaic
1267 system: A case study in japan. *Environments* 8, 65. doi:10.3390/
1268 environments8070065.

1269 Hassanpour Adeg, E., Selker, J.S., Higgins, C.W., 2018. Remarkable agri-
1270 voltaic influence on soil moisture, micrometeorology and water-use effi-
1271 ciency. *PLOS ONE* 13, e0203256. doi:10.1371/journal.pone.0203256.

1272 IPCC, 2021. *Climate Change 2021: The Physical Science Basis*. Contribution
1273 of Working Group I to the Sixth Assessment Report of the Intergovern-
1274 mental Panel on Climate Change. Cambridge University Press, Cambridge.
1275 doi:10.1017/9781009157896.

- 1276 IRENA, 2025. Renewable Capacity Statistics 2025. Technical Report. Inter-
1277 national Renewable Energy Agency. Abu Dhabi.
- 1278 Jägermeyr, J., Müller, C., Ruane, A.C., Elliott, J., Balkovic, J., Castillo, O.,
1279 Faye, B., Foster, I., Folberth, C., Franke, J.A., Fuchs, K., Guarin, J.R.,
1280 Heinke, J., Hoogenboom, G., Iizumi, T., Jain, A.K., Kelly, D., Khabarov,
1281 N., Lange, S., Lin, T.S., Liu, W., Mialyk, O., Minoli, S., Moyer, E.J.,
1282 Okada, M., Phillips, M., Porter, C., Rabin, S.S., Scheer, C., Schneider,
1283 J.M., Schyns, J.F., Skalsky, R., Smerald, A., Stella, T., Stephens, H.,
1284 Webber, H., Zabel, F., Rosenzweig, C., 2021. Climate impacts on global
1285 agriculture emerge earlier in new generation of climate and crop models.
1286 *Nature Food* 2, 873–885. doi:10.1038/s43016-021-00400-y.
- 1287 Karniadakis, G.E., Kevrekidis, I.G., Lu, L., Perdikaris, P., Wang, S., Yang,
1288 L., 2021. Physics-informed machine learning. *Nature Reviews Physics* 3,
1289 422–440. doi:10.1038/s42254-021-00314-5.
- 1290 Kingma, D.P., Ba, J., 2015. Adam: A method for stochastic optimization,
1291 in: *International Conference on Learning Representations*.
- 1292 Laub, M., Pataczek, L., Feuerbacher, A., Zikeli, S., Högy, P., 2022. Contrast-
1293 ing yield responses at varying levels of shade suggest different suitability
1294 of crops for dual land-use systems: a meta-analysis. *Agronomy for Sus-
1295 tainable Development* 42, 51. doi:10.1007/s13593-022-00783-7.
- 1296 Liu, B.Y., Jordan, R.C., 1960. The interrelationship and characteristic dis-
1297 tribution of direct, diffuse and total solar radiation. *Solar Energy* 4, 1–19.
1298 doi:10.1016/0038-092x(60)90062-1.
- 1299 Long, S.P., Ainsworth, E.A., Leakey, A.D.B., Nösberger, J., Ort, D.R., 2006.
1300 Food for thought: Lower-than-expected crop yield stimulation with ris-
1301 ing co₂ concentrations. *Science* 312, 1918–1921. doi:10.1126/science.
1302 1114722.
- 1303 Lorenzo, E., Narvarte, L., Muñoz, J., 2011. Tracking and back-tracking.
1304 *Progress in Photovoltaics: Research and Applications* 19, 747–753. doi:10.
1305 1002/pip.1085.
- 1306 Mamun, M.A.A., Dargusch, P., Wadley, D., Zulkarnain, N.A., Aziz, A.A.,
1307 2022. A review of research on agrivoltaic systems. *Renewable and Sustain-
1308 able Energy Reviews* 161, 112351. doi:10.1016/j.rser.2022.112351.

- 1309 Marchau, V.A.W.J., Walker, W.E., Bloemen, P.J.T.M., Popper, S.W. (Eds.),
1310 2019. Decision Making under Deep Uncertainty: From Theory to Practice.
1311 Springer. doi:10.1007/978-3-030-05252-2.
- 1312 Marrou, H., Dufour, L., Wery, J., 2013a. How does a shelter of solar panels
1313 influence water flows in a soil-crop system? *European Journal of Agron-*
1314 *omy* 50, 38–51. doi:10.1016/j.eja.2013.05.004.
- 1315 Marrou, H., Guilioni, L., Dufour, L., Dupraz, C., Wery, J., 2013b. Microcli-
1316 mate under agrivoltaic systems: Is crop growth rate affected in the partial
1317 shade of solar panels? *Agricultural and Forest Meteorology* 177, 117–132.
1318 doi:10.1016/j.agrformet.2013.04.012.
- 1319 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen,
1320 T., Lin, Z., Gimelshein, N., Antiga, L., et al., 2019. PyTorch: An impera-
1321 tive style, high-performance deep learning library, in: *Advances in Neural*
1322 *Information Processing Systems*, pp. 8024–8035.
- 1323 Ramos-Fuentes, I.A., Elamri, Y., Cheviron, B., Dejean, C., Belaud, G.,
1324 Fumey, D., 2023. Effects of shade and deficit irrigation on maize growth
1325 and development in fixed and dynamic agrivoltaic systems. *Agricultural*
1326 *Water Management* 280, 108187. doi:10.1016/j.agwat.2023.108187.
- 1327 Reichstein, M., Camps-Valls, G., Stevens, B., Jung, M., Denzler, J., Car-
1328 valhais, N., Prabhat, 2019. Deep learning and process understanding for
1329 data-driven Earth system science. *Nature* 566, 195–204. doi:10.1038/
1330 s41586-019-0912-1.
- 1331 Sekiyama, T., Nagashima, A., 2019. Solar sharing for both food and
1332 clean energy production: Performance of agrivoltaic systems for corn,
1333 a typical shade-intolerant crop. *Environments* 6, 65. doi:10.3390/
1334 environments6060065.
- 1335 Shen, C., Appling, A.P., Gentine, P., Bandai, T., Gupta, H., Tartakovsky, A.,
1336 Baity-Jesi, M., Fenicia, F., Kifer, D., Li, L., Liu, X., Ren, W., Zheng, Y.,
1337 Harman, C.J., Clark, M., Farthing, M., Feng, D., Kumar, P., Aboelyazeed,
1338 D., Rahmani, F., Song, Y., Beck, H.E., Bindas, T., Dwivedi, D., Fang,
1339 K., Höge, M., Rackauckas, C., Mohanty, B., Roy, T., Xu, C., Lawson,
1340 K., 2023. Differentiable modelling to unify machine learning and physical

- 1341 models for geosciences. *Nature Reviews Earth & Environment* 4, 552–567.
1342 doi:10.1038/s43017-023-00450-9.
- 1343 Siebert, S., Ewert, F., Eyshi Rezaei, E., Kage, H., Graß, R., 2014. Impact
1344 of heat stress on crop yield—on the importance of considering canopy
1345 temperature. *Environmental Research Letters* 9, 044012. doi:10.1088/
1346 1748-9326/9/4/044012.
- 1347 Sparks, A., 2018. nasapower: A nasa power global meteorology, surface solar
1348 energy and climatology data client for r. *Journal of Open Source Software*
1349 3, 1035. doi:10.21105/joss.01035.
- 1350 Trommsdorff, M., Kang, J., Reise, C., Schindele, S., Bopp, G., Ehmann, A.,
1351 Weselek, A., Högy, P., Obergfell, T., 2021. Combining food and energy
1352 production: Design of an agrivoltaic system applied in arable and vegetable
1353 farming in germany. *Renewable and Sustainable Energy Reviews* 140,
1354 110694. doi:10.1016/j.rser.2020.110694.
- 1355 Valle, B., Simonneau, T., Sourd, F., Pechier, P., Hamard, P., Frisson, T.,
1356 Ryckewaert, M., Christophe, A., 2017. Increasing the total productivity of
1357 a land by combining mobile photovoltaic panels and food crops. *Applied*
1358 *Energy* 206, 1495–1507. doi:10.1016/j.apenergy.2017.09.113.
- 1359 Waha, K., Dietrich, J.P., Portmann, F.T., Siebert, S., Thornton, P.K., Bon-
1360 deau, A., Herrero, M., 2020. Multiple cropping systems of the world
1361 and the potential for increasing cropping intensity. *Global Environmental*
1362 *Change* 64, 102131. doi:10.1016/j.gloenvcha.2020.102131.
- 1363 Webber, H., White, J.W., Kimball, B.A., Ewert, F., Asseng, S., Eyshi Rezaei,
1364 E., Pinter, P.J., Hatfield, J.L., Reynolds, M.P., Ababaei, B., Bindi, M.,
1365 Doltra, J., Ferrise, R., Kage, H., Kassie, B.T., Kersebaum, K.C., Luig,
1366 A., Olesen, J.E., Semenov, M.A., Stratonovitch, P., Ratjen, A.M., LaM-
1367 orte, R.L., Leavitt, S.W., Hunsaker, D.J., Wall, G.W., Martre, P., 2018.
1368 Physical robustness of canopy temperature models for crop heat stress
1369 simulation across environments and production conditions. *Field Crops*
1370 *Research* 216, 75–88. doi:10.1016/j.fcr.2017.11.005.
- 1371 Weselek, A., Bauerle, A., Hartung, J., Zikeli, S., Lewandowski, I., Högy, P.,
1372 2021. Agrivoltaic system impacts on microclimate and yield of different

1373 crops within an organic crop rotation in a temperate climate. *Agronomy*
1374 for Sustainable Development 41, 59. doi:10.1007/s13593-021-00714-y.

1375 Weselek, A., Ehmann, A., Zikeli, S., Lewandowski, I., Schindele, S., Högy,
1376 P., 2019. Agrophotovoltaic systems: applications, challenges, and op-
1377 portunities. a review. *Agronomy for Sustainable Development* 39, 35.
1378 doi:10.1007/s13593-019-0581-3.

1379 Willockx, B., Lavaert, C., Cappelle, J., 2023. Performance evaluation of
1380 vertical bifacial and single-axis tracked agrivoltaic systems on arable land.
1381 *Renewable Energy* 217, 119181. doi:10.1016/j.renene.2023.119181.

1382 Woli, P., Jones, J.W., Ingram, K.T., Fraisse, C.W., 2012. Agricultural
1383 reference index for drought (arid). *Agronomy Journal* 104, 287–300.
1384 doi:10.2134/agronj2011.0286.

1385 Yu, Q., You, L., Wood-Sichra, U., Ru, Y., Joglekar, A.K.B., Fritz, S., Xiong,
1386 W., Lu, M., Wu, W., Yang, P., 2020. A cultivated planet in 2010 – part 2:
1387 The global gridded agricultural-production maps. *Earth System Science*
1388 *Data* 12, 3545–3572. doi:10.5194/essd-12-3545-2020.

1389 Zhang, L., Geng, X., Ding, H., Cao, K., Wang, L., Chen, H., Deng, L., Wu,
1390 C., Xiao, M., Bao, E., 2026. Multi-objective optimization of photovoltaic
1391 array layouts on farmland via a hybrid algorithm framework for enhanced
1392 agricultural production and energy generation. *Artificial Intelligence in*
1393 *Agriculture* doi:10.1016/j.aiia.2026.09.002.

1394 Zhao, C., Liu, B., Piao, S., Wang, X., Lobell, D.B., Huang, Y., Huang, M.,
1395 Yao, Y., Bassu, S., Ciais, P., Durand, J.L., Elliott, J., Ewert, F., Janssens,
1396 I.A., Li, T., Lin, E., Liu, Q., Martre, P., Müller, C., Peng, S., Peñuelas,
1397 J., Ruane, A.C., Wallach, D., Wang, T., Wu, D., Liu, Z., Zhu, Y., Zhu,
1398 Z., Asseng, S., 2017. Temperature increase reduces global yields of major
1399 crops in four independent estimates. *Proceedings of the National Academy*
1400 *of Sciences* 114, 9326–9331. doi:10.1073/pnas.1701762114.

1401 Zhao, C., Liu, B., Xiao, L., Hoogenboom, G., Boote, K.J., Kassie, B.T.,
1402 Pavan, W., Shelia, V., Kim, K.S., Hernandez-Ochoa, I.M., Wallach, D.,
1403 Porter, C.H., Stockle, C.O., Zhu, Y., Asseng, S., 2019. A simple crop
1404 model. *European Journal of Agronomy* 104, 97–106. doi:10.1016/j.eja.
1405 2019.01.009.

Table 2: Held-out performance (92 test sites, seasons 2015–2019, canopy cooling $\kappa = 1^\circ\text{C}$). The density (GCR) of each site is the largest value on the grid 0.05, 0.075, ..., 0.5 at which the method’s mean retention in the site’s 2001–2014 seasons is ≥ 0.9 . E : electricity ($\text{MWh ha}^{-1} \text{ yr}^{-1}$); e : electricity relative to a conventional PV plant; r : mean yield retention; r_{prod} : production-weighted retention; C_s and C_{site} : share of seasons with $r \geq 0.9$ and share of sites whose test-season mean $r \geq 0.9$ (%); $\text{LER} = r + e$.

Method	GCR	E	e	r	r_{prod}	C_s	C_{site}	LER
<i>Baseline</i>								
AV static	0.06	217	0.16	0.917	0.914	93	96	1.07
Seasonal sharing	0.17	397	0.28	0.904	0.902	69	72	1.19
Phenology rule	0.17	465	0.33	0.904	0.902	69	75	1.24
Stress rule	0.13	387	0.28	0.907	0.907	76	82	1.18
SALS (ours)	0.17	510	0.36	0.905	0.902	70	74	1.27
<i>+1.5 °C</i>								
AV static	0.06	221	0.16	0.916	0.914	93	92	1.08
Seasonal sharing	0.16	393	0.28	0.905	0.903	72	74	1.19
Phenology rule	0.17	478	0.34	0.905	0.902	68	75	1.25
Stress rule	0.12	377	0.27	0.908	0.908	73	77	1.18
SALS (ours)	0.17	518	0.37	0.906	0.903	68	75	1.28
<i>+2 °C</i>								
AV static	0.06	223	0.16	0.916	0.914	92	93	1.08
Seasonal sharing	0.17	394	0.28	0.905	0.904	72	76	1.19
Phenology rule	0.17	483	0.35	0.905	0.903	69	74	1.25
Stress rule	0.12	370	0.26	0.909	0.909	75	78	1.17
SALS (ours)	0.17	521	0.38	0.907	0.904	69	75	1.28
<i>+3 °C</i>								
AV static	0.06	231	0.17	0.915	0.914	88	89	1.08
Seasonal sharing	0.17	397	0.29	0.906	0.904	69	75	1.19
Phenology rule	0.18	499	0.36	0.905	0.902	67	72	1.26
Stress rule	0.11	366	0.26	0.908	0.908	74	77	1.17
SALS (ours)	0.18	540	0.39	0.907	0.904	68	75	1.30

Table 3: LER of the controllers at different yield-retention floors ρ (main split; the historical design step is repeated for each floor).

Method	Baseline climate				+3 °C			
	0.80	0.85	0.90	0.95	0.80	0.85	0.90	0.95
AV static	1.159	1.115	1.072	1.011	1.174	1.128	1.083	1.013
Stress rule	1.380	1.282	1.182	1.077	1.364	1.269	1.172	1.071
Phenology rule	1.467	1.356	1.236	1.108	1.513	1.395	1.265	1.122
SALS (trained at $\rho = 0.9$)	1.505	1.392	1.270	1.141	1.553	1.432	1.297	1.154
SALS conditioned on ρ	1.504	1.391	1.269	1.141	1.553	1.431	1.296	1.154

Table 4: Sensitivity of the controller comparison to the rotation limit of the trackers and to the harvest-index penalty (LER; main split, every sixth held-out pair, 21 pairs, main shade response, history-based design on a GCR grid with step 0.05). SALS was trained for the main setting. Fixed arrays (no tracking; flat and tilted by 30° from the horizontal about the north–south axis) reach LER 1.048 and 1.043 in the same sample, against 1.053 for static trackers.

Setting	Static tracker	Phenology rule	Tuned rule	SALS
Main (limit 60°, window 0.45–0.65)	1.053	1.170	1.180	1.189
Rotation limit 45°	1.057	1.138	1.144	1.144
Rotation limit 80°	1.053	1.206	1.219	1.224
Window 0.30–0.50	1.053	1.169	1.176	1.185
Window 0.60–0.80	1.053	1.171	1.181	1.182
No harvest-index penalty	1.109	1.328	1.347	1.353

Table 5: Rules and the learned controller pooled over the three spatial splits (76 distinct crop-site pairs, every third held-out pair; baseline climate; main shade response; history-based design on a GCR grid with step 0.05, coarser than the 0.025 grid of the main analysis). The tuned phenology rule shares light while thermal time lies between 0.15 and 1.0 of the season requirement (window chosen on training-block seasons of split 0; the written rule uses 0.05 to 1.0). The optimised ramp schedule has five parameters (start 0.17, end 0.93, ramp widths 0.02, level 0.99) found by CMA-ES on the same training seasons.

Method	GCR	E/E_{PV}	Retention	Seasons ≥ 0.9 (%)	Pairs ≥ 0.9 (%)	LER
AV static	0.059	0.153	0.922	95	96	1.076
Seasonal sharing	0.151	0.249	0.915	84	91	1.164
Stress rule	0.118	0.255	0.918	83	88	1.173
Phenology rule	0.154	0.292	0.916	85	91	1.207
Tuned phenology rule	0.161	0.317	0.916	84	89	1.233
Optimised ramp schedule	0.155	0.322	0.916	83	87	1.238
SALS	0.161	0.335	0.916	83	87	1.251

Table 6: Simulated relative yield at the shading of published field trials (static arrays, training-block sites). Observed values: rice, single trial; wheat and potato, the worse of the two seasons (the better season gave +2.7% and +11%); maize, mean of the three seasons of the dynamic system (DAV) and single values for the other two systems. The same trials enter the calibration of Section 2.6, so the comparison is not independent.

Trial	Shading	Observed	Conservative	Calibrated
Rice, Gonocruz et al. (2021)	0.25	0.84	0.77	0.79
Wheat, Weselek et al. (2021)	0.30	0.81	0.66	0.74
Potato, Weselek et al. (2021)	0.30	0.82	0.72	0.72
Maize, Ramos-Fuentes et al. (2023) (DAV)	0.35	0.83	0.70	0.70
Maize, Ramos-Fuentes et al. (2023) (half)	0.30	0.91	0.75	0.75
Maize, Ramos-Fuentes et al. (2023) (full)	0.50	0.70	0.54	0.54

Table 7: Sensitivity of the deployable light-sharing results to the crop shade response (rule controllers, baseline climate, floor $\rho = 0.9$, split 0, every fourth held-out site, 22 sites). Conservative: proportional response (main analysis); calibrated: radiation-use compensation fitted to Laub et al. (2022); field-fitted: fitted to the field trials, maize harvest index unaffected.

Shade response	Method	GCR	E/E_{PV}	Retention	LER
conservative	AV static	0.06	0.15	0.92	1.07
	Phenology rule	0.17	0.33	0.90	1.23
	Seasonal sharing	0.16	0.28	0.90	1.18
	Stress rule	0.11	0.24	0.91	1.15
calibrated	AV static	0.07	0.18	0.91	1.09
	Phenology rule	0.19	0.37	0.90	1.28
	Seasonal sharing	0.19	0.32	0.90	1.22
	Stress rule	0.13	0.29	0.91	1.20
field-fitted	AV static	0.16	0.41	0.91	1.32
	Phenology rule	0.37	0.70	0.92	1.62
	Seasonal sharing	0.37	0.61	0.92	1.53
	Stress rule	0.29	0.63	0.90	1.53

Table 8: Predicting relative yield under shade from 428 observations of 58 studies (Laub et al., 2022): leave-one-study-out cross-validation (every study held out once; observations with $s > 0$). Covariates: irradiation, climate zone, shade type, intercropping.

Model	RMSE	MAE	Median AE	Mean MAE per study
Proportional loss, $y = 1 - s$	0.430	0.305	0.210	0.224
One curve for all crops	0.372	0.265	0.193	0.208
Curve per crop type, no pooling	0.313	0.208	0.141	0.165
Curve per crop type, partial pooling	0.312	0.210	0.150	0.162
Partial pooling + covariates	0.311	0.206	0.148	0.161
Gradient-boosted trees	0.321	0.219	0.162	0.184

Table 9: Leave-one-study-out prediction error (RMSE of relative yield, observations with $s > 0$) by crop type: proportional loss, hierarchical curve with partial pooling, and gradient-boosted trees. The last columns give the difference hierarchical minus proportional with a 95% cluster-bootstrap interval over studies (negative: hierarchical better).

Group	Obs./studies	Proportional	Hierarchical	Boosted t
C3 cereals	30/10	0.214	0.187	0.205
Maize	40/10	0.225	0.235	0.248
Grain legumes	58/14	0.215	0.225	0.234
Tubers/root crops	10/2	0.250	0.254	0.257
Four crop types used by the simulator	138/30	0.220	0.222	0.234
Other types (berries, forages, fruits, vegetables)	201/29	0.528	0.361	0.369
All types	339/58	0.430	0.312	0.321

Table 10: Relative yield at 20% and 40% shading: data-driven model (median and 80% bootstrap interval over studies; crop-type curve, covariates at their reference values: mean irradiation, temperate zone, no solar-panel shade, no intercropping) and simulator under the two calibrated ensembles (mean of members; E1 fitted to hand-read curves, E2 fitted to the data-driven curves). The data-driven curve of rice is that of C3 cereals (no rice in the data set). Observations and studies: C3 cereals 41/10, maize 51/10, grain legumes 76/14, tubers 12/2.

Crop	Shading	Data-driven	Simulator, E1	Simulator, E2	Proportional
Wheat	20%	0.93 [0.88, 0.96]	0.88	0.92	0.80
Wheat	40%	0.72 [0.59, 0.78]	0.70	0.77	0.60
Rice	20%	0.93 [0.88, 0.96]	0.85	0.90	0.80
Rice	40%	0.72 [0.59, 0.78]	0.66	0.73	0.60
Maize	20%	0.73 [0.58, 0.86]	0.86	0.90	0.80
Maize	40%	0.42 [0.27, 0.59]	0.69	0.75	0.60
Soybean	20%	0.71 [0.65, 0.79]	0.77	0.74	0.80
Soybean	40%	0.40 [0.34, 0.53]	0.53	0.48	0.60
Potato	20%	0.93 [0.92, 0.97]	0.88	0.92	0.80
Potato	40%	0.71 [0.67, 0.83]	0.73	0.79	0.60

Table 11: Shade-response parameters of ensemble E2, calibrated by gradient descent through the simulator to the data-driven curves and the field trials (12 randomised fits; mean $\pm$ s.d.): radiation-use compensation c and multiplier h on the flowering-window harvest-index penalty (1 = assumed value).

Crop	c	h multiplier
Wheat	0.91 ± 0.31	0.74 ± 0.53
Rice	0.61 ± 0.25	0.83 ± 0.33
Maize	0.14 ± 0.11	0.16 ± 0.13
Soybean	0.05 ± 0.09	1.66 ± 0.38
Potato	0.62 ± 0.20	0.90 ± 0.50

Table 12: Land-equivalent ratio of rules and learned controllers under four true shade responses, with the array density designed under the true response (split 0, baseline climate, floor $\rho = 0.9$, every sixth held-out pair, 21 pairs, coarse GCR grid, one training seed; season means, so differences differ slightly from the pair-level intervals in the text). SALS is trained on the proportional response; robust controllers are trained on random members of ensemble E1, E2 or E3 (adaptive: with the measured biomass ratio). The proportional and harsh responses lie outside the ensembles.

Method	Proportional	Calibrated	Field-fitted	Harsh
AV static	1.052	1.059	1.254	1.037
Seasonal sharing	1.133	1.156	1.385	1.097
Stress rule	1.129	1.160	1.406	1.086
Feedback rule	1.139	1.164	1.392	1.089
Phenology rule	1.164	1.190	1.456	1.122
Tuned phenology rule	1.170	1.199	1.472	1.129
Optimised ramp schedule	1.170	1.199	1.477	1.122
SALS	1.178	1.209	1.489	1.132
SALS-robust, E1	1.179	1.204	1.491	1.127
SALS-adaptive, E1	1.176	1.202	1.495	1.124
SALS-robust, E2	1.165	1.192	1.483	1.126
SALS-robust, E3	1.163	1.191	1.481	1.121
SALS-adaptive, E3	1.164	1.191	1.475	1.118

Table 13: Array design under shade-response uncertainty, phenology rule, three spatial splits pooled (76 distinct crop–site pairs, every third held-out pair). Truth: each of five members of the ensemble in turn; the mean response and the member-based rules use the other four members (leave-one-out, so the truth never enters the design). E2: ensemble calibrated to the learned curves and field trials; E3: same procedure with target errors of 0.2 on the learned curves. Pairs: share with mean retention ≥ 0.9 ; shortfall: mean of $\max(0, 0.9 - \text{retention})$ per pair.

Design	E2		LER	E3		
	Pairs ≥ 0.9 (%)	Shortfall		Pairs ≥ 0.9 (%)	Shortfall	
oracle (true response)	87	0.0025	1.486	85	0.0029	1
mean response	70	0.0062	1.493	71	0.0093	1
floor met in 50% of other members	65	0.0084	1.510	67	0.0125	1
floor met in 75% of other members	78	0.0045	1.462	79	0.0083	1
floor met in all other members	89	0.0019	1.412	88	0.0032	1
floor 0.92 under mean response	88	0.0022	1.426	85	0.0051	1
floor 0.94 under mean response	97	0.0007	1.345	92	0.0027	1
floor 0.96 under mean response	99	0.0003	1.251	96	0.0015	1

Table 14: Design rules outside the ensemble: share of crop–site pairs (%) with mean retention ≥ 0.9 under four fixed responses (P proportional, C calibrated to the meta-analysis, F field-fitted, H harsh), phenology rule, three splits pooled (76 distinct pairs). Designs use the whole ensemble E2 or E3 (five members and their mean); the oracle uses the true response. LER in the last column is the mean over the four responses.

Design	E2				E3				LER	
	P	C	F	H	P	C	F	H	E2	E3
oracle (true response)	91	89	87	84	91	89	87	84	1.283	1.283
mean response	13	13	68	12	13	17	79	7	1.435	1.448
floor met in 75% of members	16	16	80	12	22	25	93	13	1.384	1.397
floor met in all members	17	17	80	12	29	32	93	16	1.348	1.336
floor 0.92 under mean response	17	20	86	12	29	33	91	20	1.376	1.399
floor 0.94 under mean response	45	47	97	34	45	45	97	29	1.303	1.345
floor 0.96 under mean response	63	78	99	54	58	63	99	51	1.220	1.264

Table 15: Assumed economic parameters (illustrative, not data). Ranges used in the sensitivity analysis are in parentheses.

Parameter	Value
Electricity price	60 USD MWh ⁻¹ (30–120)
Capex, AV single-axis tracker plant	1.0 USD Wp ⁻¹ (0.7–1.5)
Capex, conventional tracker plant	0.75 USD Wp ⁻¹
Discount rate, lifetime, O&M	6% (4–8%), 25 years, 1.5% of capex per year
Installed capacity	2.1 MWp per hectare of module area share g (0.21 kWp m ⁻²)
Crop price per tonne of dry matter	wheat 290, rice 400, maize 250, soybean 520, potato 1000 USD (0.

Table 16: Illustrative annual economics per hectare (76 crop–site pairs, three splits, main shade response, baseline climate; parameters in Table 15). Net value is relative to the open field (crop costs unchanged) in kUSD ha⁻¹ yr⁻¹ at three electricity prices; break-even is the price at which the net value is zero. The conventional plant replaces the crop.

Method	E (MWh)	Capex annuity + O&M	Crop loss	Net at 60	Net at 90	Net
AV static	200	11.5	0.27	0.3	6.3	1
Seasonal sharing	326	29.6	0.31	-10.3	-0.5	9
Stress rule	336	23.2	0.30	-3.3	6.8	1
Phenology rule	382	30.1	0.30	-7.5	4.0	1
Tuned phenology rule	415	31.6	0.31	-7.0	5.5	1
SALS	438	31.6	0.31	-5.6	7.6	2
Conventional PV plant	1306	58.7	3.60	16.0	55.2	9

Table 17: One-at-a-time sensitivity of the net value (kUSD ha⁻¹ yr⁻¹, relative to the open field) at 90 USD MWh⁻¹, densities designed for the floor as in the main analysis. In the capex rows the capex of the conventional plant is set to 0.75 times that of the agrivoltaic plant.

Parameter	Value	Static arrays	Tuned rule	SALS	Conventional plant
Base		6.3	5.5	7.6	55.2
AV capex (USD/Wp)	0.7	9.7	14.9	17.1	72.8
AV capex (USD/Wp)	1.5	0.6	-10.3	-8.2	25.8
Discount rate	4%	8.0	10.3	12.4	64.1
Discount rate	8%	4.4	0.2	2.4	45.4
Crop prices	0.5×	6.4	5.6	7.7	57.0
Crop prices	2.0×	6.0	5.2	7.3	51.6
Electricity price	60	0.3	-7.0	-5.6	16.0
Electricity price	120	12.3	17.9	20.7	94.4

Table A.18: Paired comparison of SALS with each baseline on the 92 held-out sites of the main split (unit of analysis: site mean over test seasons). ΔLER : mean difference (SALS minus baseline) with cluster-bootstrap 95% CI; p : two-sided Wilcoxon signed-rank test; % better: share of sites where SALS has the higher LER; Δr : difference in mean yield retention with its Wilcoxon p -value.

Climate	Baseline	ΔLER [95% CI]	p	% better	Δr	p	Sites
Baseline	AV static	+0.205 [+0.176, +0.236]	$< 10^{-4}$	100	-0.011	$< 10^{-4}$	92
Baseline	Seasonal sharing	+0.088 [+0.068, +0.109]	$< 10^{-4}$	97	+0.001	0.922	92
Baseline	Phenology rule	+0.038 [+0.025, +0.053]	$< 10^{-4}$	84	+0.001	0.723	92
Baseline	Stress rule	+0.096 [+0.071, +0.124]	$< 10^{-4}$	93	-0.002	0.055	92
+1.5 °C	AV static	+0.208 [+0.180, +0.237]	$< 10^{-4}$	100	-0.010	$< 10^{-4}$	92
+1.5 °C	Seasonal sharing	+0.095 [+0.077, +0.114]	$< 10^{-4}$	98	+0.001	0.812	92
+1.5 °C	Phenology rule	+0.033 [+0.022, +0.045]	$< 10^{-4}$	86	+0.002	0.855	92
+1.5 °C	Stress rule	+0.108 [+0.083, +0.134]	$< 10^{-4}$	98	-0.001	0.041	92
+2 °C	AV static	+0.209 [+0.181, +0.238]	$< 10^{-4}$	100	-0.010	$< 10^{-4}$	92
+2 °C	Seasonal sharing	+0.097 [+0.079, +0.115]	$< 10^{-4}$	97	+0.001	0.876	92
+2 °C	Phenology rule	+0.032 [+0.022, +0.044]	$< 10^{-4}$	85	+0.002	0.496	92
+2 °C	Stress rule	+0.115 [+0.089, +0.142]	$< 10^{-4}$	98	-0.003	0.016	92
+3 °C	AV static	+0.219 [+0.190, +0.250]	$< 10^{-4}$	100	-0.009	$< 10^{-4}$	92
+3 °C	Seasonal sharing	+0.108 [+0.089, +0.129]	$< 10^{-4}$	96	+0.001	0.935	92
+3 °C	Phenology rule	+0.034 [+0.022, +0.046]	$< 10^{-4}$	86	+0.002	0.330	92
+3 °C	Stress rule	+0.131 [+0.102, +0.162]	$< 10^{-4}$	98	-0.002	0.142	92